



Oracle (Active) Data Guard

Master Slide Deck - Updated 2024.09.02

Ludovico Caldara

Senior Principal Product Manager

Oracle Database High Availability (HA), Scalability and
Maximum Availability Architecture (MAA) Team

@ludodba



<http://www.linkedin.com/in/ludovicocaldara>

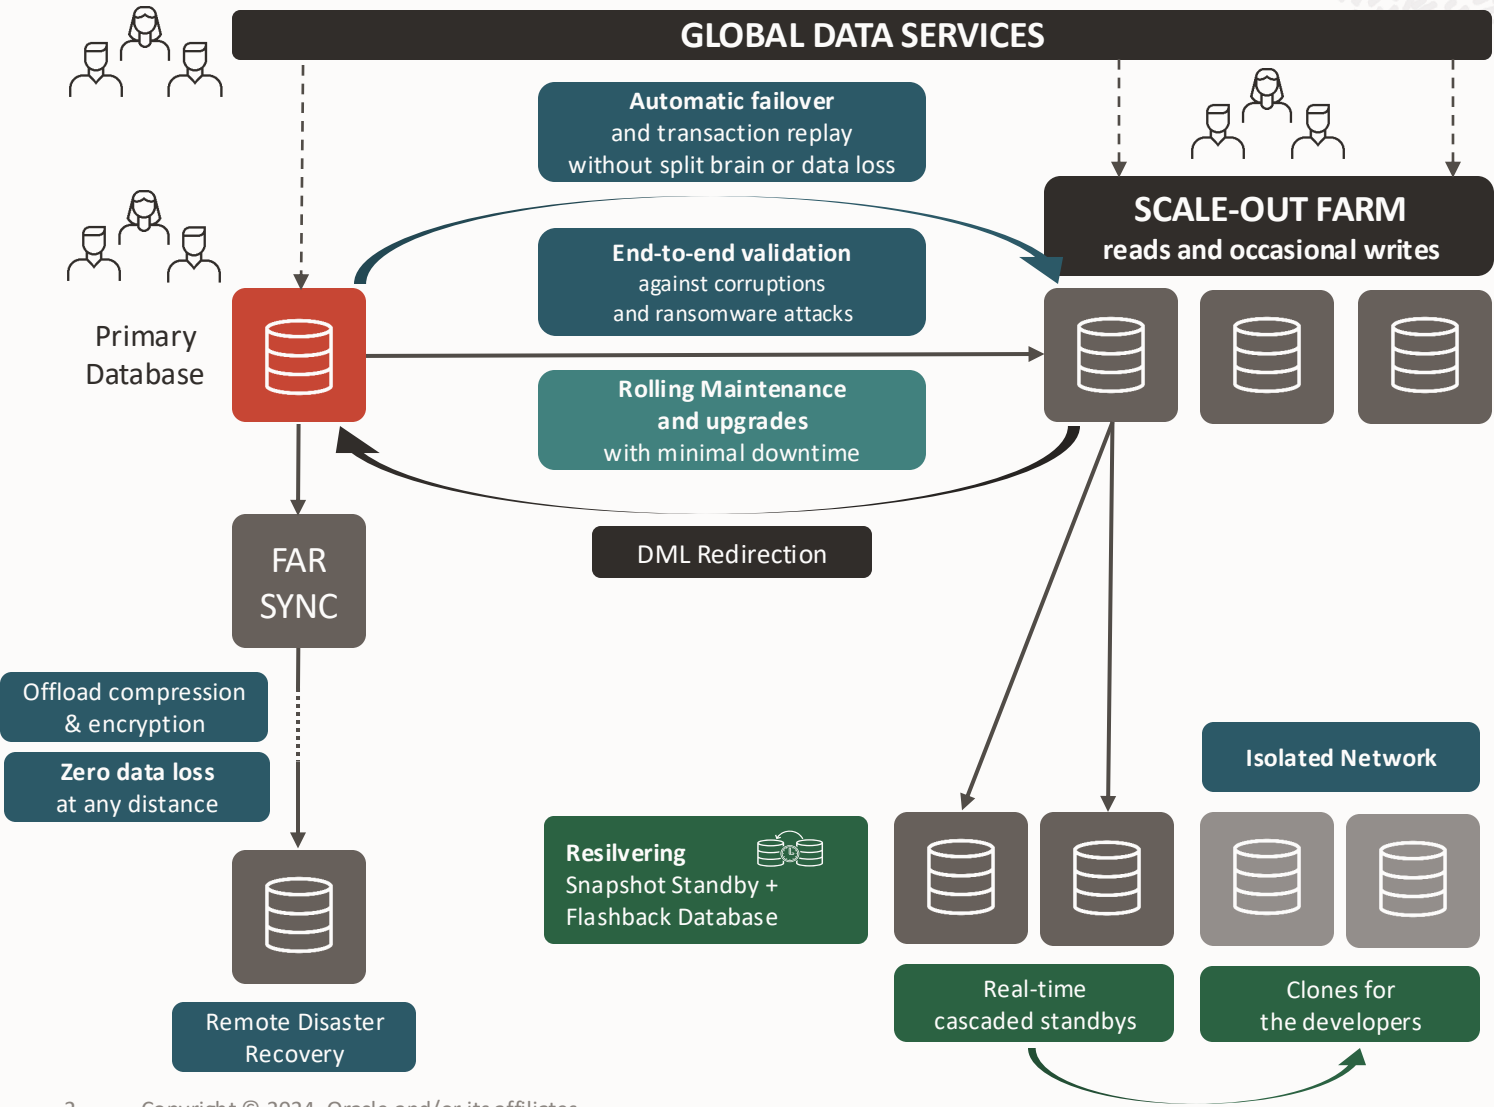


www.ludovicocaldara.net



Oracle Active Data Guard

Come for the data protection, stay for the data management












- UNMATCHED DATA PROTECTION
- OPERATIONAL CONTINUITY
- MODERN DEVELOPMENT PLATFORM
- SCALE AS YOU GROW





Active Data Guard 23^{ai} furtherly enhances the experience

Enhanced performance, observability, and manageability



-  Automatic preparation of the primary
-  PL/SQL APIs for better automation
-  SQLcl command-line integration
-  REST APIs for easier DevOps integration
-  4 new SQL views to monitor Data Guard
-  Support for mixed Transparent Data Encryption
-  Automatic temp file creation on the standby
-  Simplified AWR snapshots on the standby DB



Oracle (Active) Data Guard & MAA

Challenges of deploying highly available systems



Cost and complexity



Lack of skills



Risk of failure

Impact of database downtime



\$350K

average cost of downtime
per hour



\$10M

average cost of unplanned data center
outage or disaster



87 hours

average amount of downtime
per year



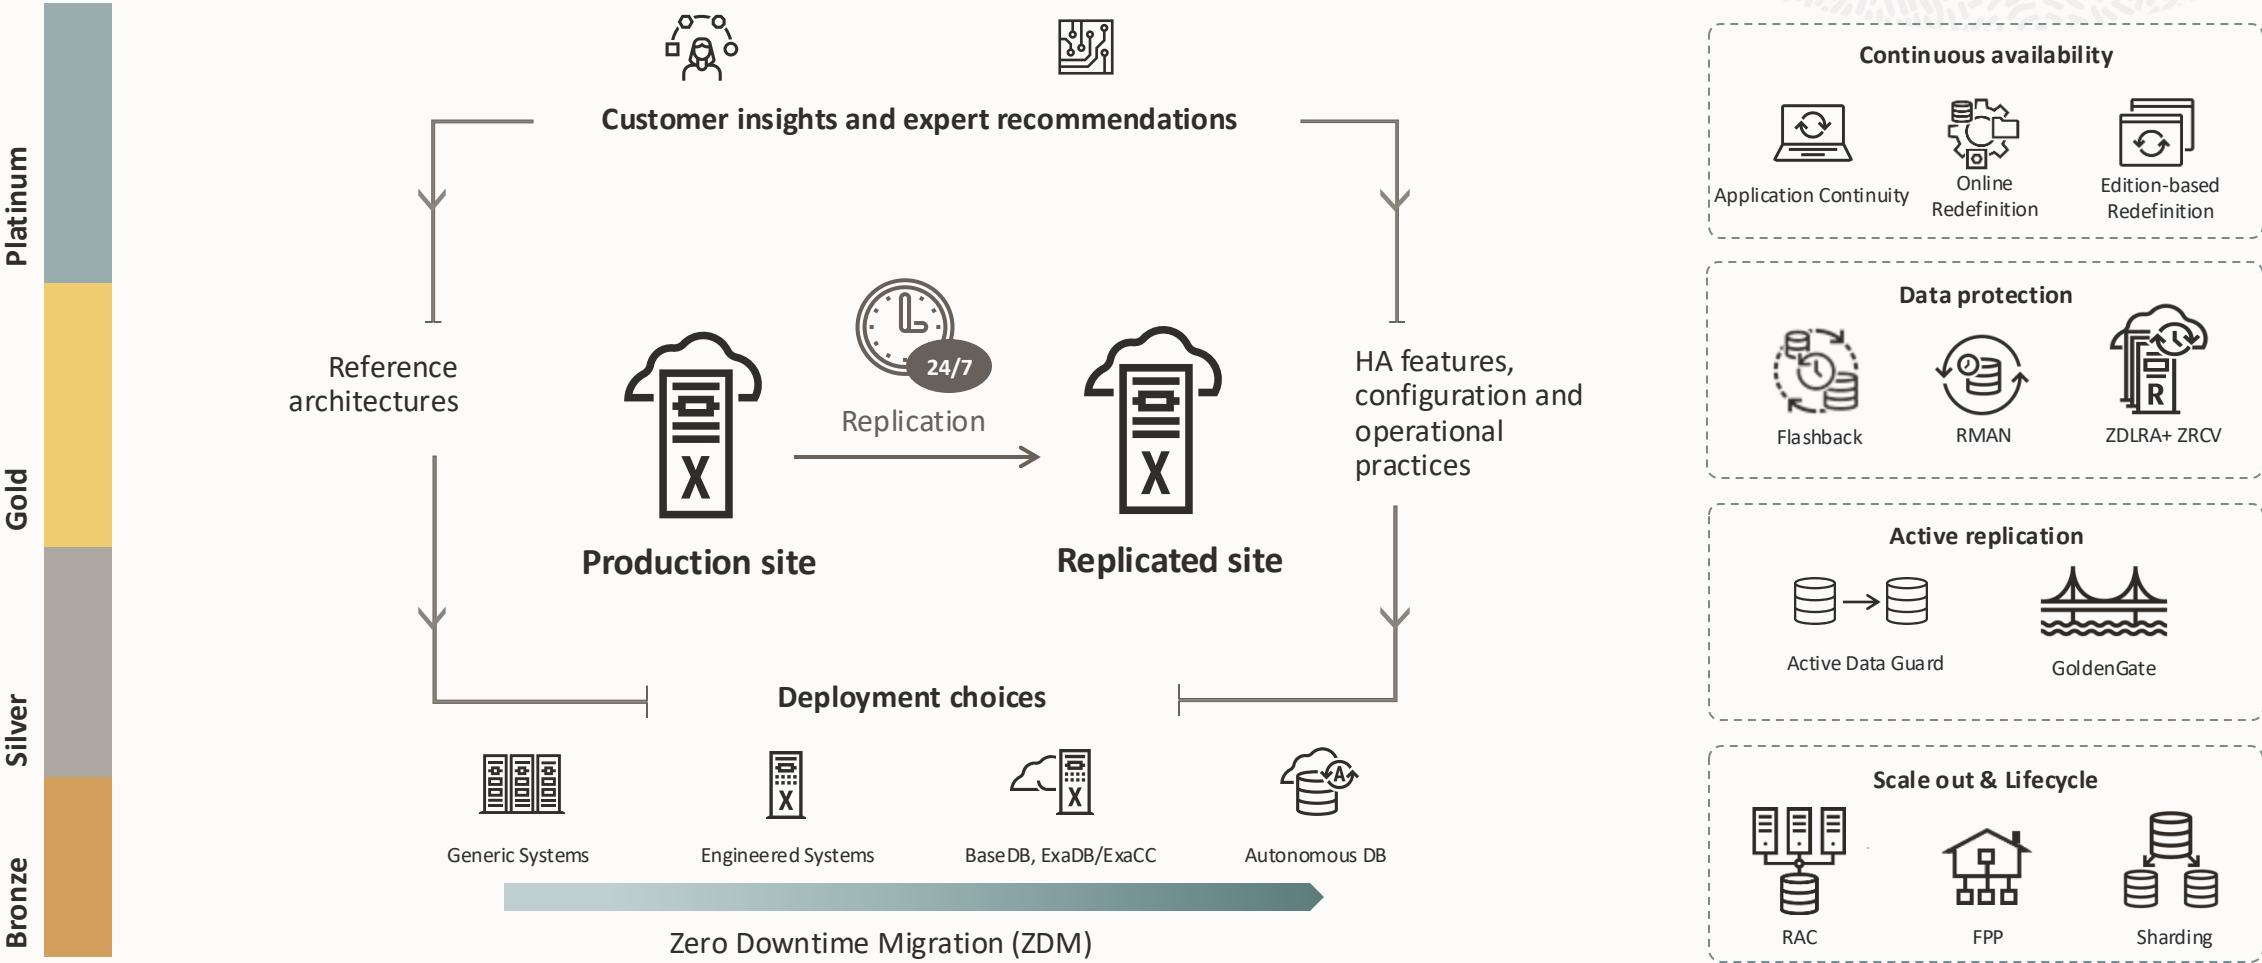
91%

percentage of companies that have
experienced an unplanned data center
outage in the last 24 months



Oracle Maximum Availability Architecture (MAA)





Standardized Reference Architectures for Never-Down Deployments





MAA reference architectures

Availability service levels




Bronze	Silver	Gold	Platinum
<div>Dev, test, prod</div> <div>Single instance DB</div> <div>Restartable</div> <div>Backup/restore</div> <div></div>	<div>Prod/departamental</div> <div>Bronze +</div> <div>Database HA with RAC</div> <div>Application continuity</div> <div>Sharding (optional)</div> <div></div>	<div>Business critical</div> <div>Silver +</div> <div>DB replication with Active Data Guard</div> <div></div>	<div>Mission critical</div> <div>Gold +</div> <div>GoldenGate</div> <div>Edition-based redefinition</div> <div></div>





Oracle Data Guard Overview

Oracle Data Guard

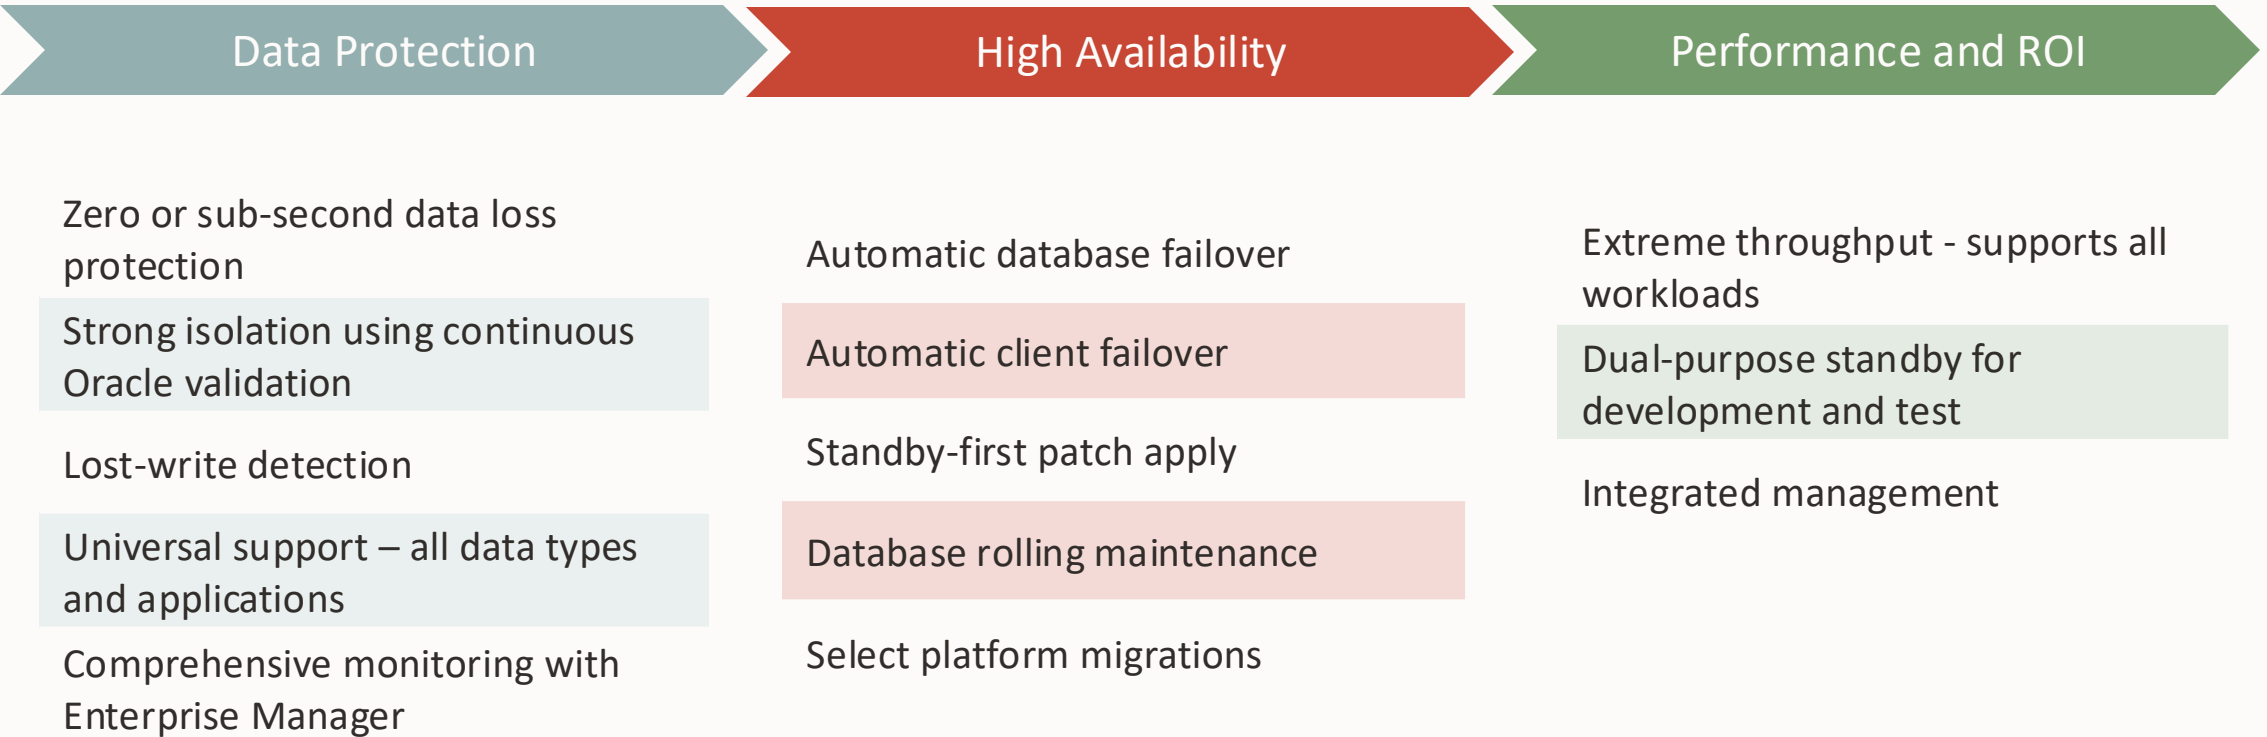


- **Disaster Recovery (included with DB EE)**
 - License primary and secondary sites
- **Active-passive**
 - Standby is used only for failovers
- **Automatic failover to Standby site**
- **Zero / near-zero data loss**
- **Continuous data validation**
- **Simple migrations and upgrades**

<https://www.oracle.com/database/technologies/high-availability/dataguard-activedataguard-demos.html>

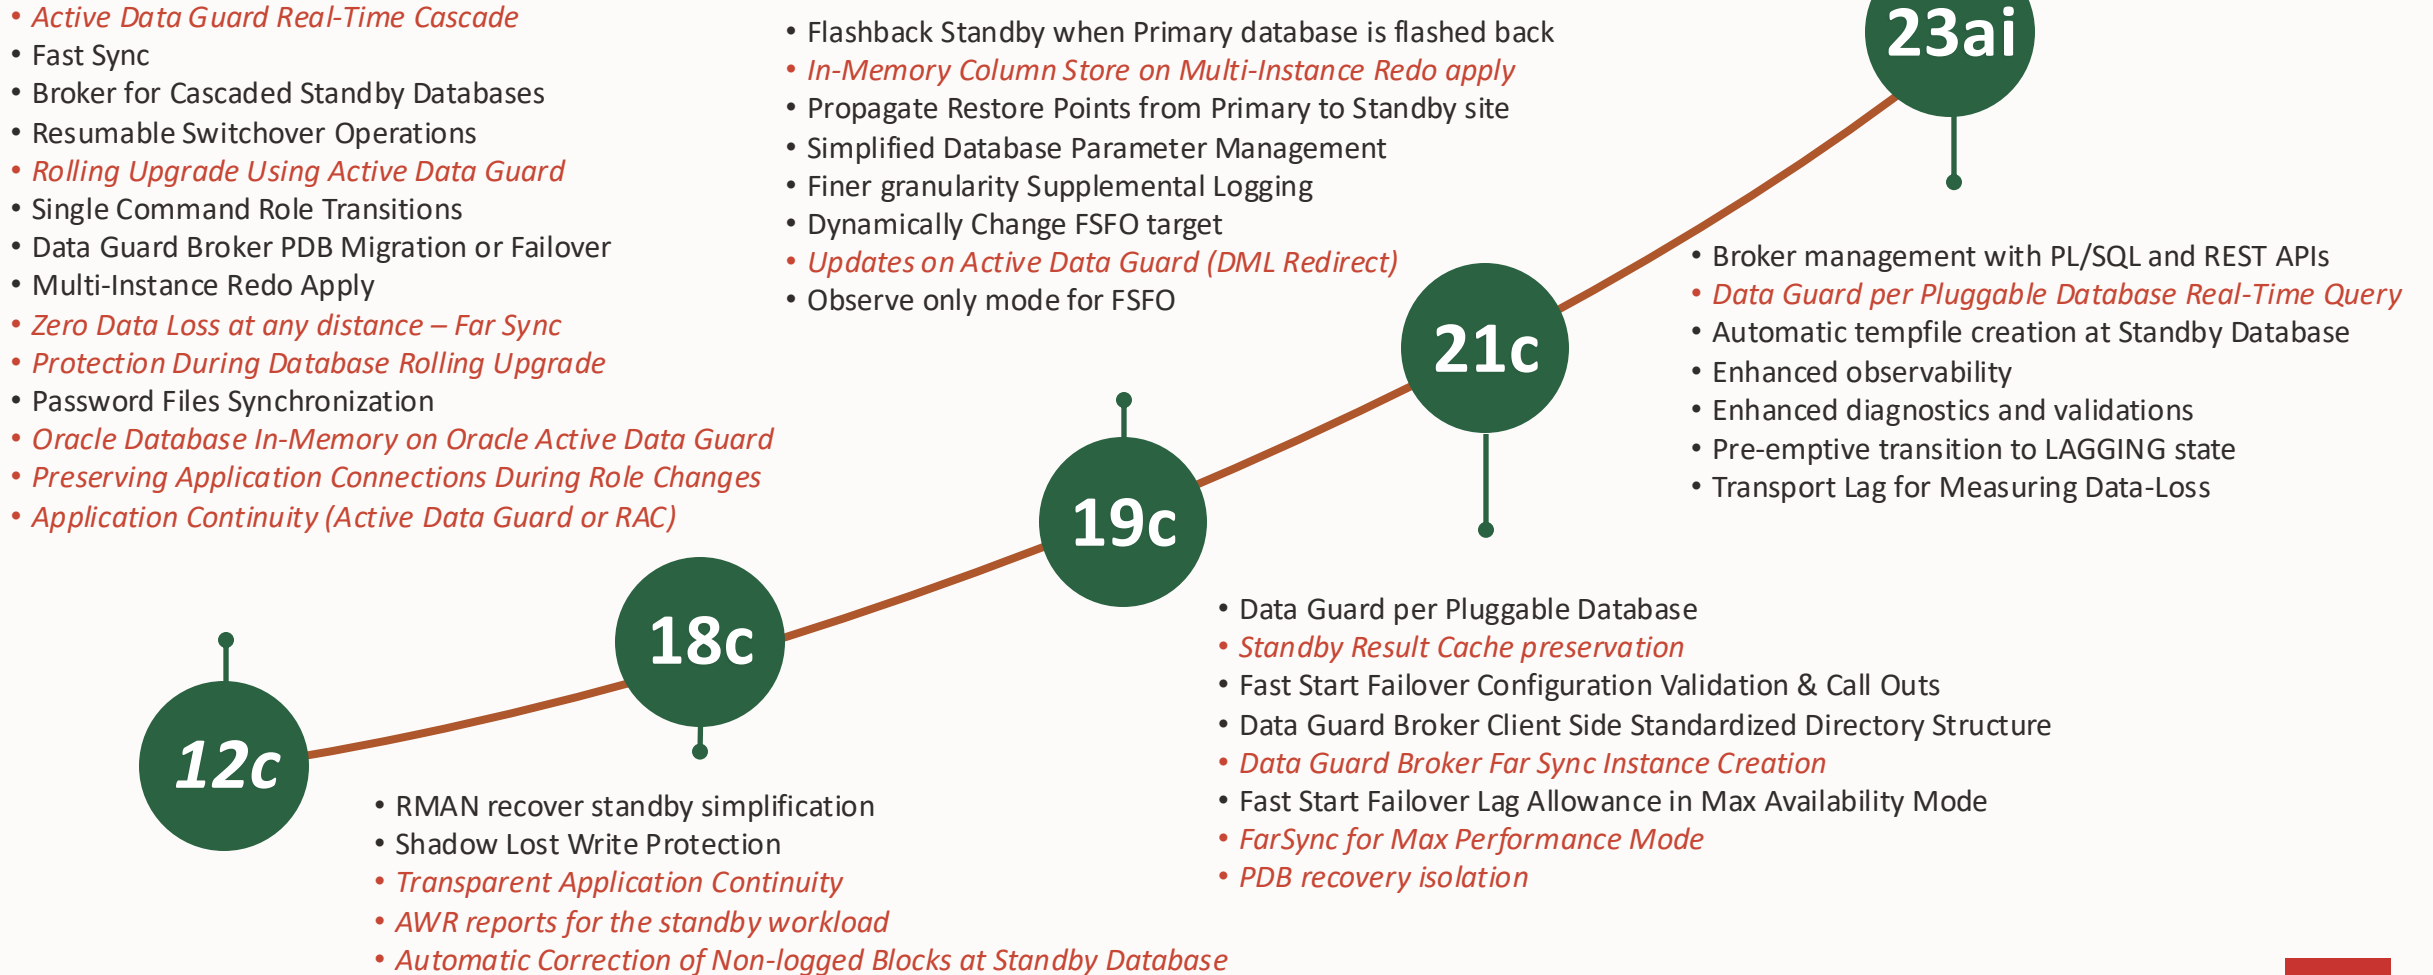
Data Guard

Capabilities Included with Oracle Database Enterprise Edition (EE)




Oracle *Active* Data Guard








Actively protecting data for the future *both* on-premises and in the cloud




Oracle Data Guard 23ai is Simple

Enhanced manageability and observability for CI/CD, DevOps, and traditional operations



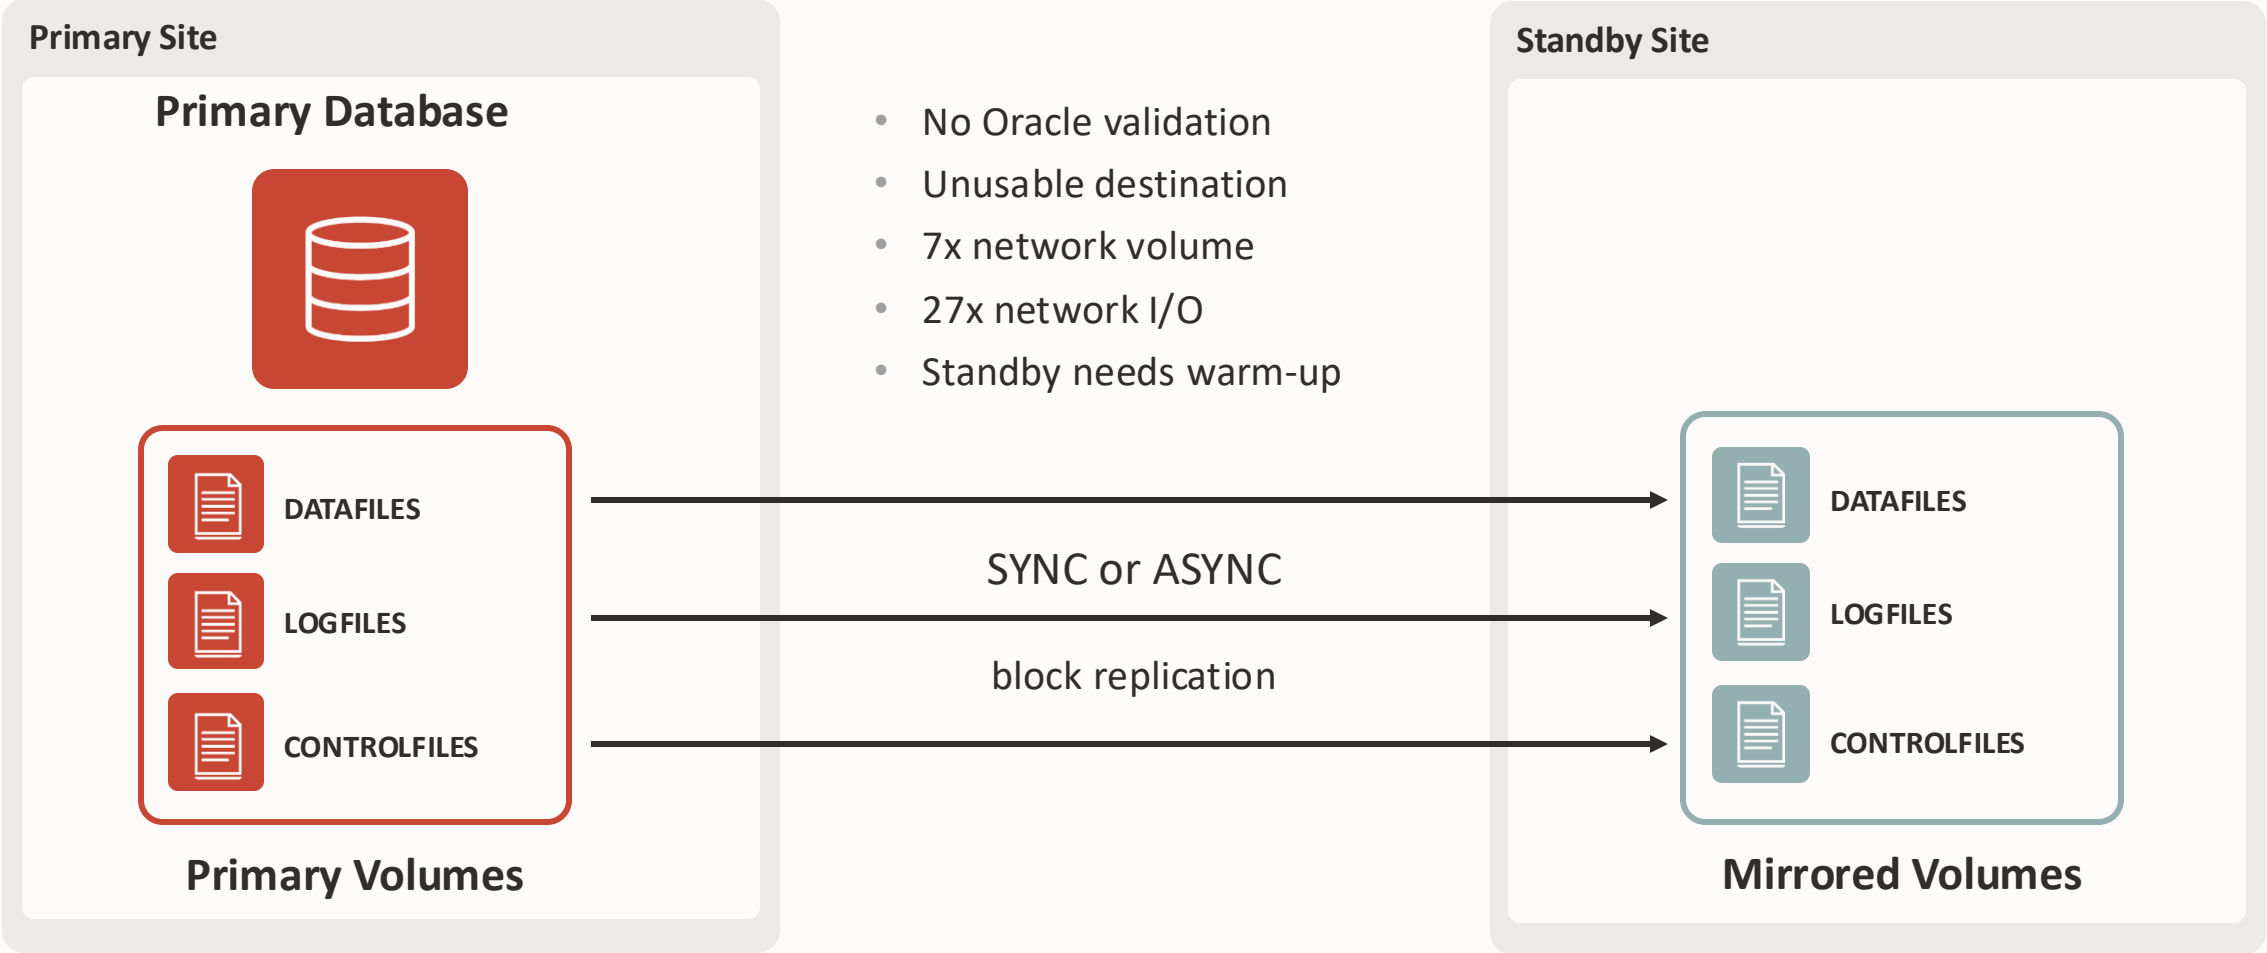
-  Automatic preparation of the primary
-  PL/SQL APIs for better automation
-  REST APIs for easier DevOps integration
-  6 new SQL views to monitor Data Guard
-  Support for mixed Transparent Data Encryption
-  Automatic temp file creation on the standby
-  Simplified AWR snapshots on the standby DB*



Oracle Data Guard vs Storage Replication

Storage Remote Mirroring Architecture

Mirrors every write to every file including those that are **corrupted or encrypted by ransomware**



Data Guard Does What Storage Mirroring Can't

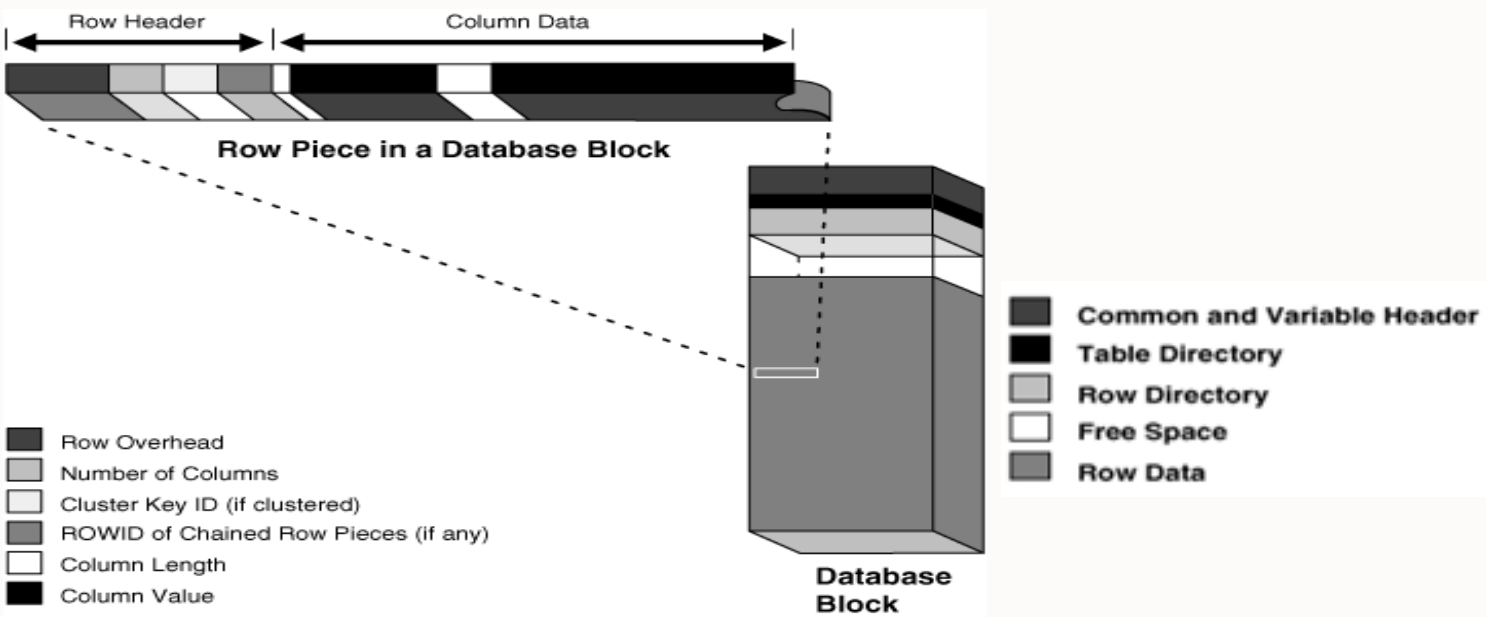
Isolate Corruption, Protect Data, Maintain Availability

Storage Remote Mirroring...
blocks are just bits on a disk

Data Guard uses physical and logical
data consistency checks for end-to-end data integrity

Block 3941 (0x0f65)

	0	1	2	3	4	5	6	7	8	9	a	b	c	d	e	f
000	20	20	20	54	77	61	73	20	74	68	65	20	6e	69	67	68
010	74	20	62	65	66	6f	72	65	20	73	74	61	72	74	2d	75
020	70	20	61	6e	64	20	61	6c	6c	20	74	68	72	6f	75	67
030	68	20	74	68	65	20	6e	65	74	2c	0a	20	20	20	20	20
040	6e	6f	74	20	61	20	70	61	63	6b	65	74	20	77	61	73
050	20	6d	6f	76	69	6e	67	3b	20	6e	6f	20	62	69	74	20
060	6e	6f	72	20	6f	63	74	65	74	2e	0a	20	20	20	54	68
070	65	20	65	6e	67	69	6e	65	65	72	73	20	72	61	74	74
080	6c	65	64	20	74	68	65	69	72	20	63	61	72	64	73	20
090	69	6e	20	64	65	73	70	61	69	72	2c	0a	20	20	20	20
0a0	20	68	6f	70	69	6e	67	20	61	20	62	61	64	20	63	68
0b0	69	70	20	77	6f	75	6c	64	20	62	6c	6f	77	20	77	69
0c0	74	68	20	61	20	66	6c	61	72	65	2e	0a	20	20	20	54
0d0	68	65	20	73	61	6c	65	73	6d	65	6e	20	77	65	72	65
0e0	20	6e	65	73	74	6c	65	64	20	61	6c	6c	20	73	6e	75
0f0	67	20	69	6e	20	74	68	65	69	72	20	62	65	64	73	2c
100	0a	20	20	20	20	20	20	77	68	69	6c	65	20	76	69	73
110	6f	6e	73	20	6f	66	20	64	61	74	61	20	6e	65	74	73
120	20	64	61	6e	63	65	64	20	69	6e	20	74	68	65	69	72
130	20	68	65	61	64	73	2e	0a	20	20	20	20	41	6e	64	20
140	20	77	69	74	68	20	6d	79	20	64	61	74	61	73	63	6f
150	70	65	20	74	72	61	63	69	6e	67	73	20	61	6e	64	20
160	64	75	6d	70	73	0a	20	20	20	20	20	20	70	72	65	70
170	72	65	64	20	66	6f	72	20	73	6f	6d	65	20	70	72	65
180	74	74	79	20	62	61	64	20	62	72	75	69	73	65	73	20
190	61	6e	64	20	6c	75	6d	70	73	2e	0a	20	20	20	57	68
1a0	65	6e	20	6f	75	74	20	69	6e	20	74	68	65	20	68	61
1b0	6c	6c	20	74	68	65	72	65	20	61	72	6f	73	65	20	73
1c0	75	63	68	20	61	20	63	6c	61	74	74	65	72	2c	0a	20
1d0	20	20	20	20	49	20	73	70	72	61	6e	67	20	66	72	6f
1e0	6d	20	6d	79	20	64	65	73	6b	20	74	6f	20	73	65	65
1f0	20	77	68	61	74	20	77	61	73	20	74	68	65	20	6d	61




See My Oracle Support Note 1302539.1 for details

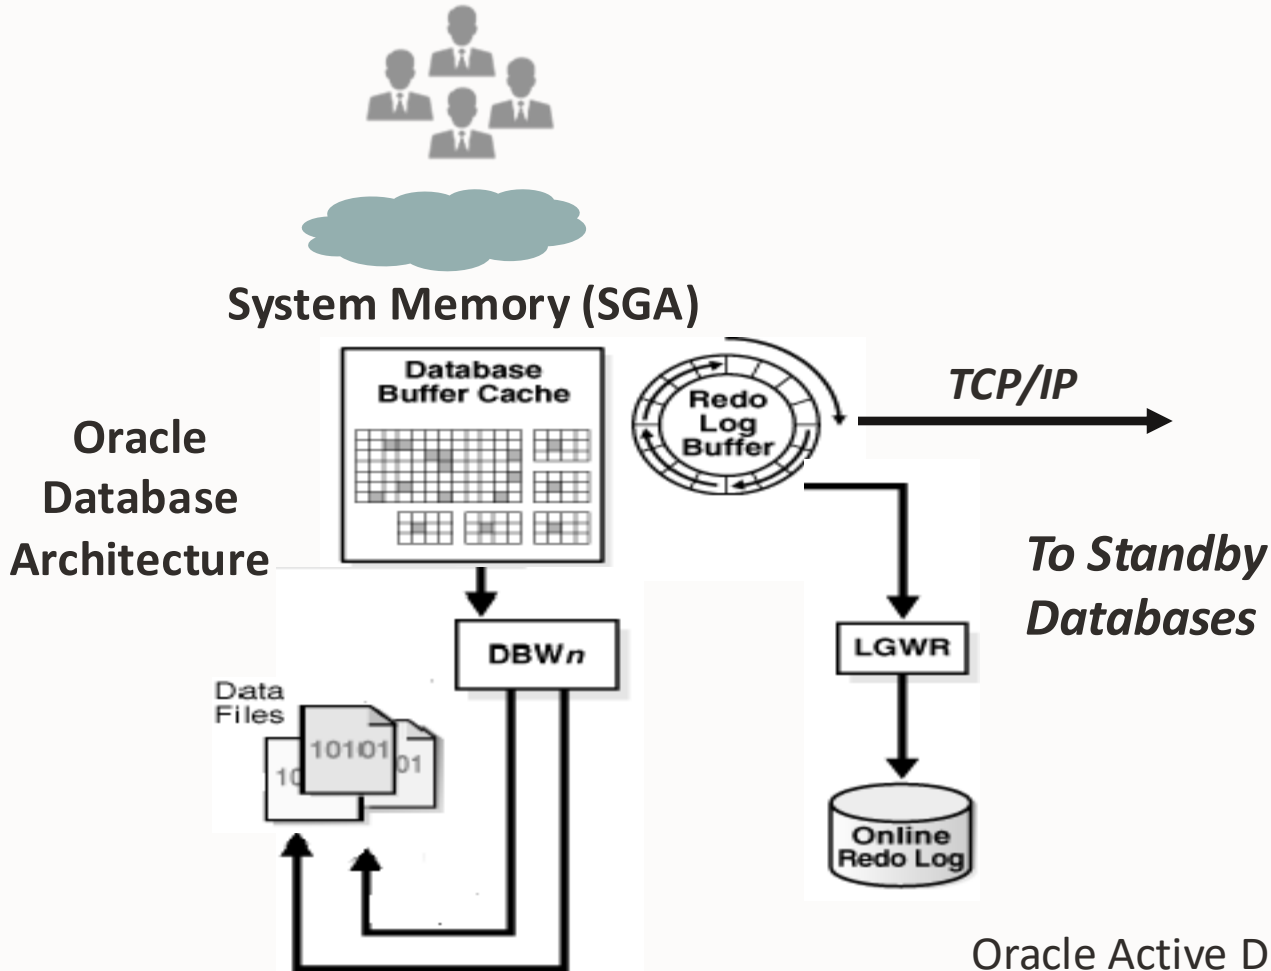


Data Guard is optimized for the database

It efficiently maintains a **physical copy** of production and **guarantees its integrity**



Data Guard Provides Strongest Fault Isolation and Best Performance



Data Guard transmits redo blocks directly from SGA: like a *memcpy* over the network

Redo received / applied by running Oracle instance: continuous Oracle-integrated data validation

- Best isolation from lower layer faults
- Best performance since no disk I/O
- Best network utilization: only redo sent
- Transactional consistency: always
- Corrupted blocks auto-repaired *
- Database-integrated application failover


* Requires Active Data Guard License

Oracle Active Data Guard Compared to Storage Remote Mirroring

<https://www.oracle.com/a/tech/docs/adg-vs-storage-mirroring.pdf>

Oracle Replication done right

<https://blogs.oracle.com/maa/replication-done-right>



Oracle Data Guard Redo Transport

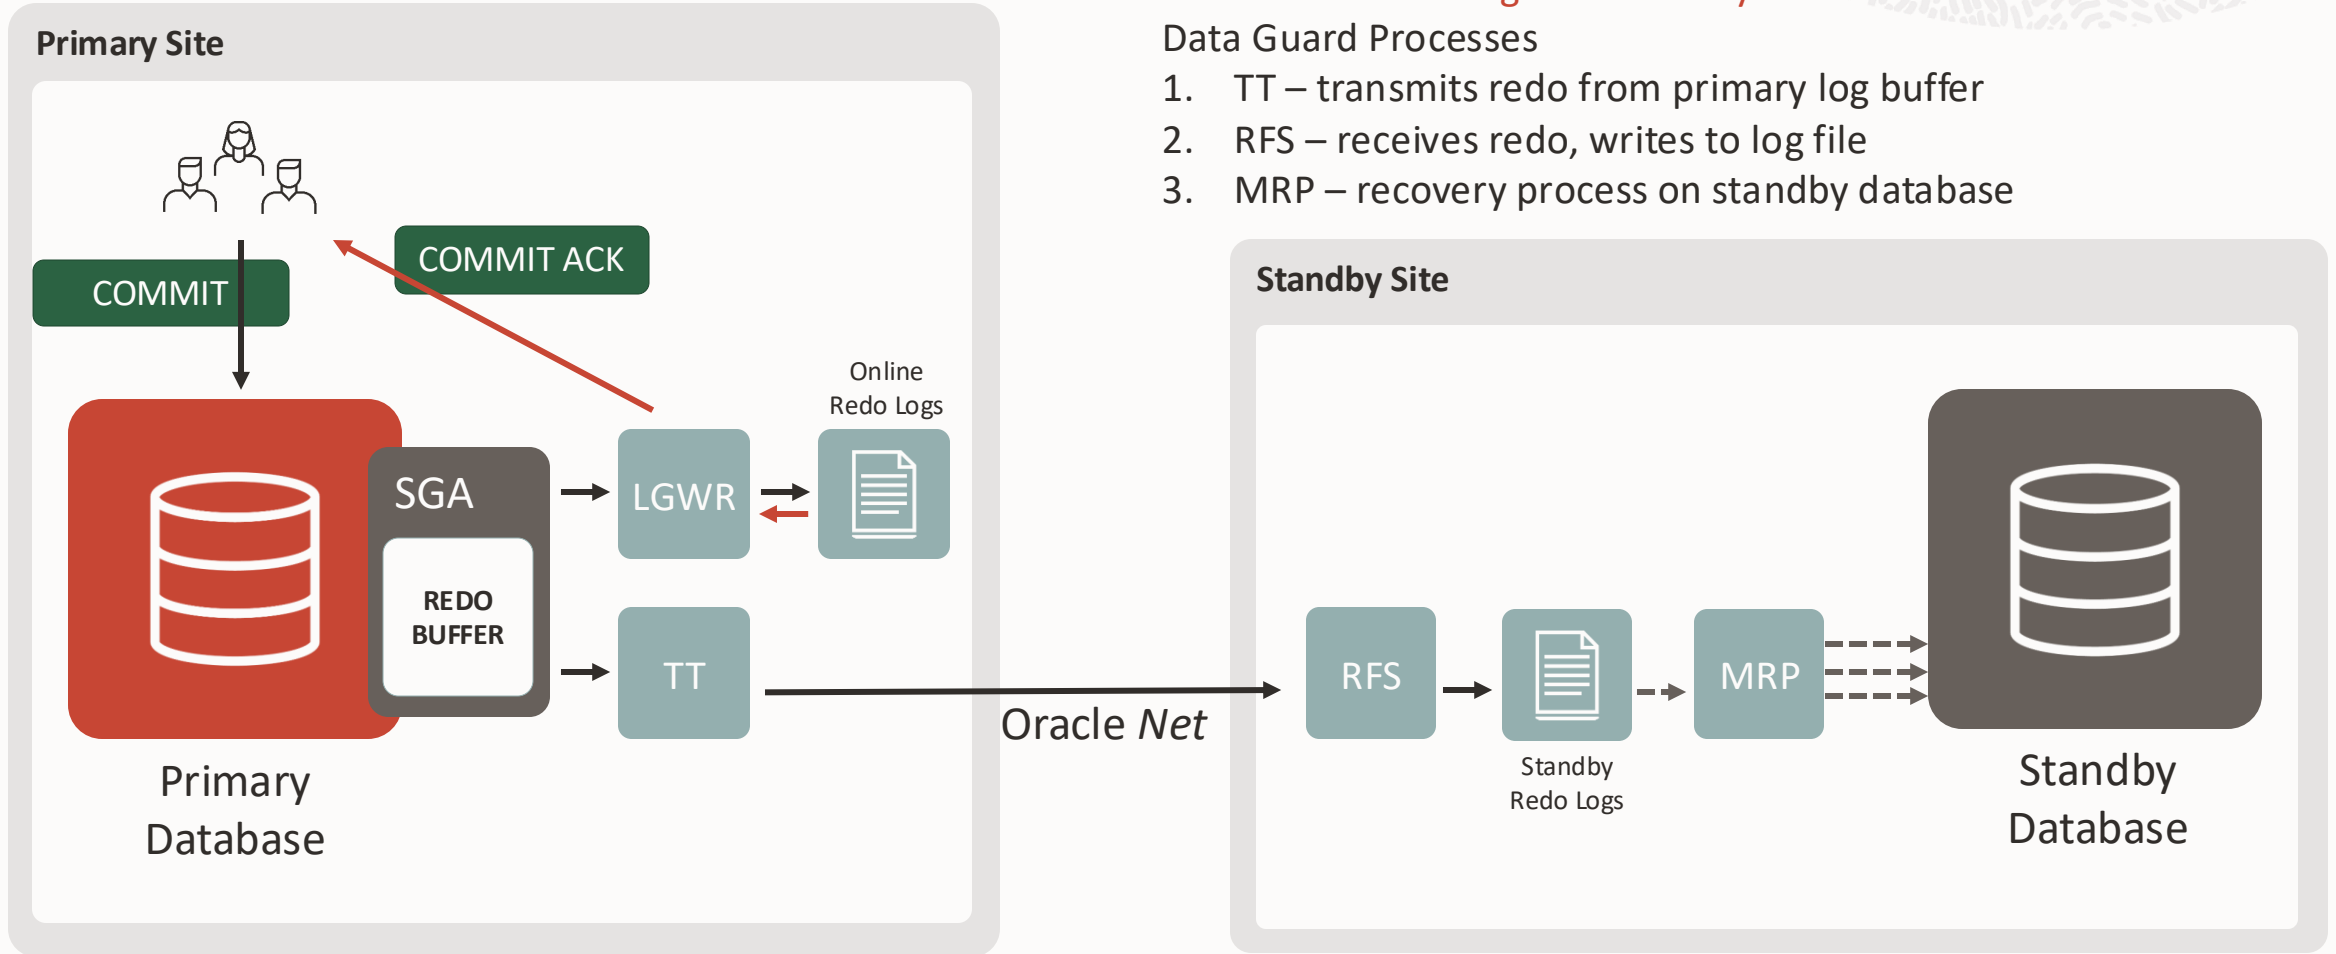
Data Guard Transport for Best Performance

Data Guard **ASYNC** Process Architecture

Commit Acknowledge is local-only

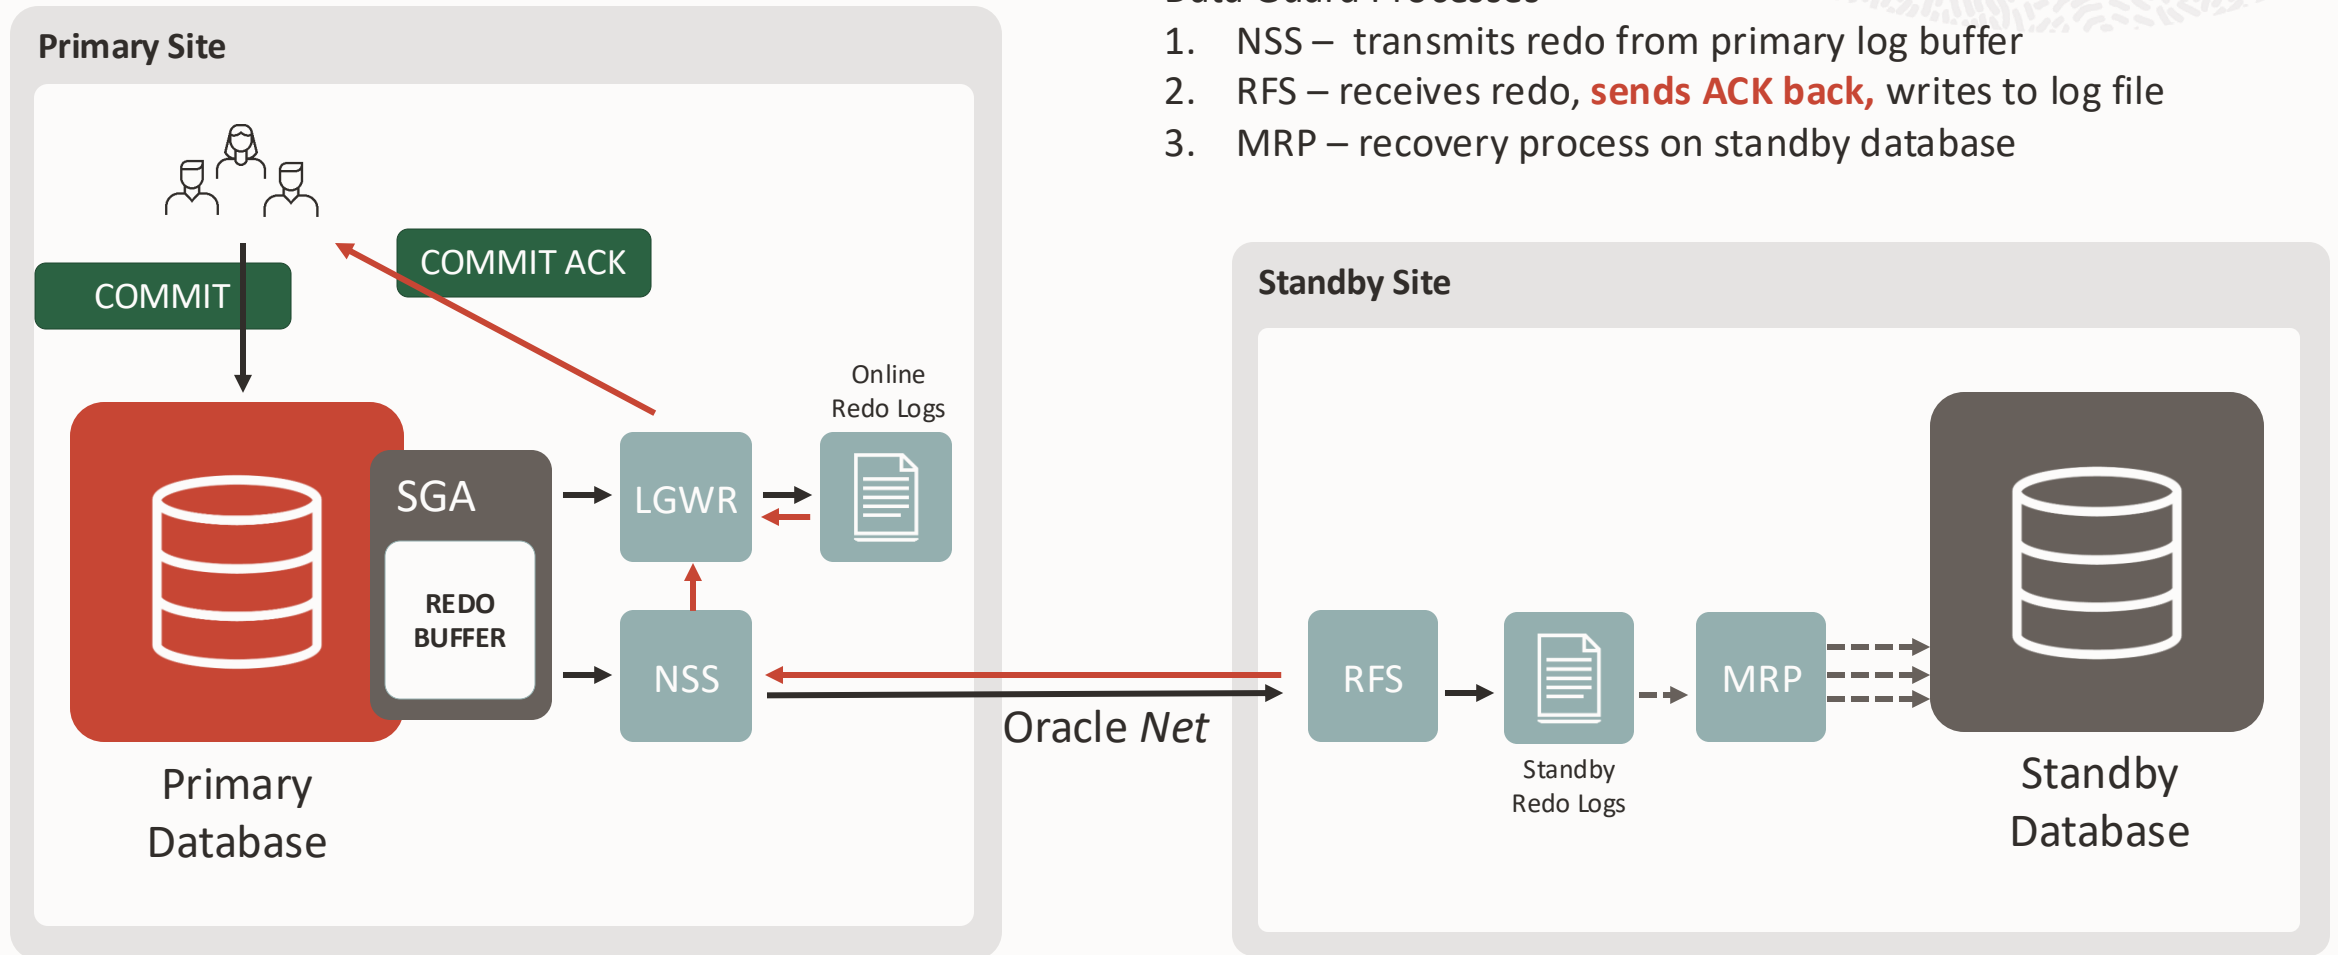
Data Guard Processes

1. TT – transmits redo from primary log buffer
2. RFS – receives redo, writes to log file
3. MRP – recovery process on standby database



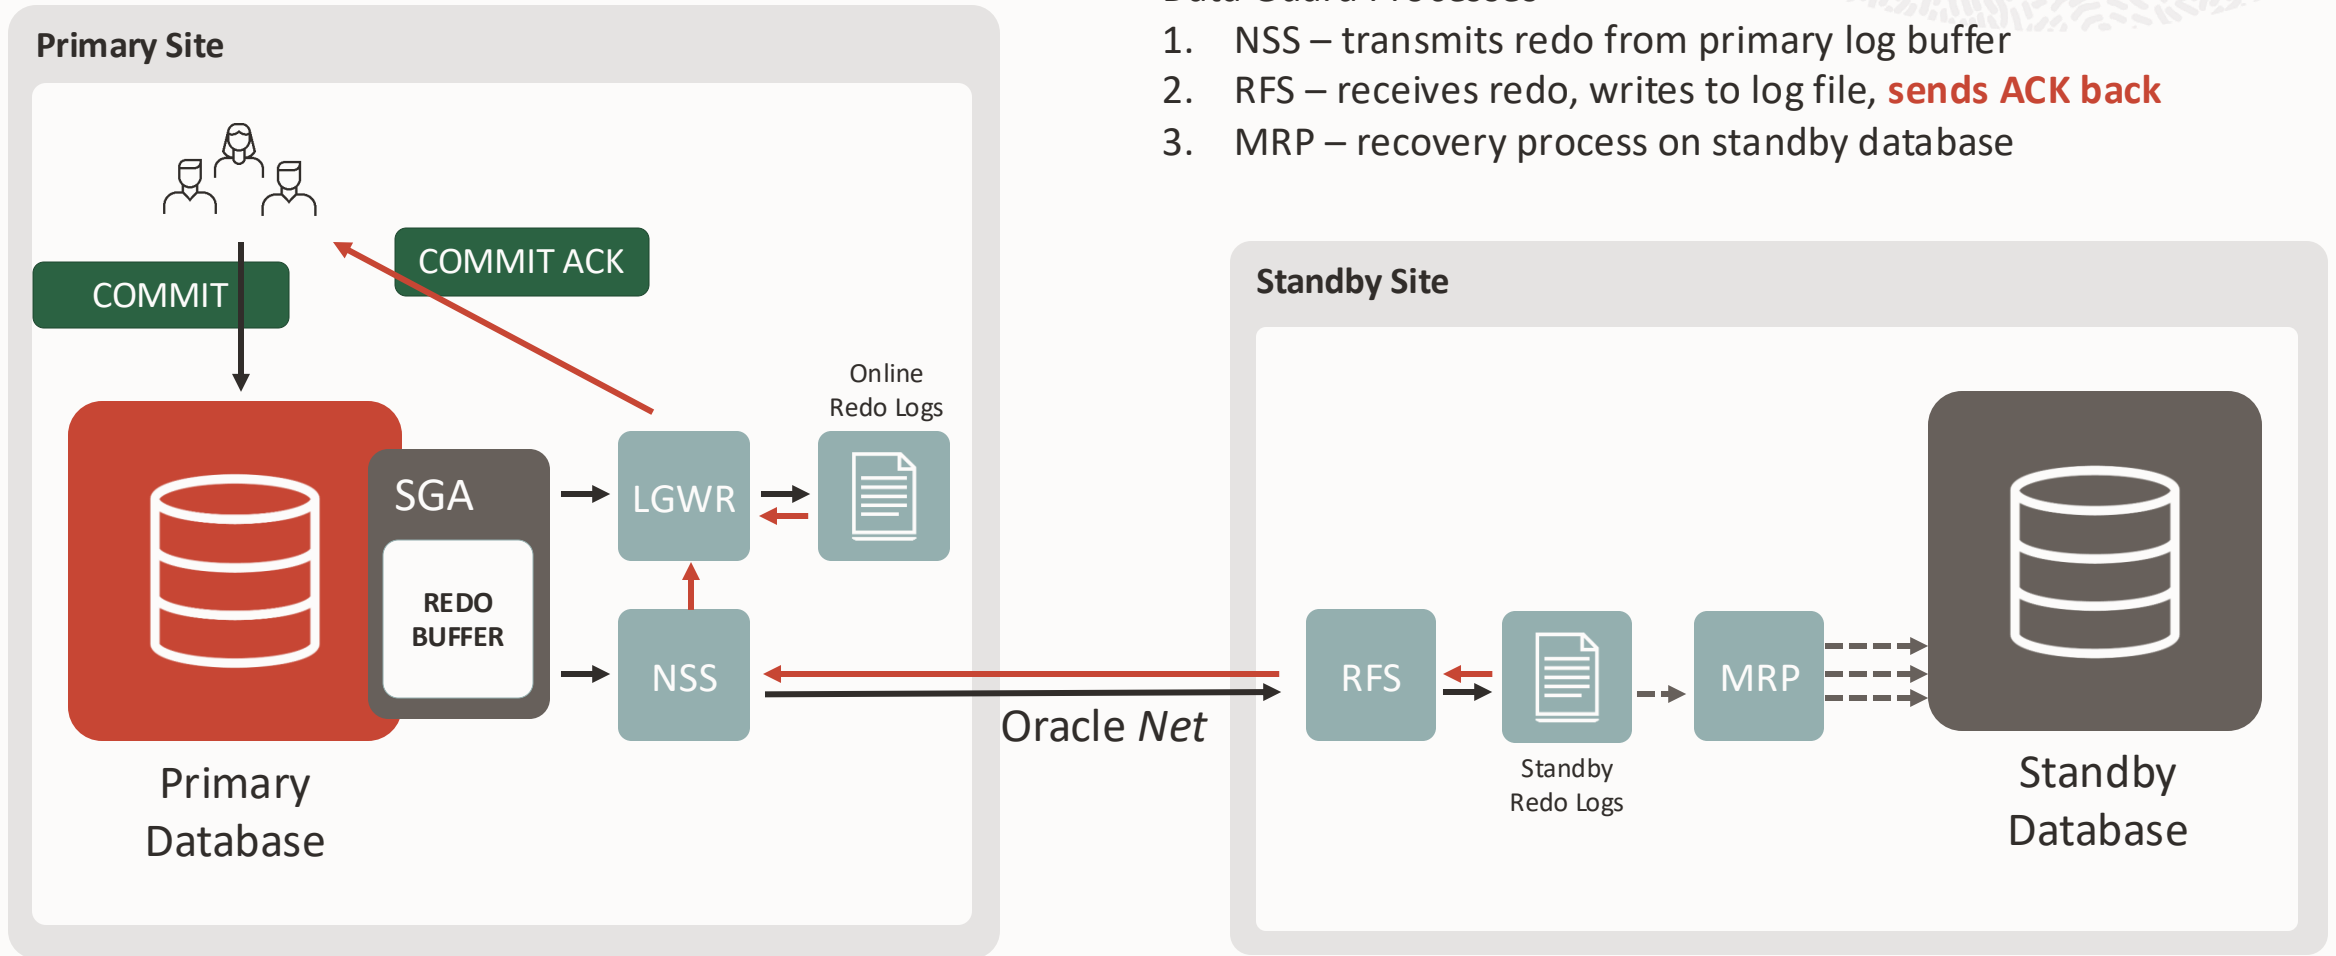
Data Guard Transport for Zero Data Loss

Data Guard FASTSYNC Process Architecture



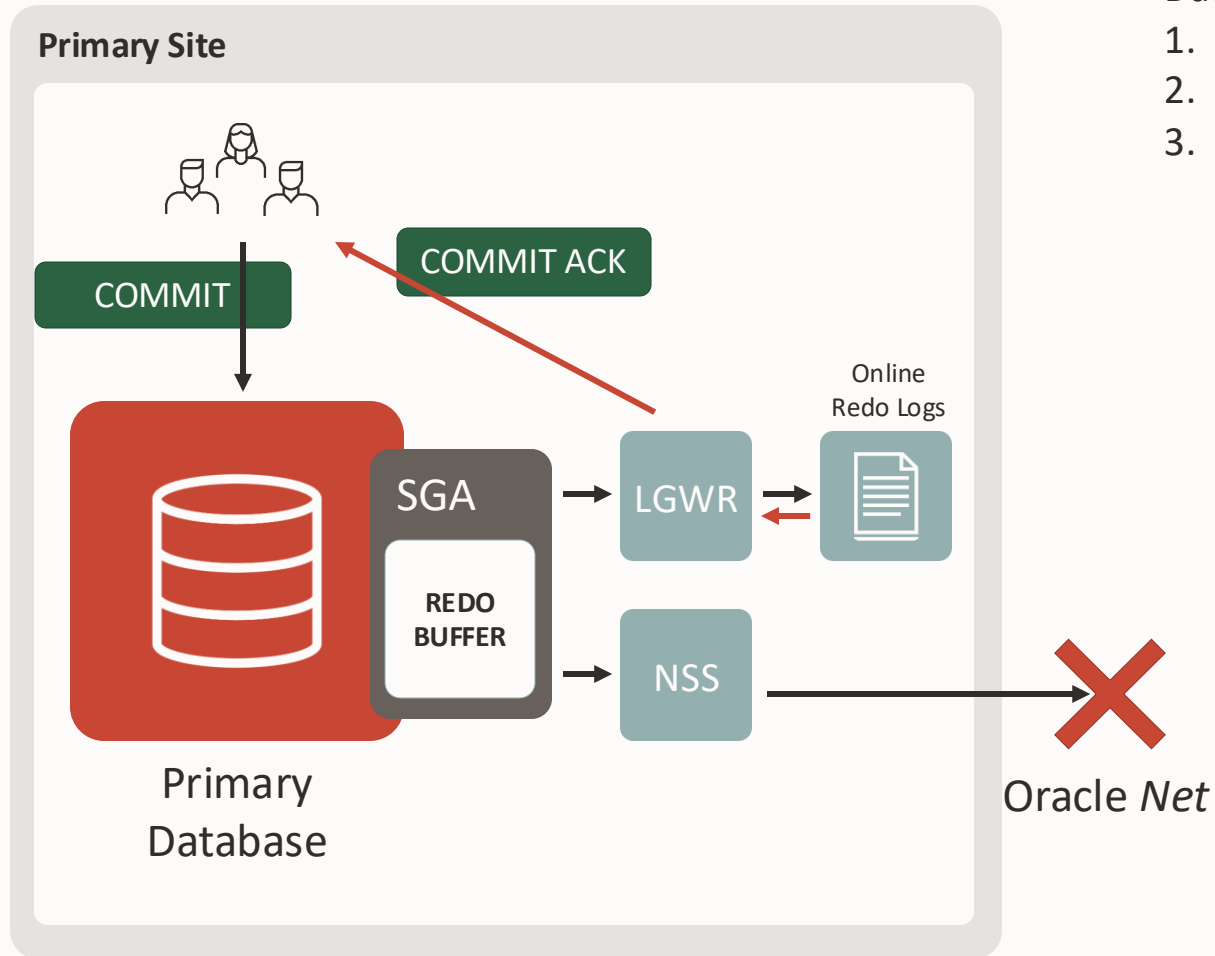
Data Guard Transport for Zero Data Loss

Data Guard SYNC Process Architecture



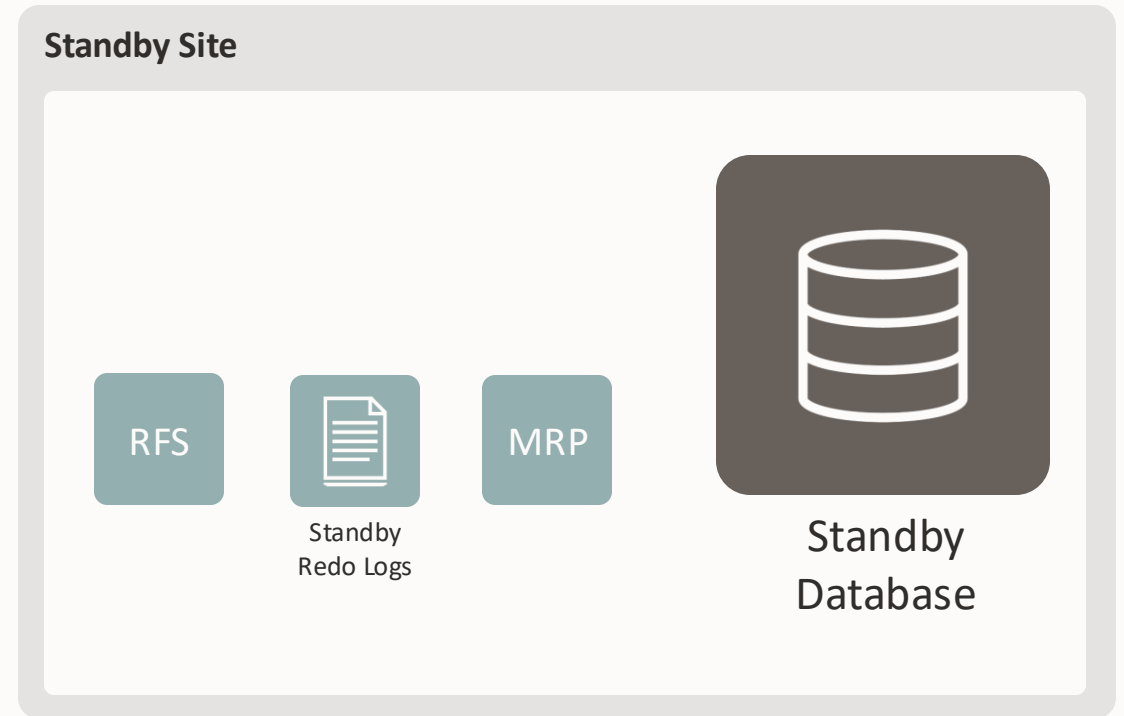
Stalling Synchronous destinations

Data Guard FASTSYNC/SYNC Process Architecture



Data Guard Processes

1. NSS – tries to send the redo to the remote destination
2. The commits stall for NetTimeout seconds
3. The destination is **abandoned**, the commits resume



The difference between receiving the redo late and not receiving it

DATUM_TIME vs TRANSPORT LAG vs LAST_TIME

Standby **not receiving the redo** from the primary:

```
SQL> select value, datum_time, from v$dataguard_stats where name='transport lag';
```

VALUE	DATUM_TIME
-----	-----
+00 00:00:00	07/11/2022 08:28:46

Standby **receiving old redo** from the primary:

```
SQL> select value, datum_time, from v$dataguard_stats where name='transport lag';
```

VALUE	DATUM_TIME
-----	-----
+01 13:50:54	07/12/2022 21:48:10

The last redo written in the standby logs:


```
SQL> select max(last_time) from v$standby_log where status='ACTIVE';
```

MAX(LAST_TIME)

07/11/2022 08:28:46

High Performance – Synchronous Redo Transport

Mixed OLTP workload with Metro-Area Network Latency



Note: 0ms latency on graph represents values <1ms

Workload profile

- Simulated OLTP with large inserts
- **112 MB/s redo**

3% impact at < 1ms RTT

5% impact at 2ms RTT

6% impact at 5ms RTT

Use **oracptest** to assess your network bandwidth and latency



Oracle Data Guard and Database Nologging



Enabling the FORCE LOGGING mode is a must for non-Engineered Systems.

On Exadata and Oracle Cloud Infrastructure, two modes are alternative to the FORCE LOGGING mode:

1. Standby Nologging for **Load Performance**
 - Ensures that standbys will receive the non-logged data changes with minimum impact on the speed of loading at the primary.
 - The standby can transiently have non-logged blocks. These non-logged blocks will be automatically resolved by managed standby recovery.
 2. Standby Nologging for **Data availability**
 - Ensures that all standbys have the data when the primary load commits, but at the cost of throttling the speed of loading data at the primary.
 - The standbys will never have any non-logged blocks.
- **These new modes cause Multi-Instance Redo Apply to return an error**
 - Single Instance Redo apply must be manually enabled to proceed past the nologging operation.

Oracle Data Guard Best Practices – Transport and Apply Tuning

Redo Apply Best Practices

<https://docs.oracle.com/en/database/oracle/oracle-database/19/haovw/tune-and-troubleshoot-oracle-data-guard.html#GUID-E8C27979-9D37-4899-9306-A5AE2B5CF6C0>

Best Practices for Redo Transport Tuning

<https://docs.oracle.com/en/database/oracle/oracle-database/19/haovw/tune-and-troubleshoot-oracle-data-guard.html#GUID-A6963335-8C5A-4DD0-AD3F-22F4CBCE3DD0>

Assessing Synchronous Redo Transport


<https://docs.oracle.com/en/database/oracle/oracle-database/19/haovw/tune-and-troubleshoot-oracle-data-guard.html#GUID-4C3E0CC9-3E54-48C4-8DD6-AB4EC0C51696>

How To Calculate The Required Network Bandwidth Transfer Of Redo In Data Guard (Doc ID 736755.1)

<https://support.oracle.com/rs?type=doc&id=736755.1>

Assessing and Tuning Network Performance for Data Guard and RMAN (Doc ID 2064368.1)

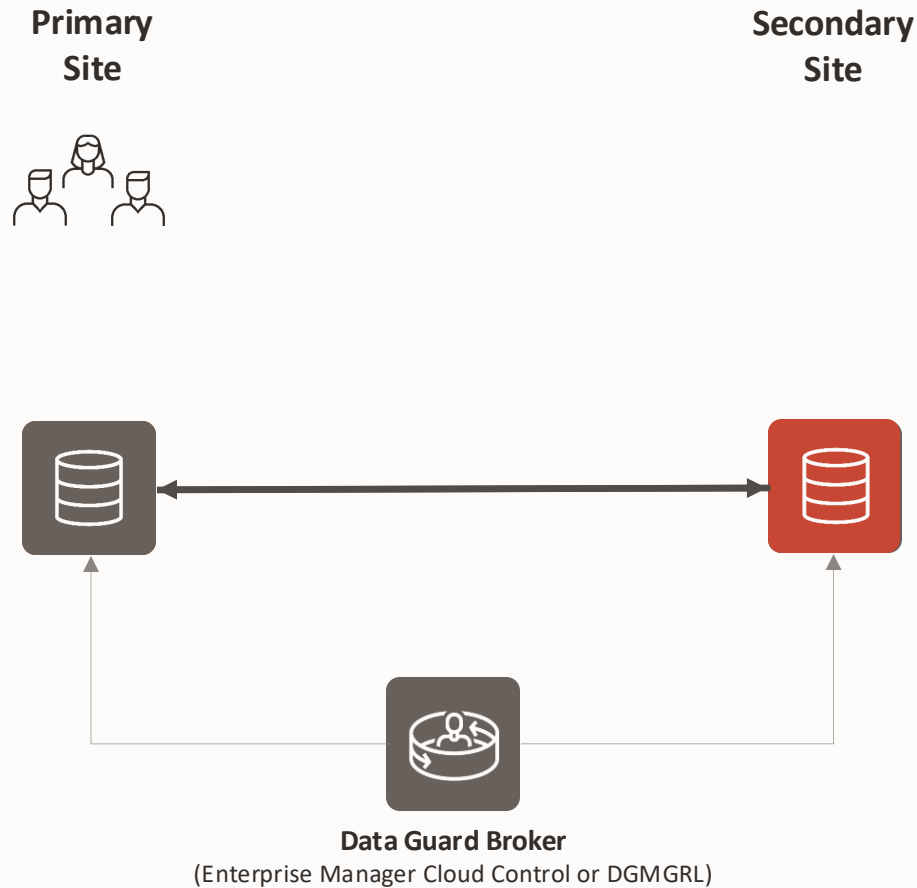
<https://support.oracle.com/rs?type=doc&id=2064368.1>



Oracle Data Guard Role Transitions

Oracle Data Guard Planned Role Transition


Switchover: Planned role transition with Zero Data Loss



- **Switchover initiated**
- **The primary ends the transactions and stops the services**
- **All the transaction are synced to the standby**
- **The standby is converted to primary and the services are started**
 - The replication starts again
- **The applications reconnect transparently to the new primary**
 - If properly configured, the application experience just a freeze for 1-2 minutes or less

Oracle Data Guard Unplanned Role Transition


Failover: In case of failure the role transition *can* be without data loss



- **The observer detects the failure of the primary**
 - Depending on the protection mode and situation, the observer initiates the failover after `FastStartFailoverThreshold` seconds
- **The standby is converted to primary and the services are started**
 - Depending on the protection mode and situation, there might be some data loss (the tolerated amount is configurable)
- **The applications reconnect to the new primary**
 - The reinstatement of the primary requires a single broker command
- The **failover** can be initiated also **manually** (DGMGRL) or by the application (DBMS_DG.INITIALIZE_FS_FAILOVER). The amount of data loss is customer's responsibility in this case.


"FAILOVER TO" vs "DBMS_DG.INITIALIZE_FS_FAILOVER"

FAILOVER TO stdby



- 1 The standby database becomes the new primary
- 2 The former primary shuts down if FSFO is in place
- 3 The former primary is automatically reinstated


DBMS_DG.INITIALIZE_FS_FAILOVER



- 1 The procedure shuts down the primary first
- 2 The standby database becomes the new primary
- 3 The former primary is not automatically reinstated


Data Guard Snapshot Standby

Standby database temporarily in Read Write



- **The standby is converted to Snapshot Standby**
 - Standby open read write
- **Users and DBAs perform tests (Upgrade, Performance, etc)**
 - The primary is still protected by the redo transfer
- **When the tests are over, the standby is flashed back and converted to physical standby again**
- Note: the snapshot standby cannot relay the redo to a cascaded standby

Tune the Switchover and Failover Operations



ADMINISTRATION

- Stop any long running operations before switchover
- Reduce `FAST_START_MTTR_TARGET`
- Switchover to a MOUNTED database whenever possible

DEVELOPMENT

- Use Oracle Connection Pools
- Make your transactions as light as possible
- Use FAN notifications
- Implement session draining
- Use the recommended connection strings

IMPLEMENTATION

- Reduce the number of data files
- Reduce the workload on the database
- Don't over-consolidate PDBs

Easier tracking of role transitions



The new fixed view **V\$DG_BROKER_ROLE_CHANGE** tracks the last 10 role transitions

```
SQL> select * from V$DG_BROKER_ROLE_CHANGE;
```

EVENT	STANDBY_TYPE	OLD_PRIMARY	NEW_PRIMARY	FS_FAILOVER_REASON	BEGIN_TIME	END_TIME

Failover	Physical	mydb1	mydb1b	Manual Failover	30-SEP-2022 19:01:14	30-SEP-2022 19:01:35
Switchover	Physical	mydb1b	mydb1		30-SEP-2022 19:04:53	30-SEP-2022 19:05:15
Switchover	Physical	mydb1	mydb1b		30-SEP-2022 20:51:38	30-SEP-2022 20:52:03
Failover	Physical	mydb1b	mydb1	Manual Failover	30-SEP-2022 20:52:46	30-SEP-2022 20:53:04
Switchover	Physical	mydb1	mydb1c		30-SEP-2022 19:53:14	30-SEP-2022 19:54:14
Switchover	Physical	mydb1c	mydb1		30-SEP-2022 20:03:14	30-SEP-2022 20:04:04
Switchover	Logical	mydb1	mydb1d		30-SEP-2022 20:24:46	30-SEP-2022 20:26:32
Switchover	Logical	mydb1d	mydb1		30-SEP-2022 20:35:27	30-SEP-2022 20:35:48
Fast-Start Failover	Physical	mydb1	mydb1b	Primary Disconnected	30-SEP-2022 20:13:51	30-SEP-2022 20:14:53



Strict validation of switchover readiness

New command `VALIDATE DATABASE STRICT`

```
DGMGRL> VALIDATE DATABASE mydb_site2
```

```
Database Role:      Physical standby database  
Primary Database:   mydb_site1
```

```
Ready for Switchover: Yes
```

```
Ready for Failover:  Yes (Primary Running)
```

```
Flashback Database Status:
```

Database	Status	Retention Target
mydb_site1	Off	1440
mydb_site2	Off	1440

```
...
```

No strict validation

The DB shows as Ready for Switchover if no conditions would prevent the switchover from working.

Syntax:

```
VALIDATE DATABASE [VERBOSE] <database> STRICT
```

```
{TEMP_FILES | FLASHBACK | LOG_FILES_CLEARED | LOG_FILE_CONFIGURATION | APPLY_PROPERTY | TRANSPORT_PROPERTY | ALL}
```

```
DGMGRL> VALIDATE DATABASE mydb_site2 STRICT ALL
```

```
Database Role:      Physical standby database  
Primary Database:   mydb_site1
```

```
Ready for Switchover: No
```

```
Ready for Failover:  Yes (Primary Running)
```

```
Flashback Database Status:
```

Database	Status	Retention Target
mydb_site1	Off	1440
mydb_site2	Off	1440


```
...
```

Strict validation

No conditions would prevent the switchover from working, but the new primary would miss some important configurations.

Faster Role Transitions in Oracle Data Guard 23.5


Between 50% and 85% faster role transition in Oracle Data Guard 23ai



2-node Exadata RAC
5 PDBs
50 data files

10 services

60MB redo/sec



4-node Exadata RAC
100 PDBs
10k data files


1200 services

100MB redo/sec




Oracle Data Guard 23ai Provides Lower RTO

Faster role transitions. Pre-emptive actions to prevent stalls.



Faster role transitions
compared to 19c*

Up to 3 seconds faster observer acknowledgment



19c

23ai

in Fast-Start Failover Max Performance



* Figures vary with the workload and the environment

Oracle Data Guard Role Transitions – Read More



Role Transition Assessment and Tuning

<https://docs.oracle.com/en/database/oracle/oracle-database/19/haovw/tune-and-troubleshoot-oracle-data-guard.html#GUID-CBA9FC61-9894-4D62-9569-EFBD7960267F>




Client Failover and Application Continuity


Services for Location Transparency and High Availability

Services provide a “dial in number” for your application


- Use Custom services with FAN notifications and Application Continuity
- Regardless of location, application keeps the name!
- Client failover best practices across the Oracle technology stack




Real Application Clusters



Active Data Guard



GoldenGate



PDB Relocation

Connections Appear Continuous

Standard for All Drivers from 12.2

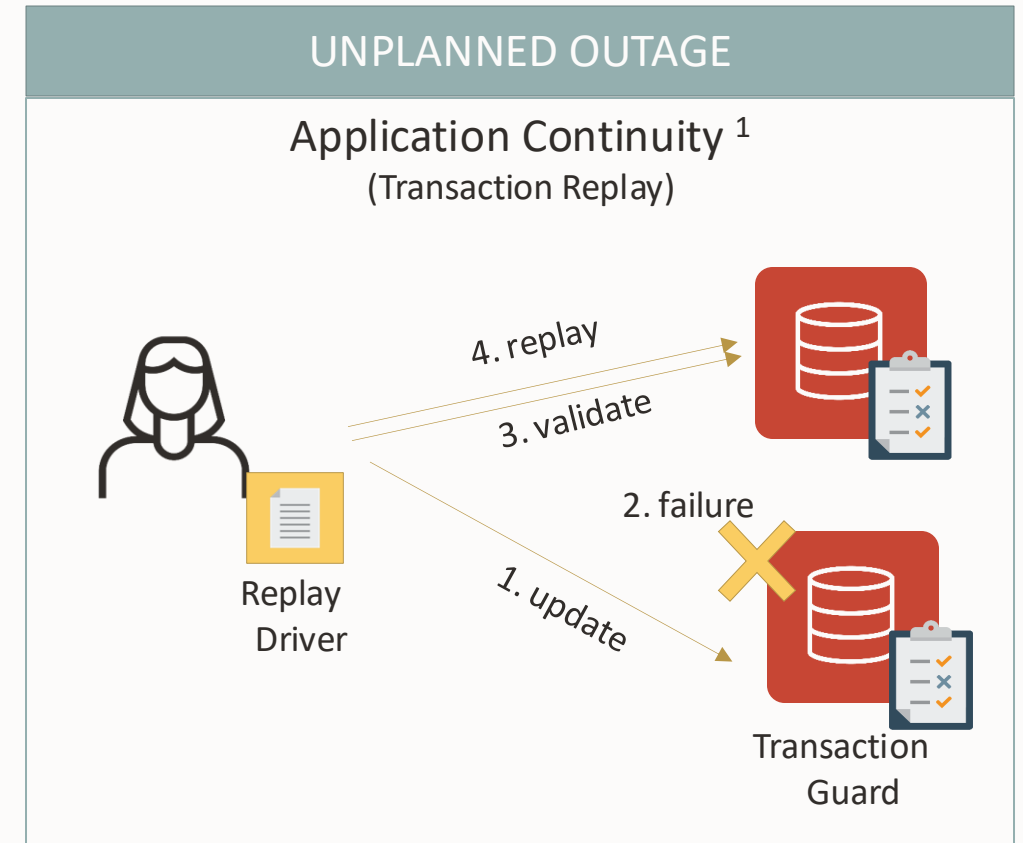
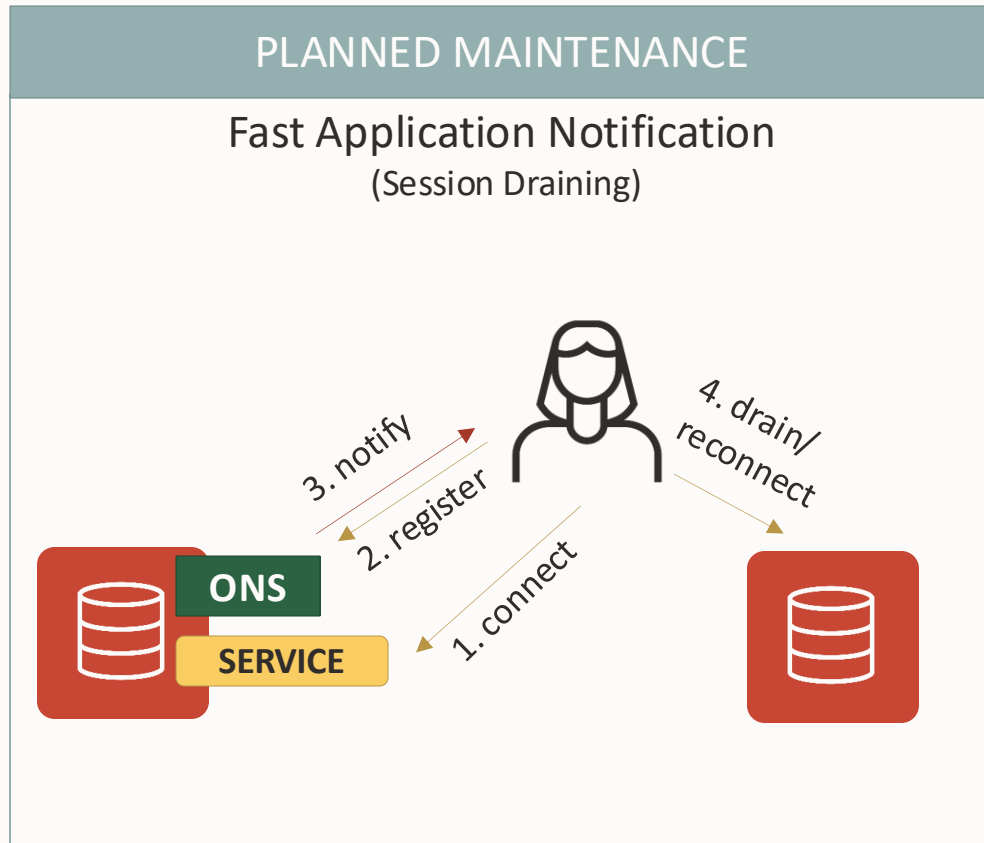
Automatic retries until the service is available

```
HR = (DESCRIPTION =  
  (CONNECT_TIMEOUT=120)(RETRY_COUNT=50)(RETRY_DELAY=3)  
  (TRANSPORT_CONNECT_TIMEOUT=3)  
  (ADDRESS_LIST =  
    (LOAD_BALANCE=on)  
    (ADDRESS=(PROTOCOL=TCP)(HOST=cluster1-scan)(PORT=1521)))  
  (ADDRESS_LIST =  
    (LOAD_BALANCE=on)  
    (ADDRESS=(PROTOCOL=TCP)(HOST=cluster2-scan)(PORT=1521)))  
  (CONNECT_DATA=(SERVICE_NAME = HR.oracle.com)))
```

Always use a custom service!
Do NOT use PDB or DB Name

Client-side required technologies

Client draining/failover is a crucial part of high availability for applications connecting to the database.




¹ Application Checklist for Continuous Service for MAA Solutions

<https://www.oracle.com/technetwork/database/clustering/checklist-ac-6676160.pdf>

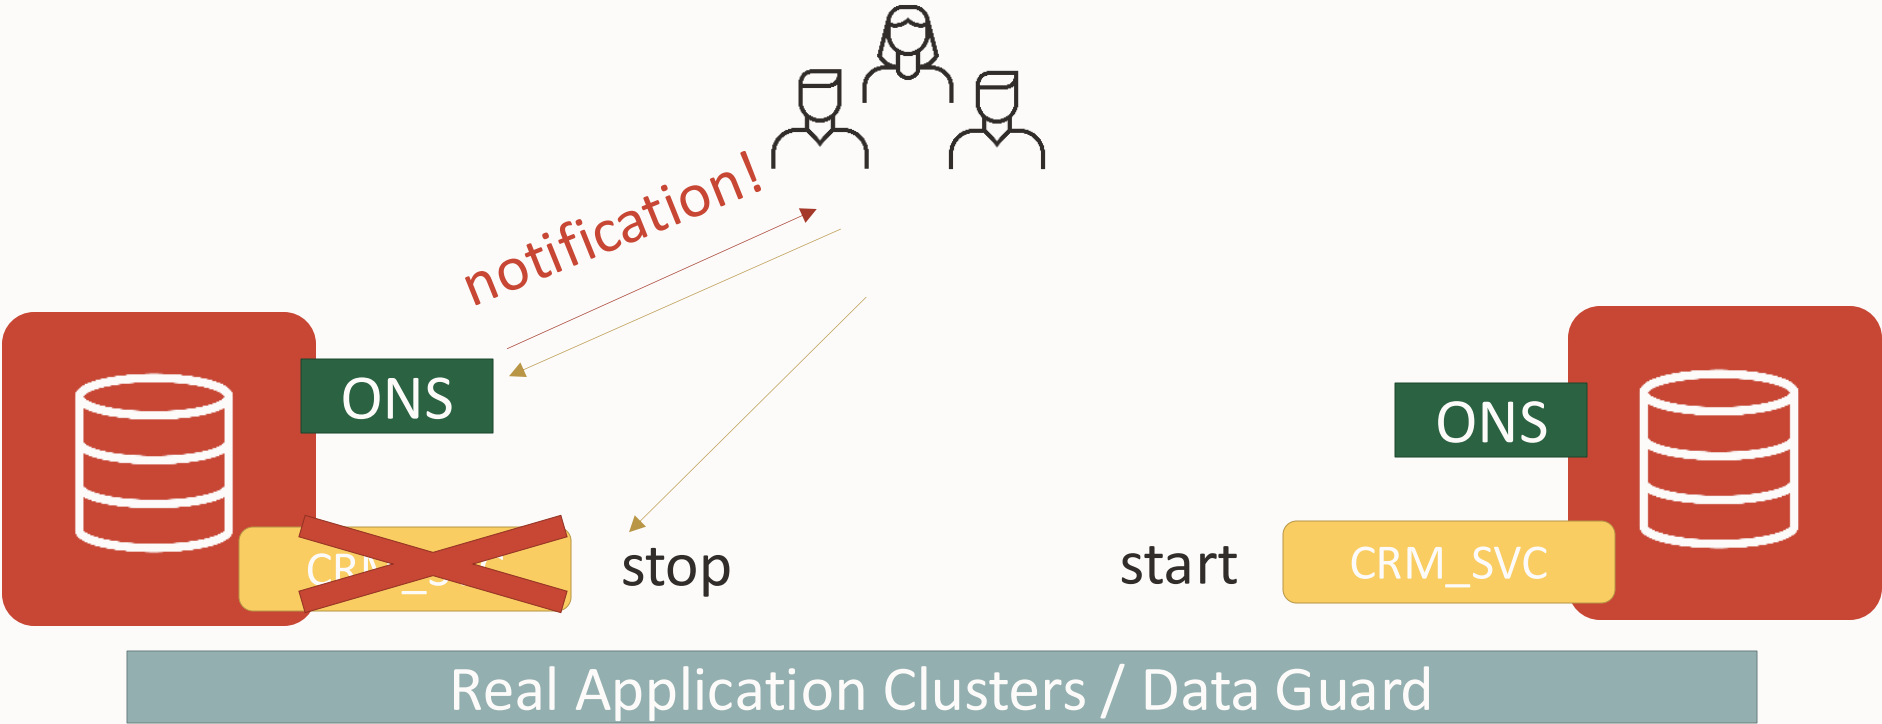
Fast Application Notification

Session Draining for planned maintenance



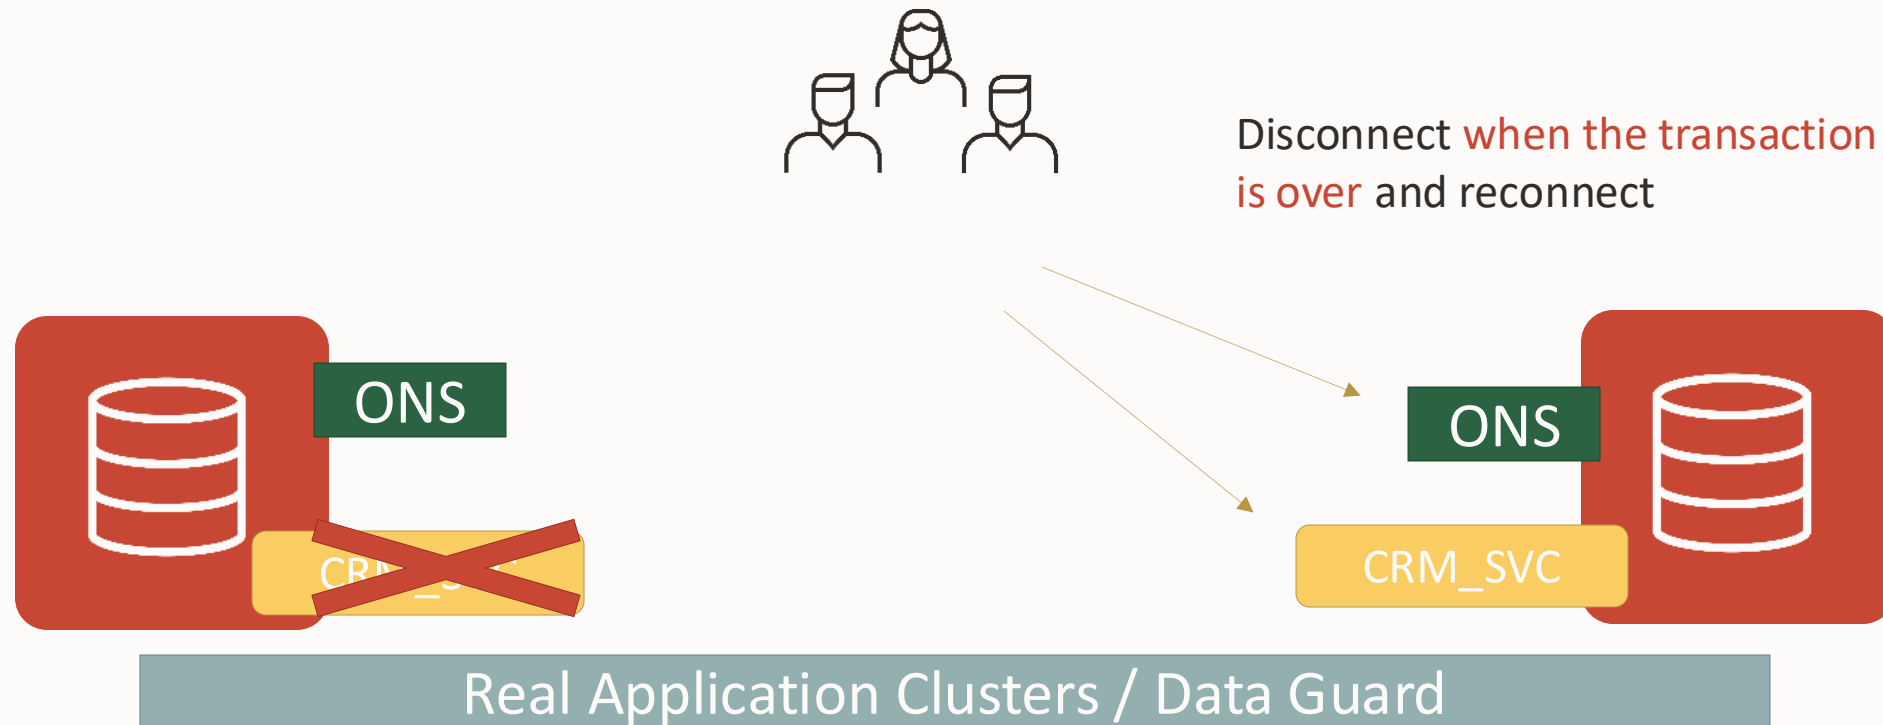
Fast Application Notification

Session Draining for planned maintenance



Fast Application Notification

Session Draining for planned maintenance



Fast Connection Failover (FCF)

FAN integrated in connection pools



- Pre-configured FAN integration
- Uses connection pools
- The application must be pool aware
 - (borrow/release)
- The connection pool leverages FAN events to:
 - Remove quickly dead connections on a DOWN event
 - (opt.) Rebalance the load on a UP event

Fast Connection Failover (FCF)

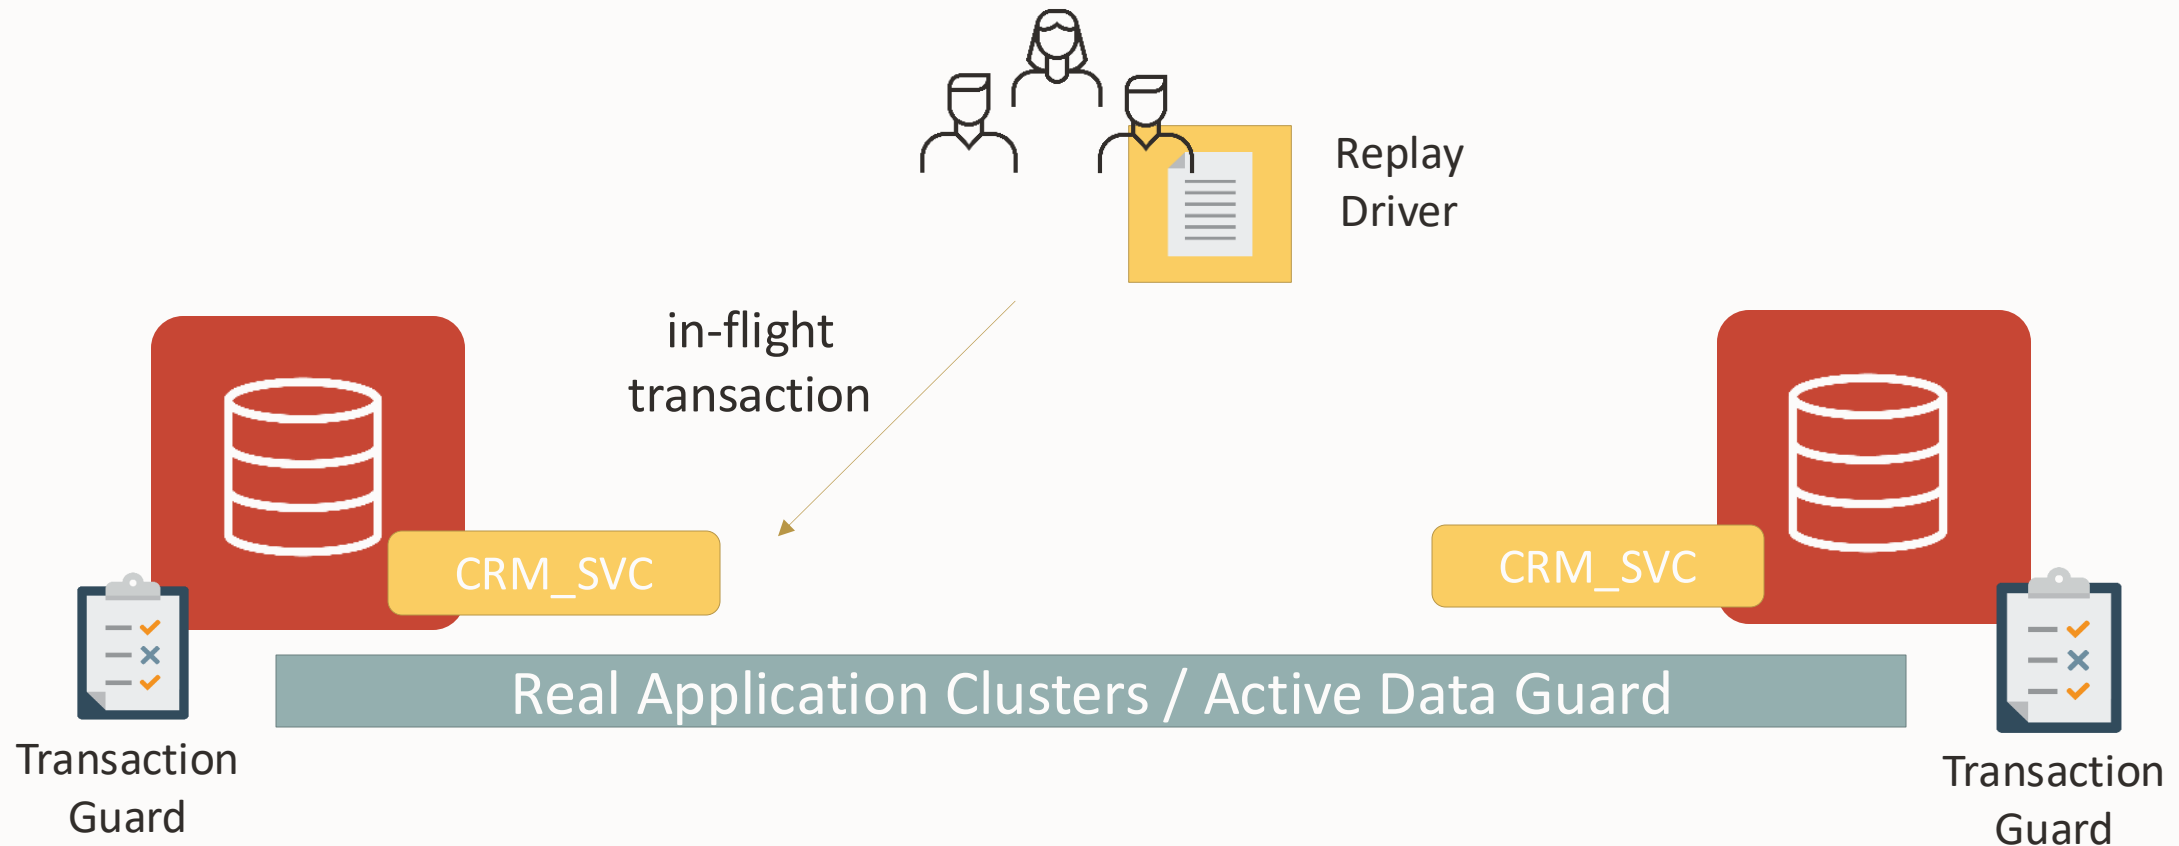
FAN integrated in connection pools



- **UCP** (Universal Connection Pool, ucp.jar) and WebLogic **Active GridLink** handle FAN out of the box.
No code changes! Just enable **FastConnectionFailoverEnabled**
- Third-party connection pools can implement FCF
 - If JDBC driver version ≥ 12.2
 - simplefan.jar and ons.jar in CLASSPATH
 - Connection validation options are set in pool properties
 - Connection pool can plug **javax.sql.ConnectionPoolDataSource**
 - Connection pool checks connections at borrow/release

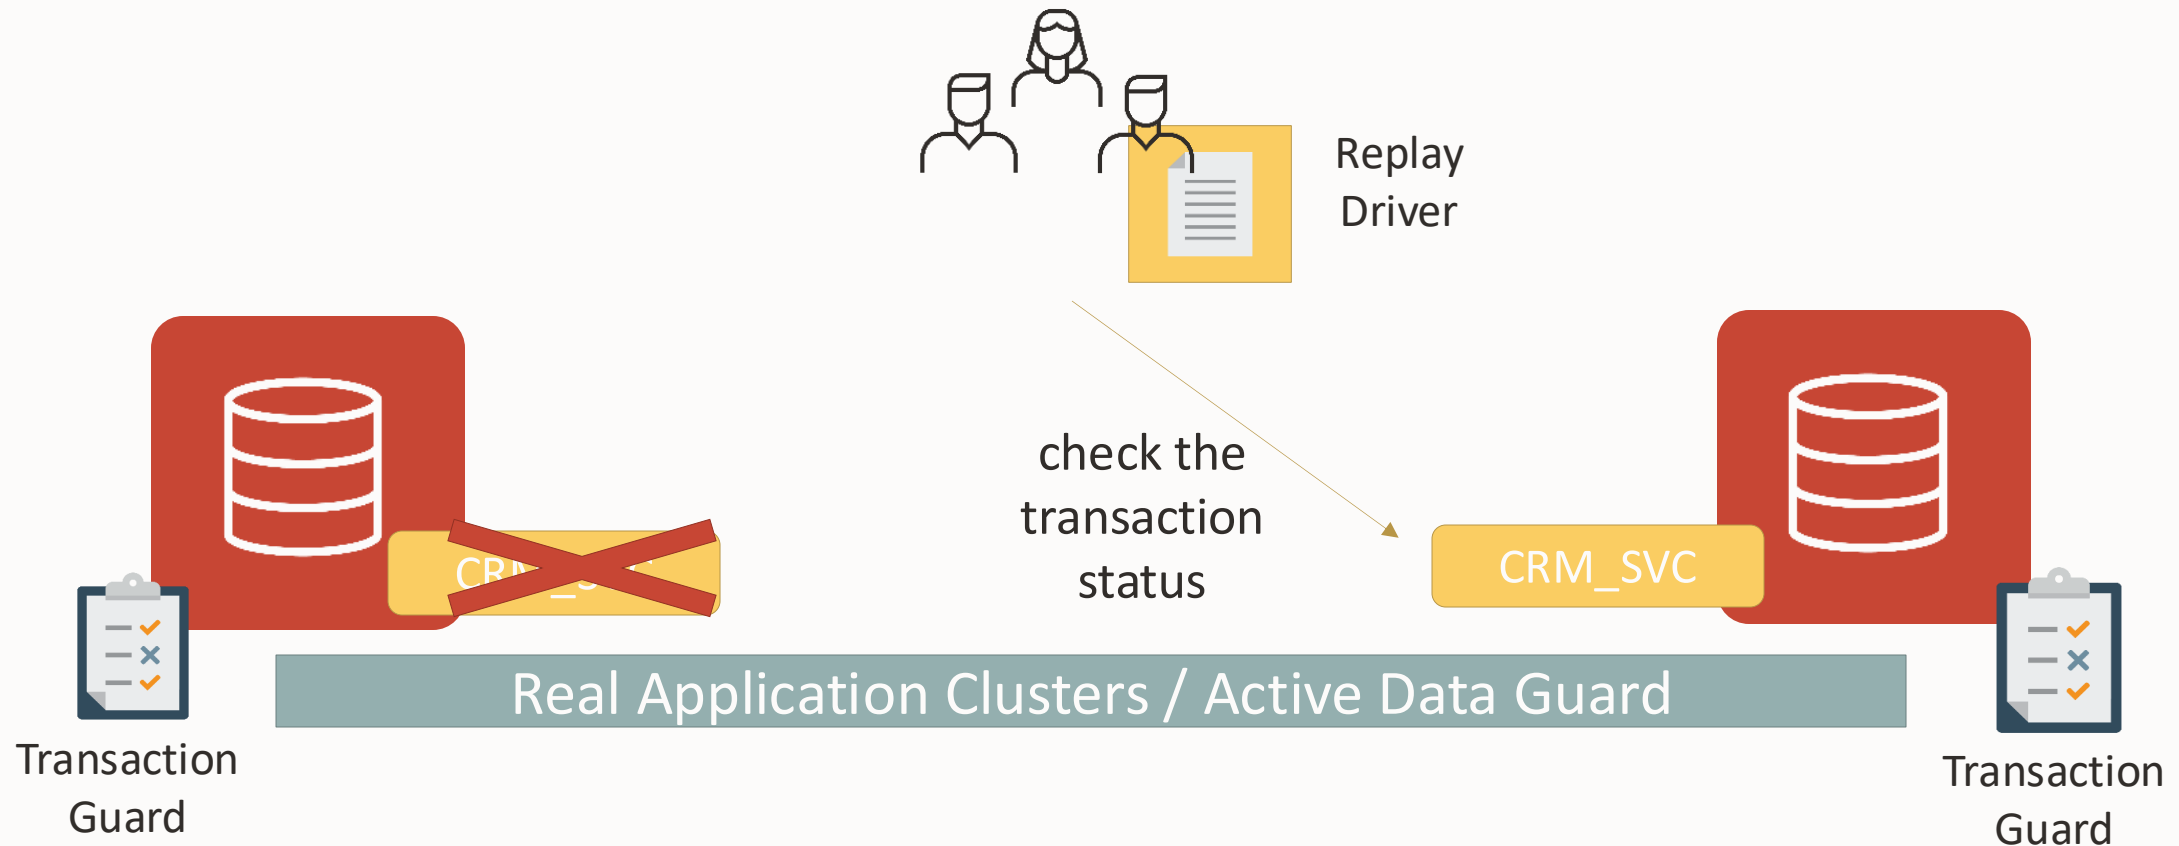
Application Continuity (AC)

Protects the in-flight transaction from failures and disconnections



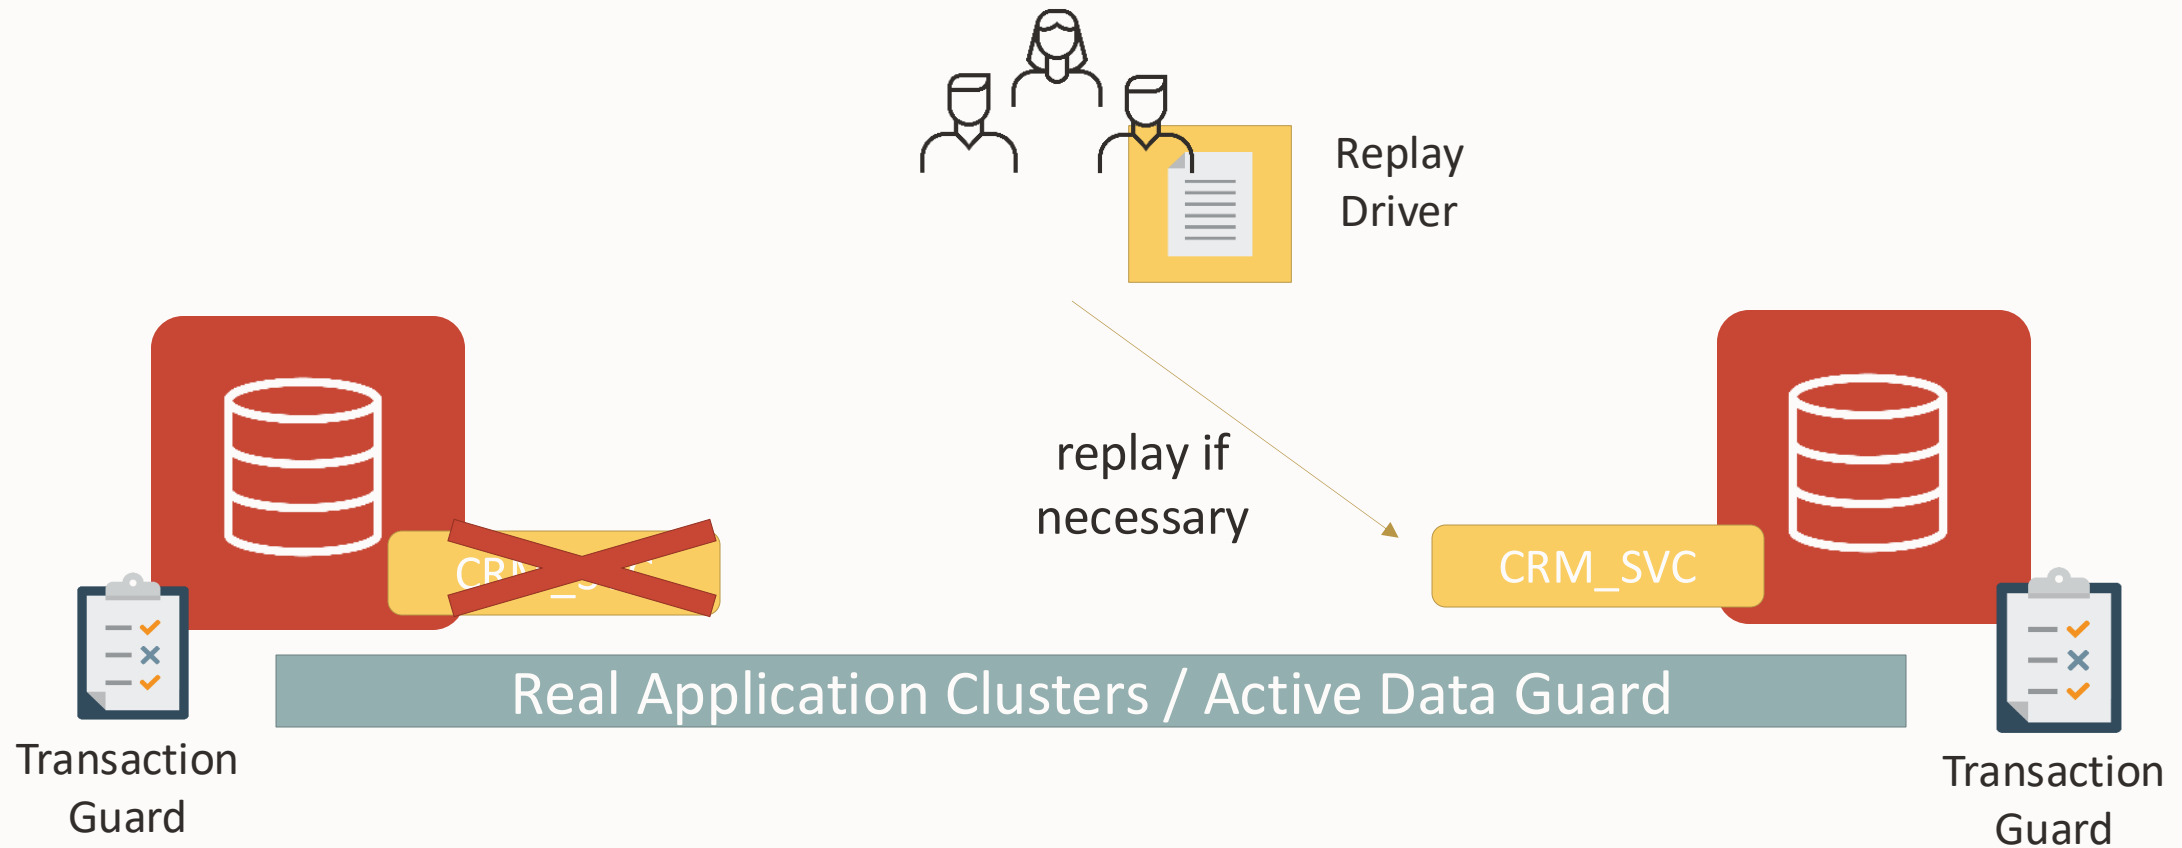
Application Continuity (AC)

Protects the in-flight transaction from failures and disconnections



Application Continuity (AC)

Protects the in-flight transaction from failures and disconnections



Application Continuity (AC)

Protects the in-flight transaction from failures and disconnections

- AC with UCP: no code change

```
PoolDataSource pds = PoolDataSourceFactory.getPoolDataSource();
pds.setConnectionFactoryClassName("oracle.jdbc.replay.OracleDataSourceImpl");
...
conn = pds.getConnection();           // Implicit database request begin
// calls protected by Application Continuity
conn.close();                         // Implicit database request end
```

- AC without connection pool: code change

```
OracleDataSourceImpl ods = new OracleDataSourceImpl();
conn = ods.getConnection();
...
((ReplayableConnection) conn).beginRequest(); // Explicit database request begin
// calls protected by Application Continuity
((ReplayableConnection) conn).endRequest();   // Explicit database request end
```

Transparent Application Continuity (TAC)

Application Continuity for every connection and application type

- Introduced in **18c** for JDBC thin, **19c** for OCI (Oracle Call Interface)
- Records session and transaction state server-side
- No application change
- Works without connection pools (although they are still recommended)
- Replayable transactions are replayed
- Non-replayable transactions raise exception
- **Good driver coverage but check the doc!**
- **Side effects are never replayed**

Key Differences between FAN, AC, and TAC



	Best for	Since Version	Application Changes	Requires Connection Pool	Replay Side Effects	JDBC/OCI
FAN	Planned Maintenance	10g	Catch FAN events (or use UCP)	No, but recommended (FCF)	N/A	Both
AC	Unplanned Outage	12c	Use explicit boundaries (or use UCP)	Yes	Yes (Choose)	Both
TAC (Recommended)	Unplanned Outage	19c	No	No, but recommended	Never	Both



Help Center

Database / Oracle / Oracle Database / Release 19

JDBC Developer's Guide

- Table of Contents
- List of Tables
- Title and Copyright Information
- + Preface
- + Changes in This Release for Oracle Database JDBC Developer's Guide
- + Part I Overview
- + Part II Oracle JDBC
- + Part III Connection and Security
- + Part IV Data Access and Manipulation
- + Part V Performance and Scalability
- Part VI High Availability
 - + 27 Transaction Guard for Java
 - + [28 Application Continuity for Java](#)
 - + 29 Oracle JDBC Support for FAN Events
 - + 30 Transparent Application Failover
 - + 31 Single Client Access Name
- + Part VII Transaction Management
- + Part VIII Manageability
- + Appendixes
- Index

Documentation!

JDBC

OCI



Help Center

Database / Oracle / Oracle Database / Release 19

Programmer's Guide

- + 8 Describing Schema Metadata
- + 9 LOB and BFILE Operations
- + 10 Managing Scalable Platforms
- + 11 Session Pooling and Connection Pooling in OCI
- [12 High Availability in OCI](#)
 - 12.1 Runtime Connection Load Balancing
 - + 12.2 HA Event Notification
 - + 12.3 Transparent Application Failover in OCI
 - + 12.4 OCI and Transaction Guard
 - + 12.5 OCI and Application Continuity
- + 13 Notification Methods and Database Advanced Queuing
- + 14 User-Defined Callback Functions in OCI
- + 15 Performance Topics
- + 16 Database Startup and Shutdown
- + 17 Support for Pluggable Databases






Fast-Start Failover: an overview

Network partitioning

When did the primary disconnect?



```
-- THE LAST TIME THE STANDBY HEARD FROM THE PRIMARY (1 second tolerance)
SQL> select datum_time, (sysdate-to_date(datum_time,'MM/DD/YYYY HH24:MI:SS'))*86400 secs_ago
2>   from v$dataguard_stats where name='transport lag';
```

DATUM_TIME	SECS_AGO
07/11/2022 08:28:46	90361




The columns might be null in some cases



Do not rely on the transport lag value

Network partitioning

Up to which point can the standby recover?



```
-- THE TIMESTAMP OF THE LAST REDO ENTRY RECEIVED FROM THE PRIMARY
SQL> alter session set nls_date_format='MM/DD/YYYY HH24:MI:SS';
SQL> select coalesce(max(s.last_time), max(a.next_time)) as last_redo_from_prim,
2>    (sysdate-coalesce(max(s.last_time), max(a.next_time)))*86400 as secs_ago
3>    from v$standby_log s , v$archived_log a ;
```

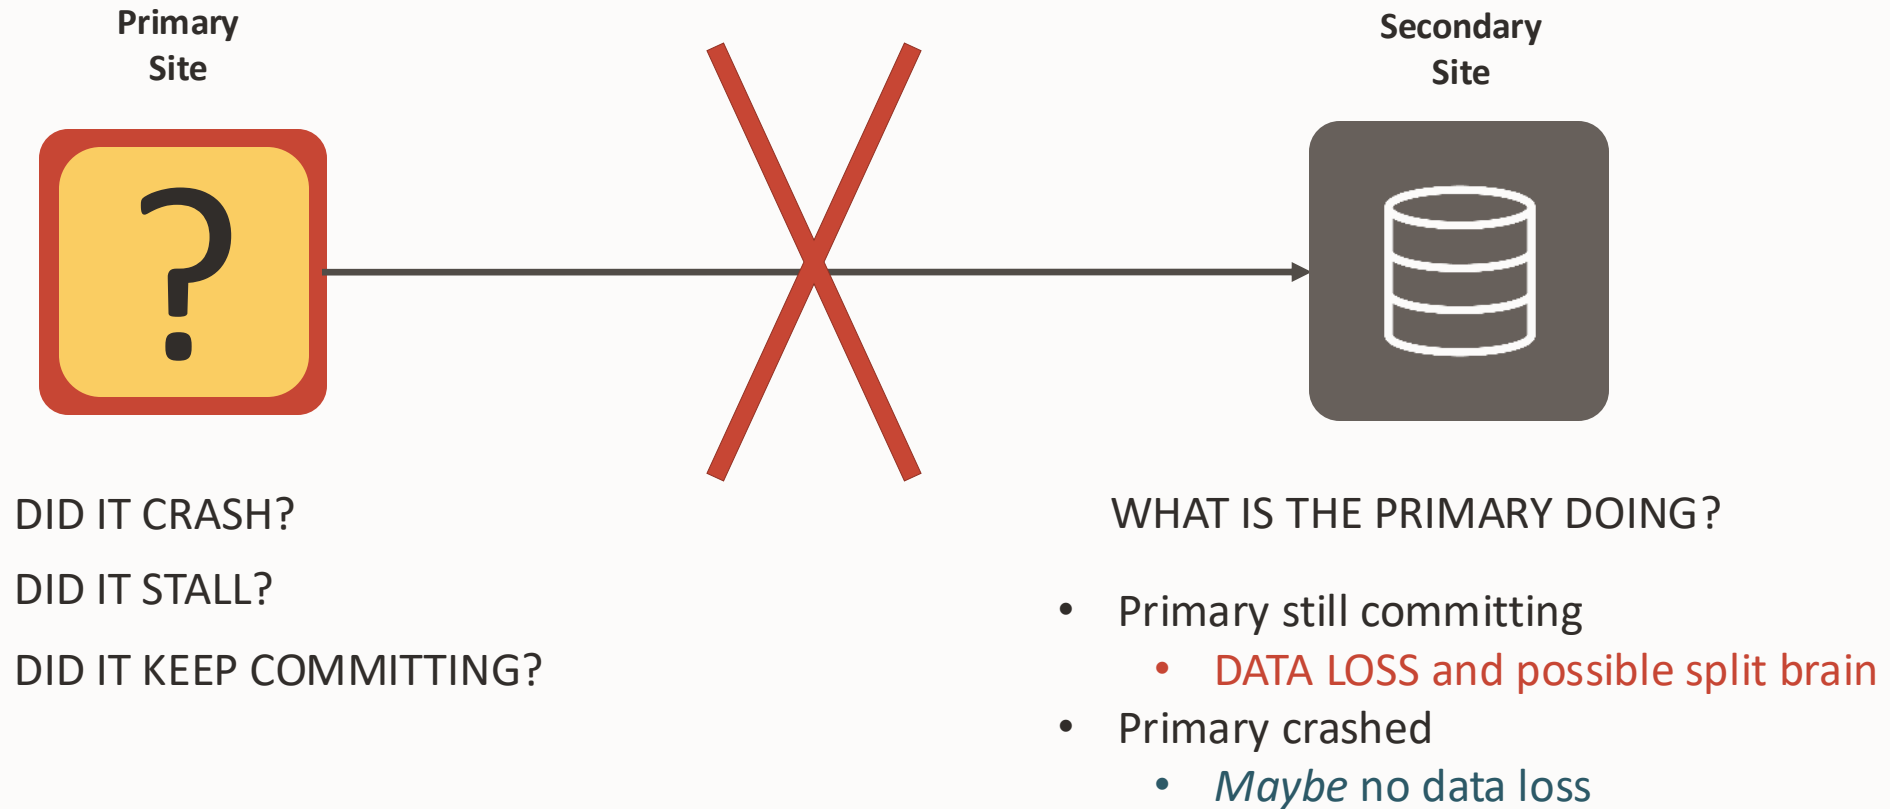
LAST_REDO_FROM_PRIM	SECS_AGO
07/11/2022 08:28:46	91061



Is it a reliable way to calculate data loss?


Network partitioning

Is there a way to calculate the data loss upon failover?



Network partitioning

When did the primary disconnect?



```
-- THE LAST TIME THE STANDBY HEARD FROM THE PRIMARY (1 second tolerance)
SQL> select datum_time, (sysdate-to_date(datum_time,'MM/DD/YYYY HH24:MI:SS'))*86400 secs_ago
2>   from v$dataguard_stats where name='transport lag';
```

DATUM_TIME	SECS_AGO
07/11/2022 08:28:46	90361




The columns might be null in some cases



Do not rely on the transport lag value

Network partitioning

Up to which point can the standby recover?



```
-- THE TIMESTAMP OF THE LAST REDO ENTRY RECEIVED FROM THE PRIMARY
SQL> alter session set nls_date_format='MM/DD/YYYY HH24:MI:SS';
SQL> select coalesce(max(s.last_time), max(a.next_time)) as last_redo_from_prim,
2>    (sysdate-coalesce(max(s.last_time), max(a.next_time)))*86400 as secs_ago
3>   from v$standby_log s , v$archived_log a ;
```

LAST_REDO_FROM_PRIM	SECS_AGO
07/11/2022 08:28:46	91061



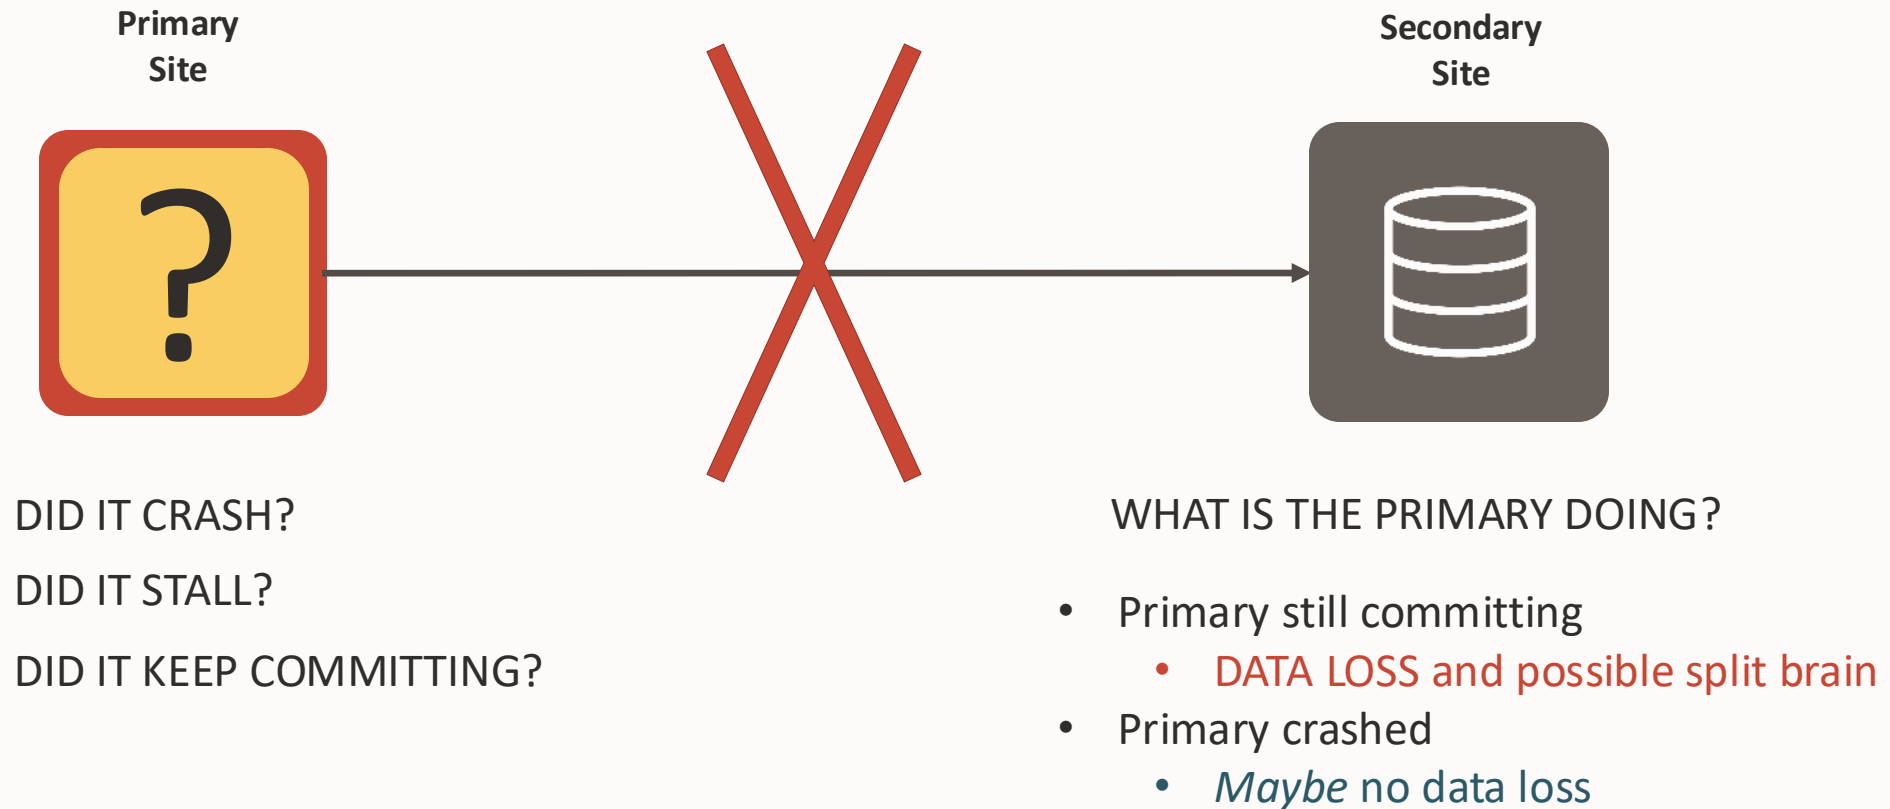
Is it a reliable way to calculate data loss?

Network partitioning

Is there a way to calculate the data loss upon failover?




Can I safely fail over to the standby site?



Oracle Fast-Start Failover introduces a quorum

Automatic failover when the Primary Database is unavailable



- The observer monitors both primary and standby
- **Primary has the quorum** (standby is isolated):
 - The primary keeps writing
- **Primary and standby have the quorum** (observer is isolated):
 - The configuration keeps working unobserved
- **Standby has the quorum** (primary is isolated):
 - Failover!
The primary loses the quorum and stops committing
- The observer can work in “**OBSERVE ONLY**” mode
 - Reports a failure without failing over

Fast-Start Failover automates the failover and solves important problems



Recovery Point Objective is honored

No automatic failover if the data loss breaches your RPO.

RPO can be set to ZERO (no data loss).

Recovery Time Objective improves


Automatic failover begins immediately after the primary is not reachable for more than a specified threshold.

Split-brain protection

The quorum mechanism intrinsically prevents having two primary databases after an automatic failover.

Fast-Start Failover callouts

Execute custom actions before and after the automatic failover occurs




```
$ cat $DG_ADMIN/config_ConfigName/callout/fsfocallout.ora

# The pre-callout script is run before failover
FastStartFailoverPreCallout=fsfo_precallout.sh
FastStartFailoverPreCalloutTimeout=1200
FastStartFailoverPreCalloutSucFileName=fsfo_precallout.suc
FastStartFailoverPreCalloutErrorFileName=precallout.err
FastStartFailoverActionOnPreCalloutFailure=STOP

# The post-callout script is run after failover succeeds
FastStartFailoverPostCallout=fsfo_postcallout.sh
$
```

Fast Start Failover Configuration Validation

Ensure everything is configured properly for the automatic failover



```
DGMGRL> VALIDATE FAST_START FAILOVER;
Fast-Start Failover: Enabled in Potential Data Loss Mode
Protection Mode:      MaxPerformance
Primary:               North_Sales
Active Target:         South_Sales
```

Fast-Start **Failover Not Possible**:
Fast-Start Failover observer not started


Post Fast-Start **Failover Issues**:
Flashback database disabled for database 'dgv1'

Other issues:
FastStartFailoverThreshold may be too low for RAC databases.

Fast-start failover **callout configuration** file "fsfocallout.ora" has the following **issues**:
Invalid lines
foo=foo
The specified file "./precallout" contains a path.

Easier checking of Fast-Start Failover configurations

The new fixed view **V\$FAST_START_FAILOVER_CONFIG** shows the Fast-Start Failover settings and status



```
SQL> desc V$FAST_START_FAILOVER_CONFIG;
```

Name	Null?	Type
-----	-----	-----
FSFO_MODE		VARCHAR2(19)
STATUS		VARCHAR2(22)
CURRENT_TARGET		VARCHAR2(30)
THRESHOLD		NUMBER
OBSERVER_PRESENT		VARCHAR2(7)
OBSERVER_HOST		VARCHAR2(512)
PING_INTERVAL		NUMBER
PING_RETRY		NUMBER
PROTECTION_MODE		VARCHAR2(30)
LAG_LIMIT		NUMBER
AUTO_REINSTATE		VARCHAR2(5)
OBSERVER_RECONNECT		NUMBER
OBSERVER_OVERRIDE		VARCHAR2(5)
SHUTDOWN_PRIMARY		VARCHAR2(5)
CON_ID		NUMBER

```
SQL> SELECT fsfo_mode, status, current_target, threshold, observer_present, observer_host,
2> protection_mode, lag_limit, auto_reinstate, observer_override, shutdown_primary FROM V$FAST_START_FAILOVER_CONFIG;
```

FSFO_MODE	STATUS	CURRENT_TARGET	THRESHOLD	OBSERVE	OBSERVER_HOST	PROTECTION_MODE	LAG_LIMIT	AUTO_	OBSER	SHUTD
-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----
POTENTIAL DATA LOSS	TARGET UNDER LAG LIMIT	mydb_site2	180	YES	mydb-obs	MaxPerformance	300	TRUE	FALSE	TRUE

Note: V\$DATABASE columns starting with FS_FAILOVER_ are therefore deprecated.





Fast-Start Failover Lag Histogram

The view **V\$FS_LAG_HISTOGRAM** displays the frequency of Fast-Start Failover lags.


```
SQL> select * from v$fs_lag_histogram;
```

THREAD#	LAG_TYPE	LAG_TIME	LAG_COUNT	LAST_UPDATE_TIME	CON_ID
1	APPLY	5	122	01/23/2023 10:46:07	0
1	APPLY	10	5	01/02/2023 16:12:42	0
1	APPLY	15	2	12/25/2022 12:01:23	0
1	APPLY	30	0		0
1	APPLY	60	0		0
1	APPLY	120	0		0
1	APPLY	180	0		0
1	APPLY	300	0		0
1	APPLY	65535	0		0

- Useful to calculate the optimal **FastStartFailoverLagTime** property.
- It shows also the most recent occurrence for each bucket.
- **LAG_TIME** is the upper bound of the bucket:
 - 5 -> between 0 and 5 seconds
 - 10 -> between 5 and 10 seconds
 - etc.
- It's calculated every minute, only when Fast-Start Failover is enabled (also in observe-only mode)

Data Guard Broker Client Side Standardized Directory Structure

A single environment variable to define all the locations



New configuration property. It defaults to the Configuration Name

Location of Client-side Broker Files


<https://docs.oracle.com/en/database/oracle/oracle-database/21/dgbkr/using-data-guard-broker-to-manage-switchovers-failovers.html#GUID-0C8473F6-33B5-479F-9208-9CA651F1B483>

Enhanced observer diagnostic

New columns in V\$FS_FAILOVER_OBSERVERS with additional details

New columns:



- LAST_PING_PRIMARY
- LAST_PING_TARGET
- LOG_FILE
- STATE_FILE
- CURRENT_TIME



```
SQL> select name, registered, host, ismaster, pinging_primary, pinging_target ,
2> last_ping_primary, last_ping_target, log_file, state_file, current_time
3> from V$FS_FAILOVER_OBSERVERS where host is not null;
```

NAME	REGI	HOST	ISMA	PING	PING	LAST_PING_PRIMARY	LAST_PING_TARGET	LOG_FILE	STATE_FILE	CURRENT_TIME
host-obs	YES	host-obs	YES	YES	YES	0	2	/.../observer.lst	/.../observer.dat	06-OCT-22 06.38.14.000000000 AM






Fast-Start Failover: Oracle Data Guard Protection Modes

Data Guard Fast-Start Failover Protection Modes

Balance Data Protection with Performance and Availability



What does the primary do if the standby does not acknowledge the transaction?

MAXIMUM PERFORMANCE



Never waits for acknowledge (**ASYNC**).
Some transactions might get lost when primary fails.

MAXIMUM AVAILABILITY

Waits until *NetTimeout* seconds (**SYNC** or **FASTSYNC**), then continue without standby. Data loss possible with manual failover or within specified limit (21c).

MAXIMUM PROTECTION

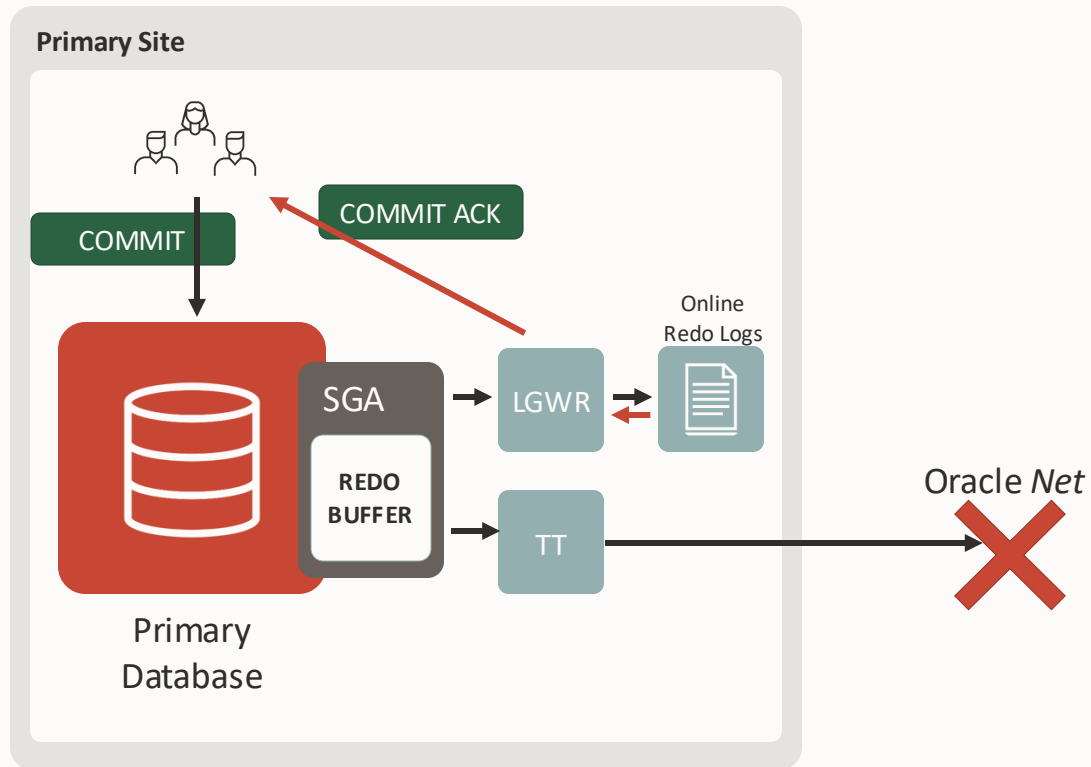
Waits until the standby is available again (**SYNC only**).
No transactions are lost, ever.



Fast-Start Failover: Maximum Performance

the standby is not reachable or is slow:

Data Guard **ASYNC** redo transport

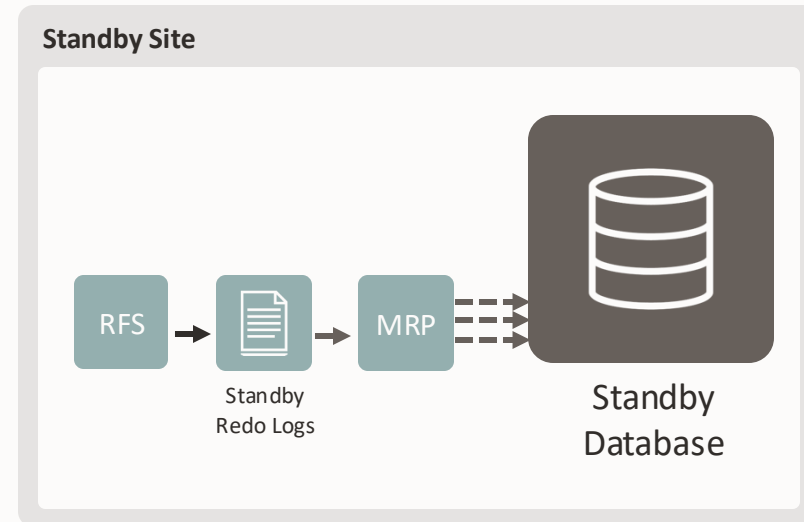


If the standby is not reachable or is slow:

- The primary keeps writing at its pace
- The lag (data loss exposure) increases

If the primary fails:

- The standby requires manual failover

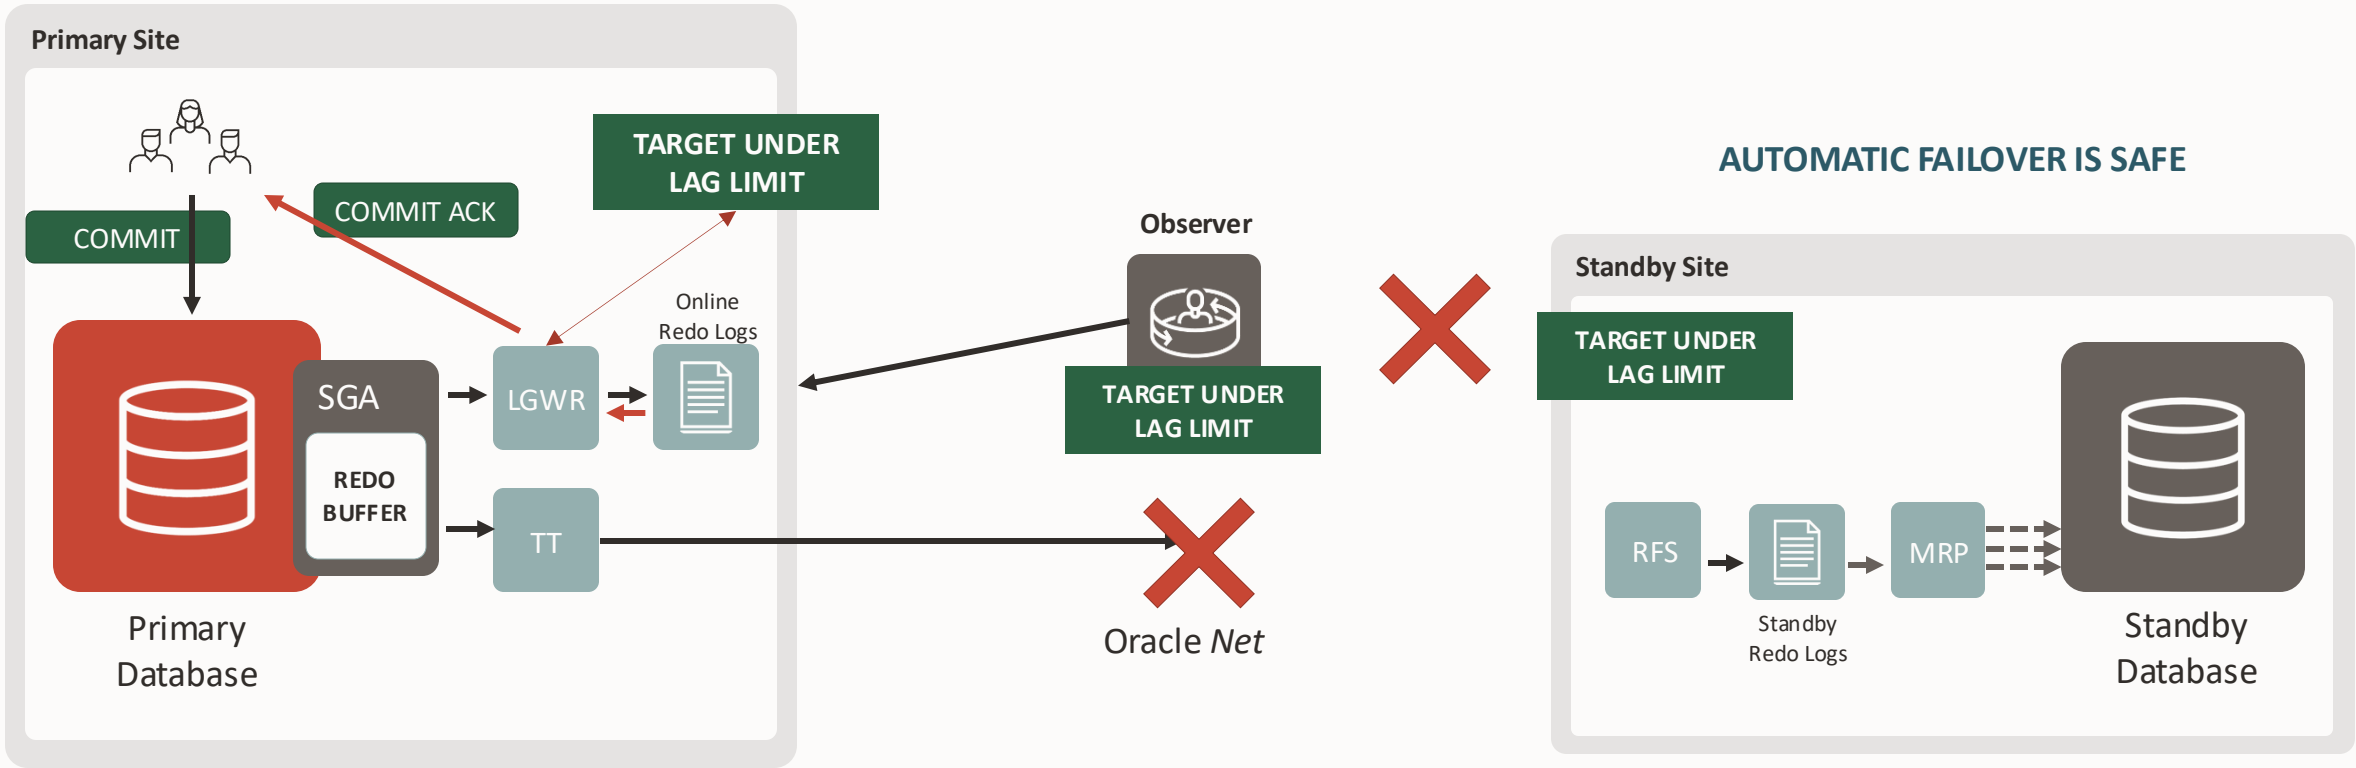


Max Performance with Fast-Start Failover

The primary continuously computes the lag with the Fast-Start Failover Target

If the standby is not reachable or is slow:

- The primary keeps committing until reaching the lag limit

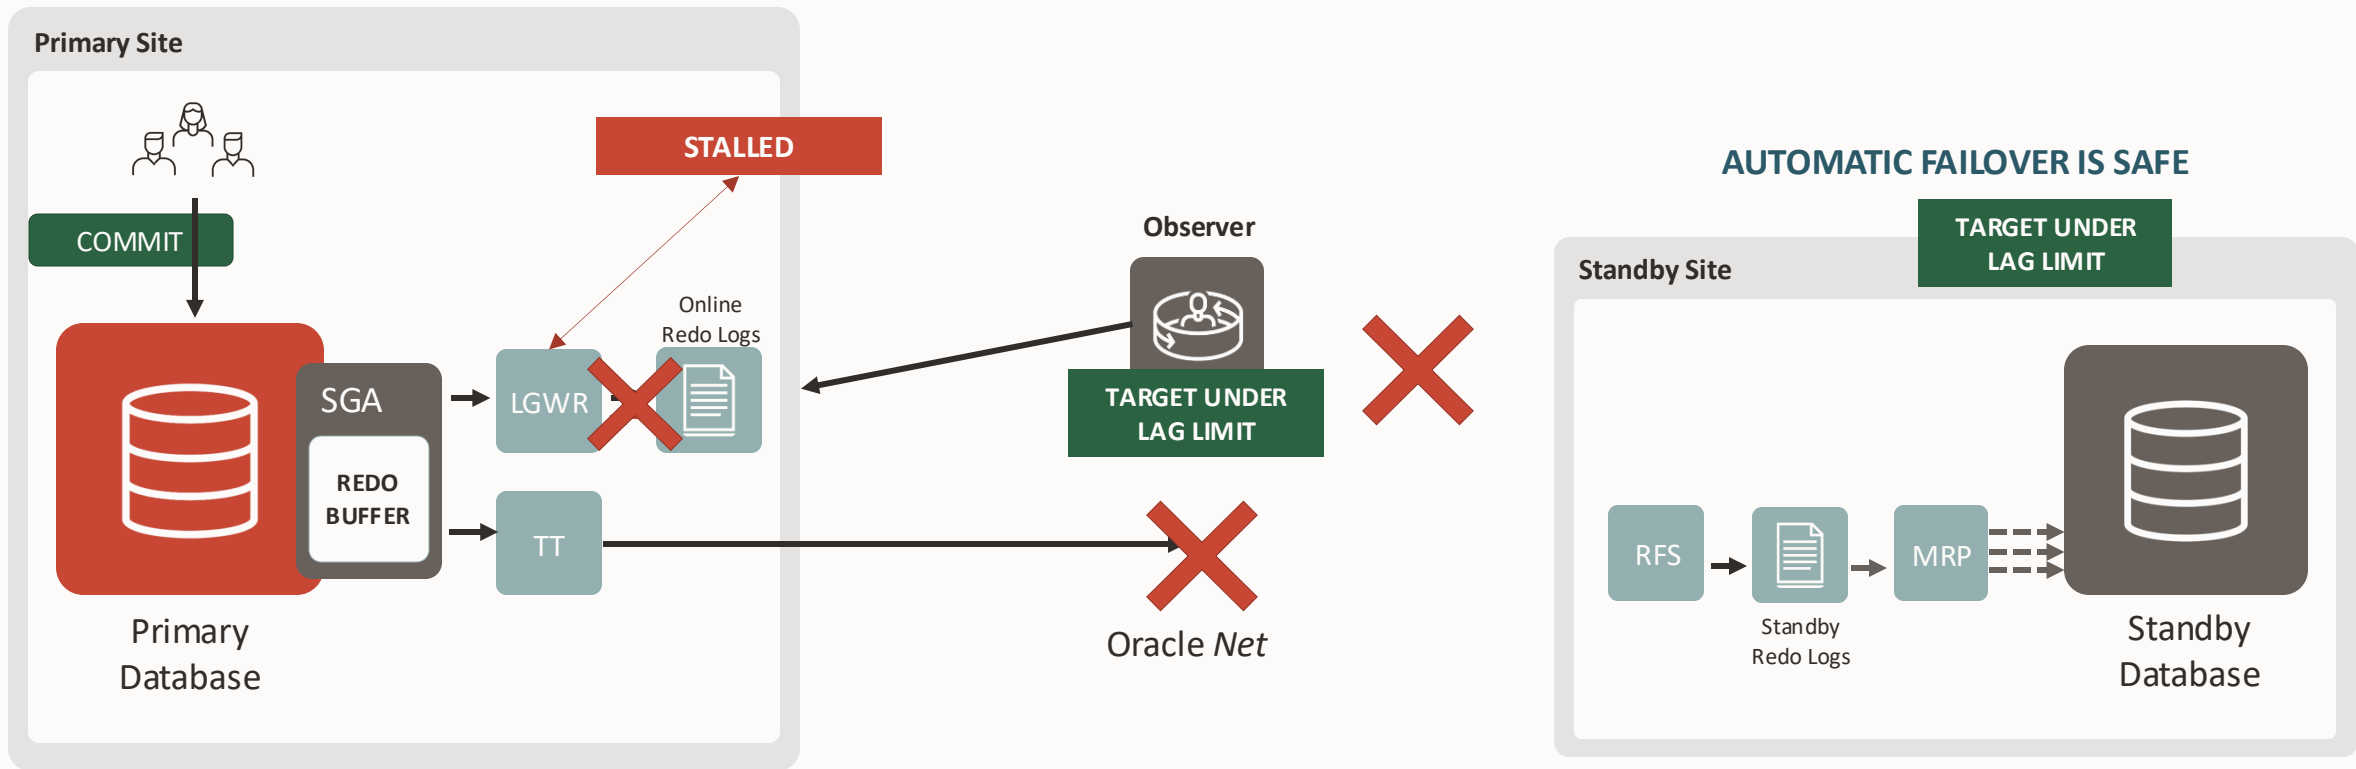


Max Performance with Fast-Start Failover

The primary continuously computes the lag with the Fast-Start Failover Target

If the standby is not reachable or is slow:

- The primary keeps committing until reaching the lag limit
- Then it goes to STALLED mode (no commits possible)

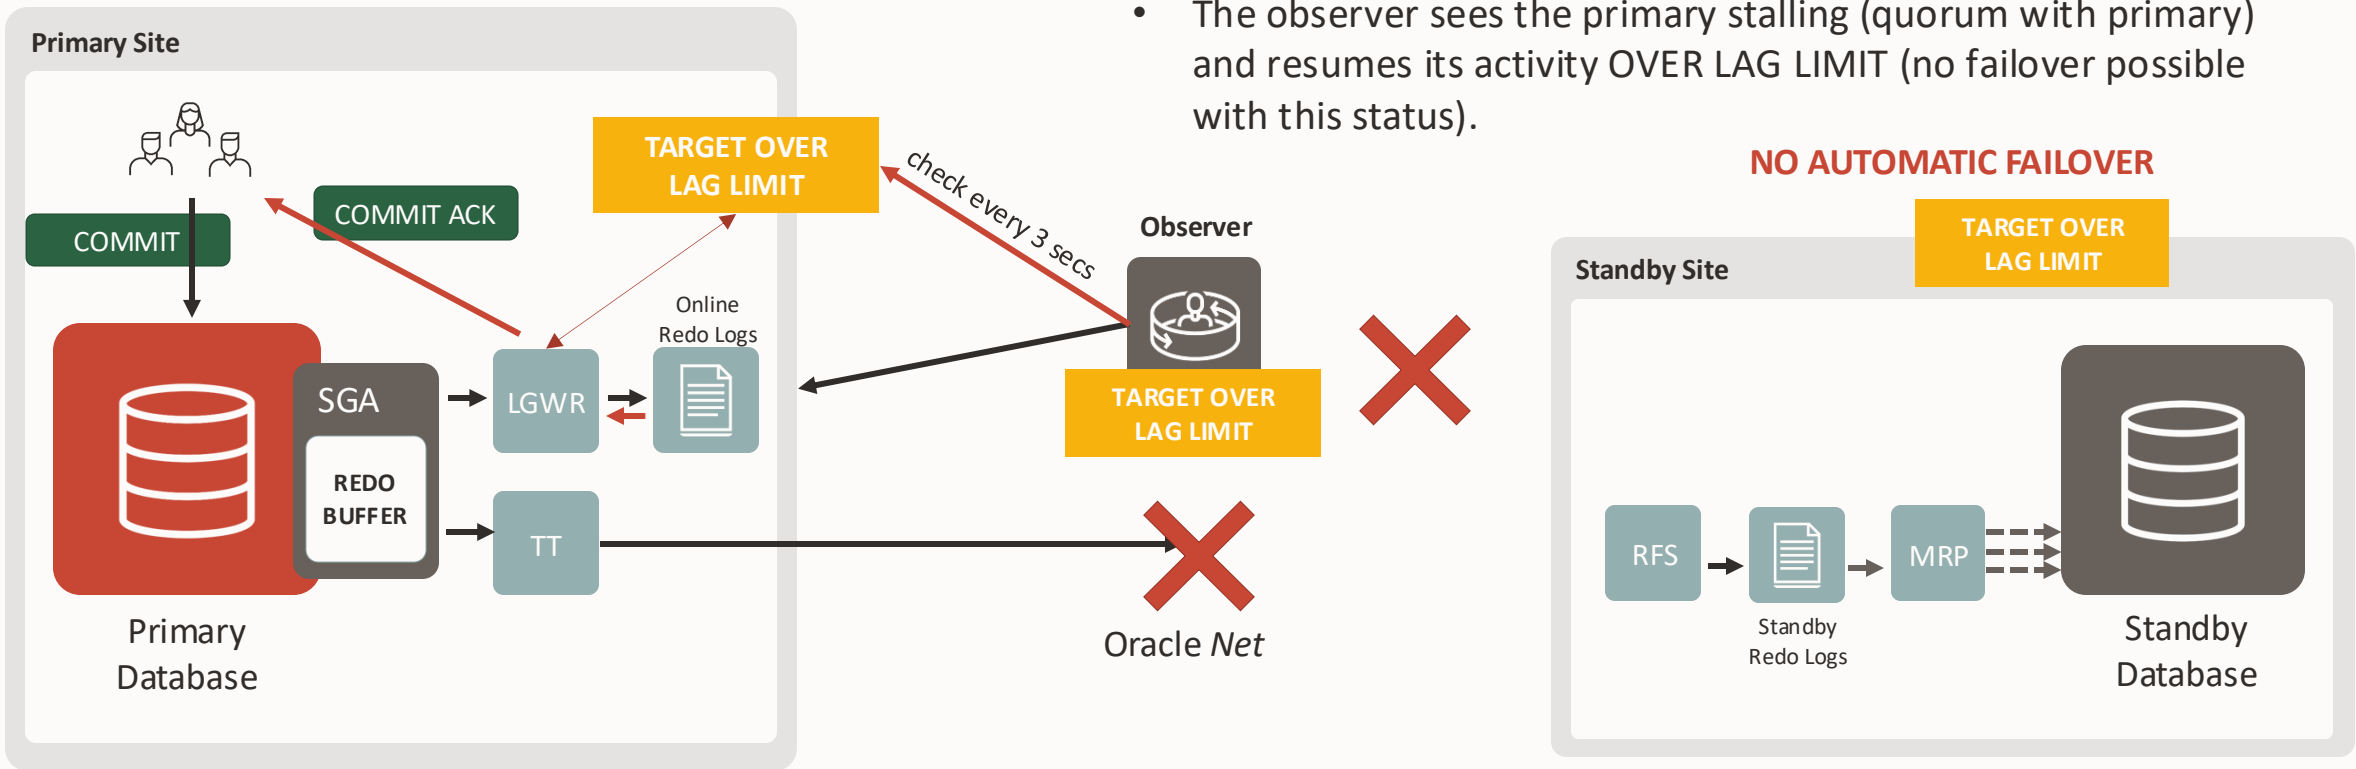


Max Performance with Fast-Start Failover

The primary continuously computes the lag with the Fast-Start Failover Target

If the standby is not reachable or is slow:

- The primary keeps committing until reaching the lag limit
- Then it goes to STALLED mode (no commits possible)
- The observer sees the primary stalling (quorum with primary) and resumes its activity OVER LAG LIMIT (no failover possible with this status).

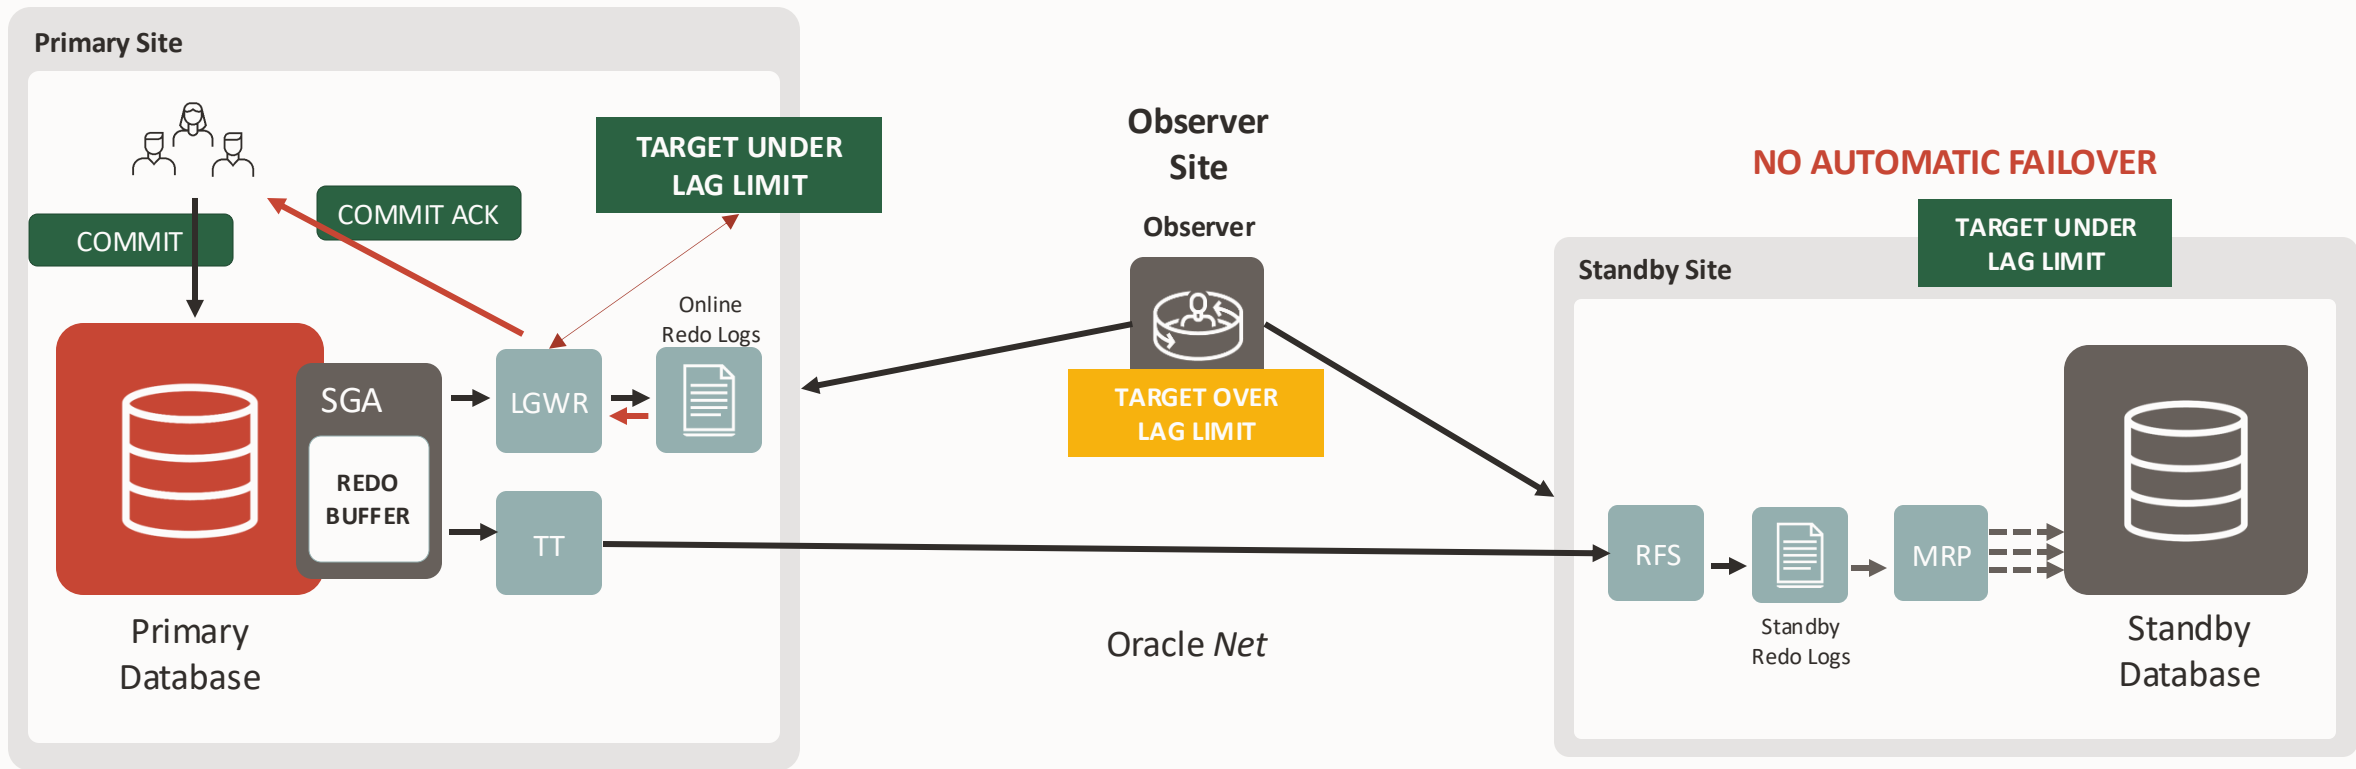


Max Performance with Fast-Start Failover

The primary continuously computes the lag with the Fast-Start Failover Target

When the standby catches up with the primary

- The primary goes autonomously under lag limit

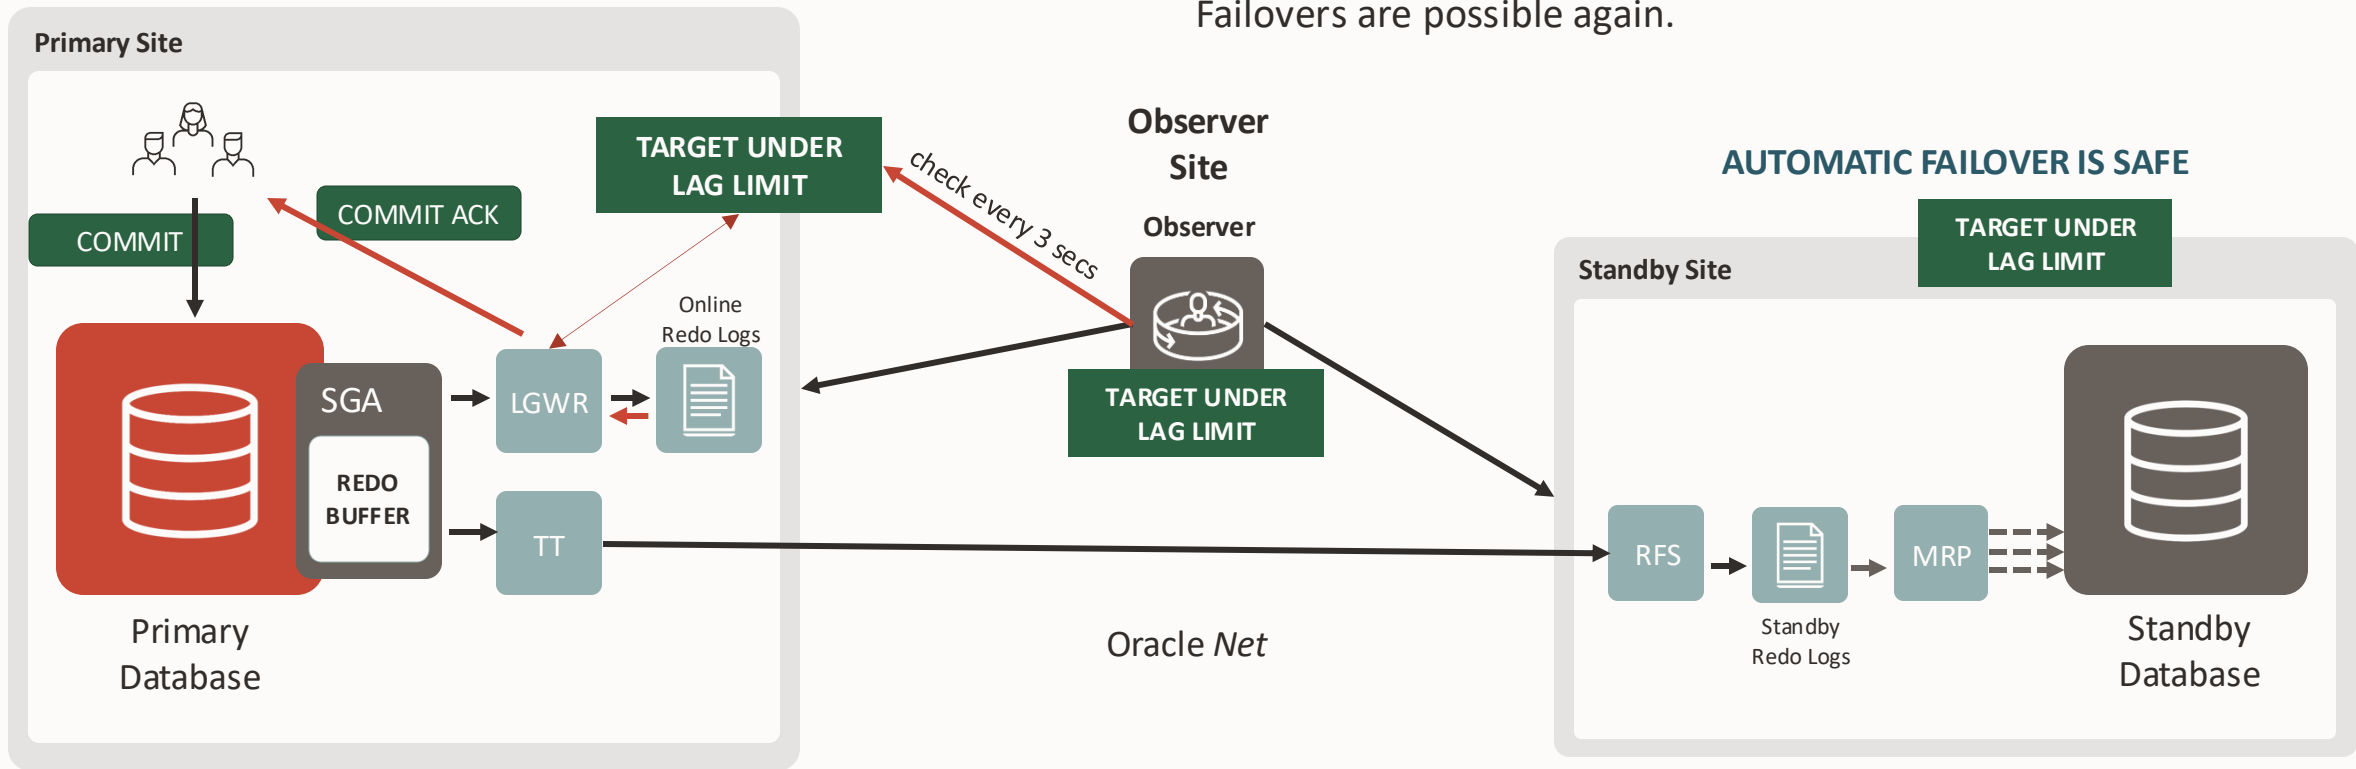


Max Performance with Fast-Start Failover

The primary continuously computes the lag with the Fast-Start Failover Target


When the standby catches up with the primary

- The primary goes autonomously under lag limit
- At the next check, the observer acknowledges the change. Failovers are possible again.



Automatic Failover with MaxPerformance


Choose how much data loss you can tolerate



- 1 ASYNC Transport. The Standby has a residual lag
Status: TARGET UNDER LAG LIMIT
- 2 At t_0 + ping time, the primary cannot contact the standby. The Primary keeps committing, the Standby lag increases.
Status: TARGET UNDER LAG LIMIT
- 3 The primary reaches FastStartFailoverLagLimit. It temporarily stalls (~3 seconds) until it gets permission from the observer to continue.
Status: STALLED
- 4 After the observer pings the primary and gives permission to continue, the primary resumes the commit activity. Status: TARGET OVER LAG LIMIT
- 5 The observer acknowledges that the lag is above the limit and will not permit a failover in case it loses connectivity with the primary. The standby is declared out of sync.
Status: TARGET OVER LAG LIMIT

Automatic Failover with MaxPerformance

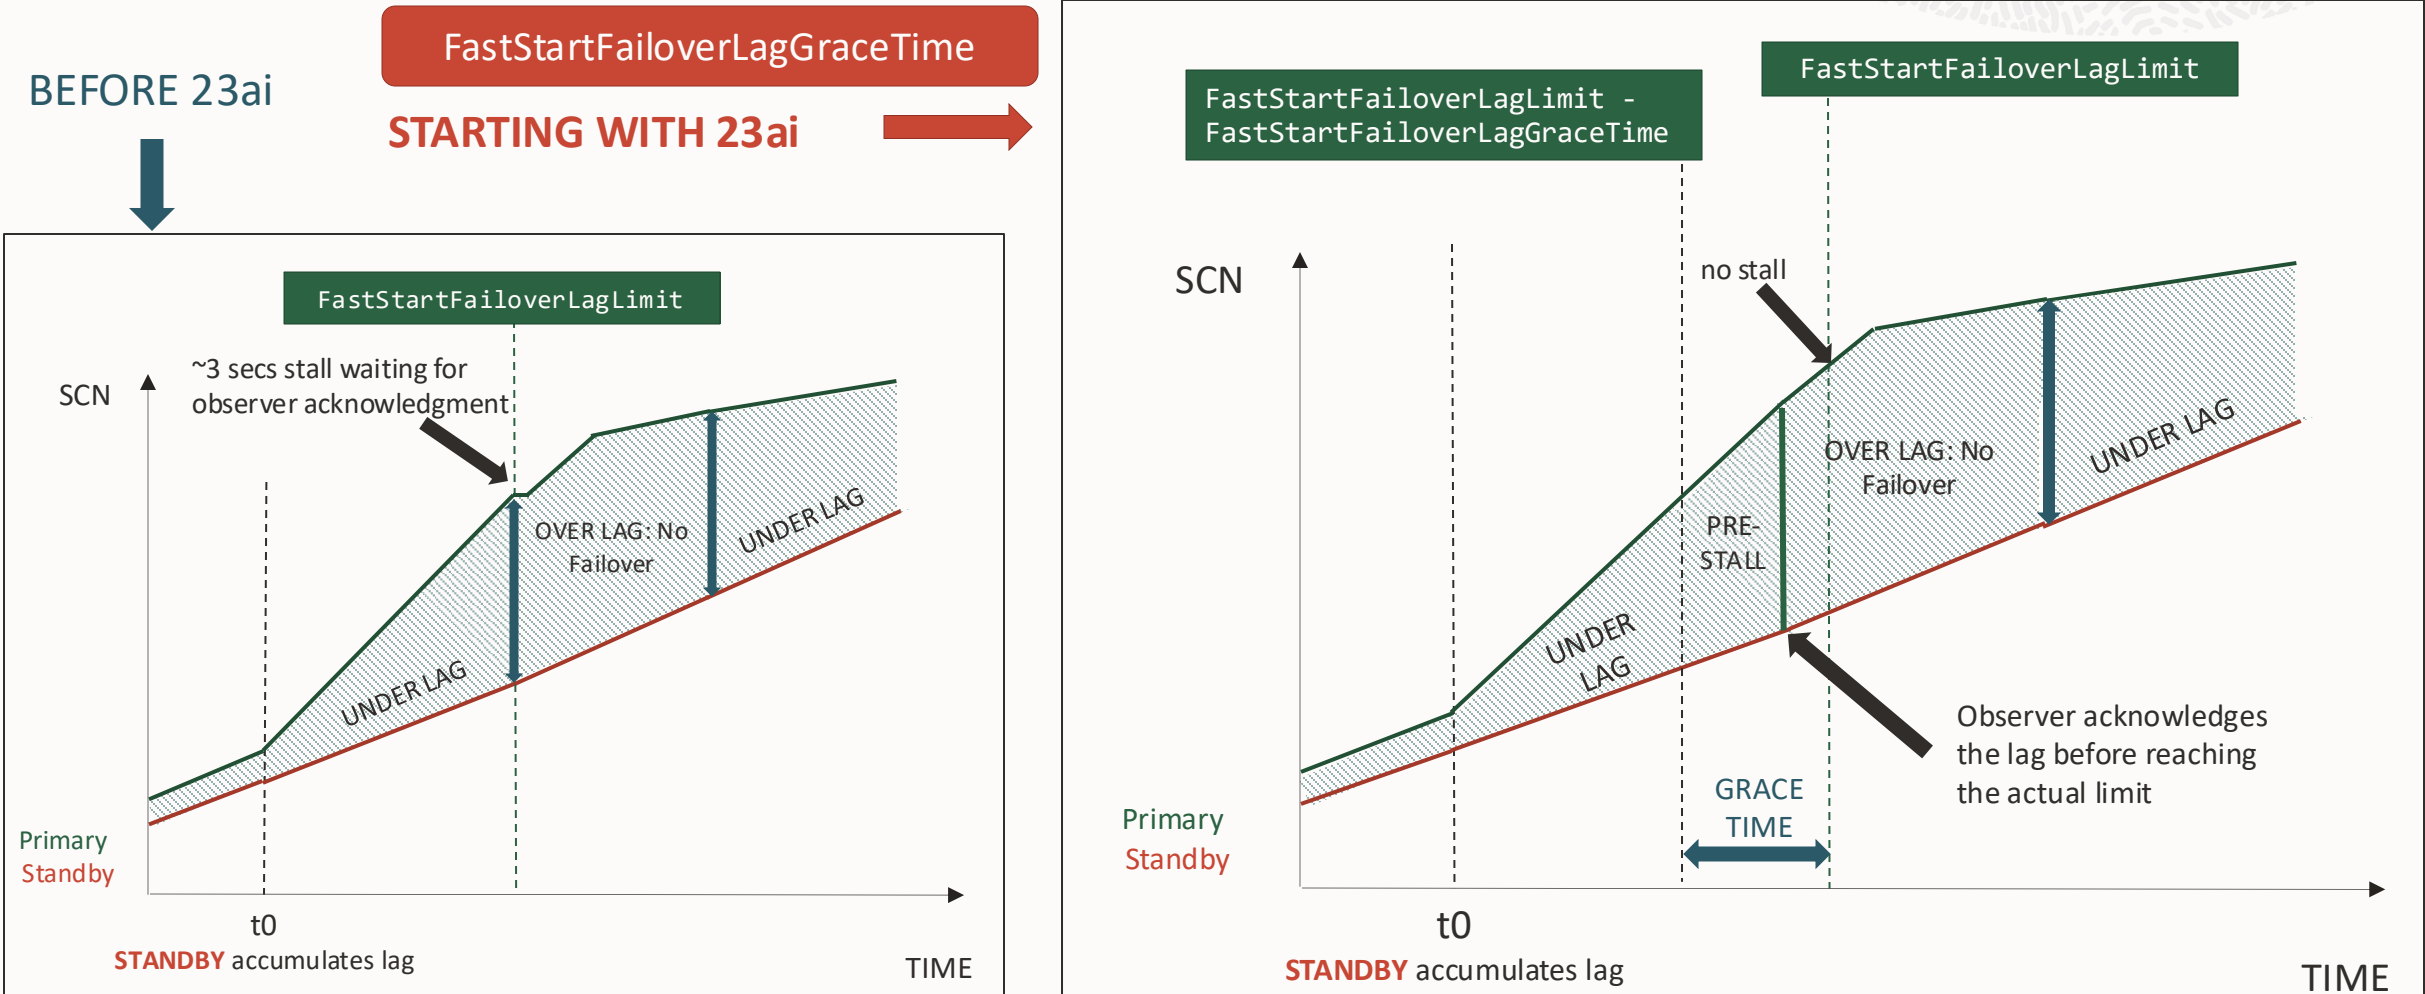
Choose how much data loss you can tolerate



- 1 ASYNC Transport. The Standby has a residual lag
Status: TARGET UNDER LAG LIMIT
- 2 At t_0 the primary increases the activity rate. The Standby lag increases.
Status: TARGET UNDER LAG LIMIT
- 3 The primary reaches `FastStartFailoverLagLimit`. It temporarily stalls (~3 seconds) until it gets permission from the observer to continue.
Status: STALLED
- 4 After the observer pings the primary and gives permission to continue, the primary resumes the commit activity. Status: TARGET OVER LAG LIMIT
- 5 The observer acknowledges that the lag is above the limit and will not permit a failover in case it loses connectivity with the primary. The standby is declared out of sync.
Status: TARGET OVER LAG LIMIT
- 6 The standby catches up with the primary and the lag goes under the limit. The observer can now failover if the primary fails.
Status: TARGET UNDER LAG LIMIT

Minimized Stall in Fast-Start Failover Maximum Performance

Reduce/avoid delays during FSFO state transition to "OVER LAG LIMIT"





* or 19c Patch 34995066: MINIMIZE STALL IN DATA-LOSS FAST-START FAILOVER



Minimized Stall in Fast-Start Failover Maximum Performance

Reduce/avoid delays during FSFO state transition to "OVER LAG LIMIT"



Choose how much data loss you can tolerate

The diagram illustrates the FastStartFailover mechanism. The Y-axis represents the Sequence Number (SCN), and the X-axis represents time. A green line represents the Primary database, and a red line represents the Standby database. At time t_0 , the Primary disconnects. The Standby starts applying redo, creating a 'Potential LAG' (shaded area). A 'STALL' occurs when the lag reaches 'FastStartFailoverLagLimit' (marked with a black X at point 3). The Standby then becomes the 'New primary' (point 4) after a 'DATA LOSS' event (marked with a red arrow at point 4). The 'FastStartFailoverThreshold' is indicated by a red dashed line.

- 83 Copyright © 2024, Oracle and/or its affiliates



Set the FastStartFailoverLagLimit wisely to avoid split-brain conditions



- 84 Copyright © 2024, Oracle and/or its affiliates





Choose the Lag Type for Maximum Performance Mode

New Property FastStartFailoverLagType

The broker can now use the standby's transport lag to determine whether a data loss situation exists. The amount of tolerated data loss is still set with **FastStartFailoverLagLimit**.

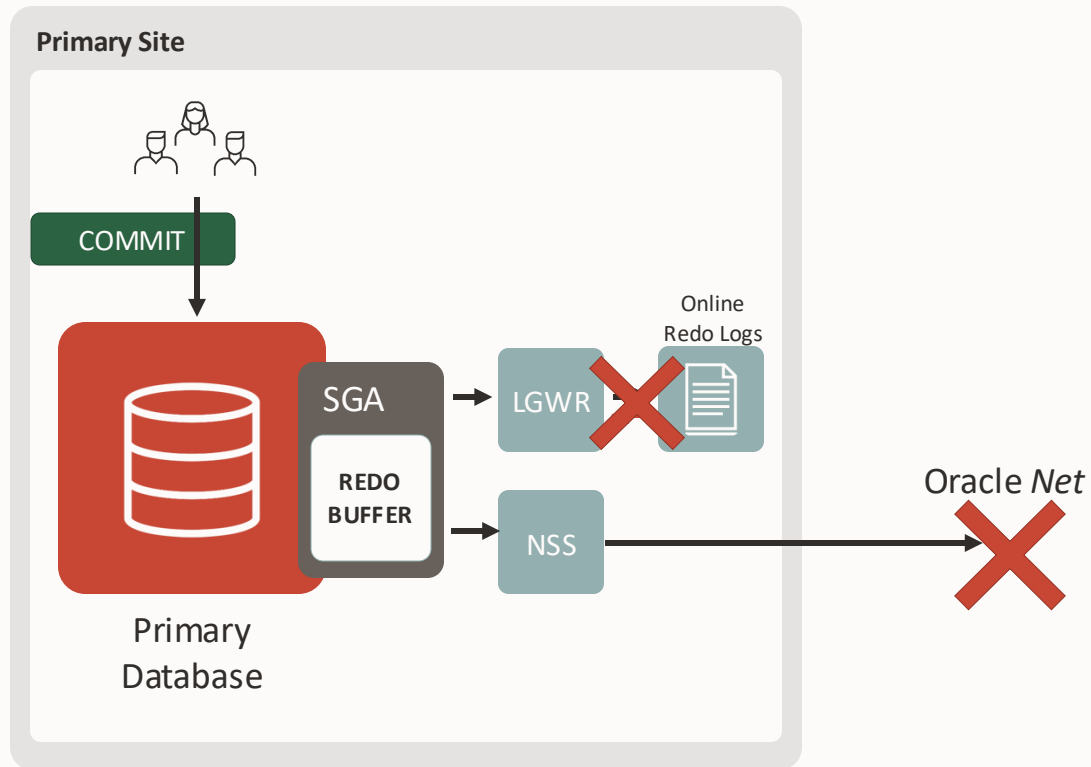
Before 23ai	Starting with 23ai
Only the APPLY lag is used to determine the data loss exposure.	<p>The new property FastStartFailoverLagType property can be used to choose which type of lag should be used.</p> <p>It can be TRANSPORT or APPLY.</p> <p>APPLY is the default to keep the old behavior.</p>



Fast-Start Failover: Maximum Availability

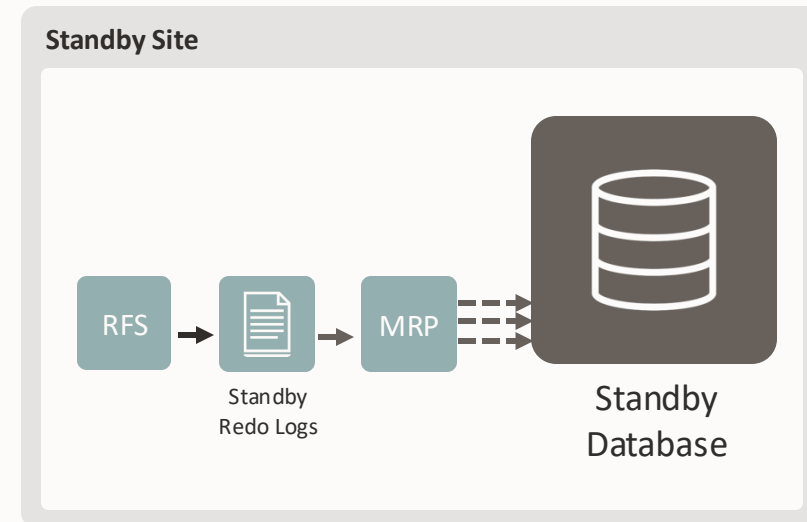
Max Availability without Fast-Start Failover

Data Guard **SYNC** redo transport



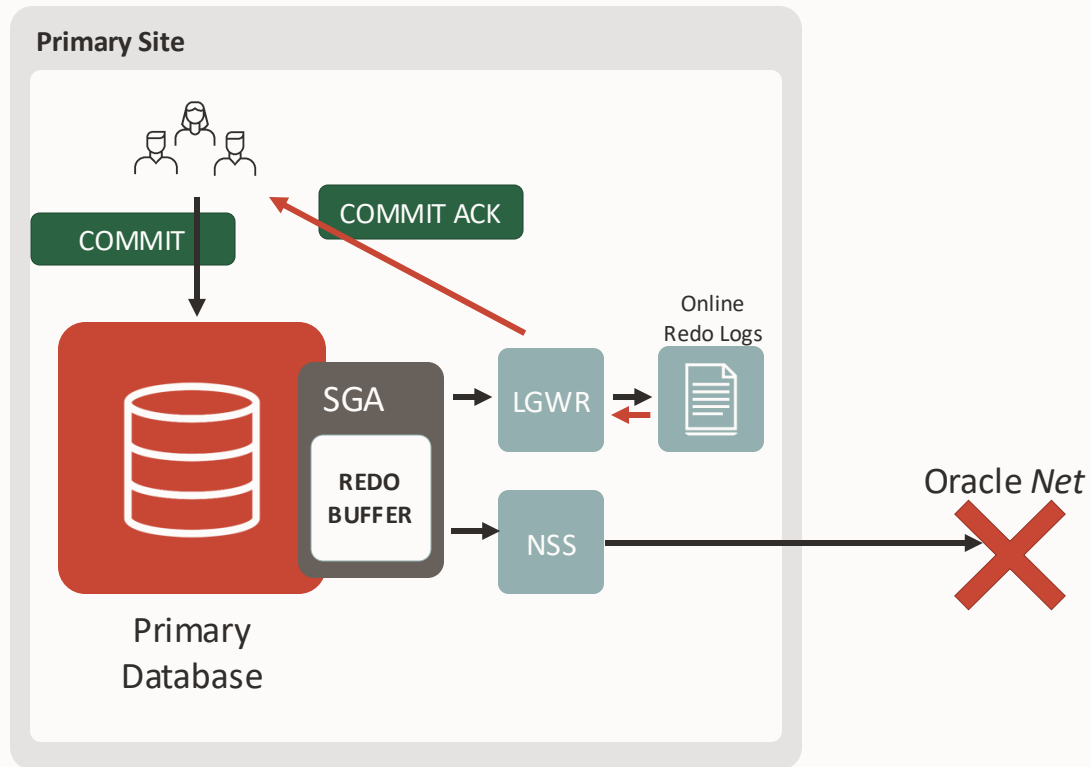
If the standby is not reachable:

- The primary stalls waiting for the SYNC destination (no commits possible)




Max Availability without Fast-Start Failover

Data Guard **SYNC** redo transport



If the standby is not reachable:

- The primary stalls waiting for the SYNC destination (no commits possible)
- After NetTimeout seconds, it resumes the activity without protection

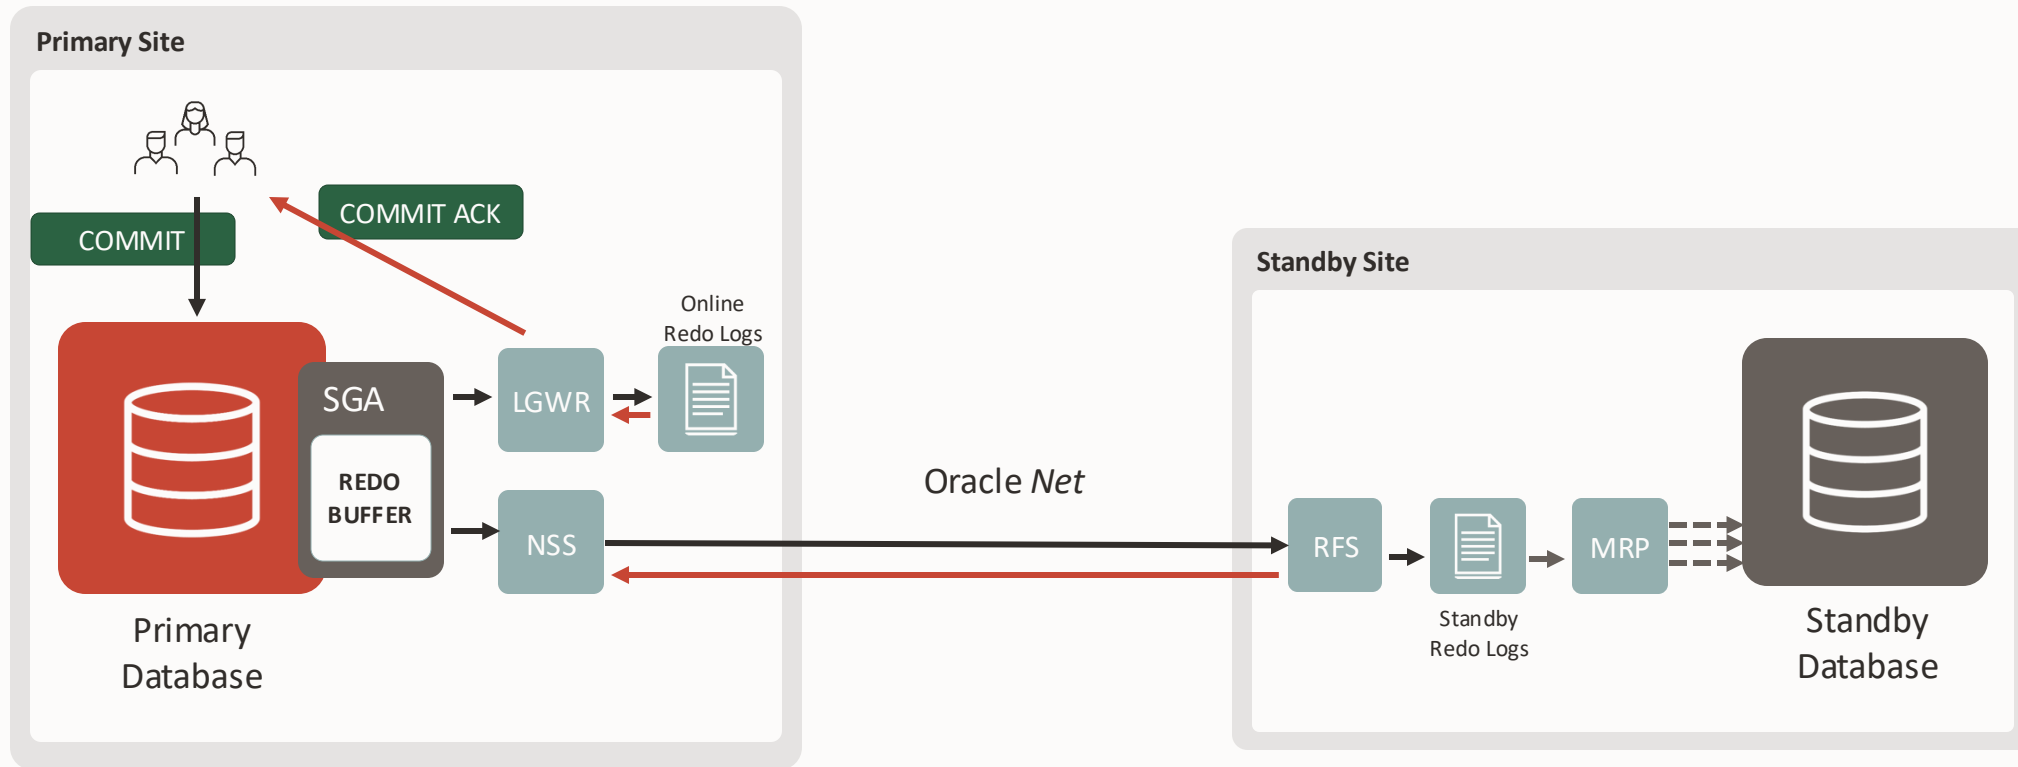


Max Availability without Fast-Start Failover

Data Guard **SYNC** redo transport

If the standby is slow:

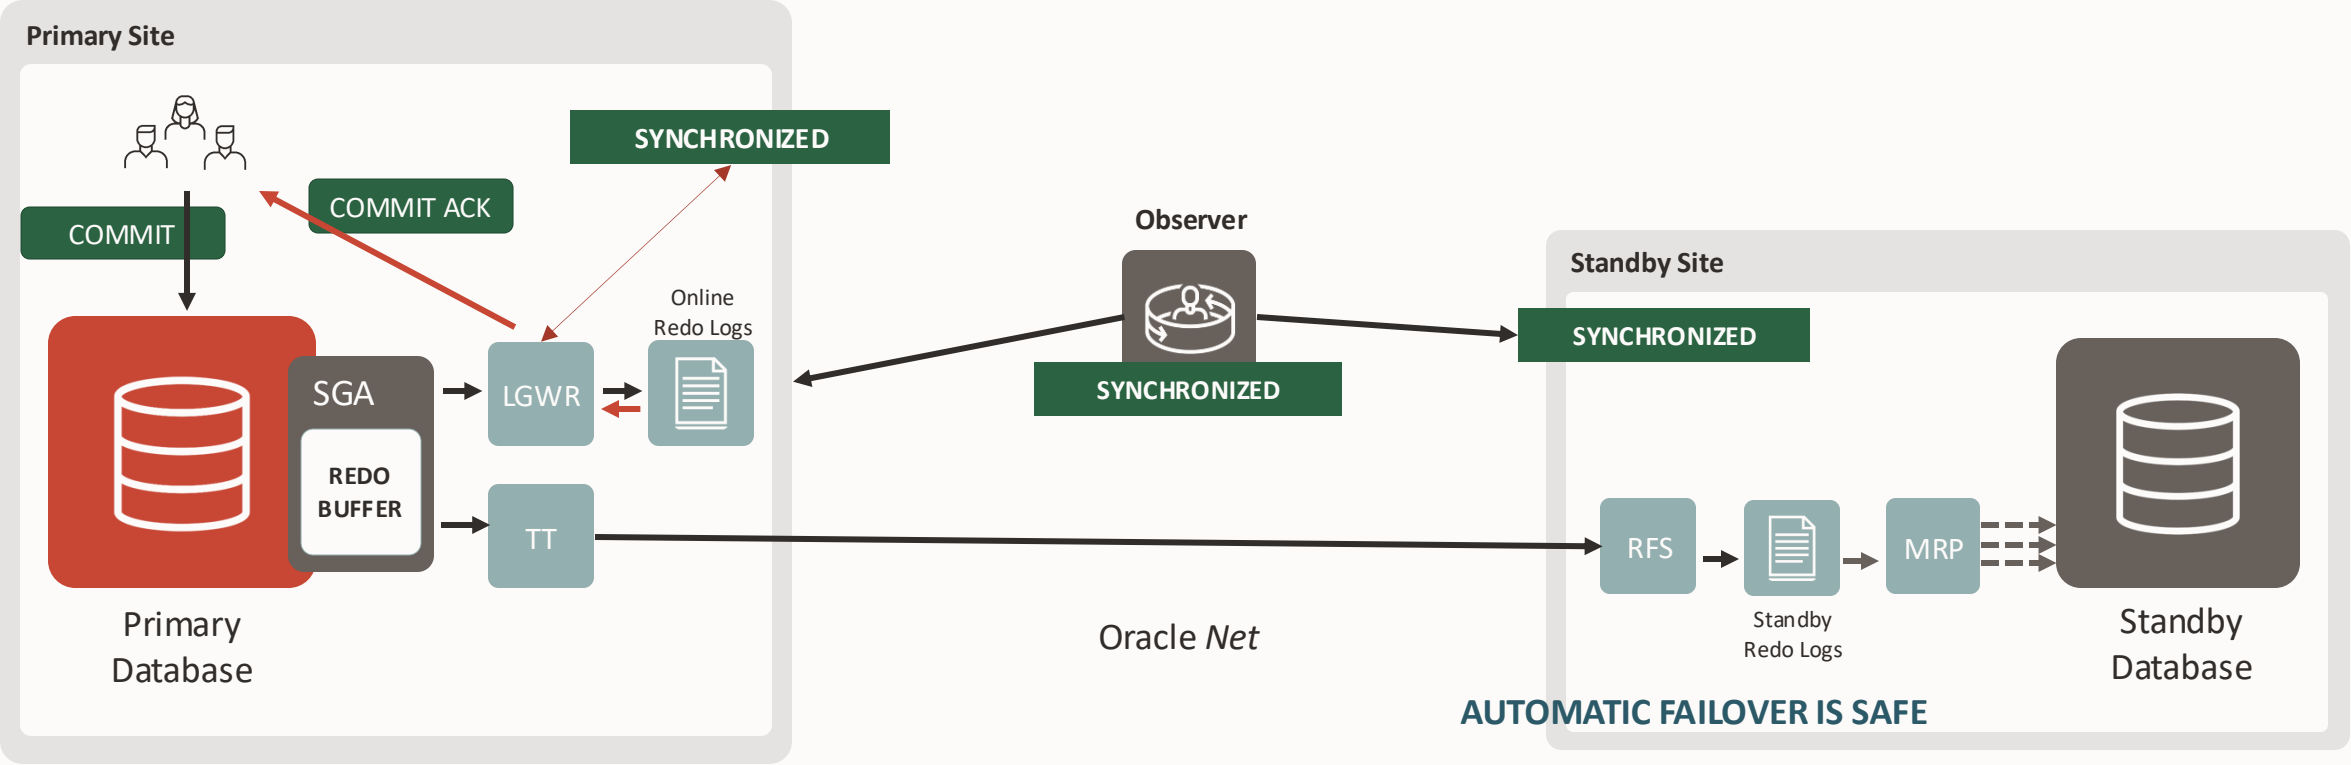
- The commits will take longer, decreasing the primary database performance
- Latencies (disk and/or network) have a crucial role



Max Availability with Fast-Start Failover

The primary never commits without the observer quorum

In a normal situation, the status is SYNCHRONIZED




Max Availability with Fast-Start Failover

The primary never commits without the observer quorum

If the standby is not reachable:

- The primary waits for the SYNC destination (no commits possible) but keeps the status synchronized

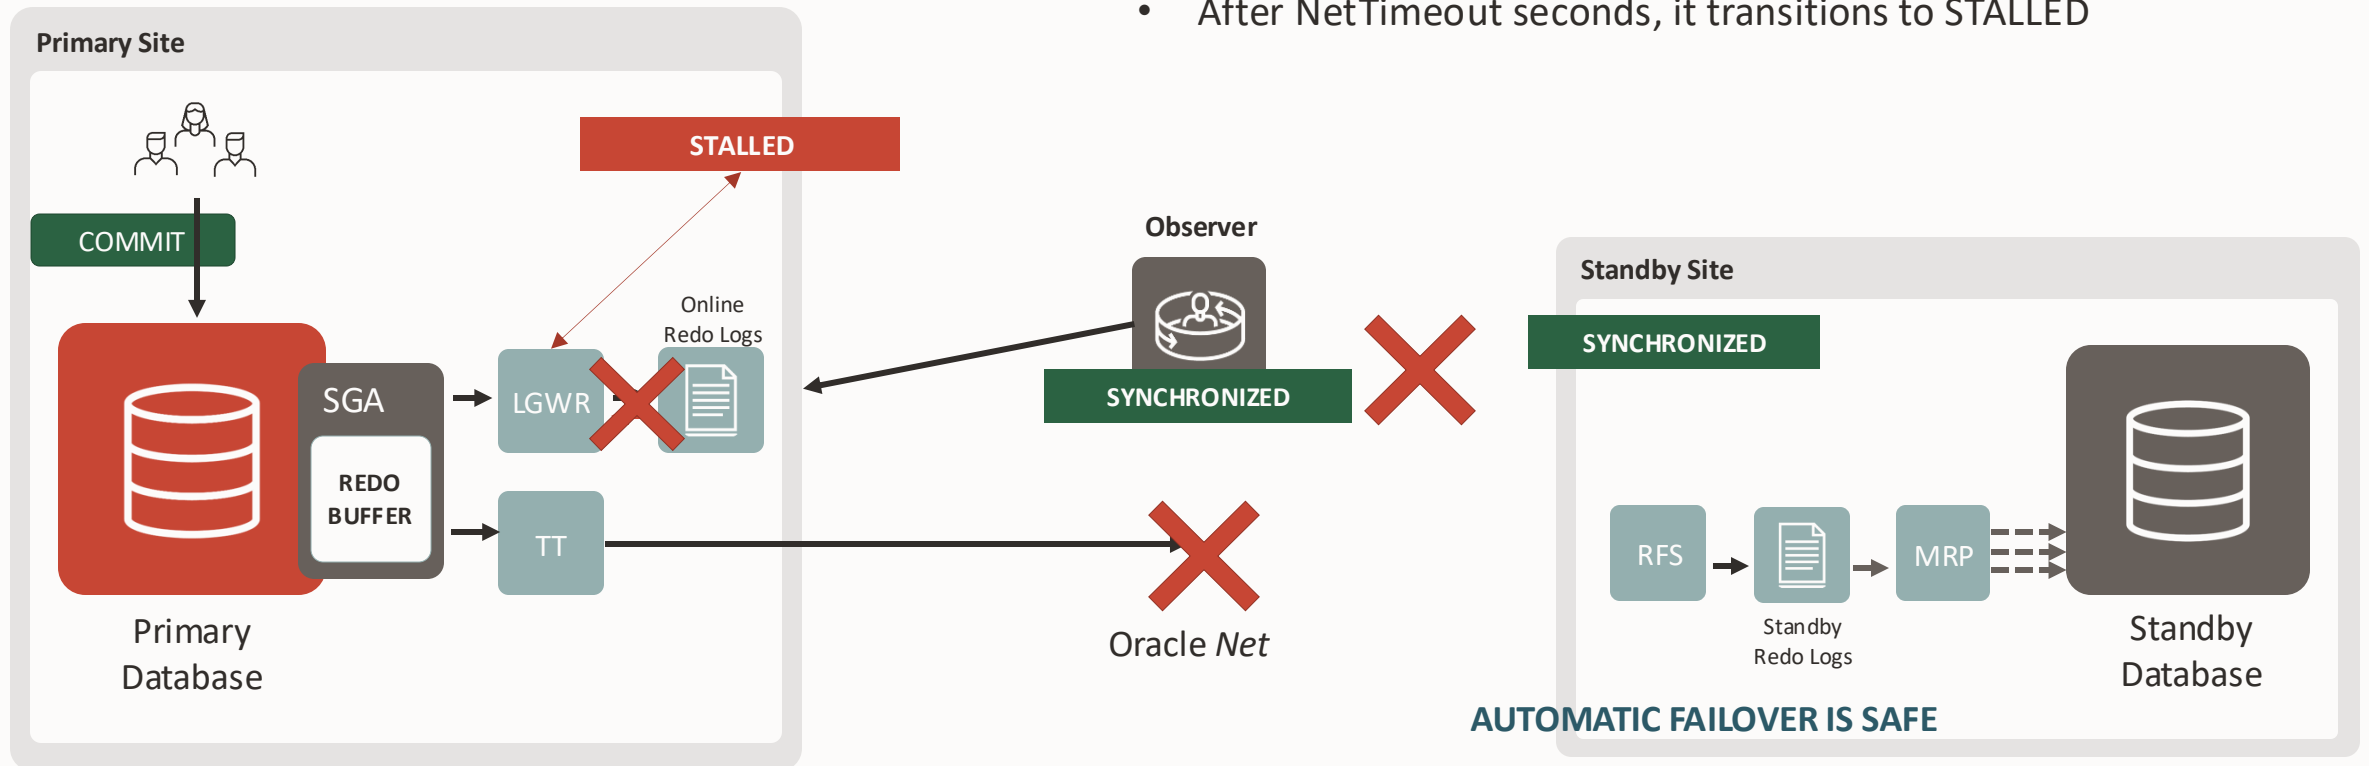


Max Availability with Fast-Start Failover

The primary never commits without the observer quorum

If the standby is not reachable:

- The primary waits for the SYNC destination (no commits possible) but keeps the status synchronized
- After NetTimeout seconds, it transitions to STALLED

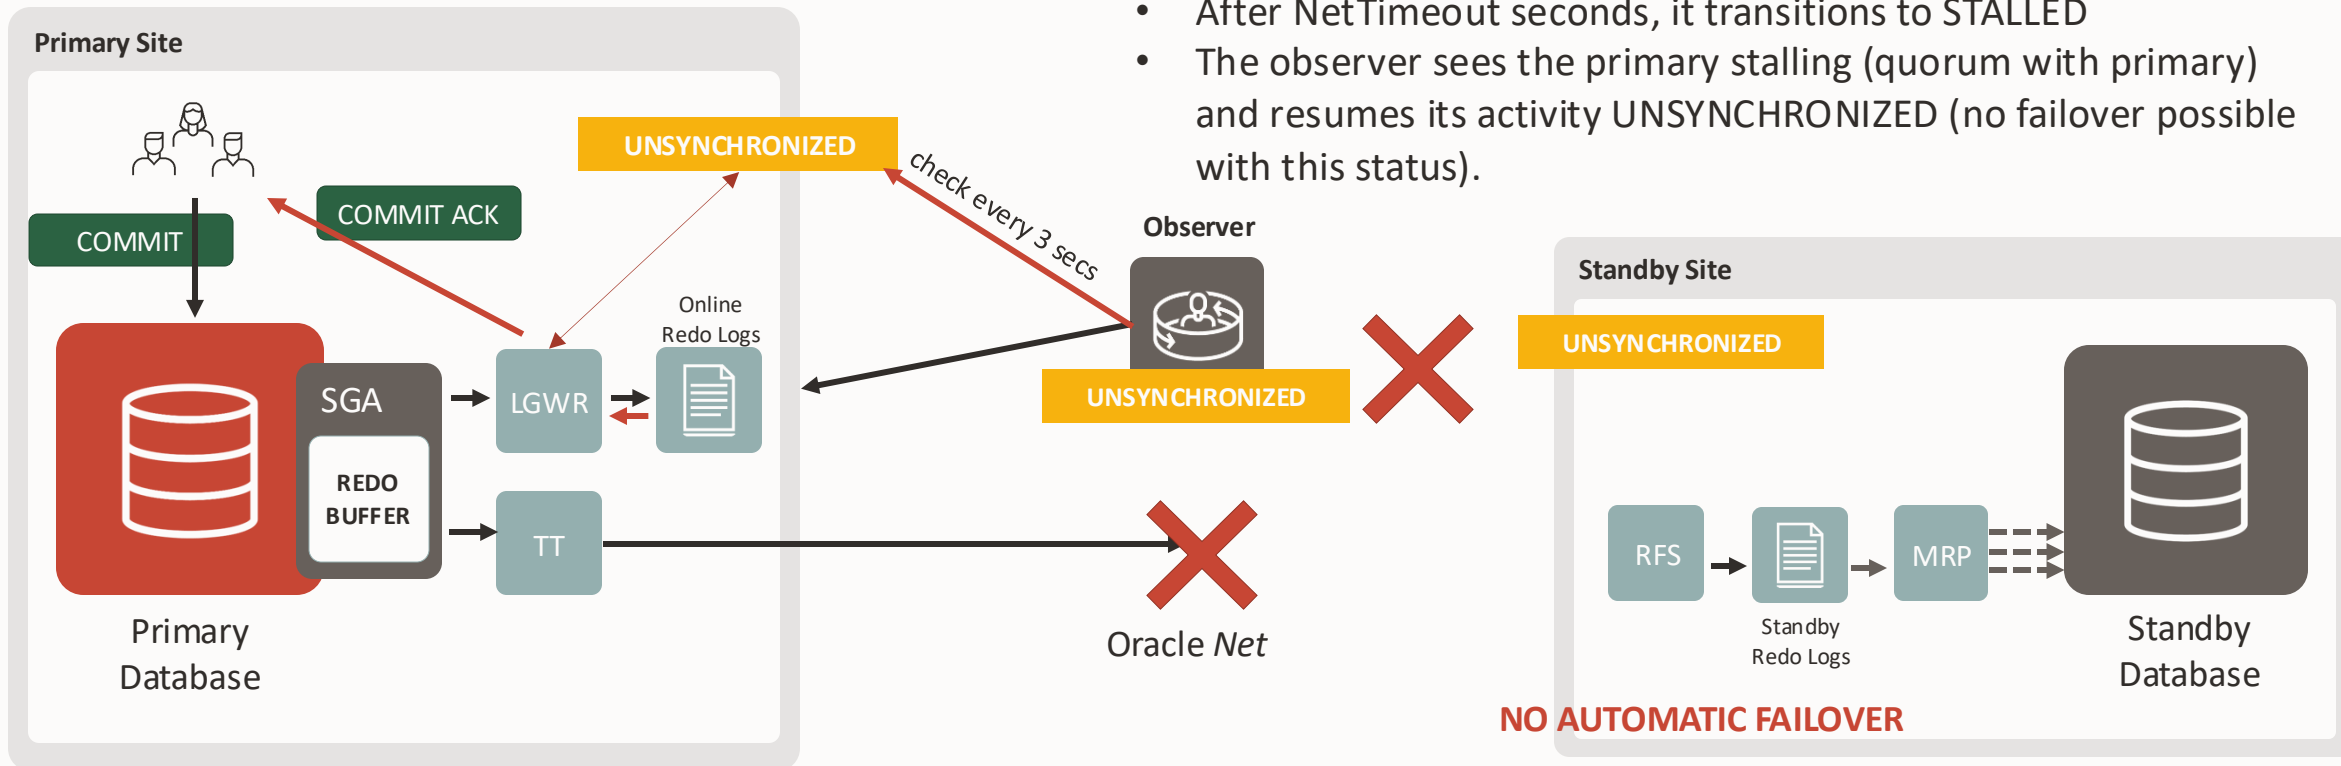


Max Availability with Fast-Start Failover

The primary never commits without the observer quorum

If the standby is not reachable:

- The primary waits for the SYNC destination (no commits possible) but keeps the status synchronized
- After NetTimeout seconds, it transitions to STALLED
- The observer sees the primary stalling (quorum with primary) and resumes its activity UNSYNCHRONIZED (no failover possible with this status).

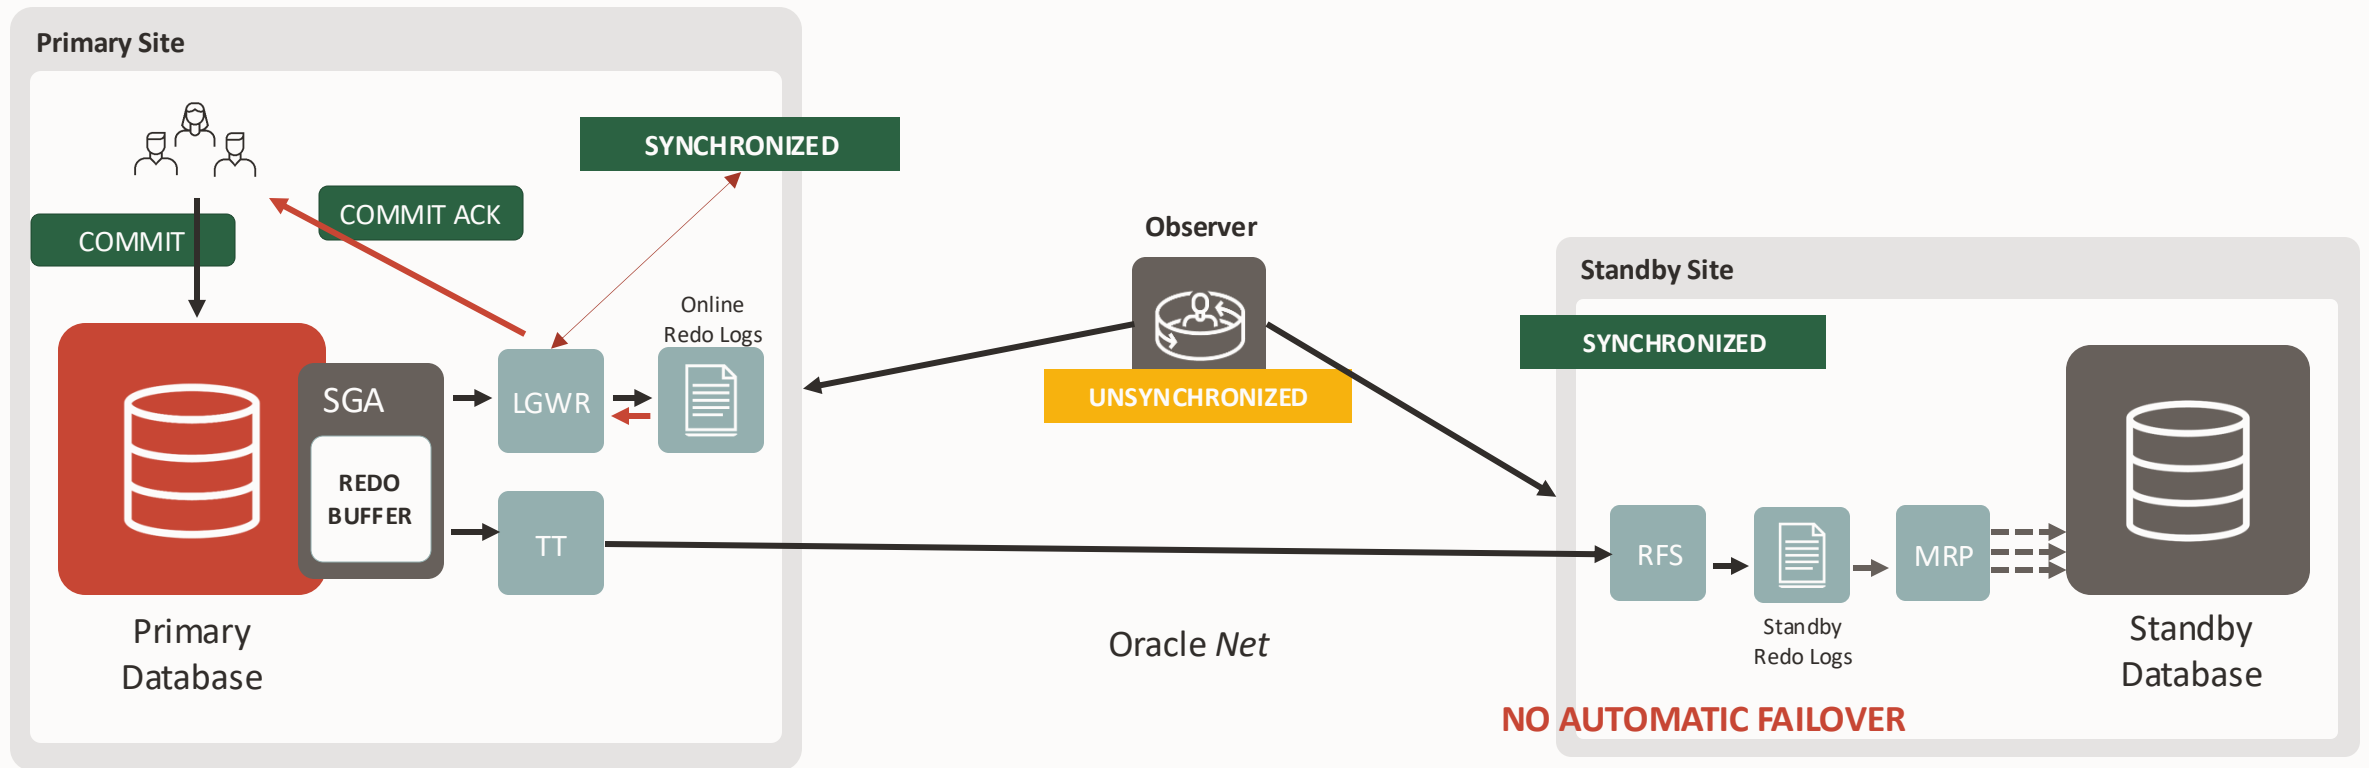


Max Availability with Fast-Start Failover

The primary never commits without the observer quorum

When the standby is in SYNC again

- The primary autonomously change the status to SYNCHRONIZED

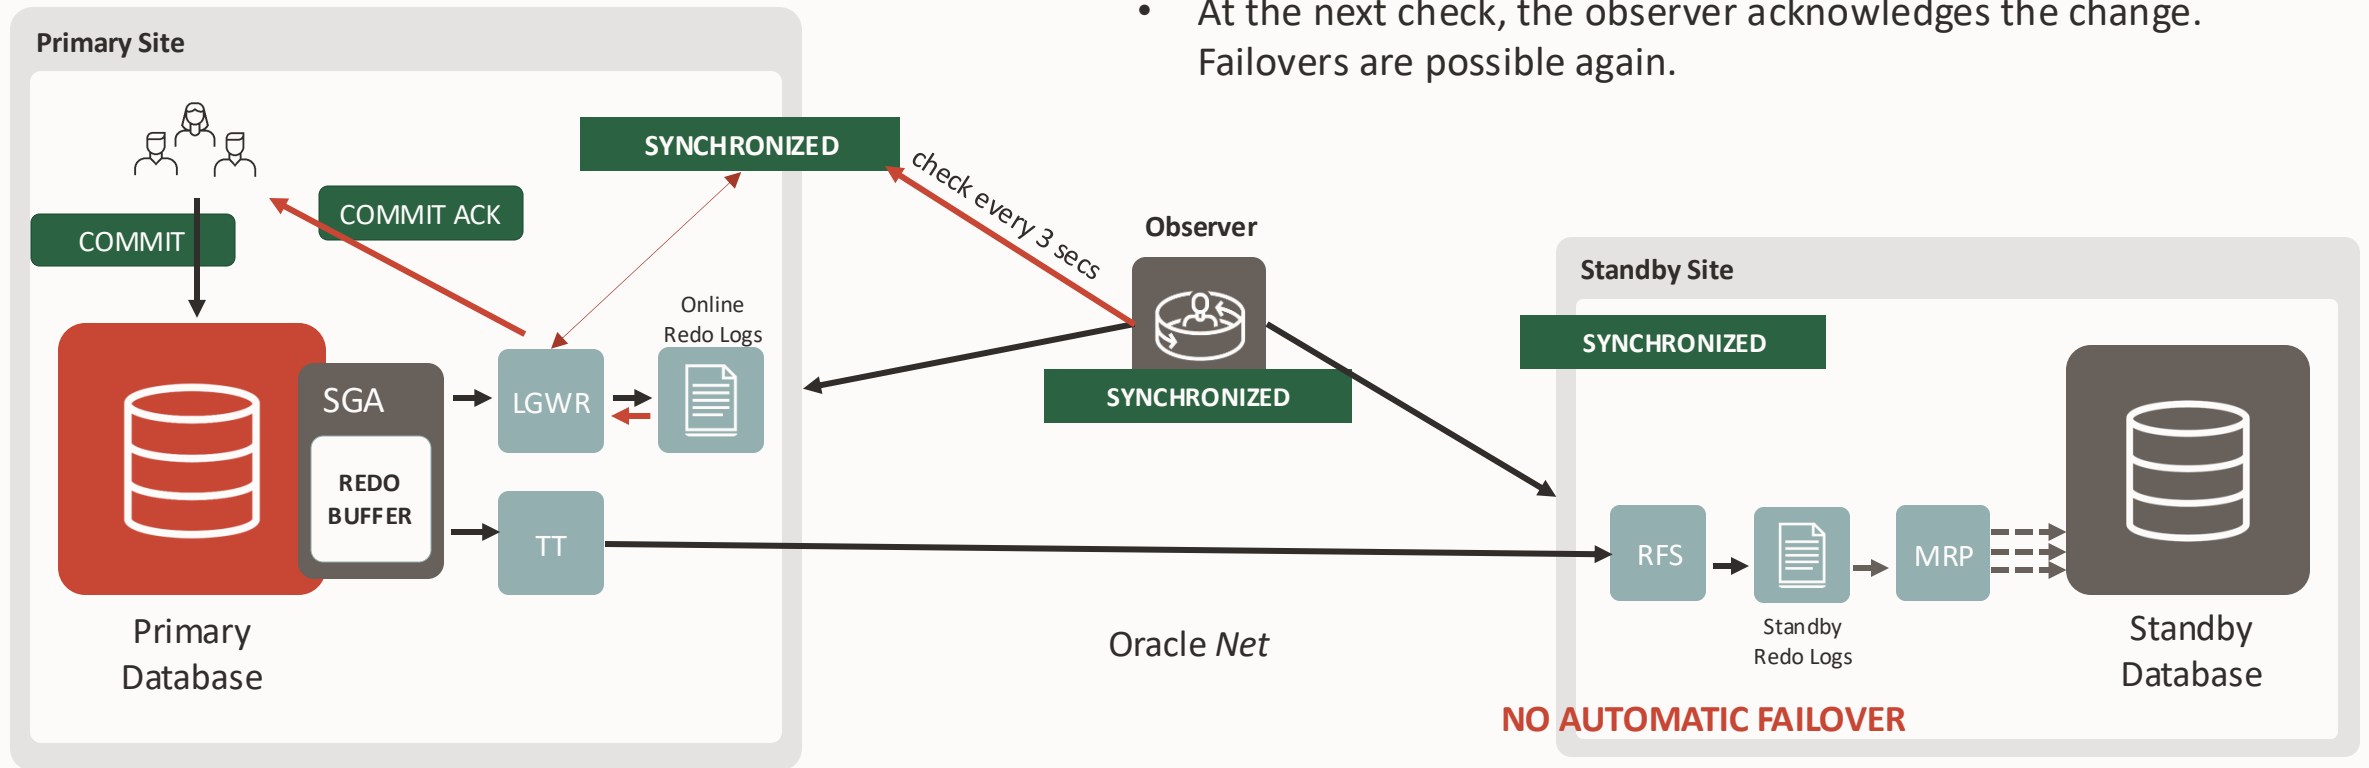


Max Availability with Fast-Start Failover

The primary never commits without the observer quorum

When the standby is in SYNC again

- The primary autonomously change the status to SYNCHRONIZED
- At the next check, the observer acknowledges the change. Failovers are possible again.



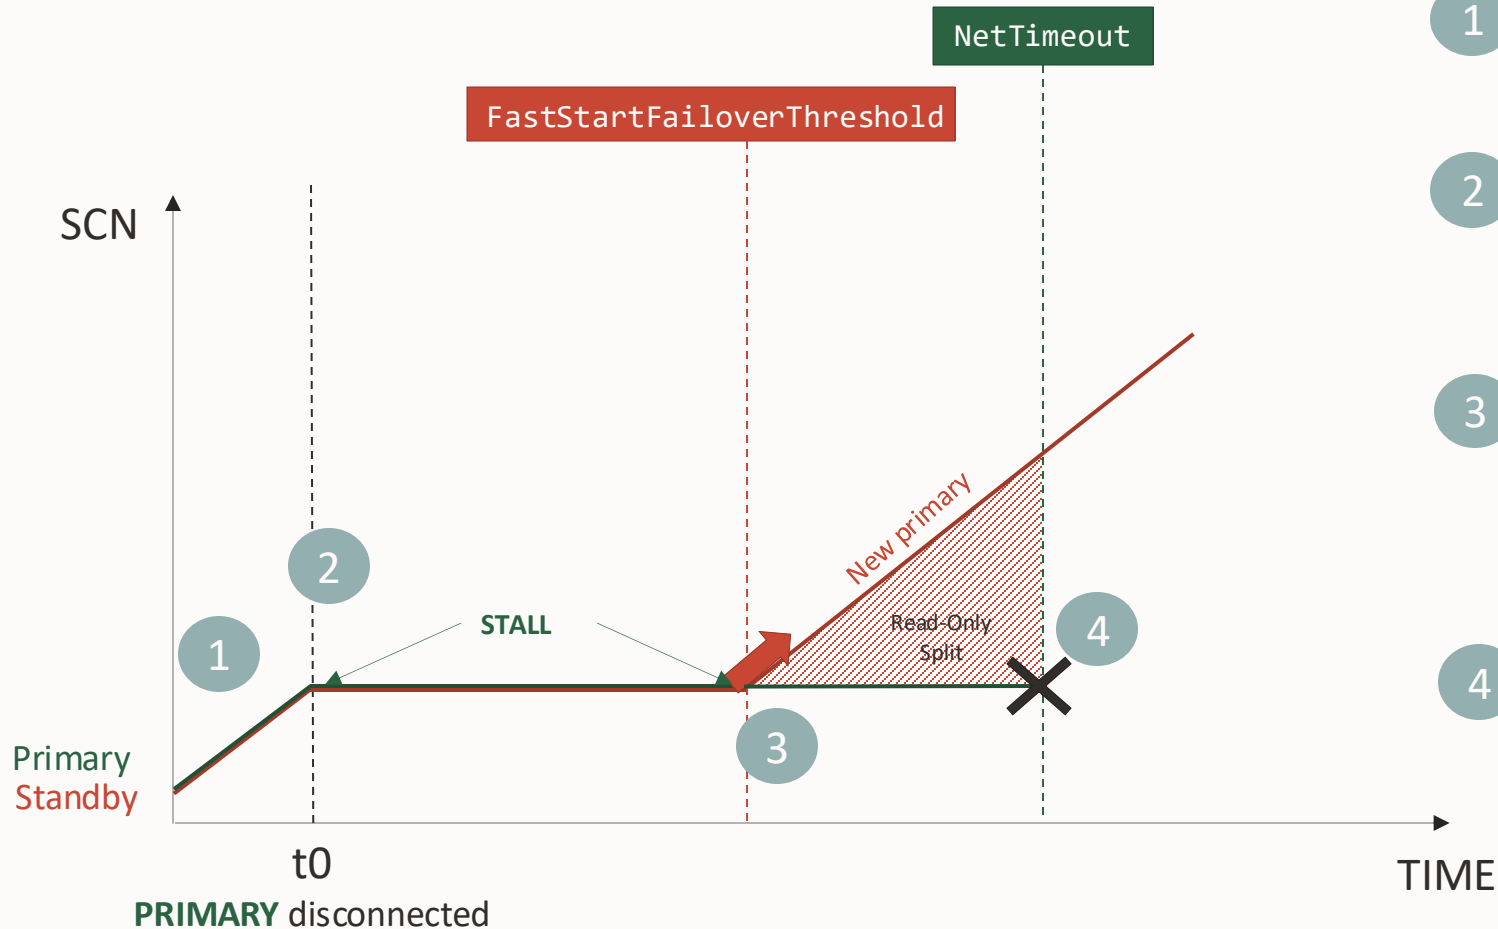
Automatic Failover with MaxAvailability

Failover only if Zero Data Loss is guaranteed



Automatic Failover with MaxAvailability


Failover only if Zero Data Loss is guaranteed



- 1 SYNC Transport. The Standby is synched.
Status: SYNCHRONIZED
- 2 At $t_0 + \text{ping time}$, the observer cannot contact the primary and starts a timer for `FastStartFailoverThreshold` seconds. The Primary stalls and keeps retrying for `NetTimeout` Seconds.
Status: STALLED
- 3 The observer timer reaches `FastStartFailoverThreshold`. Still no connection with the primary: it initiates the Failover. If `NetTimeout` is higher than the threshold, the Primary is reachable but not committing (read-only split).
Primary status: STALLED
Standby status: REINSTATE REQUIRED
- 4 The Primary keeps stalling or shuts down after `NetTimeout` because it cannot get the permission from the observer to abandon the destination. A lower `NetTimeout` reduces the potential of read-only split.
Status: REINSTATE REQUIRED

Automatic Failover with MaxAvailability

Optionally choose a limit to have an automatic failover with potential data loss



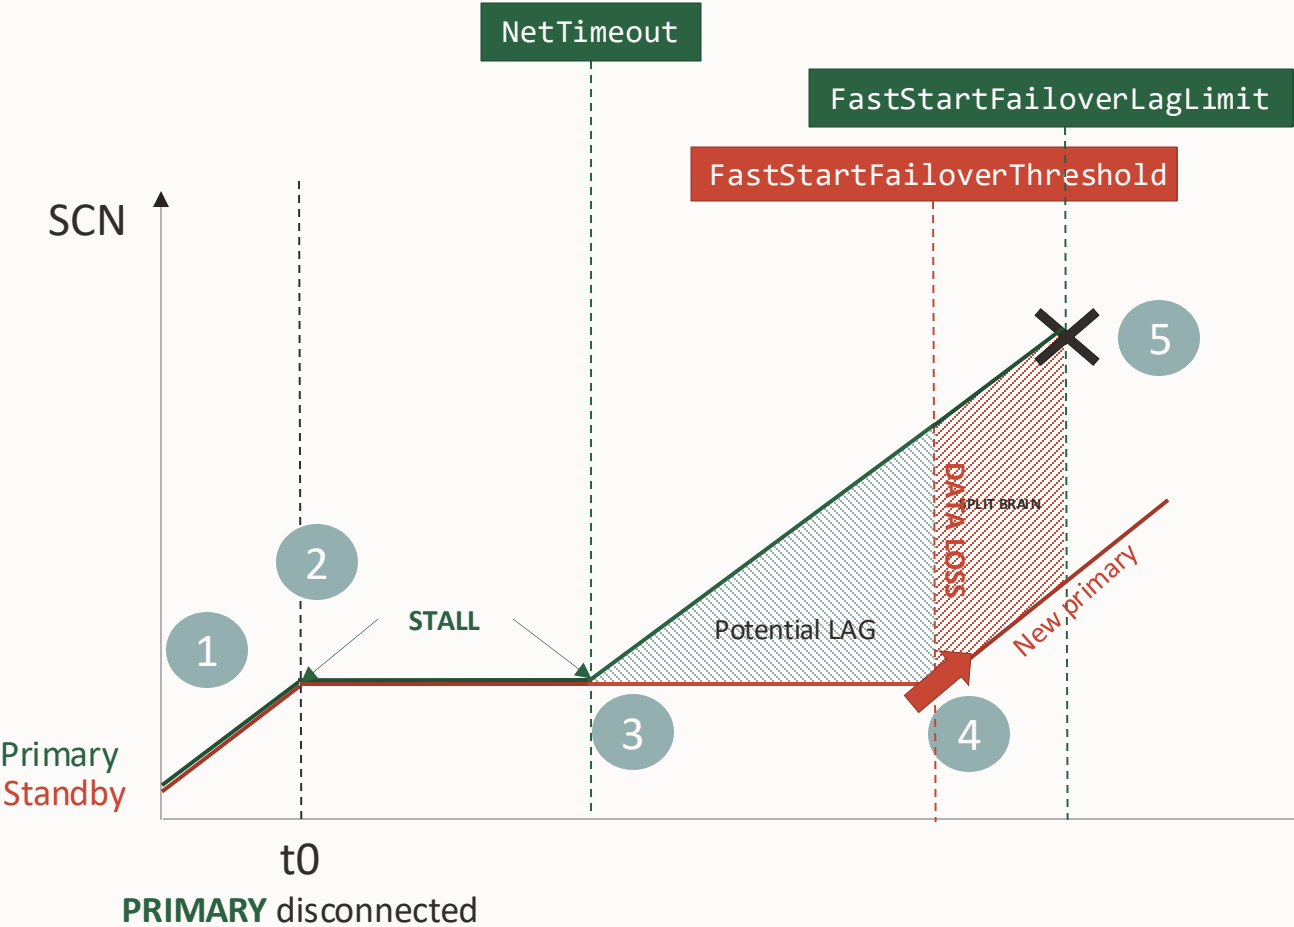
$$\text{FastStartFailoverLagLimit} \leq \text{FastStartFailoverThreshold}$$

- 1 SYNC Transport. The Standby is synched.
- 2 At $t_0 + \text{ping time}$, the primary cannot contact the standby. The primary stalls and keeps retrying for `NetTimeout` seconds. The observer starts the timer.
- 3 At `NetTimeout` seconds, the primary stops shipping the redo without asking permission to the observer, because `FastStartFailoverLagLimit` is set. The observer does not acknowledge the unsynched status.
- 4 The primary reaches `FastStartFailoverLagLimit`. It cannot obtain permission from the observer to continue. It stalls or shuts down depending on `FastStartFailoverPmyShutdown`
- 5 The observer timer reaches `FastStartFailoverThreshold`. Still no connection with the primary: it initiates the Failover.



Automatic Failover with MaxAvailability



Set the FastStartFailoverLagLimit wisely to avoid split-brain conditions



`FastStartFailoverThreshold` < `FastStartFailoverLagLimit`

- 1 SYNC Transport. The Standby is synched.
- 2 At $t_0 + \text{ping time}$, the primary cannot contact the standby. The primary stalls and keeps retrying for `NetTimeout` seconds. The observer starts the timer.
- 3 At `NetTimeout` seconds, the primary switches to ASYNC redo transport without requiring permission to the observer, because `FastStartFailoverLagLimit` is set. The observer does not acknowledge the unsynched status.
- 4 The observer timer reaches `FastStartFailoverThreshold`. Still no connection with the primary: it initiates the Failover
- 5 The Primary keeps committing until it reaches `FastStartFailoverLagLimit`, then stalls or shuts down. A Split-Brain condition may occur depending on timings and parameter values.

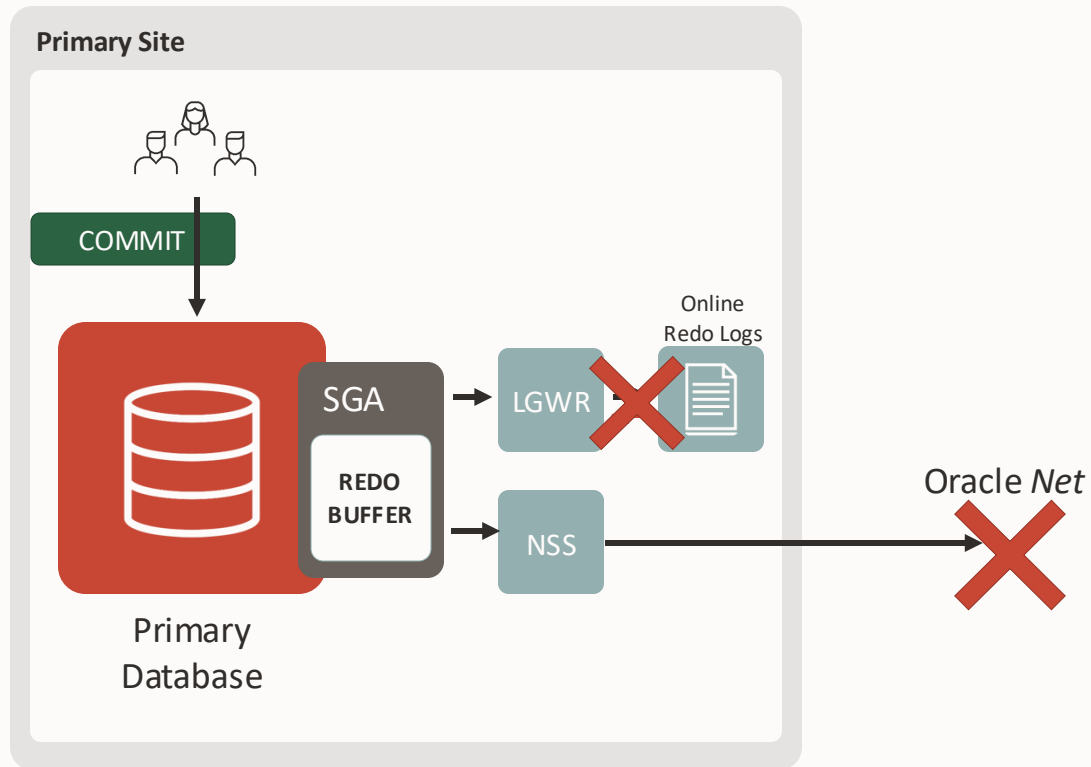




Fast-Start Failover: Maximum Protection

Max Protection without Fast-Start Failover

Data Guard **SYNC** redo transport

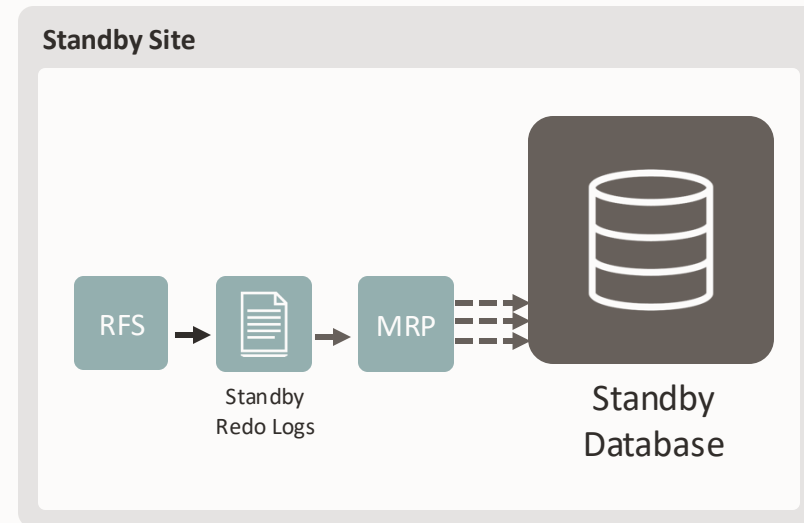


If the standby is not reachable:

- The primary stalls waiting for the SYNC destination (no commits possible)
- Even after NetTimeout, the commits cannot happen

If the primary is not reachable:

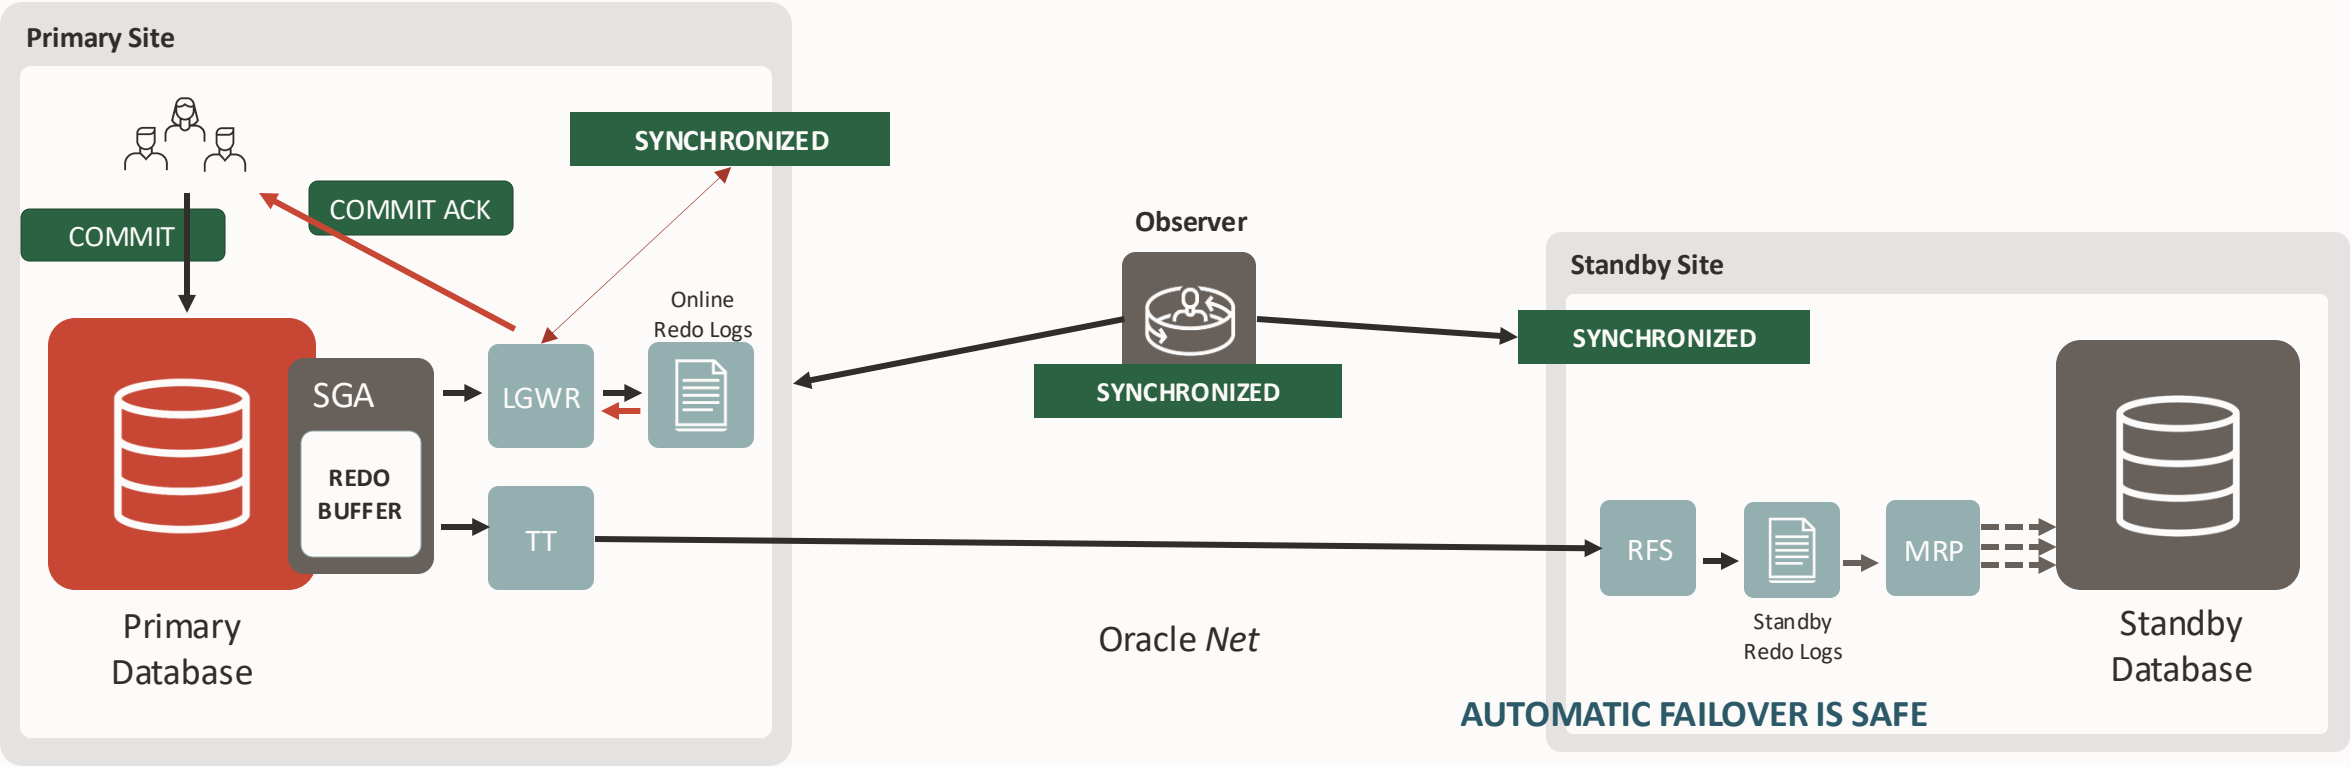
- Failovers, even manual, are safe if no other standbys are protecting the primary



Max Protection with Fast-Start Failover

The primary never commits without the observer quorum

In a normal situation, the status is **SYNCHRONIZED**

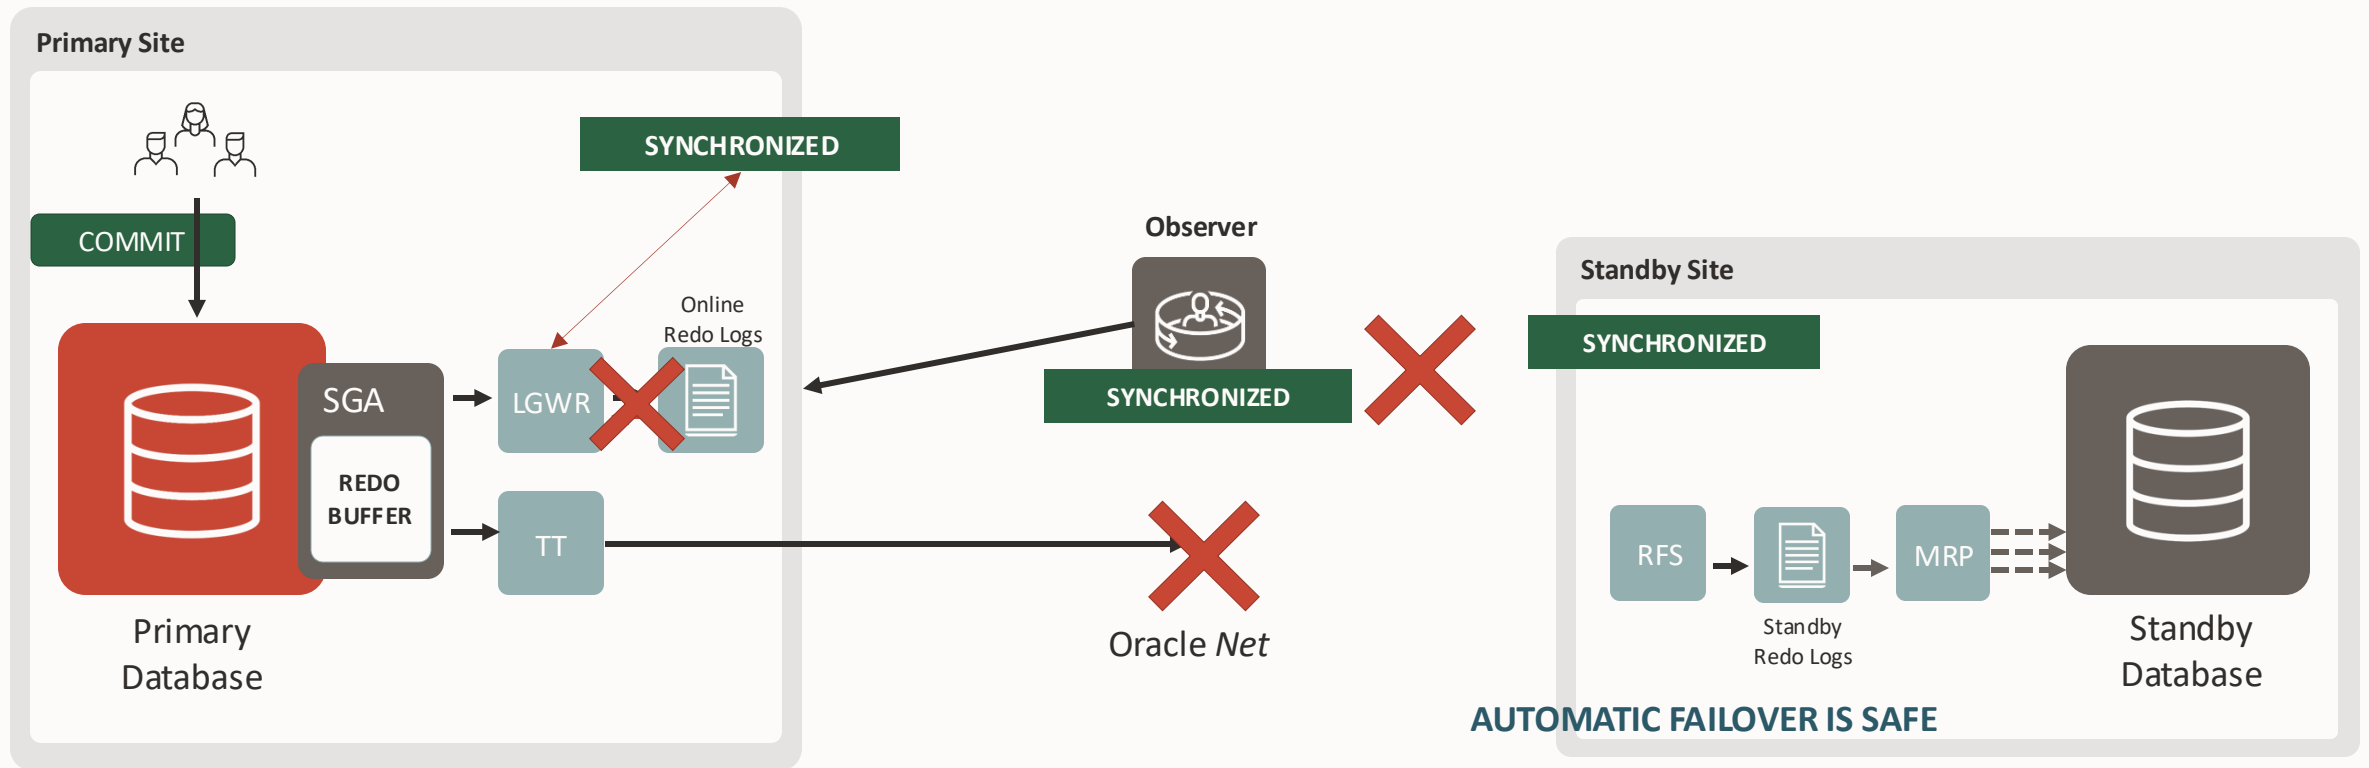


Max Protection with Fast-Start Failover

The primary never commits without the observer quorum

If the standby is not reachable:

- The primary waits for the SYNC destination (no commits possible) but keeps the status synchronized

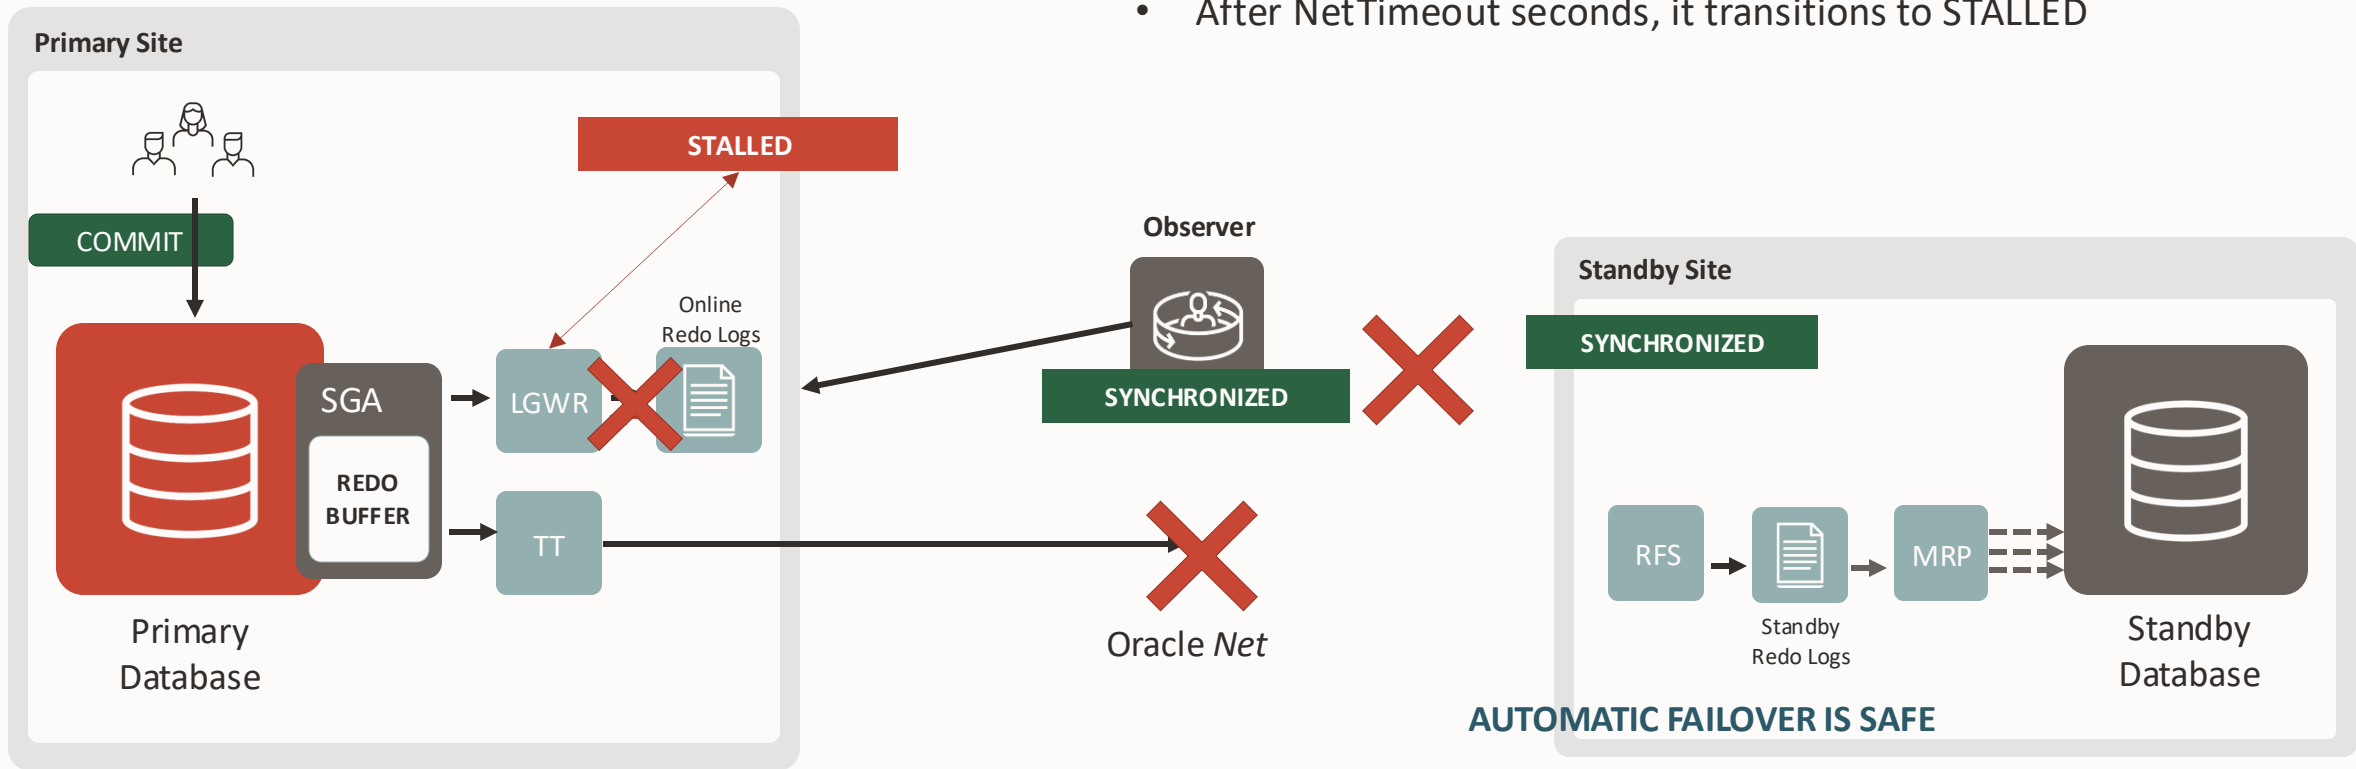


Max Protection with Fast-Start Failover

The primary never commits without the observer quorum

If the standby is not reachable:

- The primary waits for the SYNC destination (no commits possible) but keeps the status synchronized
- After NetTimeout seconds, it transitions to STALLED



- The primary waits for the SYNC destination (no commits possible) but keeps the status synchronized
- After NetTimeout seconds, it transitions to STALLED
- The observer sees the primary stalling (quorum with primary) and acknowledges the new status.

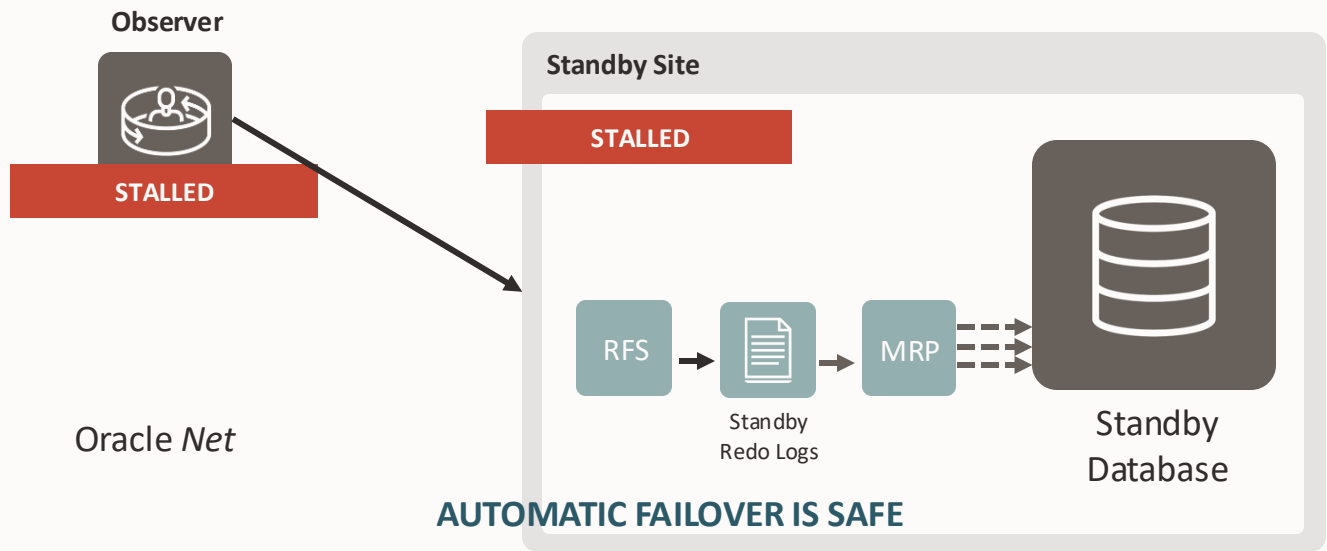
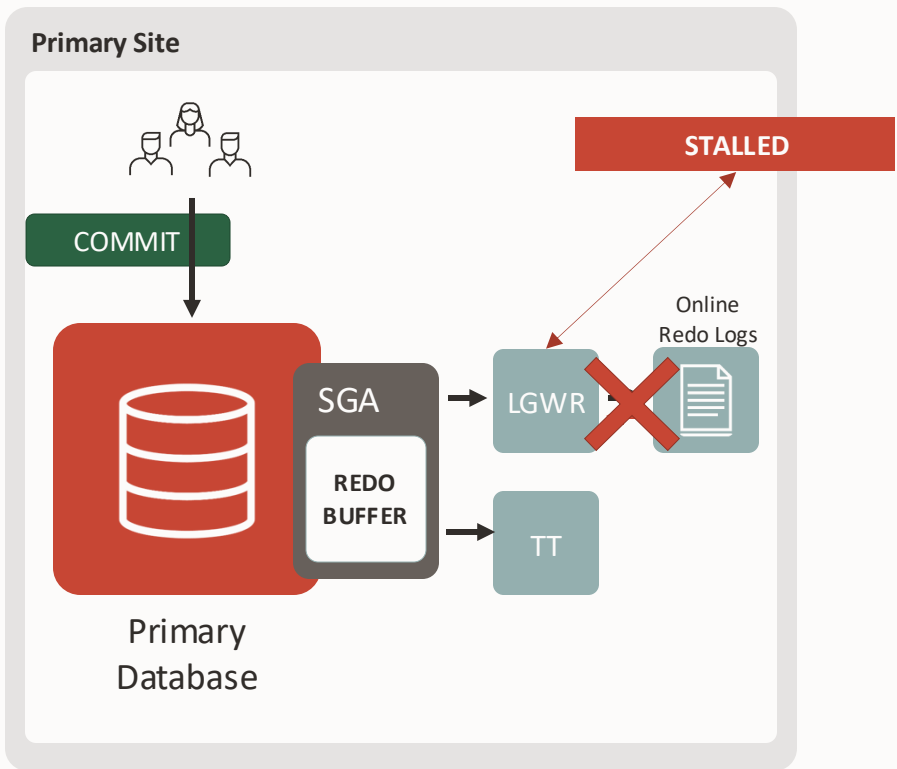


Max Protection with Fast-Start Failover

The primary never commits without the observer quorum


If the primary is not reachable:

- The primary never committed without a SYNC standby
- The automatic failover is always safe, but the new primary will require another standby to keep the current protection level



Automatic Failover with MaxProtection

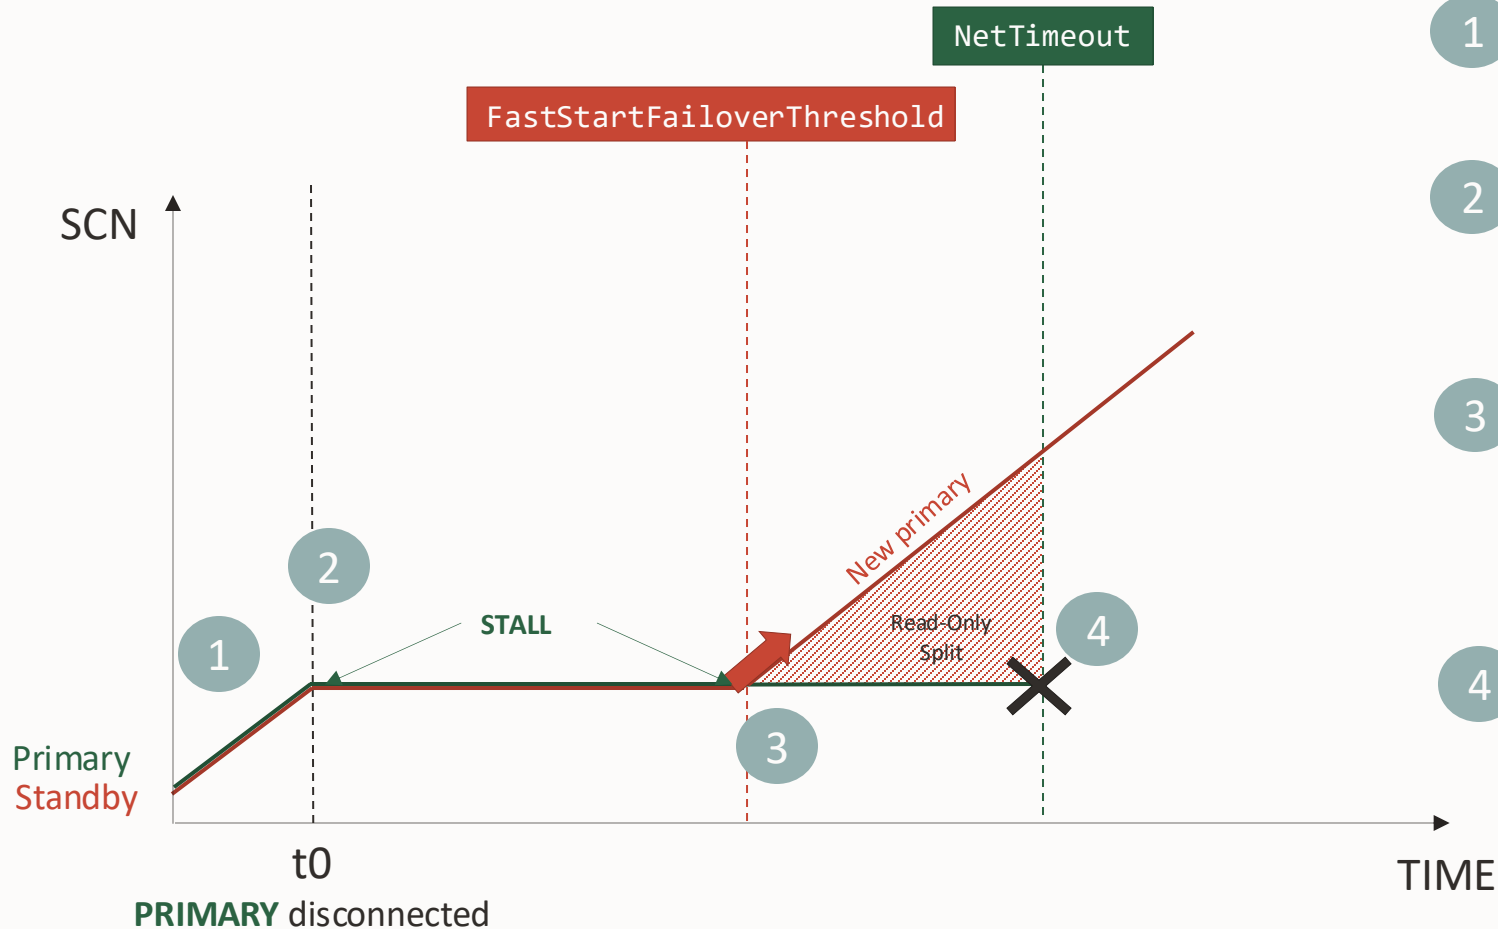
Zero Data Loss in any condition





- 1 SYNC Transport. The Standby is synched.
Status: SYNCHRONIZED
- 2 At $t_0 + \text{ping time}$, the primary cannot contact the standby and keeps retrying for NetTimeout Seconds.
Status: STALLED
- 3 After NetTimeout, the primary still has connectivity with the observer. Knowing that a failover did not occur, it keeps trying forever.
Status: STALLED

Automatic Failover with MaxProtection

Zero Data Loss in any condition



- 1 SYNC Transport. The Standby is synced.
Status: SYNCHRONIZED
- 2 At $t_0 + \text{ping time}$, the observer cannot contact the primary and starts a timer for `FastStartFailoverThreshold` seconds. The Primary stalls and keeps retrying for `NetTimeout` Seconds.
Status: STALLED
- 3 The observer timer reaches `FastStartFailoverThreshold`. Still no connection with the primary: it initiates the Failover. If `NetTimeout` is higher than the threshold, the Primary is reachable but not committing (read-only split).
Former Primary Status: STALLED
New Primary Status: REINSTATE_REQUIRED
- 4 The primary keeps stalling or shuts down after `NetTimeout` depending on `FastStartFailoverPmyShutdown`. A lower `NetTimeout` reduces the potential of read-only split.
Former Primary Status: STALLED or REINSTATE_REQUIRED
New Primary Status: REINSTATE_REQUIRED

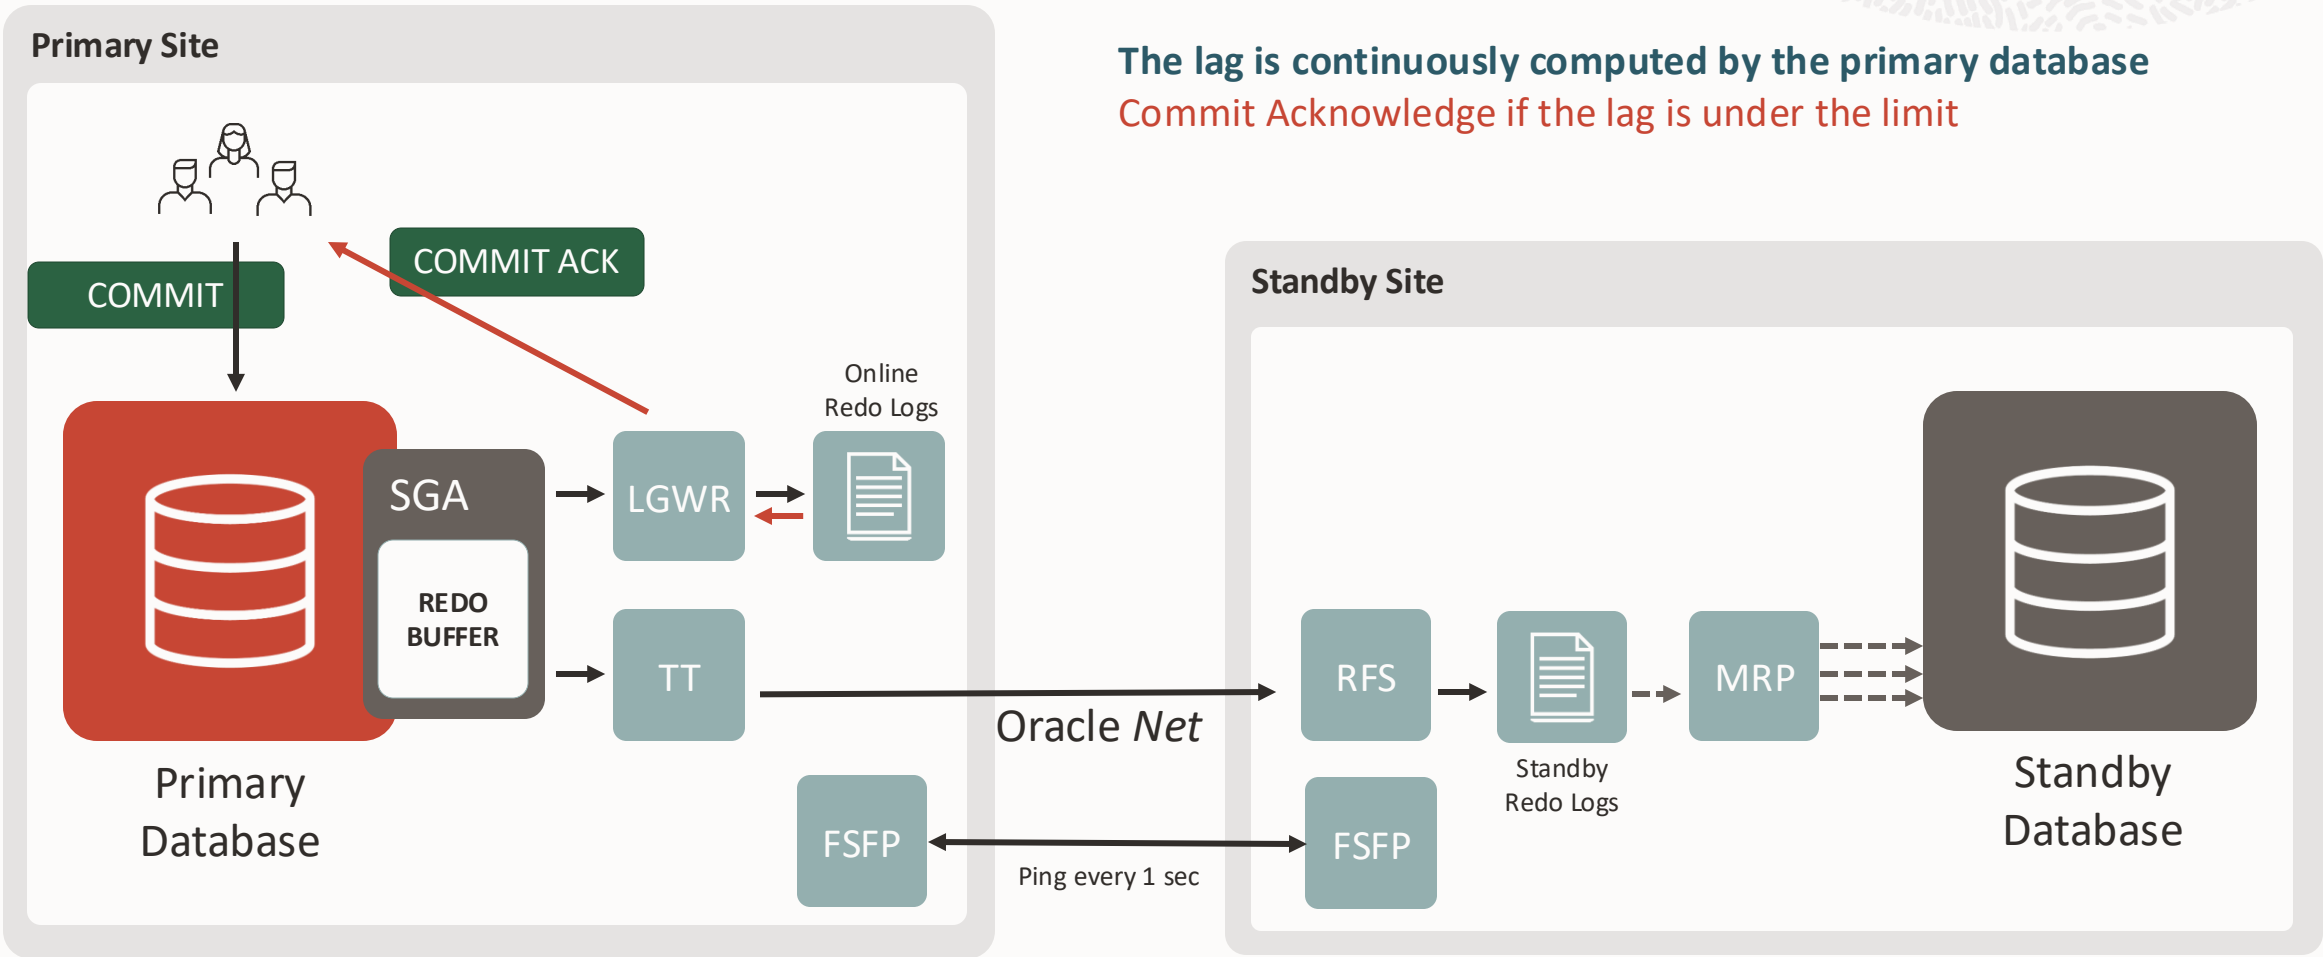


Fast-Start Failover: Oracle Data Guard Observer

How Fast-Start Failover limits data loss


Data Guard **ASYNC** Process Architecture (Possible Data Loss)

The lag is continuously computed by the primary database
Commit Acknowledge if the lag is under the limit




Data Guard Transport for Zero Data Loss

Data Guard SYNC Process Architecture



Oracle Data Guard Observer Placement

Observer at the primary site

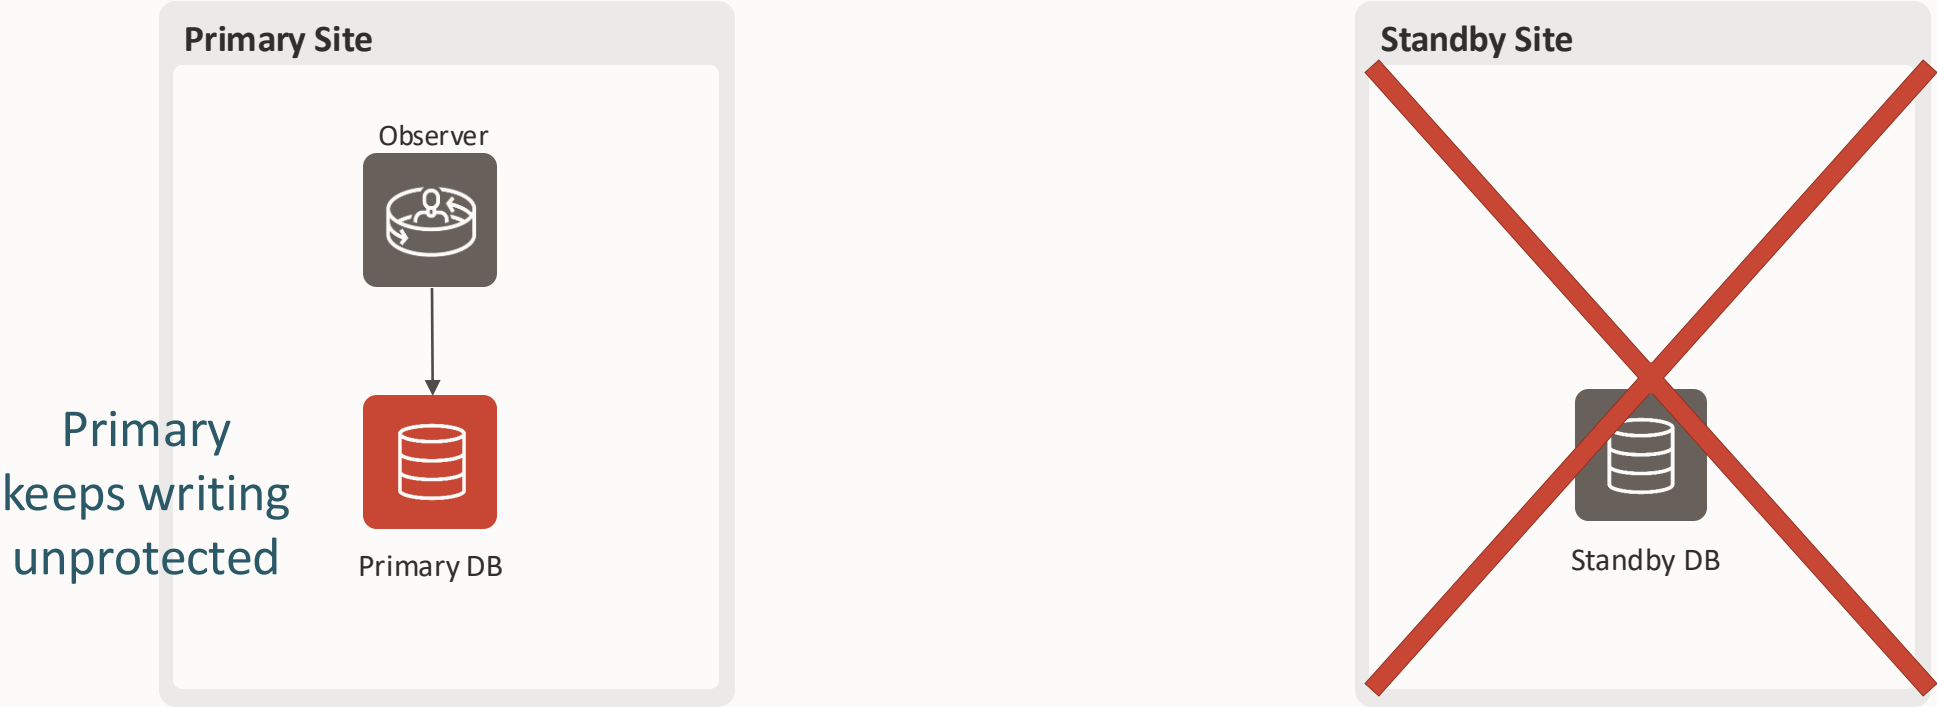


Failure:	Primary DB	Standby DB	Network	Primary Site	Standby Site	Observer
	✓	✓	✓	✗	✓	✓



Oracle Data Guard Observer Placement

Observer at the primary site




Failure:	Primary DB	Standby DB	Network	Primary Site	Standby Site	Observer
	✓	✓	✓	✗	✓	✓



Oracle Data Guard Observer Placement

Observer at the primary site

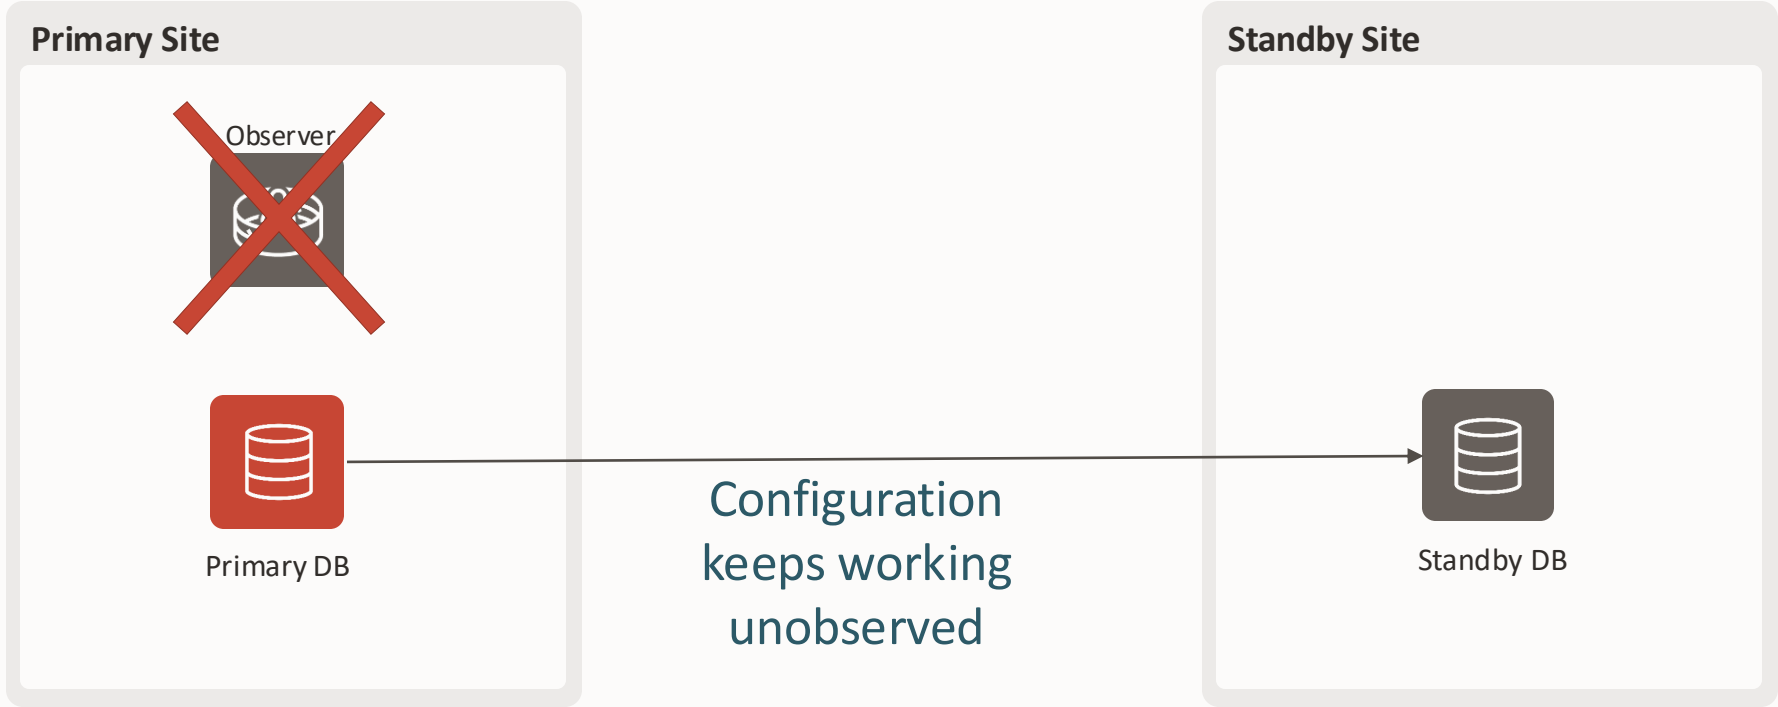


Failure:	Primary DB	Standby DB	Network	Primary Site	Standby Site	Observer
	✓	✓	✓	✗	✓	✓



Oracle Data Guard Observer Placement

Observer at the primary site

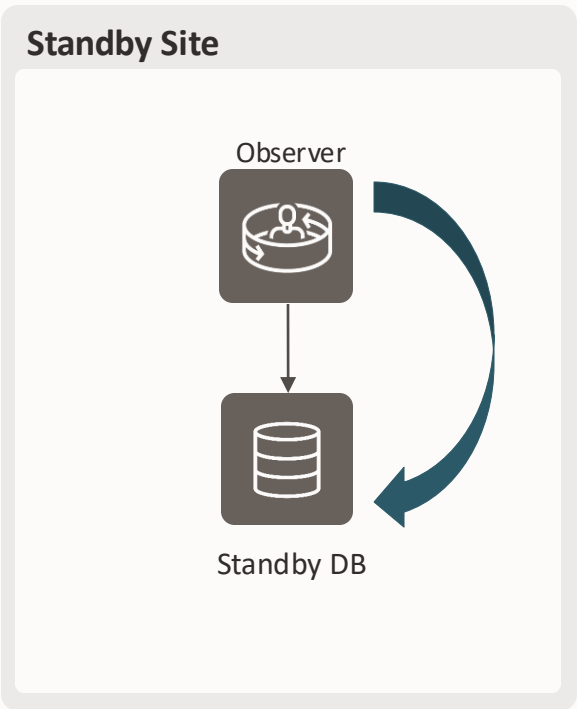



Failure:	Primary DB	Standby DB	Network	Primary Site	Standby Site	Observer
	✓	✓	✓	✗	✓	✓



Oracle Data Guard Observer Placement

Observer at the standby site




Automatic failover!

Failure:	Primary DB	Standby DB	Network	Primary Site	Standby Site	Observer
	✓	✓	✗	✓	✗	✓



Oracle Data Guard Observer Placement

Observer at the standby site




Failure:	Primary DB	Standby DB	Network	Primary Site	Standby Site	Observer
	✓	✓	✗	✓	✗	✓



Oracle Data Guard Observer Placement

Observer at the standby site

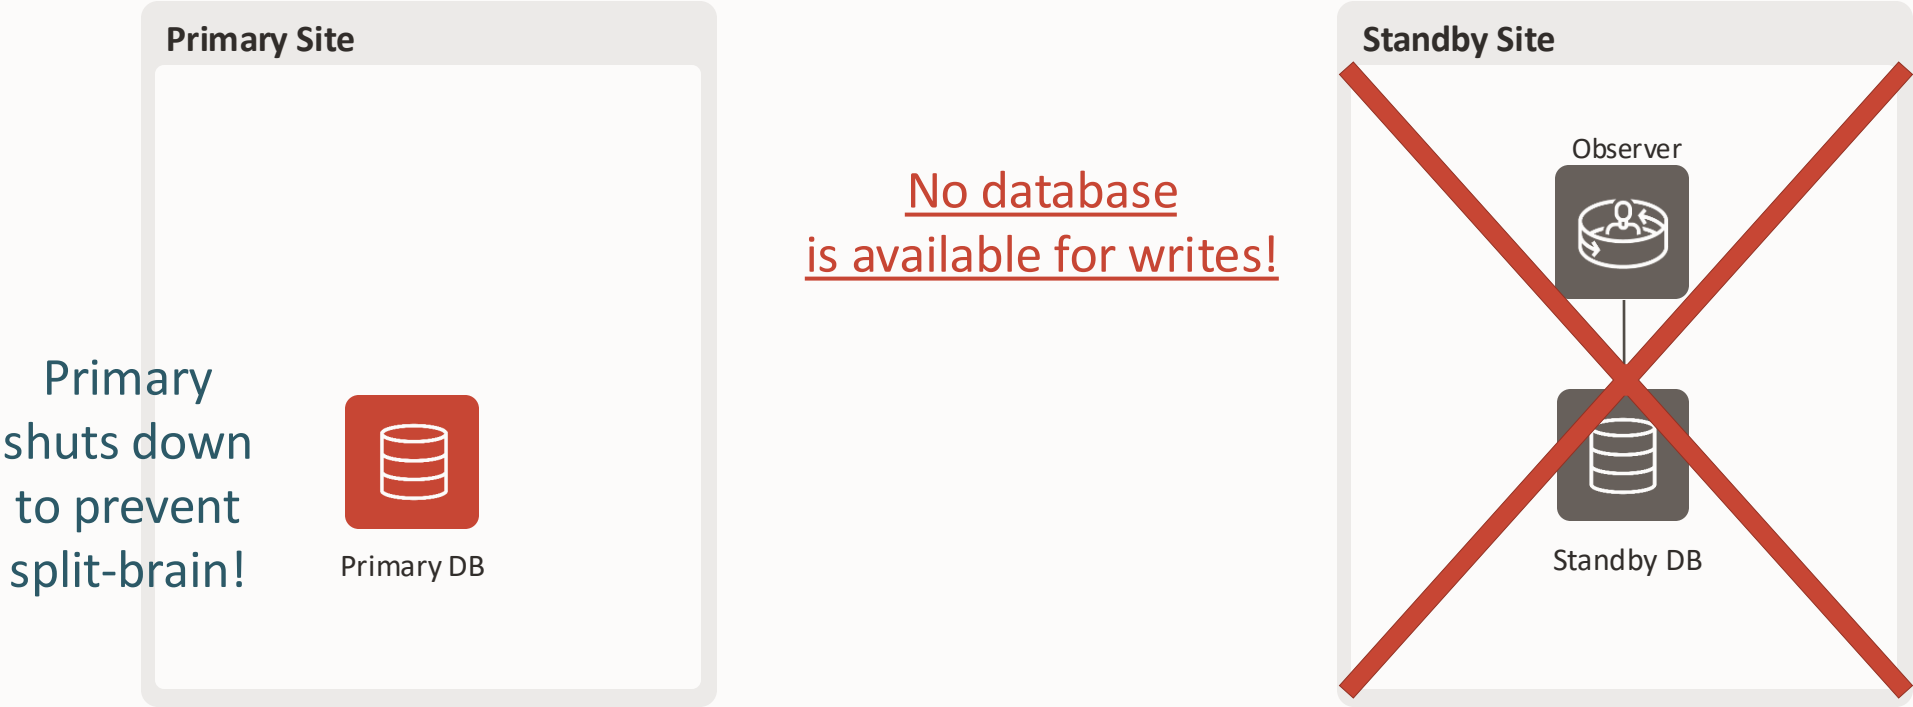


Failure:	Primary DB	Standby DB	Network	Primary Site	Standby Site	Observer
	✓	✓	✗	✓	✗	✓



Oracle Data Guard Observer Placement

Observer at the standby site




Failure:	Primary DB	Standby DB	Network	Primary Site	Standby Site	Observer
	✓	✓	✗	✓	✗	✓



Oracle Data Guard Observer Placement

Observer at the standby site




Failure:	Primary DB	Standby DB	Network	Primary Site	Standby Site	Observer
	✓	✓	✗	✓	✗	✓



Oracle Data Guard Observer Placement

Observer at an external size

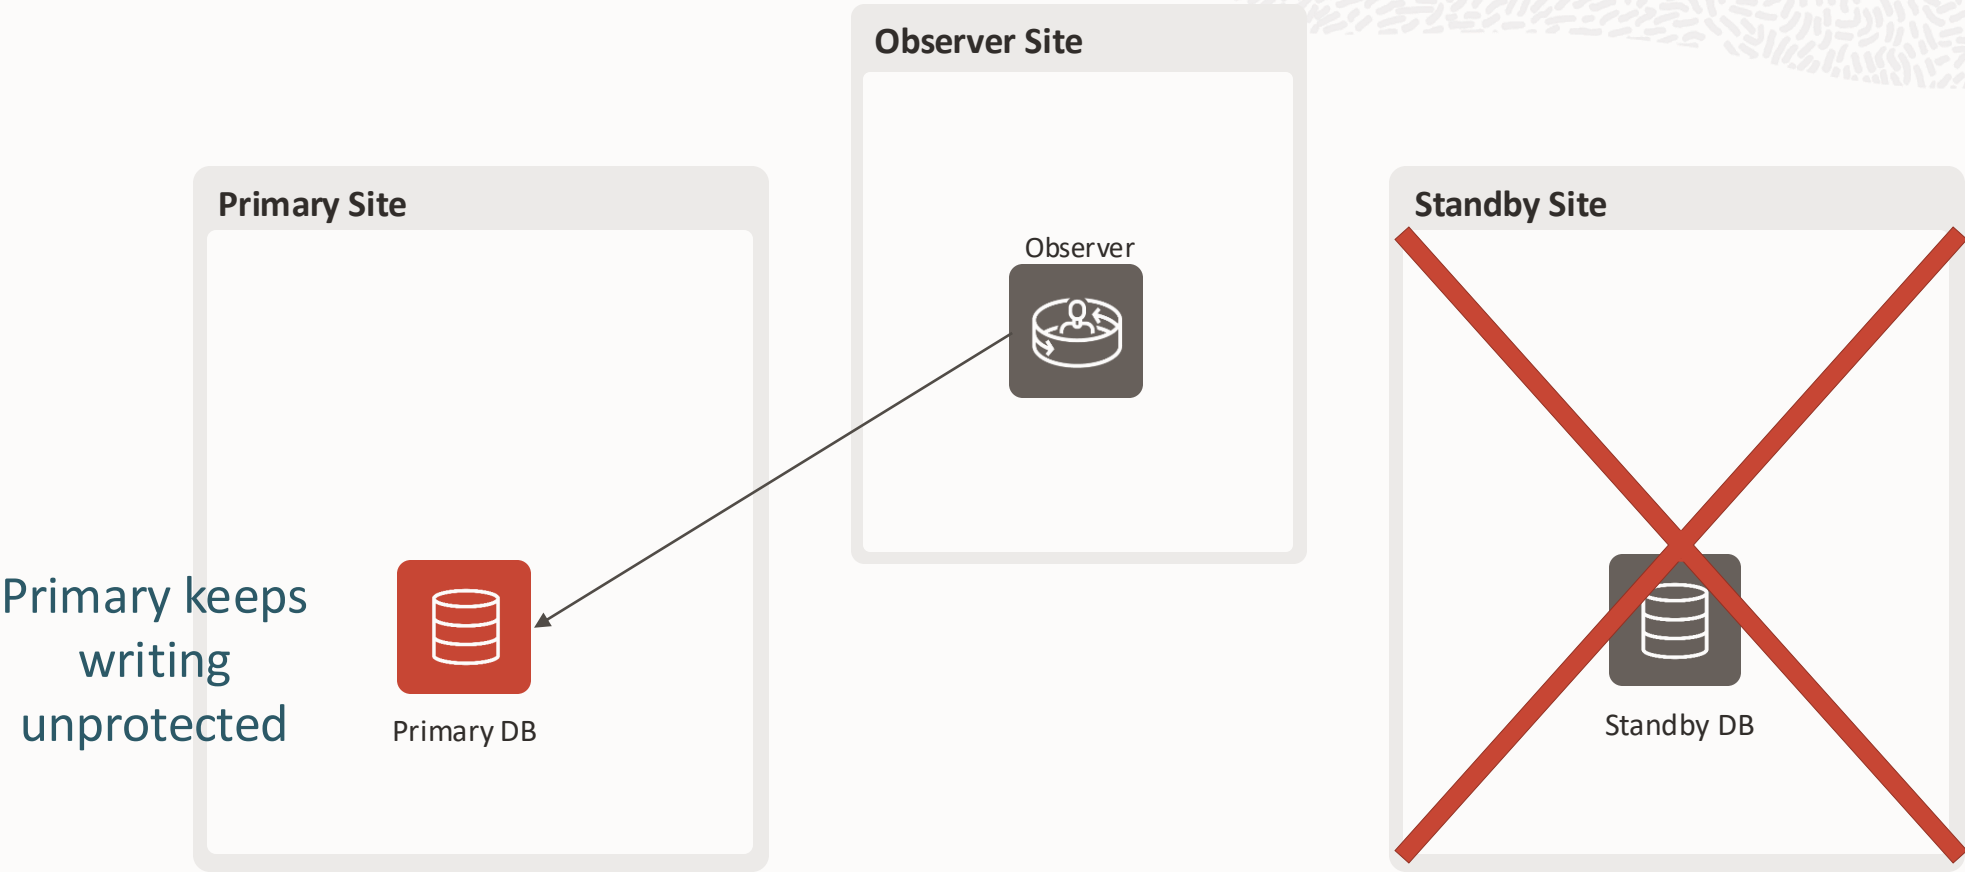


Failure:	Primary DB	Standby DB	Network	Primary Site	Standby Site	Observer
	✓	✓	✓	✓	✓	✓



Oracle Data Guard Observer Placement

Observer at an external size

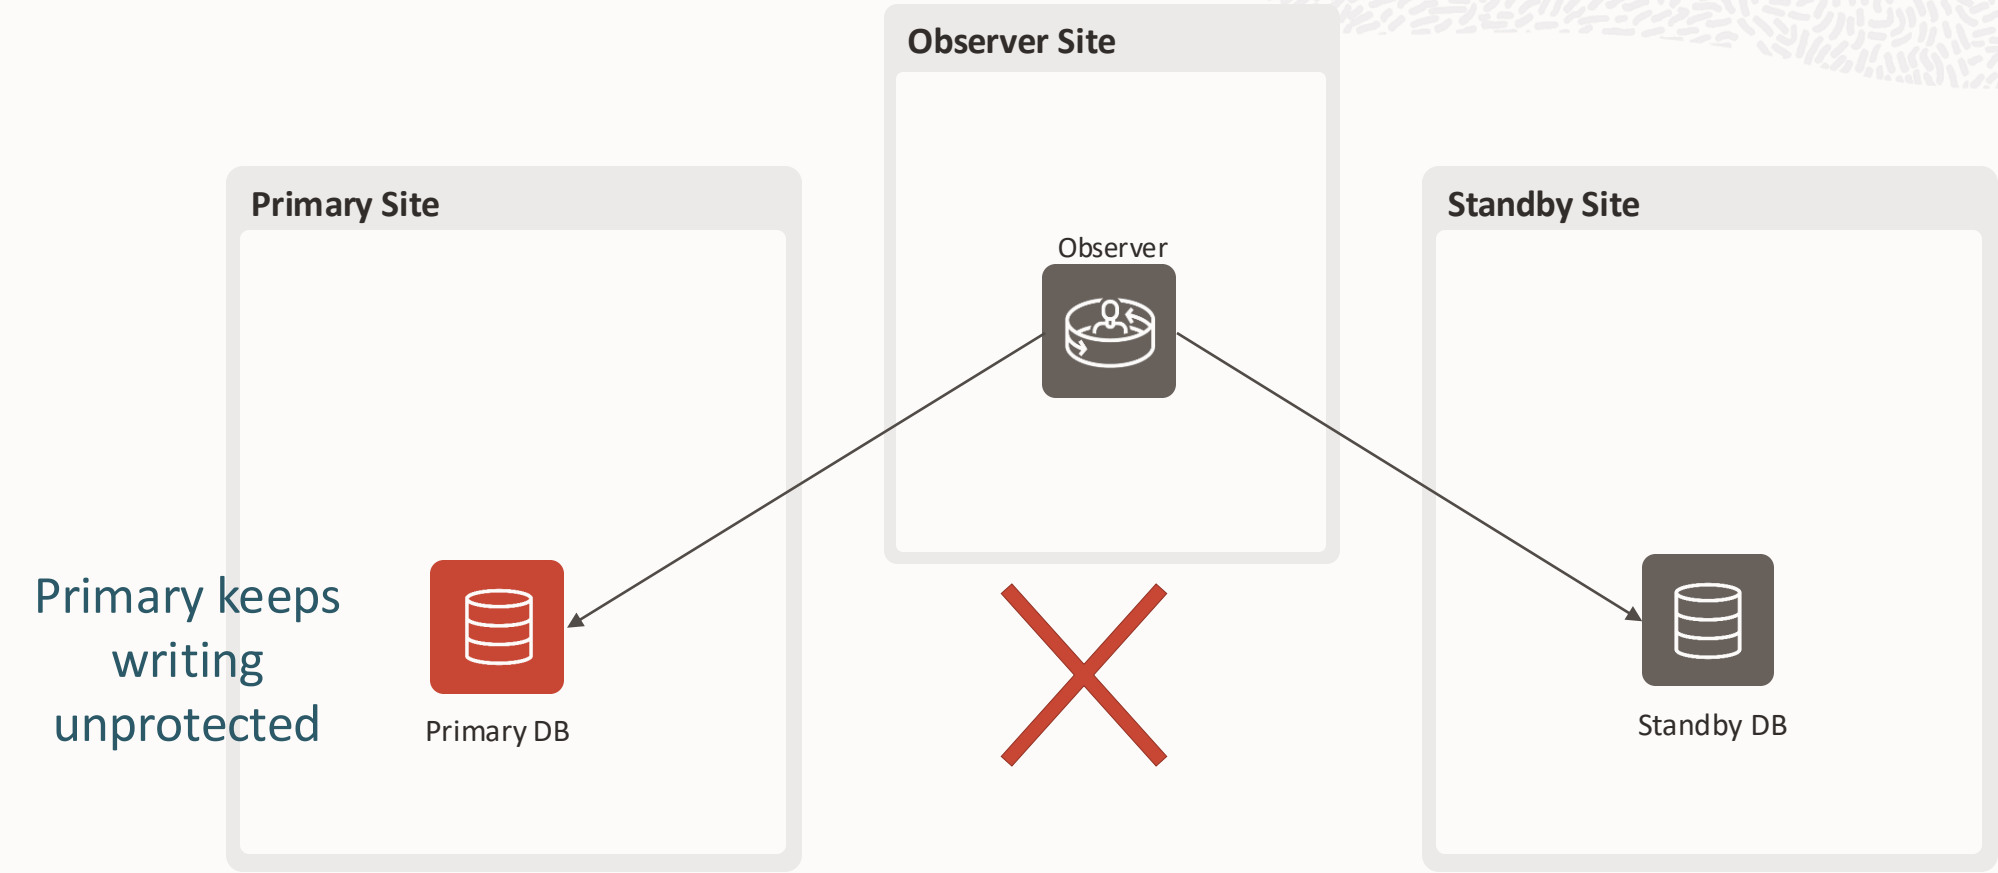


Failure:	Primary DB	Standby DB	Network	Primary Site	Standby Site	Observer
	✓	✓	✓	✓	✓	✓



Oracle Data Guard Observer Placement

Observer at an external size

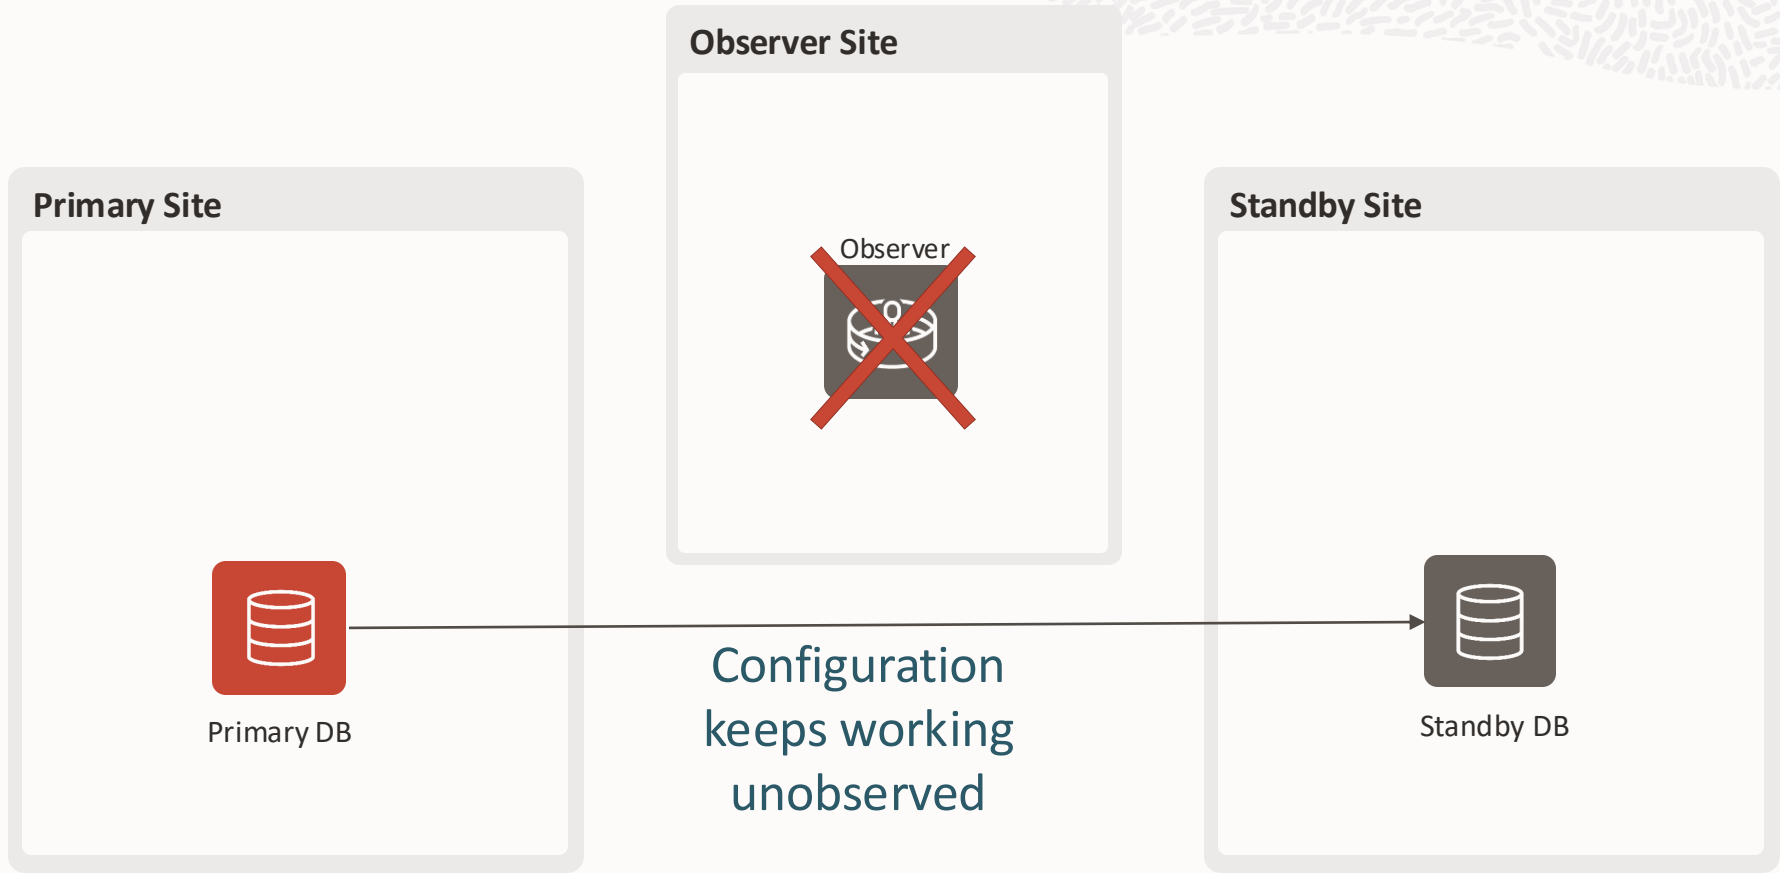


Failure:	Primary DB	Standby DB	Network	Primary Site	Standby Site	Observer
	✓	✓	✓	✓	✓	✓



Oracle Data Guard Observer Placement

Observer at an external size

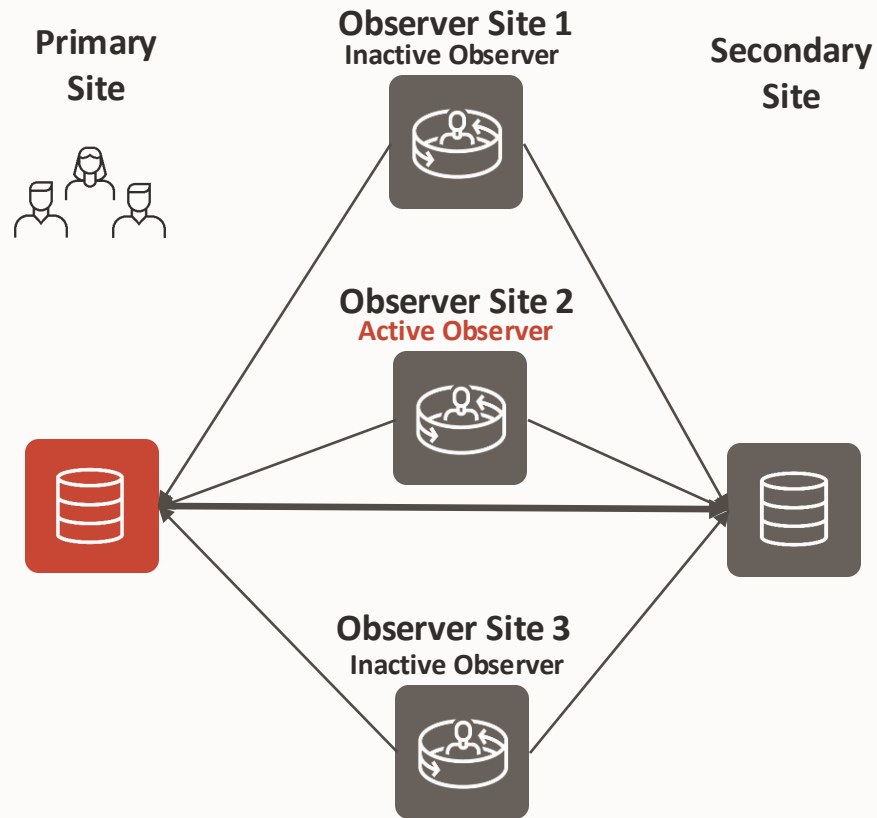


Failure:	Primary DB	Standby DB	Network	Primary Site	Standby Site	Observer
	✓	✓	✓	✓	✓	✓



Oracle Data Guard Observer High Availability

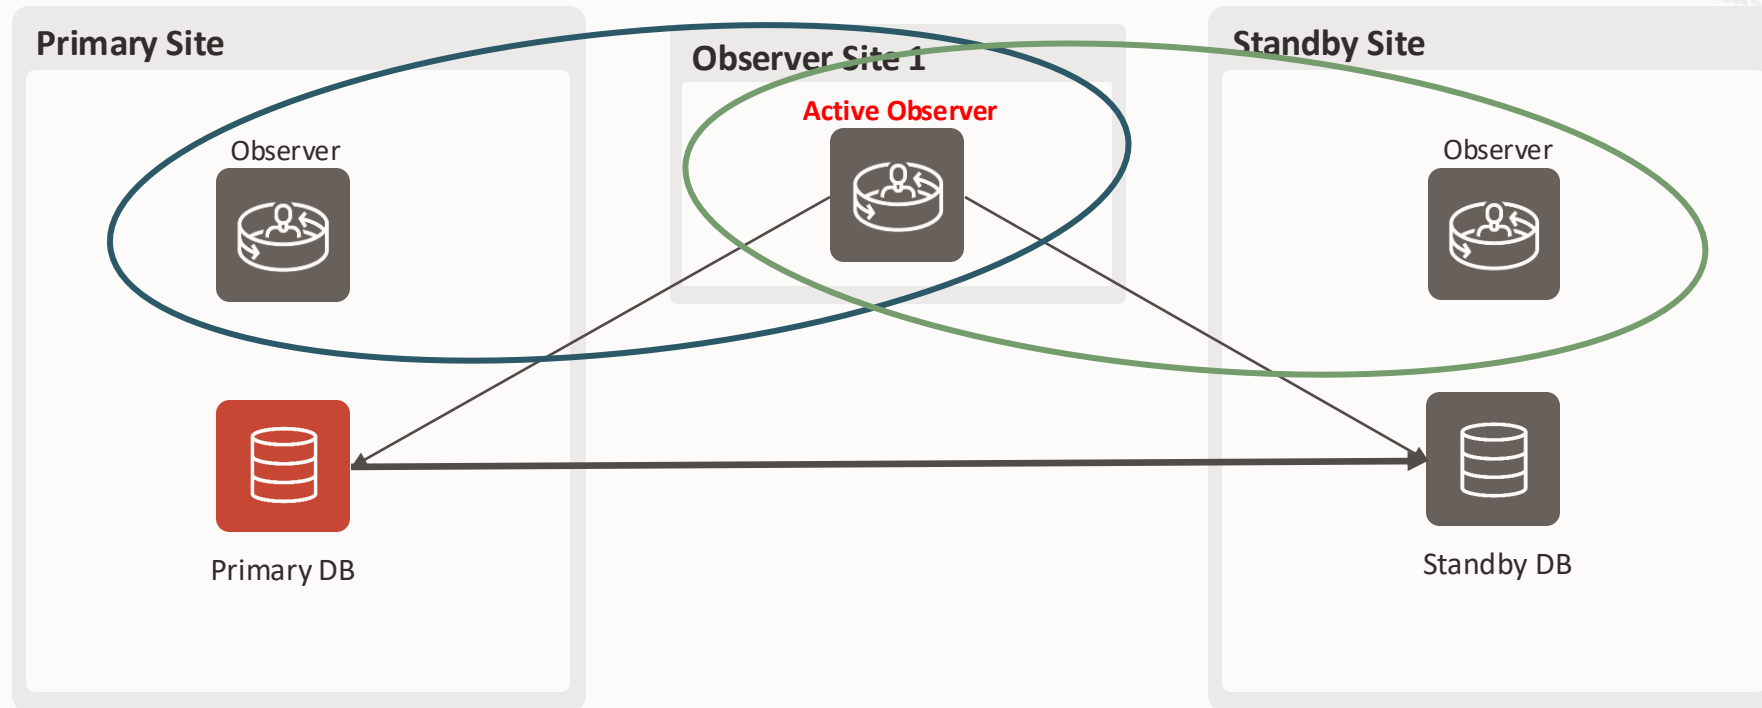
Up to four observers configured (one active at a time)



- **Optimal: 2 or 3 different Regions/Data Centers/Ads**
 - Ensure there are no SPOFs (network, power...)
- If one observer fails, another is promoted

Oracle Data Guard Observer High Availability

Tolerate observer site failure

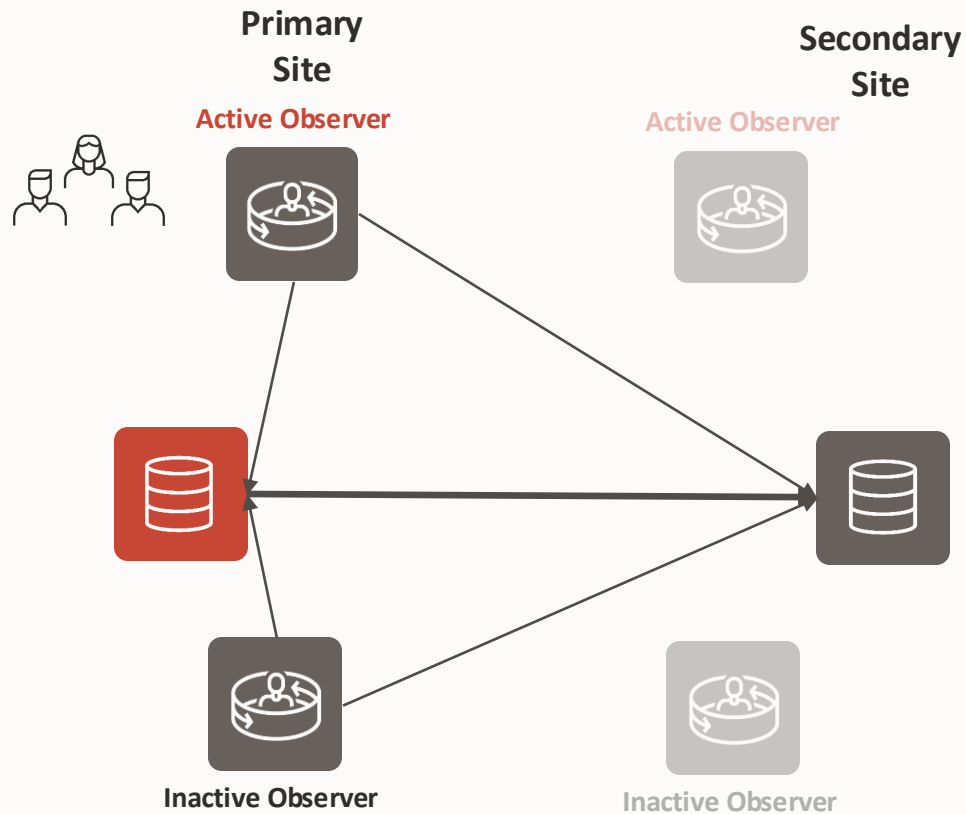


```
edit database db_site1 set property PreferredObserverHosts='obs_ext:1,obs_site1:2';
```

```
edit database db_site2 set property PreferredObserverHosts='obs_ext:1,obs_site2:2';
```

Oracle Data Guard Observer High Availability

Up to four observers configured (one active at a time)




- **No external site?**

- Configure HA observers at the primary site
- Ideally, at least one in the application network
- When role change occurs, have two observers ready at the secondary site

Oracle Data Guard Observer High Availability

Optimal configuration with two sites




```
edit database db_site1 set property PreferredObserverHosts='obs1_site1,obs2_site1';
```

```
edit database db_site2 set property PreferredObserverHosts='obs1_site2,obs2_site2';
```

Oracle Data Guard Observer High Availability

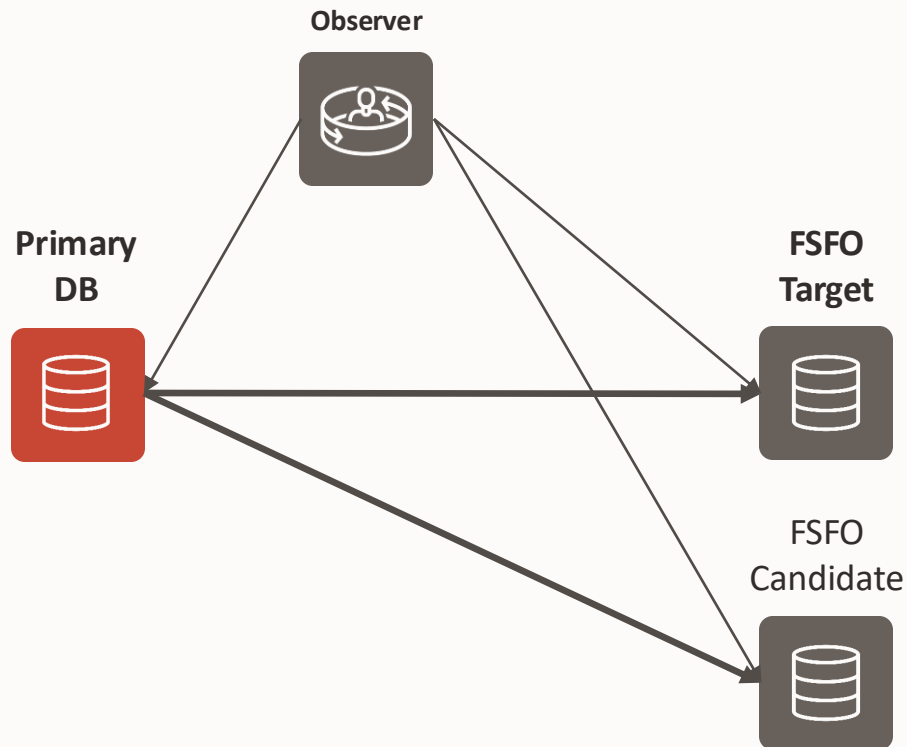
Observer promotion requires both primary and standby databases



- The surviving observer cannot tell if the configuration isn't working on the primary site: a promotion and failover might cause a split-brain

Multiple Fast-Start Failover Targets

Don't let a database failure compromise your protection



```
DGMGRL> edit database BOSTON set property
FastStartFailoverTarget='NASHUA,NEWYORK';
```

```
DGMGRL> edit database NASHUA set property
FastStartFailoverTarget='BOSTON,NEWYORK';
```

```
DGMGRL> edit database NEWYORK set property
FastStartFailoverTarget='BOSTON,NASHUA';
```

```
DGMGRL> show fast_start failover;
```

...

Active Target: NASHUA

Potential Targets: NEWYORK

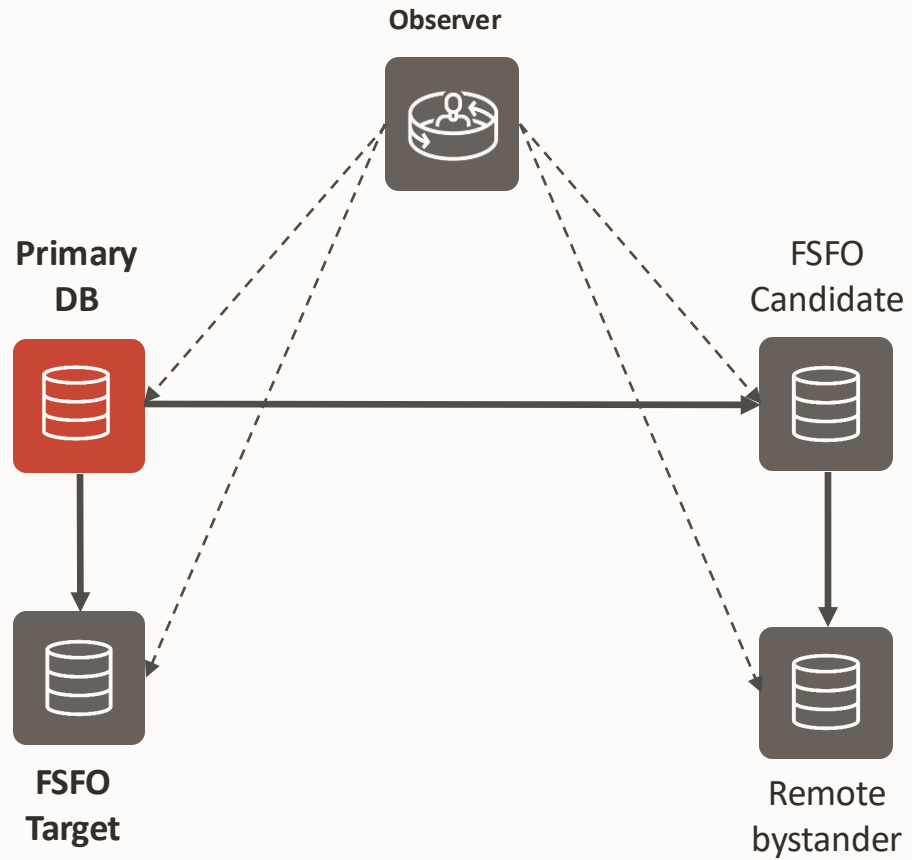
NEWYORK valid

...

Always use at least two standbys in Max Protection!


Network partitions and consistency

Multiple Fast-Start Failover Targets



Network partitions and consistency

Multiple Fast-Start Failover Targets



Primary Isolated

- The observer can still contact the FSFO target.

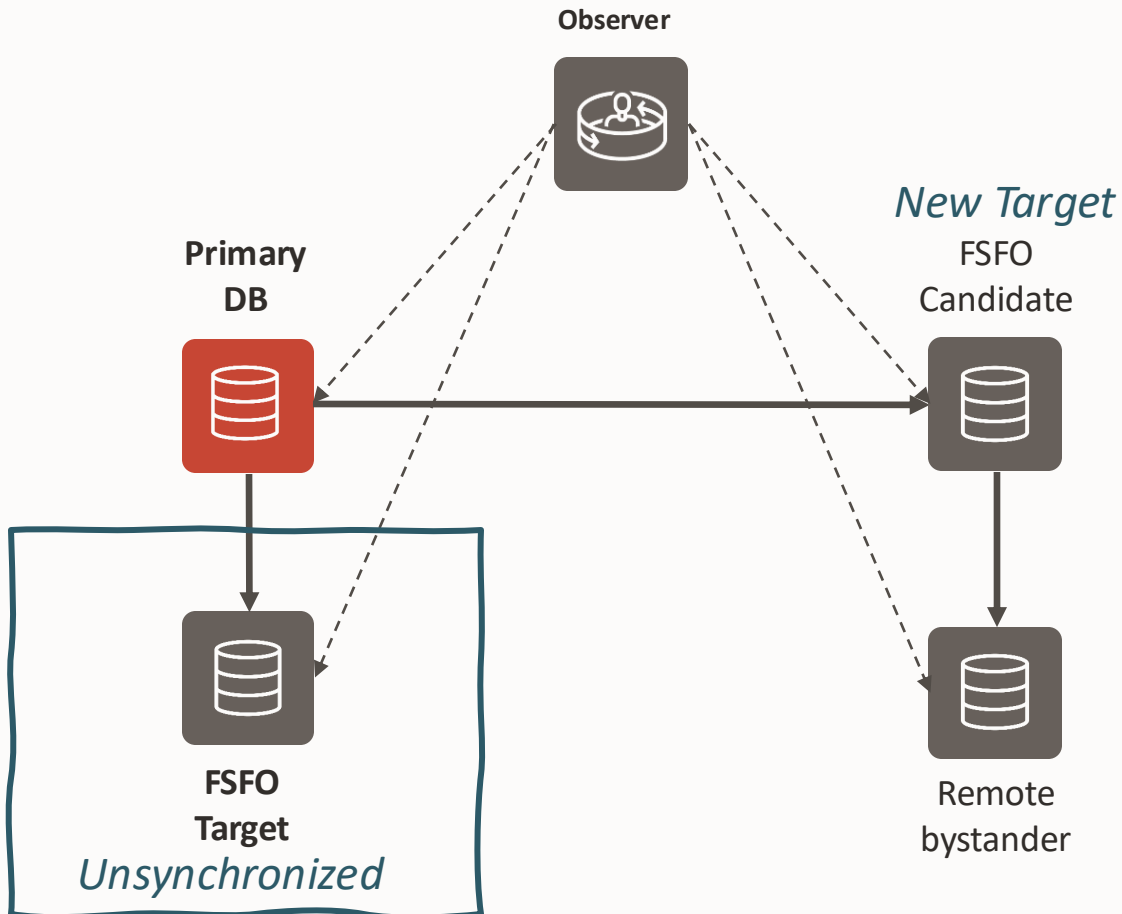
What happens:

1. The primary is STALLED.
2. The observer initiates the failover to FSFO Target.
3. The FSFO Target becomes primary.
4. The new primary and the observer agree to a new FSFO Target.
5. The former Primary DB will require a reinstate.

Execution of Fast Start Failover

Network partitions and consistency

Multiple Fast-Start Failover Targets



FSFO Target Isolated

- The observer can still contact the primary.

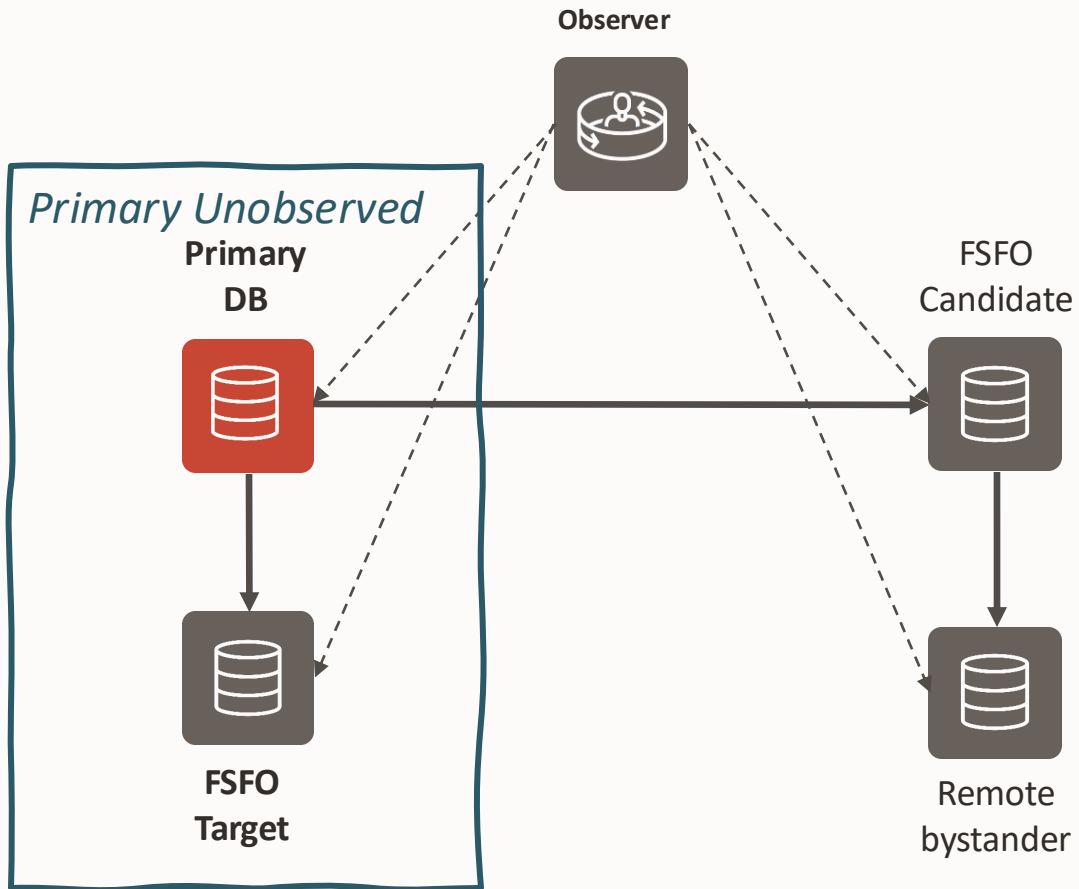
What happens:

1. The primary goes temporarily UNSYNCHRONIZED (no FSFO, unless Max Protection).
2. The new primary and the observer agree to a new FSFO Target.
3. As soon as the new FSFO target is ready, FSFO is possible again.

Fast Start Failover not possible for the time of target switch, then possible again

Network partitions and consistency

Multiple Fast-Start Failover Targets



Primary and FSFO Target Isolated

- The observer cannot contact the primary nor the FSFO target

What happens:

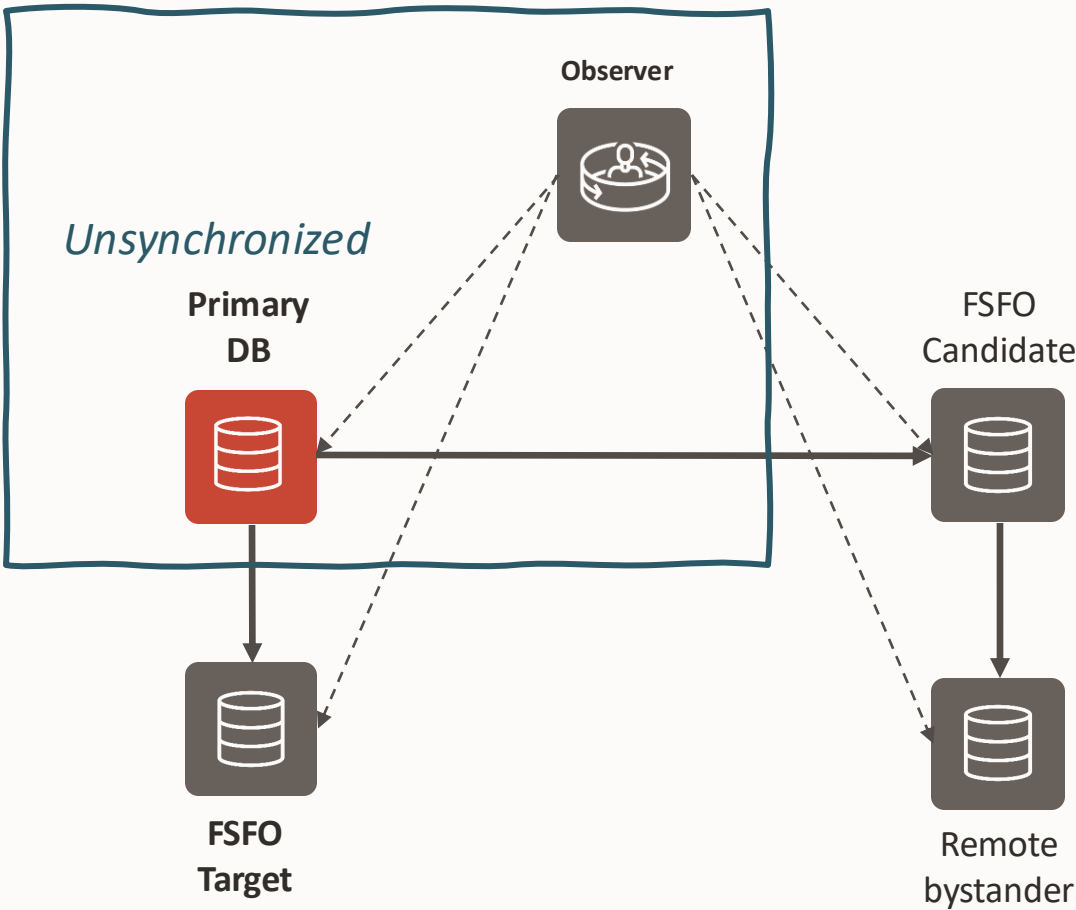
1. The observer cannot tell if the network is unreachable, or the whole site is down. The primary might still write to the standby (valid LAD destination).
2. The primary and FSFO targets keep working, the configuration is UNOBSERVED.
3. The observer cannot initiate a failover, as it would lead to split-brain and data loss.

Fast Start Failover not possible



Network partitions and consistency

Multiple Fast-Start Failover Targets



Primary and Observer Isolated



- No FSFO target or candidates can be contacted.

What happens:

1. The primary keeps working without protection (unless Max Protection).

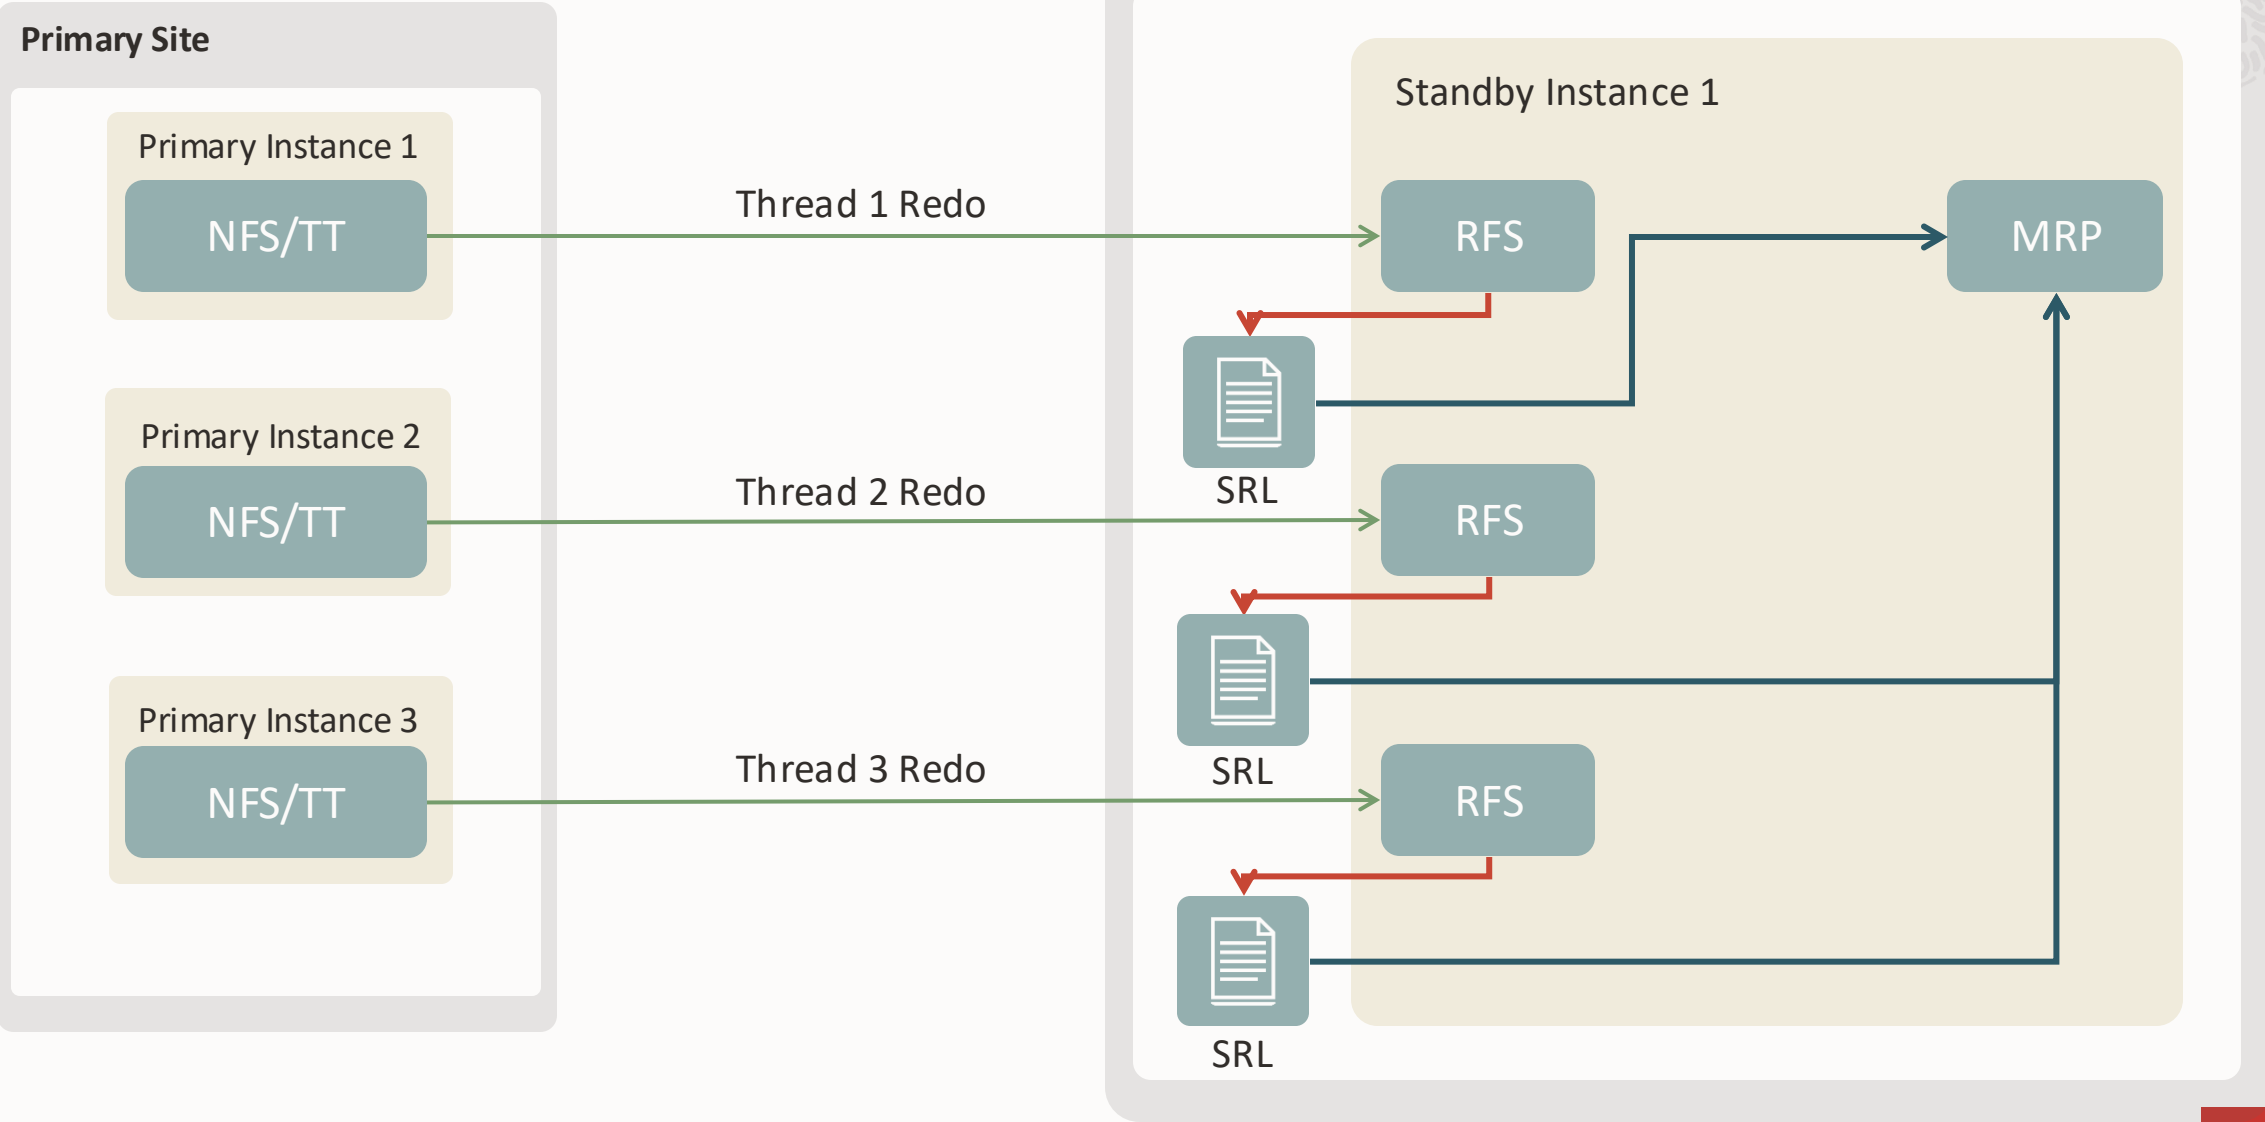
Fast Start Failover not possible






Multi-Instance Redo Apply (MIRA)

Single-Instance Redo Apply (SIRA)




Single-Instance Redo Apply (SIRA)

- The MRP and its redo apply servers run on one node of a Physical Standby RAC.
- Single-Instance Redo apply performance generally meets all use cases.
- Before considering Multi-Instance apply, make sure you apply the best practices for Redo Apply



Multi-Instance Redo Apply




Multi-Instance Redo Apply

- Utilizes all RAC nodes on the Standby database to parallelize recovery
- OLTP workloads on Exadata show great scalability
- Generally 30% improvement or more, depending on the workload



Standby
Apply
Rate
MB/sec



When to consider Multi-Instance Redo Apply



- **IO bottlenecks** and **database wait events** affecting SIRA, will affect MIRA as well!
- Consider MIRA only if SIRA *is* the bottleneck and cannot meet the SLA
- We recommend Oracle Database **19.17 or higher** (it includes critical fixes for MIRA)
- CPU-bounded apply coordinator or worker are the best indicators that MIRA is needed
- **Oracle Data Guard Configuration Best Practices**
<https://docs.oracle.com/en/database/oracle/oracle-database/19/haovw/configure-and-deploy-oracle-data-guard.html#GUID-97769612-4980-42C2-A28C-4C5E49FE2824>
- **Redo Apply Troubleshooting and Tuning**
<https://docs.oracle.com/en/database/oracle/oracle-database/19/haovw/tune-and-troubleshoot-oracle-data-guard.html#GUID-E8C27979-9D37-4899-9306-A5AE2B5CF6C0>

How to Enable Multi-Instance Redo Apply

With the Data Guard Broker:

```
DGMGRL> edit database <standby> set property ApplyInstances=<#|ALL>;
```

From SQL*Plus:

```
SQL> alter database recover managed standby database disconnect from session instances <#|ALL>;
```

Entries in the alert log:

```
ALTER DATABASE RECOVER MANAGED STANDBY DATABASE DISCONNECT FROM SESSION INSTANCES ALL
2018-05-23T11:37:09.937690+01:00
Attempt to start background Managed Standby Recovery process (<db_unique_name>)
...
2018-05-23T11:37:15.111518+01:00
Started logmerger process on instance id 1
Started logmerger process on instance id 2
Starting Multi Instance Redo Apply (MIRA) on 2 instances
...
2018-05-23T11:37:16.027775+01:00
Started 24 apply slaves on instance id 1
2018-05-23T11:37:16.545221+01:00
Started 24 apply slaves on instance id 2
```

Multi-Instance Redo Apply on Exadata



- Exadata prerequisites to enable MIRA

Exadata Systems	RDBMS version	Steps
With PMEM	19.13 and higher	No additional steps
Without PMEM	19.13 and higher	Set dynamic parameter on all instances: "_cache_fusion_pipelined_updates_enable"=FALSE (*)
Any Exadata System	19.12 and lower	Apply Patch 31962730 and set dynamic parameter on all instances: "_cache_fusion_pipelined_updates_enable"=FALSE (*)

(*) MIRA can recover only redo generated with the “_cache_fusion_pipelined_updates_enable” set to FALSE


- Using ExaWatcher Charts to Monitor Exadata Database Machine Performance
<https://docs.oracle.com/en/engineered-systems/exadata-database-machine/dbmmn/exadata-general-maintenance.html#GUID-5AEB3139-333D-453F-91D6-8EB09CB6E6EB>



Tuning Multi-Instance Redo Apply

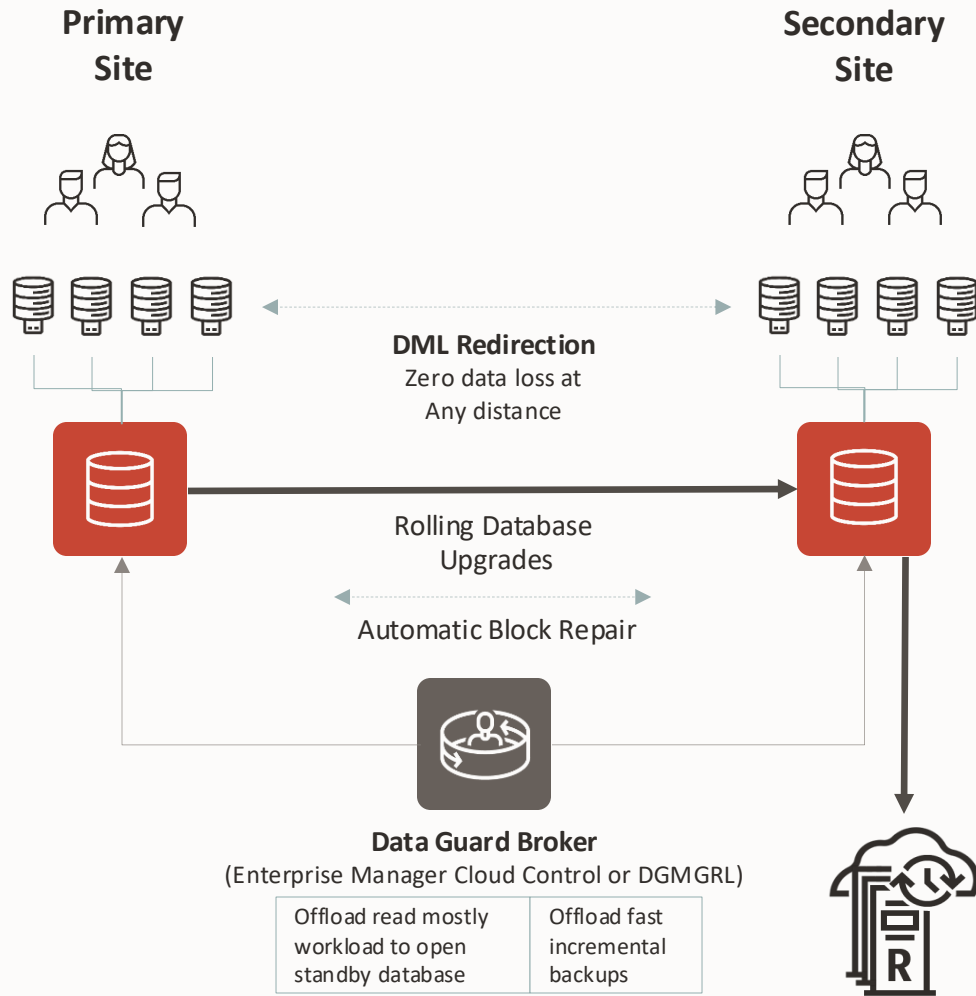


- Tune Redo Apply by evaluating Database **wait events**
- If **recovery apply pending** and/or **recovery receive buffer free** are among the top wait events:
 - Incrementally increase `_mira_num_receive_buffers` and `_mira_num_local_buffers` by 100
 - Additionally set “`_mira_rcv_max_buffers`”=10000
 - The additional memory requirements for each participating MIRA RAC instance:
 $(_mira_num_receive_buffers + _mira_num_local_buffers) * (\#instances * 2MB)$
- If **parallel recovery change buffer free** is among the top wait events:
 - Increase `_change_vector_buffers` to 2 or 4.



Oracle Active Data Guard Overview

Oracle *Active* Data Guard



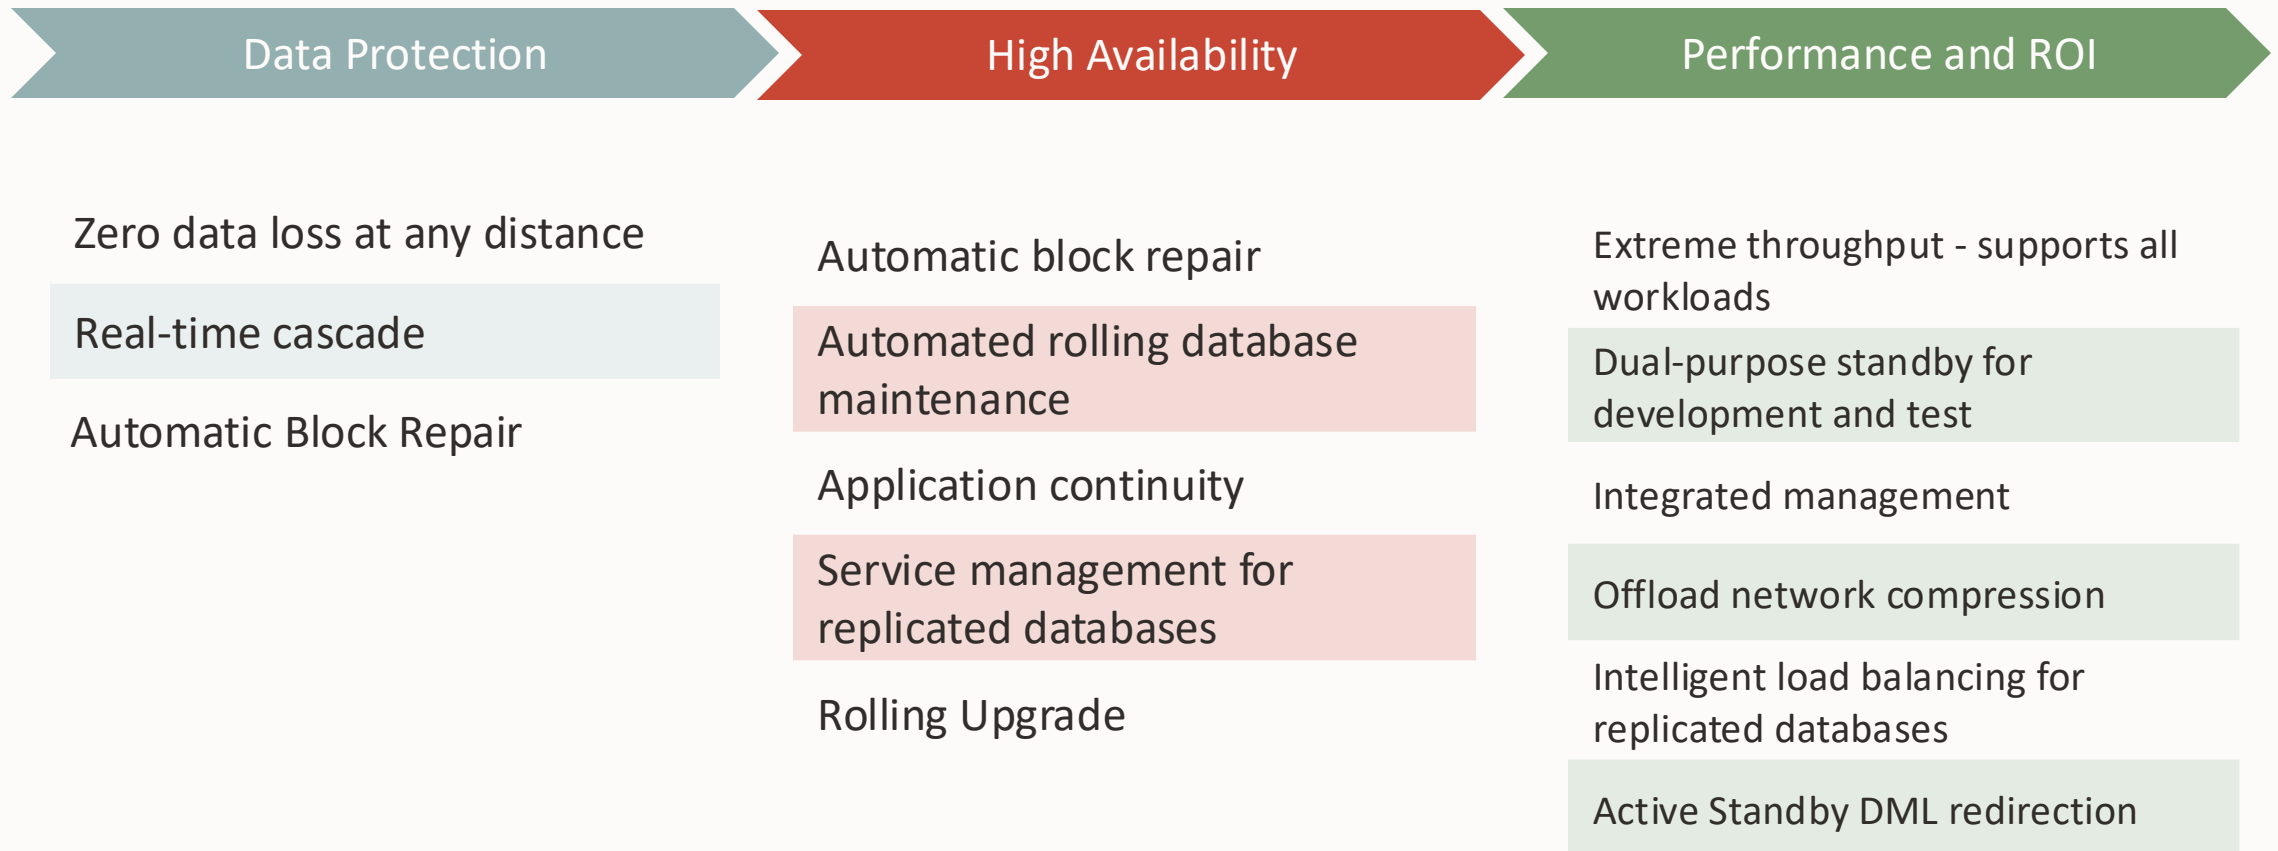
Oracle Data Guard features, plus:



- **Active-active**
 - Queries, reports, backups
 - Occasional updates (19c)
 - Assurance of knowing system is operational
- **Automatic block repair**
- **Application Continuity**
- **Zero data loss across any distance**
- **Rolling Upgrades and Maintenance**
- **Real-time cascaded standbys**
- **Global Data Services**

<https://www.oracle.com/database/technologies/high-availability/dataguard-activedataguard-demos.html>

Active Data Guard

Option of Oracle Database for Advanced Capabilities and Protection

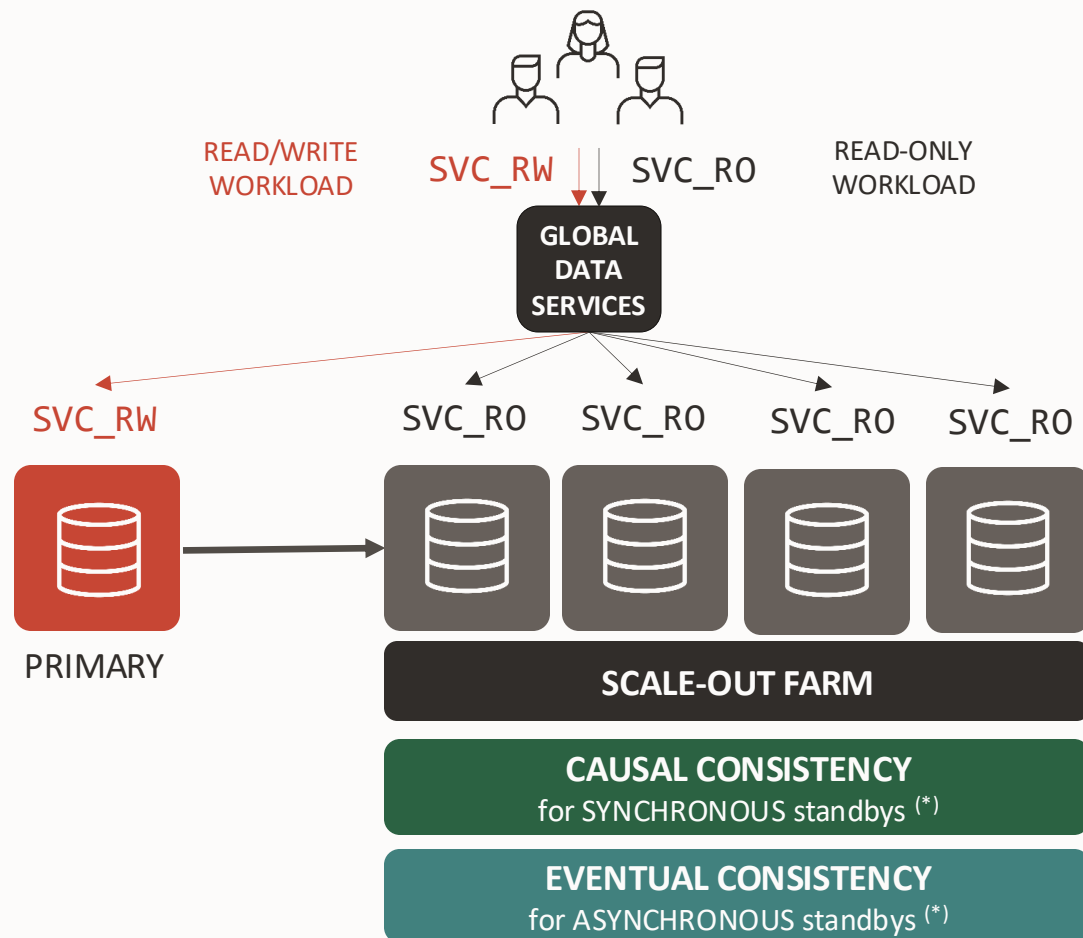




Oracle Active Data Guard Real-Time Query

Scale and Improve ROI of your Active Data Guard Environments

Scale linearly by opening standbys read-only for additional processing power



Benefits of offloading read-only workload:


- Linear scalability of read-only (RO) workloads
- Isolation of primary from heavy queries
- RO sessions uninterrupted during role changes
- Reports and data extractions off the primary
- Fast incremental backups on the standby databases

Global Data Service (optional):

- Simplifies the application configuration with a single, highly-available endpoint
- Sessions are redirected to the best standby based on locality and load
- Problematic or lagging standbys are excluded automatically

Active Data Guard Real-Time Query

Read-only Standby while Recovery is Active



Activation

With Data Guard Broker:


```
SQL> ALTER DATABASE OPEN;
```

Without Data Guard Broker:

```
ALTER DATABASE RECOVER MANAGED STANDBY DATABASE CANCEL;  
ALTER DATABASE OPEN;  
ALTER DATABASE RECOVER MANAGED STANDBY DATABASE DISCONNECT;
```

Real-Time Query Apply Lag Limit

Full read consistency at the session level



```
-- log transport to the standby must be synchronous
DGMGRL> edit database prim
  set property LogXptMode='SYNC';
```

```
-- write on the primary
insert into emp values (...);
commit;
```


```
-- read on the standby
-- wait once until current SCN is applied
alter session sync with primary;
```

```
-- or always have READ COMMITTED in the session
alter session set standby_max_data_delay=0;
```

```
select first_name from emp where ...;
```



Offload Read-Only Workloads

Increase Performance and ROI – Standby is a Production System



Standby Offload Increases Performance for all Workloads

Bring Idle Capacity Online



Double read-write throughput

Increase read-only throughput by 70%

Eliminate contention between read-write and read-only workloads



Real-Time Query

Not just Selects for your Application Workloads!



SQL Performance Analyzer

Oracle Database In-Memory *

Global Temporary Tables

R/O Connections Preserved

Sequences

Updates on Active Data Guard

Standby Result Cache preservation

Simplified AWR snapshots


NEW in 21c

NEW in 23ai

* Only on Engineered Systems or Oracle Cloud Infrastructure

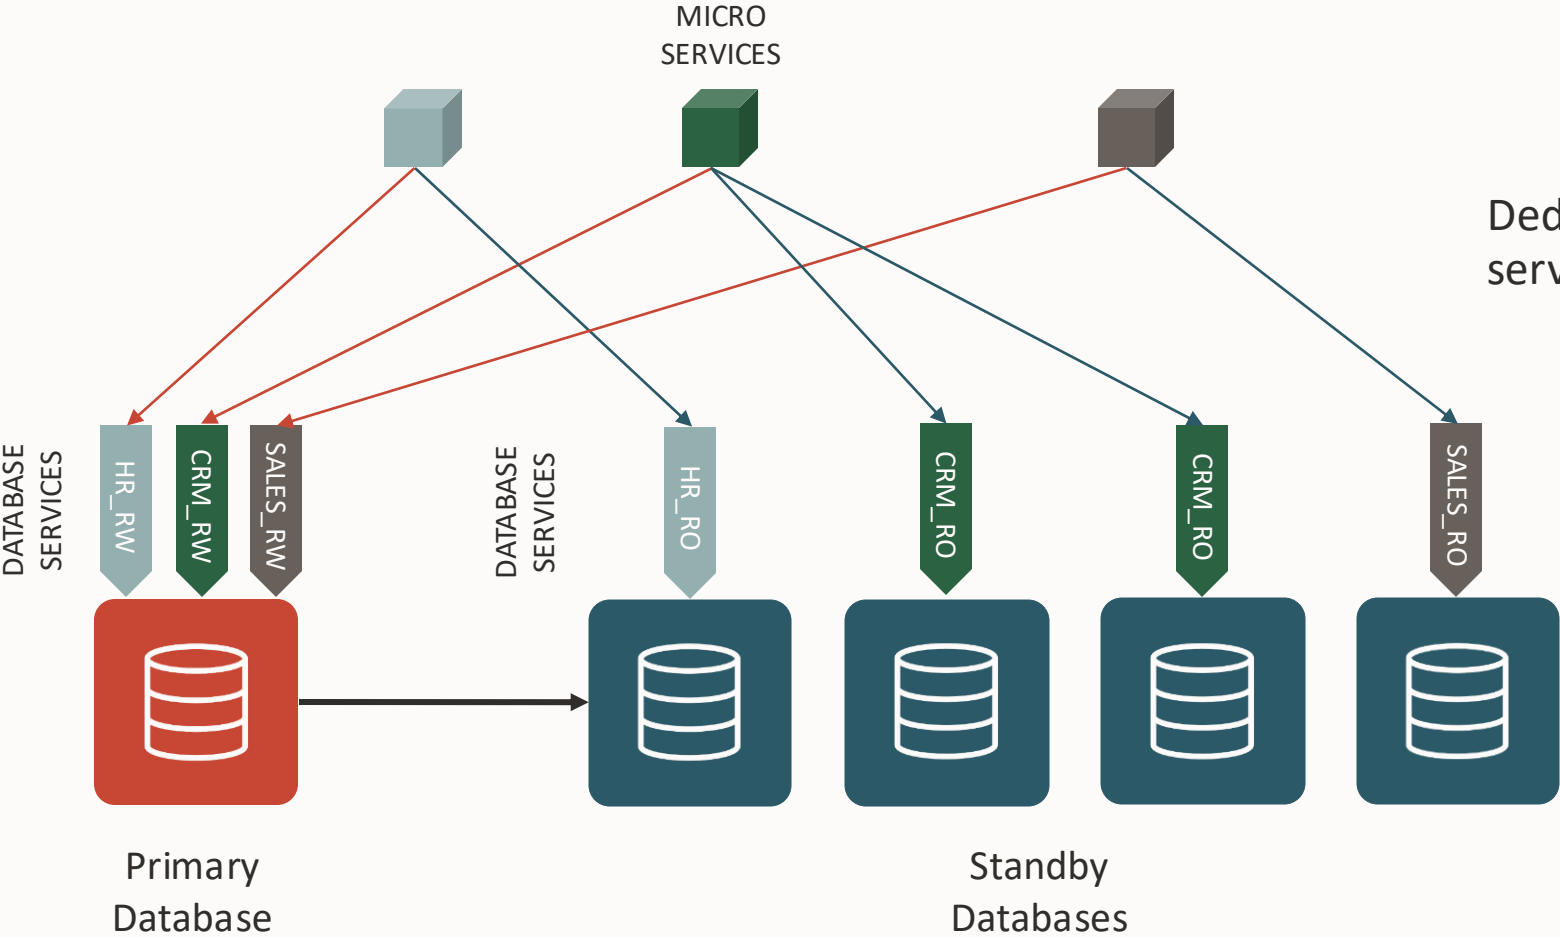


Active Data Guard and Database In-Memory



- **In-Memory queries run on standby**
 - No impact on the primary database
 - Full use of standby database resources
- Standby can have different in-memory contents from Primary
 - **DISTRIBUTE FOR SERVICE** subclause used to determine data placement
 - Increases total effective in-memory columnar **capacity**
 - Increases column store **availability**:
 - Reporting workload on standby unaffected by primary site outage

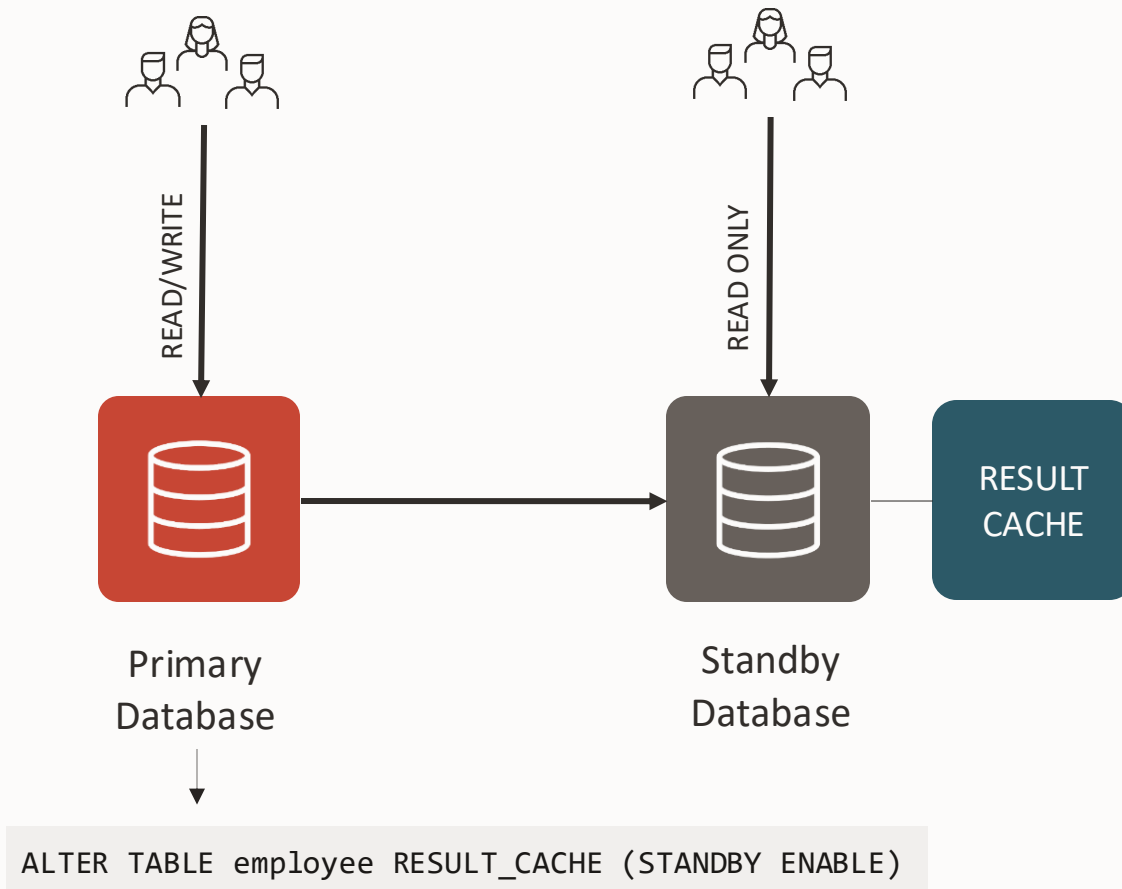
Service Distribution to Optimize the Buffer Cache



Standby Result Cache preservation

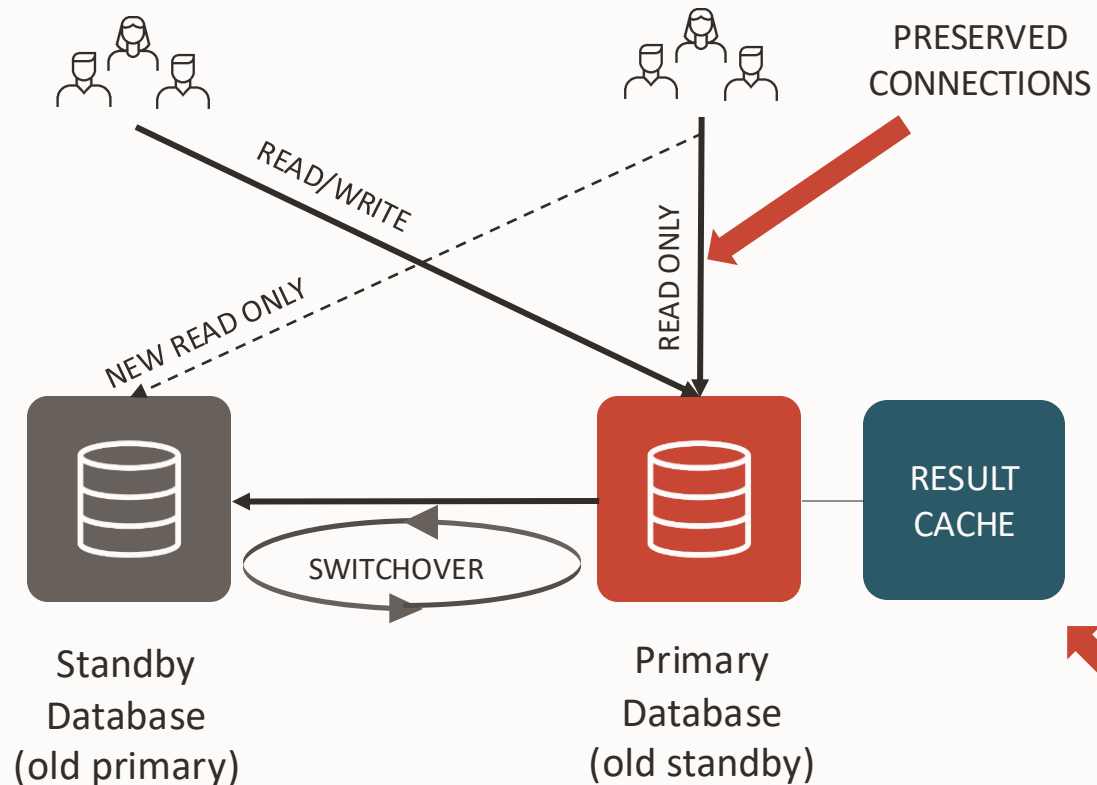
Keep the Result Cache warm after a role transition

- Real-Time Query supports the Result Cache for queries run on the standby database (tables only)
- Result Cache improves query performance for recurring queries and reduces resource usage (CPU, I/O)



Standby Result Cache preservation

Keep the Result Cache warm after a role transition



- Real-Time Query supports the Result Cache for queries run on the standby database (tables only)
- Result Cache improves query performance for recurring queries and reduces resource usage (CPU, I/O)
- In **21c**, after a role transition (switchover or failover), the Result Cache is preserved
 - Query performance not impacted
 - No cache warm-up required

Simplified AWR snapshots on Active Data Guard


Just create the snapshot on the standby, and you are ready to go.

```
-- at the PDB or CDB level
alter session set container = PDB1;


-- the snapshot can be created without additional configuration
select dbms_workload_repository.create_snapshot from dual;

-- create the report at the PDB or CDB level
@?/rdbms/admin/awrrpti
```

```
-- databases already using SYS$UMF can set this
-- to use the new framework instead:
alter system set "_umf_remote_enabled"=FALSE;
```

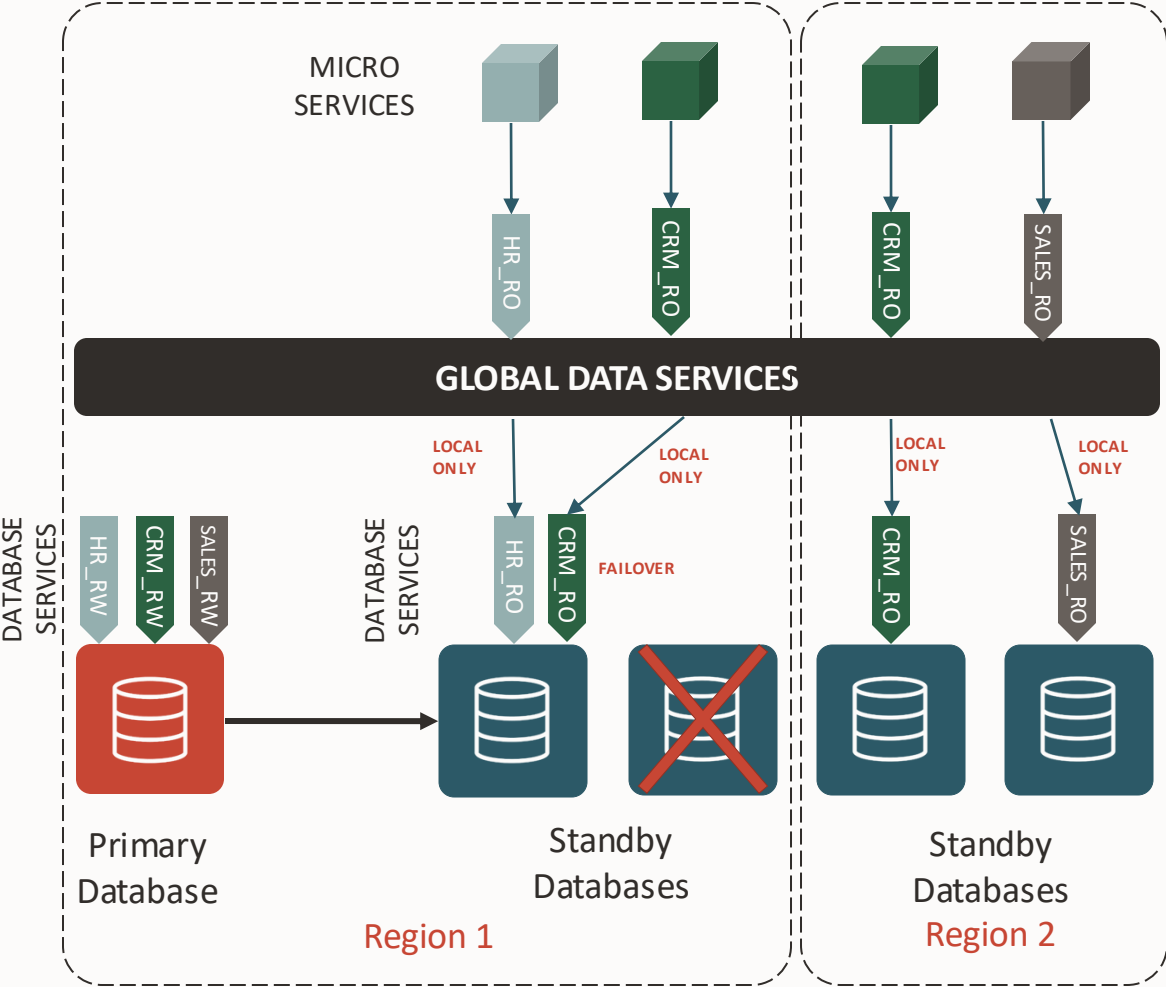


Oracle Global Data Services (GDS)



- Automatic and transparent client workload management across replicas
- Extends the concept of services across clusters
- Capabilities
 - Workload routing based on standby **load**, **locality**, or **lag**
 - Service failover across replicas
- Benefits
 - Maximize application performance
 - Mitigate downtime during planned and unplanned outages
 - Centrally manage services and resources of replicas


Oracle Global Data Services (GDS)



- **Inter-database service failover**
 - If a cluster fails, the service is restarted automatically on another cluster
 - Clients reconnect automatically where the service is available
- Workload routing (region-based and lag-based)
- Load balancing (connect-time & run-time)




Oracle Global Data Services (GDS)



- **Region-based** services can be configured to accept (fail over) deny requests from remote clients should a region become unavailable





Oracle Active Data Guard DML Redirection

Bigger Footprint of ADG Applications


DML on Active Data Guard

DML Re-direction is automatically performed from an Active Data Guard standby to the primary without compromising ACID compliance

- New documented parameter `ADG_REDIRECT_DML` controls DML Redirection
- New `alter system set ADG_REDIRECT_DML | alter session enable ADG_REDIRECT_DML`
- New `ADG_REDIRECT_PLSQL` commands

Supported with Oracle Database 19c

Targeted for “Read-Mostly,
Occasional Updates” applications



Active Data Guard DML Replication

Easy and ready to use



By default DMLs are not possible on the standby

```
SQL> update hr.employee set salary=salary+100 where employee_id=1;  
ERROR at line 1:  
ORA-16000: database or pluggable database open for read-only access
```


Enable DML redirection

```
SQL> alter session enable ADG_REDIRECT_DML;
```

DMLs work seamlessly



```
SQL> update hr.employee set salary=salary+100 where employee_id=1;  
1 row updated.  
SQL> commit;  
Commit complete.
```

Improved Performance of Redirected Transactions



- 19c redirected statements wait for the DML to be applied on the standby before returning to the application.
 - Wait event: `"standby query scn advance"`
- From **21c** onwards, the statement returns as soon as it's executed on the primary.
 - The session waits only at commits or when the modified data is needed for consistent reads.
 - The non-documented parameter: `"_alter_adg_redirect_behavior"` can be set to `"sync_each_dml"` to restore the previous behavior.

* 16 concurrent Order Entry sessions simulated with Swingbench with mixed 'NewCustomerProcess' and 'BrowseProducts' transactions.





Oracle Active Data Guard Automatic Block Repair

Oracle Active Data Guard Automatic Block Repair

Transparently repairs corrupted blocks

- Oracle detects if a block is corrupted when reading it
- The corruption is automatically repaired using a good copy
 - From the standby when the corruption is on the primary
 - From the primary when the corruption is on the standby






Oracle Active Data Guard Fast Incremental Backup on Physical Standby

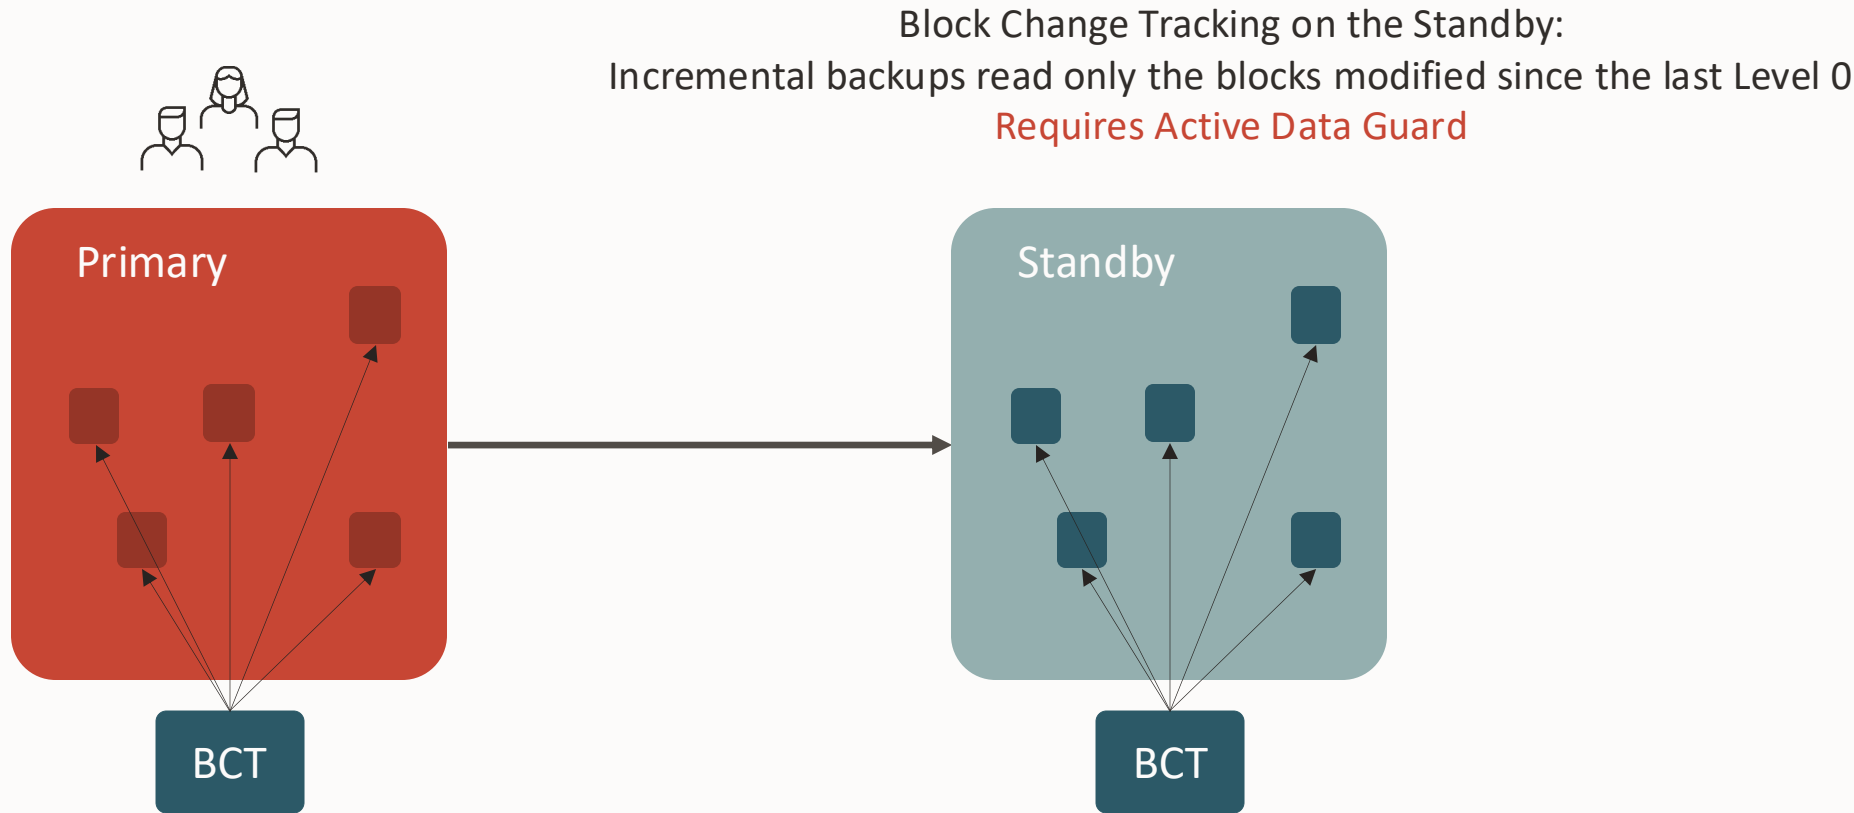
Fast Incremental Backup on Physical Standby

Enable the Block Change Tracking to speed up backups and avoid unnecessary I/O



Fast Incremental Backup on Physical Standby

Enable the Block Change Tracking to speed up backups and avoid unnecessary I/O






Oracle Active Data Guard Real-Time Cascade Standbys

Active Data Guard: up to 30 direct standbys and 253 total members

Far Sync and Cascading Standby open endless possibilities



Active Data Guard Real-Time Cascade Standby

Offload multiple redo transports to a first-level standby

BOSTON



RedoRoutes=

```
(LOCAL : ( NASHUA SYNC PRIORITY=1, NEWYORK ASYNC PRIORITY=8, NEWARK ASYNC PRIORITY=8 ))  
(NASHUA : NEWYORK ASYNC, NEWARK ASYNC ))
```

NASHUA



RedoRoutes=

```
(LOCAL : ( BOSTON SYNC PRIORITY=1, NEWYORK ASYNC PRIORITY=8, NEWARK ASYNC PRIORITY=8 ))  
(BOSTON : NEWYORK ASYNC, NEWARK ASYNC ))
```

NEWYORK



NEWARK

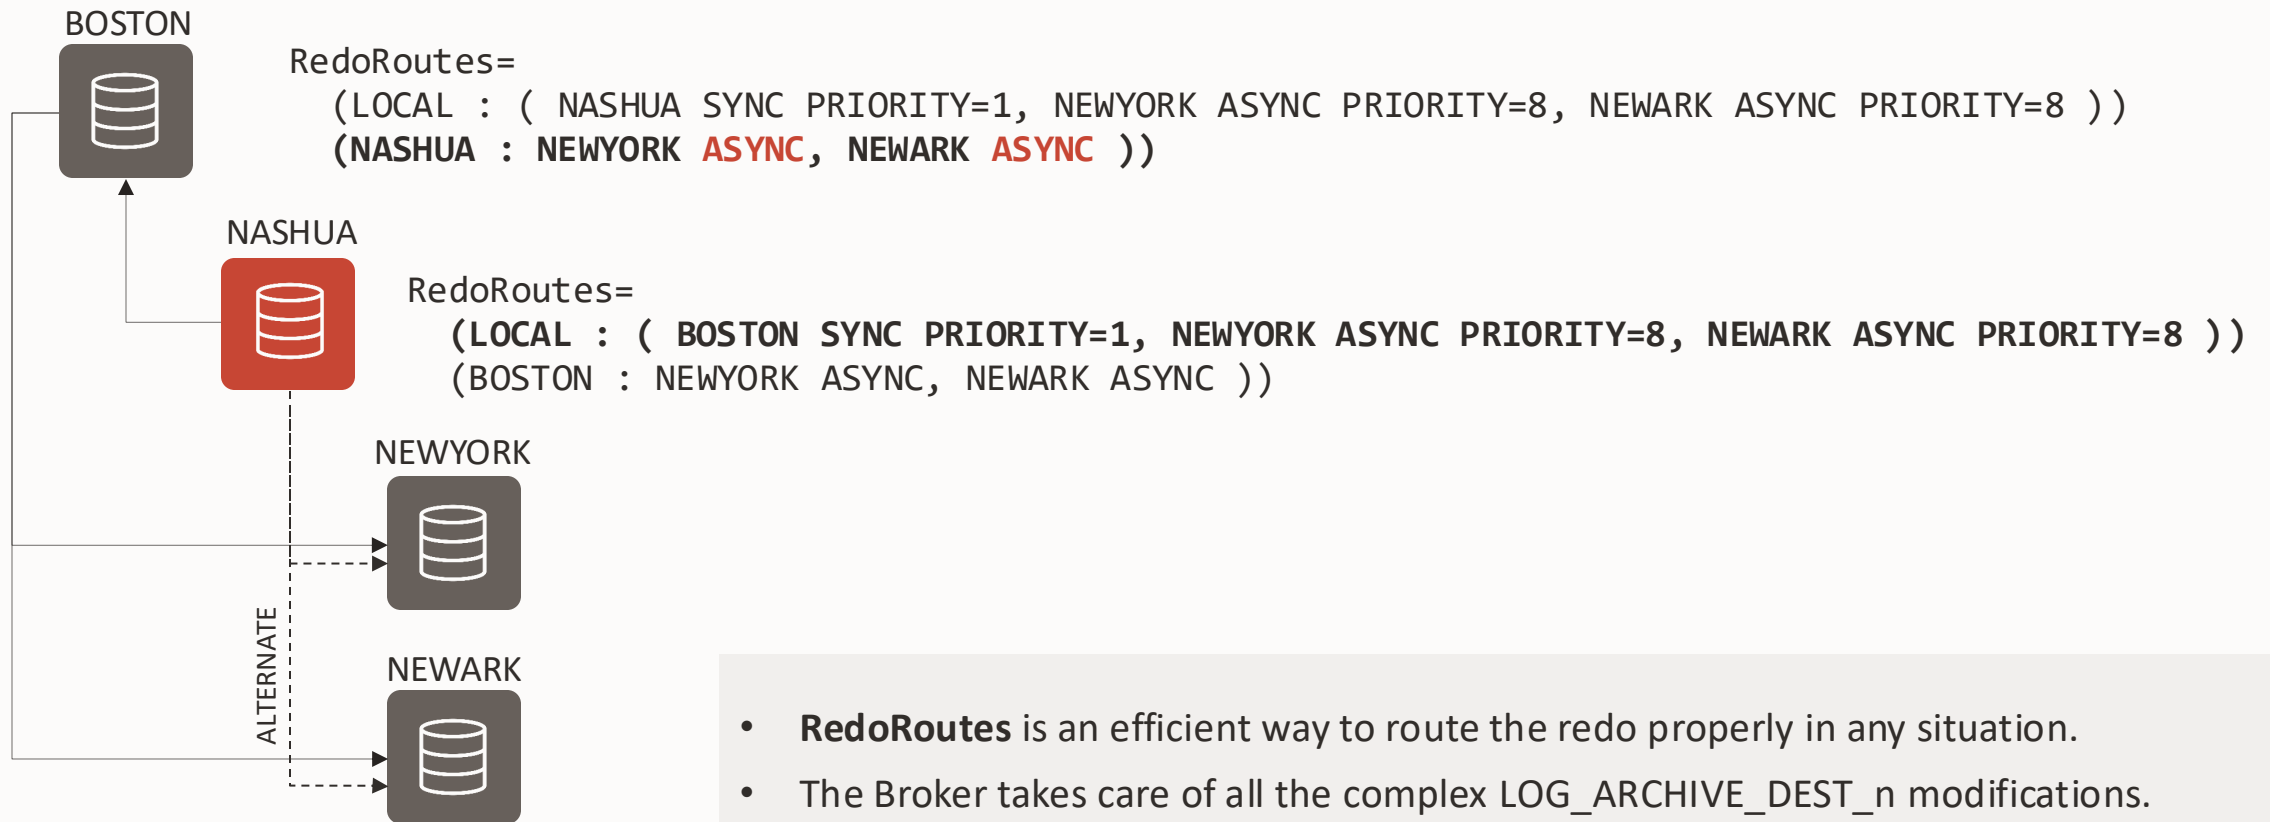


ALTERNATE

- Explicit “ASYNC” in the cascading member means “Real-Time Cascade”. Such configuration requires **Active Data Guard**.
- If not specified, the redo is shipped at log switch.

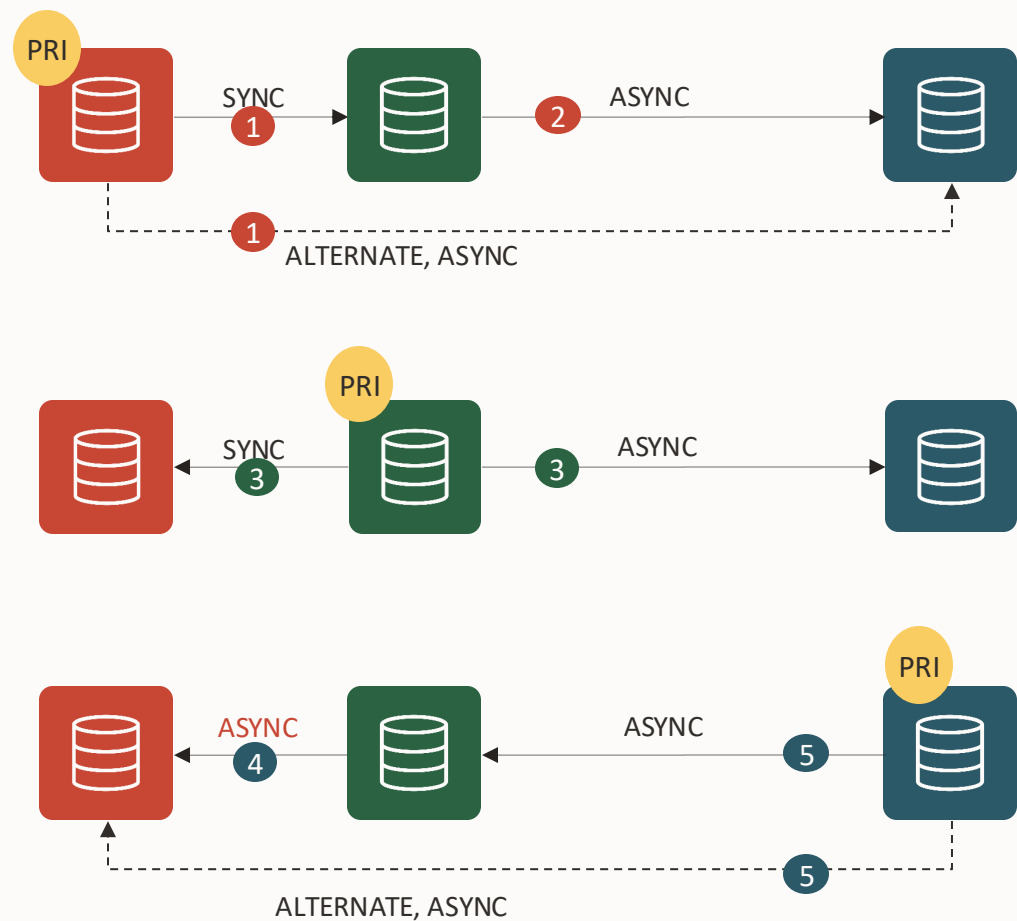
Active Data Guard Real-Time Cascade Standby

Offload multiple redo transports to a first-level standby



Understanding RedoRoutes

Start by drawing all the permutations and describe what you want



Where does RED send the redo?

- 1 When RED is primary:
 - to GREEN (SYNC)
 - or BLUE (ASYNC) if GREEN is not available

Where does GREEN send the redo?

- 2 When RED is primary:
 - to BLUE (ASYNC) (real-time cascade)
- 3 When GREEN is primary:
 - to RED (SYNC)
 - and to BLUE (ASYNC)
- 4 When BLUE is primary:
 - to RED (ASYNC) (real-time cascade)

Where does BLUE send the redo?


- 5 When BLUE is primary:
 - to GREEN (ASYNC)
 - or RED (ASYNC) if GREEN is not available


R
e
d
o
R
o
u
t
e
s






Understanding RedoRoutes


RedoRoutes tells where a DB (or Far Sync) should send the redo, depending on the primary


EDIT DATABASE RED SET PROPERTY RedoRoutes =  Where does RED send the redo?

- ① (RED: (GREEN SYNC PRIORITY=1, BLUE ASYNC PRIORITY=2))  ① When RED is primary:
- to GREEN (SYNC)
or BLUE (ASYNC) if GREEN is not available

EDIT DATABASE GREEN SET PROPERTY RedoRoutes =  Where does GREEN send the redo?

- ② (RED: BLUE ASYNC)  ② When RED is primary:
- to BLUE (ASYNC) (real-time cascade)
- ③ (GREEN: RED SYNC, BLUE ASYNC)  ③ When GREEN is primary:
- to RED (SYNC)
and to BLUE (ASYNC)
- ④ (BLUE: RED ASYNC)  ④ When BLUE is primary:
- to RED (ASYNC) (real-time cascade)

EDIT DATABASE BLUE SET PROPERTY RedoRoutes =  Where does BLUE send the redo?

- ⑤ (BLUE: (GREEN ASYNC PRIORITY=1, RED ASYNC PRIORITY=2))  ⑤ When BLUE is primary:
- to GREEN (ASYNC)
or RED (ASYNC) if GREEN is not available

Verifying the RedoRoutes configuration


```
EDIT DATABASE RED SET PROPERTY RedoRoutes = '(RED: ( GREEN SYNC PRIORITY=1, BLUE ASYNC PRIORITY=2 ))';  
EDIT DATABASE GREEN SET PROPERTY RedoRoutes = '(RED:BLUE ASYNC)(GREEN:RED SYNC,BLUE ASYNC)(BLUE:RED ASYNC)';  
EDIT DATABASE BLUE SET PROPERTY RedoRoutes = '(BLUE: ( GREEN ASYNC PRIORITY=1, RED ASYNC PRIORITY=2 ))';
```

```
DGMGRL> show configuration when primary is red ;
```

Configuration when red is primary - redoroutes_demo

Members:

- red - Primary database
- green - Physical standby database
- blue - Physical standby database (receiving current redo)
- blue - Physical standby database (alternate of green)



Verifying the RedoRoutes configuration



```
EDIT DATABASE RED SET PROPERTY RedoRoutes = '(RED: ( GREEN SYNC PRIORITY=1, BLUE ASYNC PRIORITY=2 ))';  
EDIT DATABASE GREEN SET PROPERTY RedoRoutes = '(RED:BLUE ASYNC)(GREEN:RED SYNC,BLUE ASYNC)(BLUE:RED ASYNC)';  
EDIT DATABASE BLUE SET PROPERTY RedoRoutes = '(BLUE: ( GREEN ASYNC PRIORITY=1, RED ASYNC PRIORITY=2 ))';
```

```
DGMGRL> show configuration when primary is green ;
```

Configuration when green is primary - redoroutes_demo

Members:

green - Primary database

red - Physical standby database

blue - Physical standby database



Verifying the RedoRoutes configuration


```
EDIT DATABASE RED SET PROPERTY RedoRoutes = '(RED: ( GREEN SYNC PRIORITY=1, BLUE ASYNC PRIORITY=2 ))';  
EDIT DATABASE GREEN SET PROPERTY RedoRoutes = '(RED:BLUE ASYNC)(GREEN:RED SYNC,BLUE ASYNC)(BLUE:RED ASYNC)';  
EDIT DATABASE BLUE SET PROPERTY RedoRoutes = '(BLUE: ( GREEN ASYNC PRIORITY=1, RED ASYNC PRIORITY=2 ))';
```

```
DGMGRL> show configuration when primary is blue ;
```

Configuration when blue is primary - redoroutes_demo

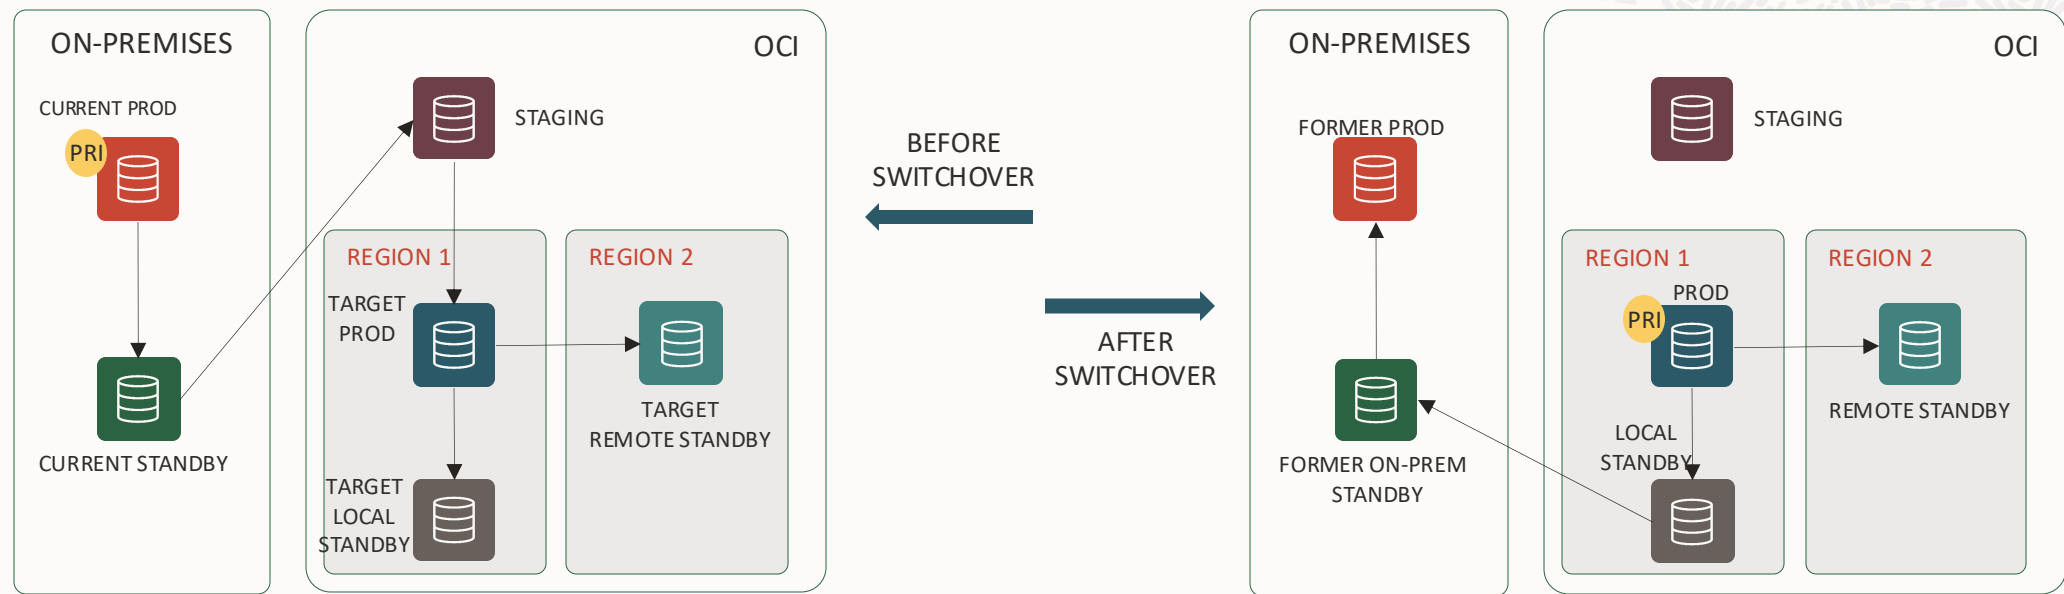
Members:

- blue - Primary database
- green - Physical standby database
- red - Physical standby database (receiving current redo)
- red - Physical standby database (alternate of green)



Real Life Use Case #1

Highly available migration to the Oracle Cloud Infrastructure



```
DGMGRL> show configuration when primary is red;
```

Configuration when blue is primary - redoroutes_demo

Members:

- red - Primary database
- green - Physical standby database
- purple - Physical standby database (receiving current redo)
- blue - Physical standby database (receiving current redo)
- gray - Physical standby database (receiving current redo)
- turquoise - Physical standby database (receiving current redo)

```
DGMGRL> show configuration when primary is blue ;
```

Configuration when blue is primary - redoroutes_demo

Members:

- blue - Primary database
- turquoise - Physical standby database
- gray - Physical standby database
- green - Physical standby database (receiving current redo)
- red - Physical standby database (receiving current redo)


Members Not Receiving Redo:

- purple - Physical standby database



Real Life Use Case #2

Read-only farm for intensive, latency-sensitive workloads

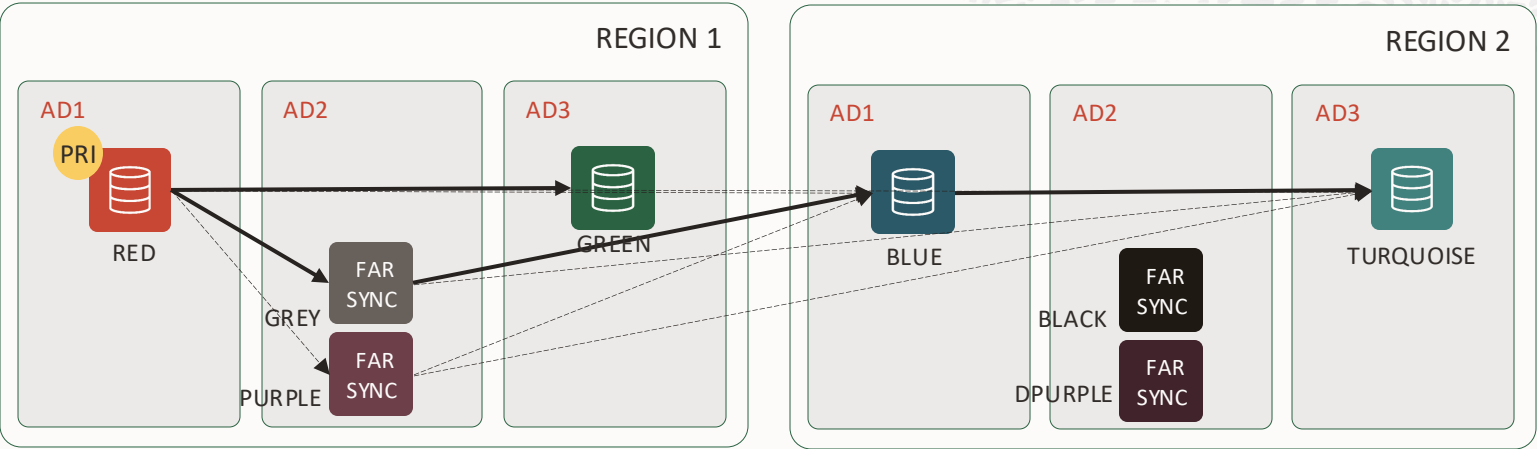


- Members:
- prod** - Primary database
 - site1** - Physical standby database
 - site101 - Physical standby database (receiving current redo)
 - site102 - Physical standby database (receiving current redo)
 - site103 - Physical standby database (receiving current redo)
 - site104 - Physical standby database (receiving current redo)
 - site2** - Physical standby database
 - site201 - Physical standby database (receiving current redo)
 - site202 - Physical standby database (receiving current redo)
 - site203 - Physical standby database (receiving current redo)
 - site204 - Physical standby database (receiving current redo)
 - site3** - Physical standby database
 - site301 - Physical standby database (receiving current redo)
 - site302 - Physical standby database (receiving current redo)
 - site303 - Physical standby database (receiving current redo)
 - site304 - Physical standby database (receiving current redo)
 - site4** - Physical standby database
 - site401 - Physical standby database (receiving current redo)
 - site402 - Physical standby database (receiving current redo)
 - site403 - Physical standby database (receiving current redo)
 - site404 - Physical standby database (receiving current redo)
 - site5** - Physical standby database
 - site501 - Physical standby database (receiving current redo)
 - site502 - Physical standby database (receiving current redo)
 - site503 - Physical standby database (receiving current redo)
 - site504 - Physical standby database (receiving current redo)
 - site6** - Physical standby database
 - site601 - Physical standby database (receiving current redo)
 - site602 - Physical standby database (receiving current redo)
 - site603 - Physical standby database (receiving current redo)
 - site604 - Physical standby database (receiving current redo)
 - site7** - Physical standby database
 - site701 - Physical standby database (receiving current redo)
 - site702 - Physical standby database (receiving current redo)
 - site703 - Physical standby database (receiving current redo)
 - site704 - Physical standby database (receiving current redo)
 - site8** - Physical standby database
 - site801 - Physical standby database (receiving current redo)
 - site802 - Physical standby database (receiving current redo)
 - site803 - Physical standby database (receiving current redo)
 - site804 - Physical standby database (receiving current redo)



Real Life Use Case #3

Highly available cloud blueprint for multi-AD regions

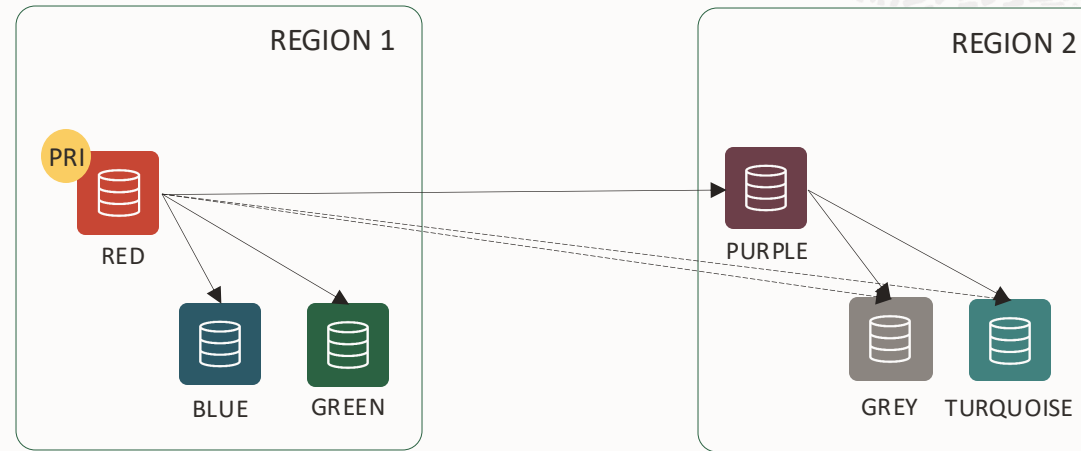


```
DGMGRL> show configuration when primary is RED;
Configuration - HADB
Protection Mode: MaxAvailability
Members:
RED      - Primary database
GREEN    - Physical standby database
GREY     - Far sync instance
BLUE     - Physical standby database
          TURQUOISE - Physical standby database (receiving current redo)
          TURQUOISE - Physical standby database (alternate of BLUE)
PURPLE   - Far sync instance (alternate of GREY)
BLUE     - Physical standby database
          TURQUOISE - Physical standby database (receiving current redo)
          TURQUOISE - Physical standby database (alternate of BLUE)
BLUE     - Physical standby database (alternate of GREY)
          TURQUOISE - Physical standby database (receiving current redo)
          TURQUOISE - Physical standby database (alternate of BLUE)
Members Not Receiving Redo:
BLACK    - Far sync instance
DPURPLE  - Far sync instance
```



Real Life Use Case #4

Two local Active Data Guard standbys and a symmetric region for DR



```
DGMGRL> show configuration
```

Configuration - HADB

Protection Mode: MaxPerformance

Members:

RED - Primary database

GREEN - (*) Physical standby database

BLUE - Physical standby database

PURPLE - Physical standby database


TURQUOISE - Physical standby database (receiving current redo)

GREY - Physical standby database (receiving current redo)

Fast-Start Failover: Enabled in Potential Data Loss Mode

Configuration Status:

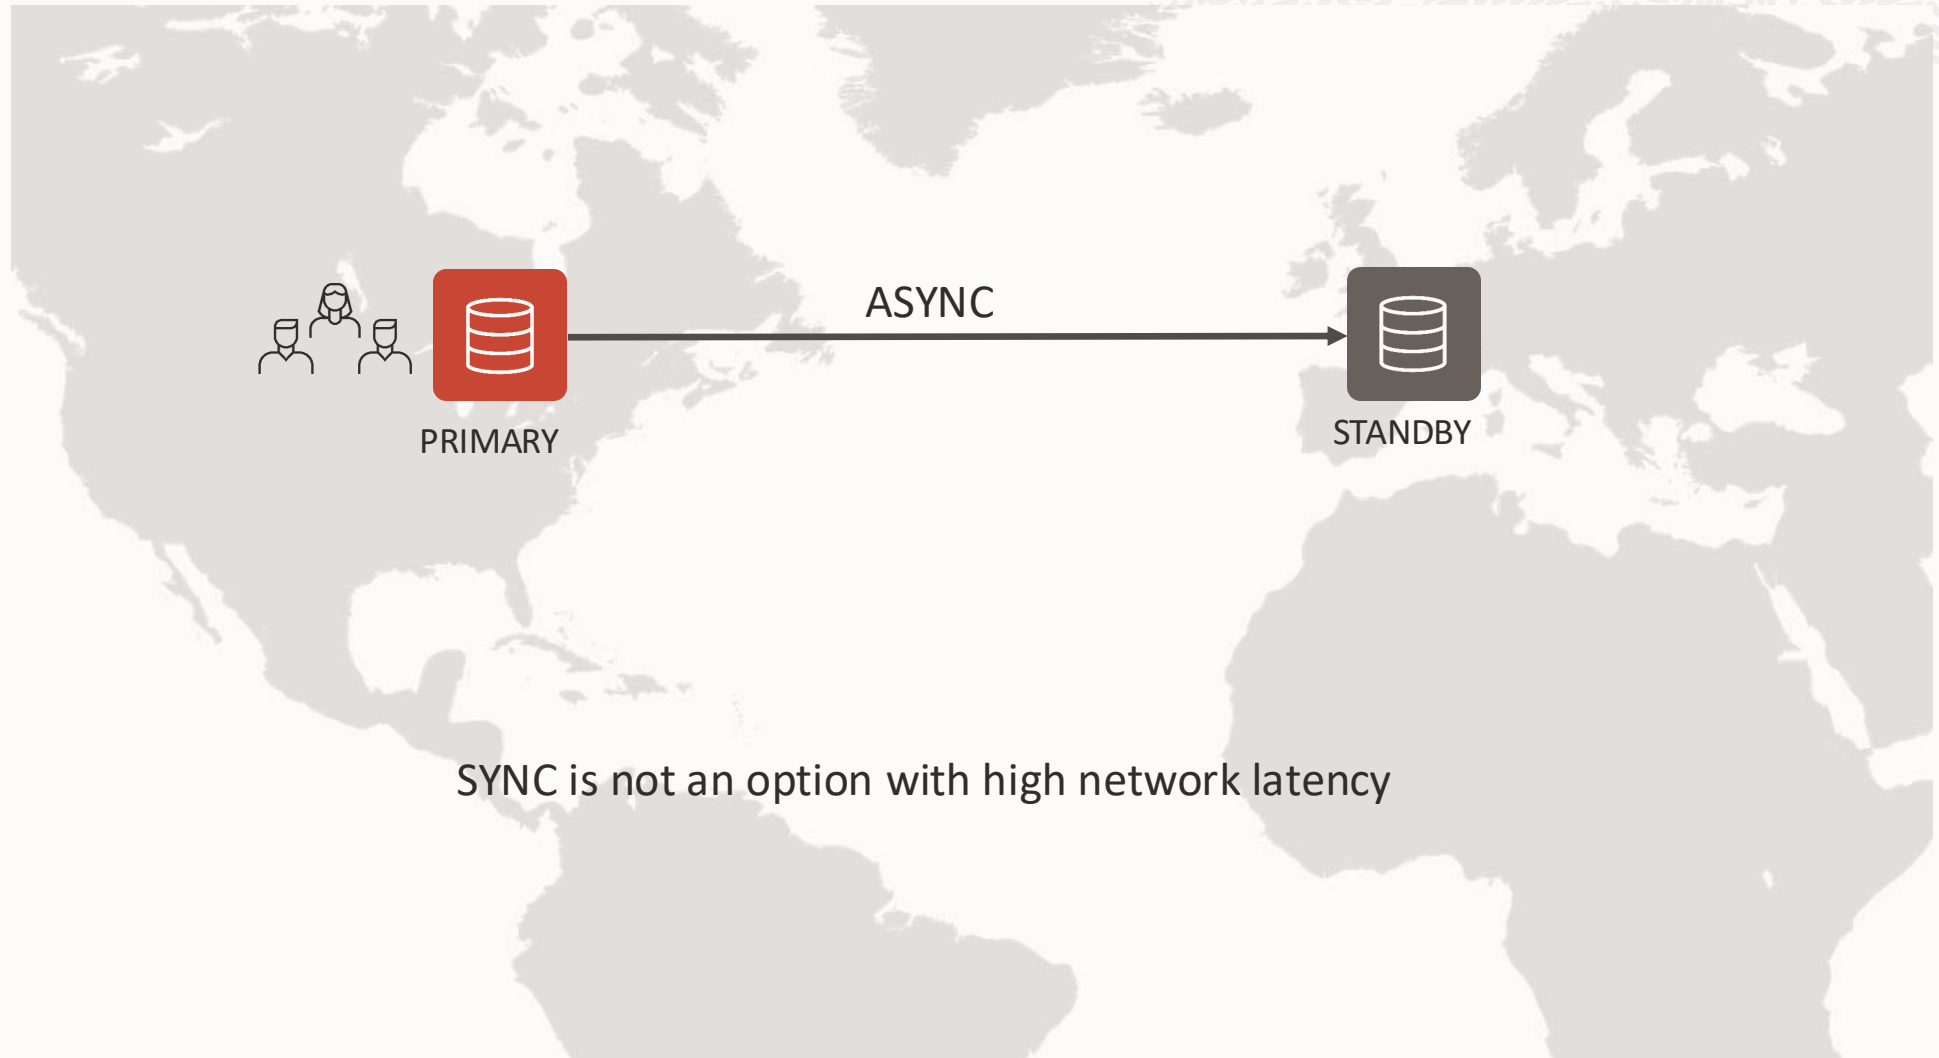
SUCCESS (status updated 29 seconds ago)



Oracle Active Data Guard Far Sync

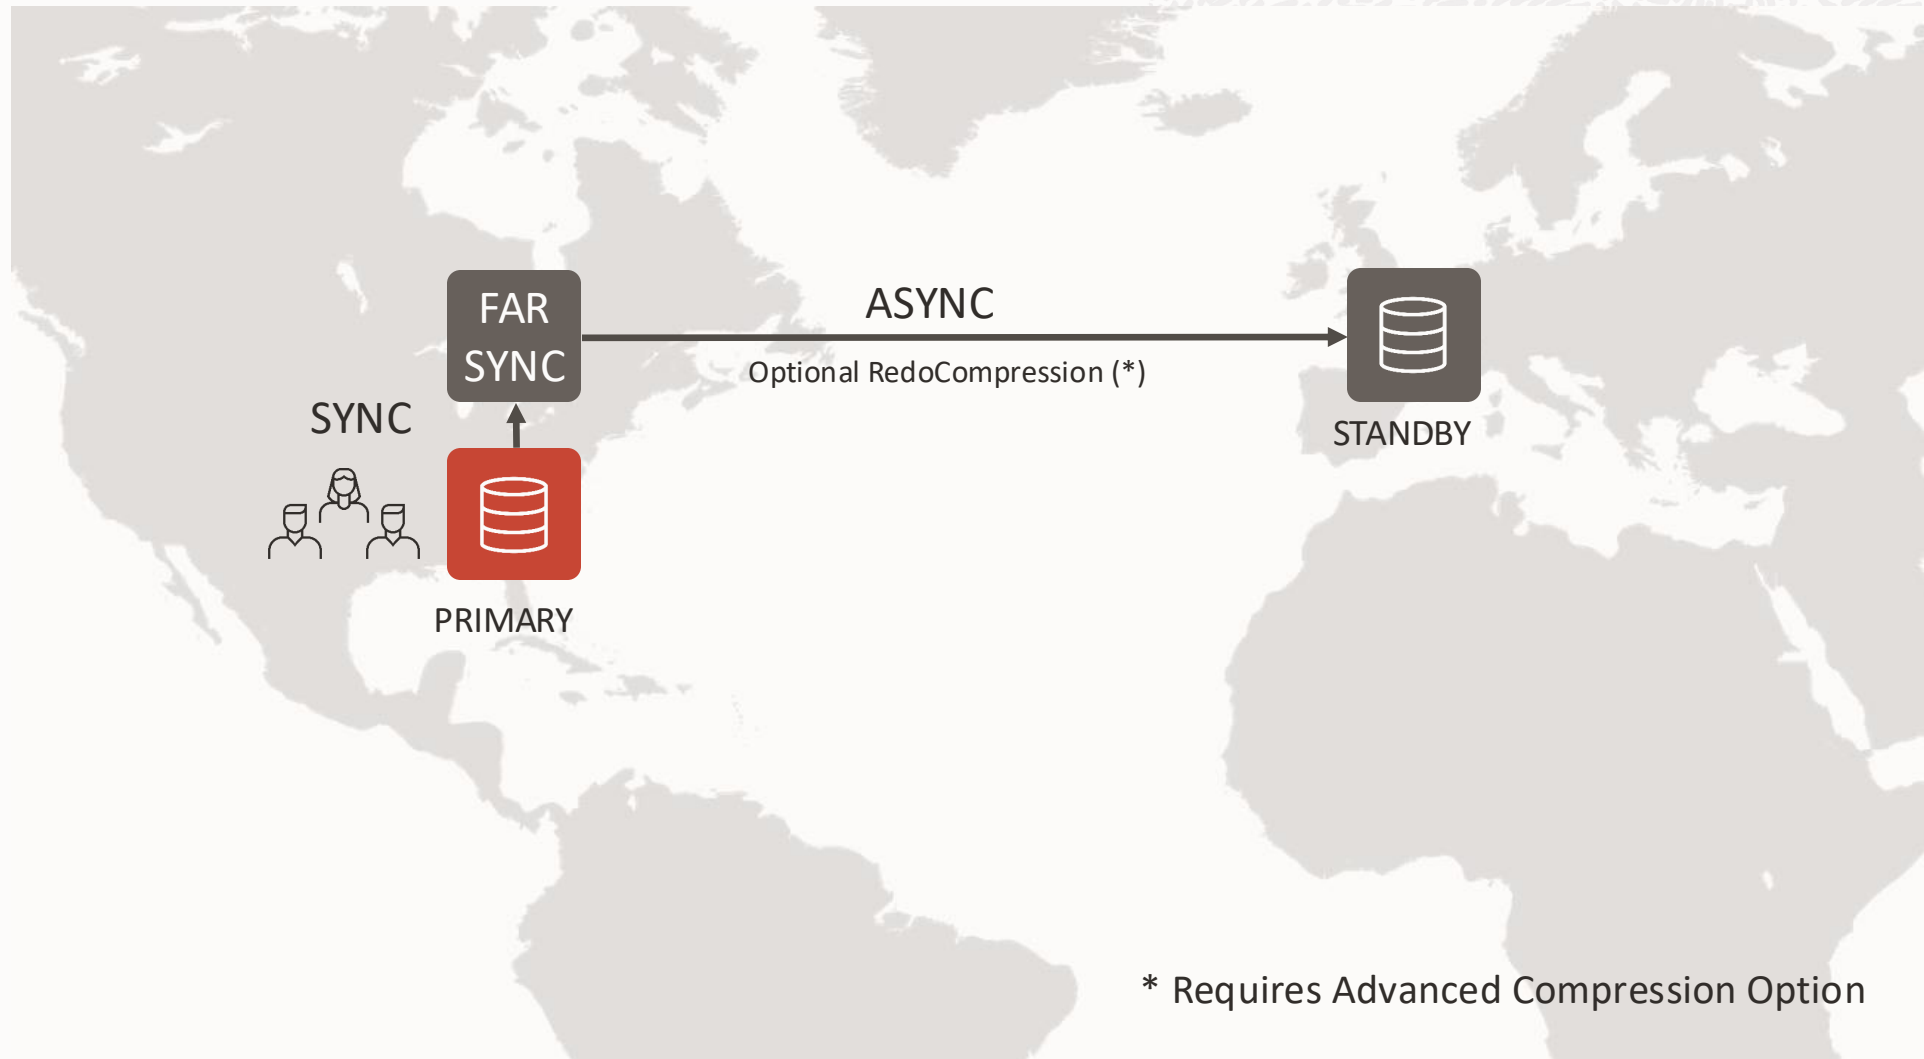
The Zero Data Loss Challenge

Trade-off Performance for Protection



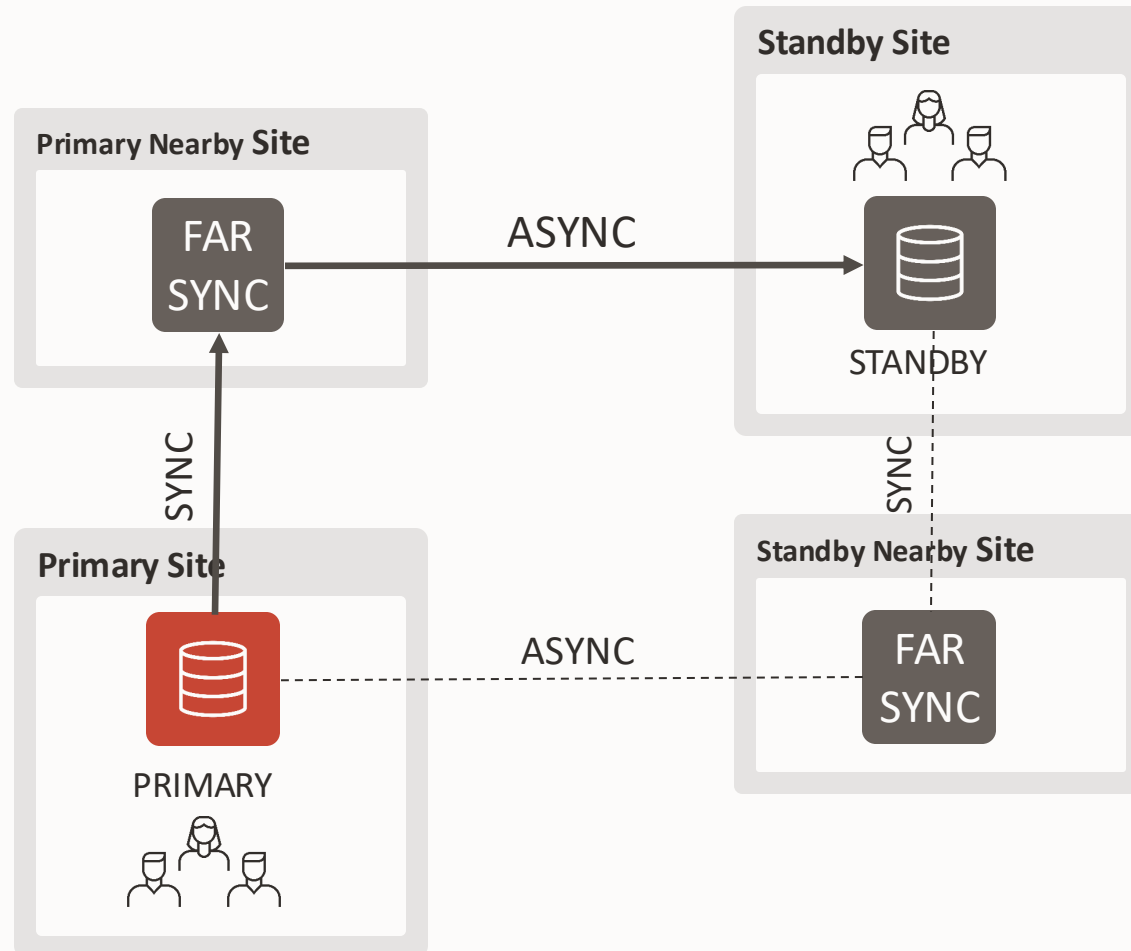
Active Data Guard Far Sync

Zero Data Loss Protection at Any Distance



Active Data Guard Far Sync

Do not Trade-off Protection for Performance



Far Sync

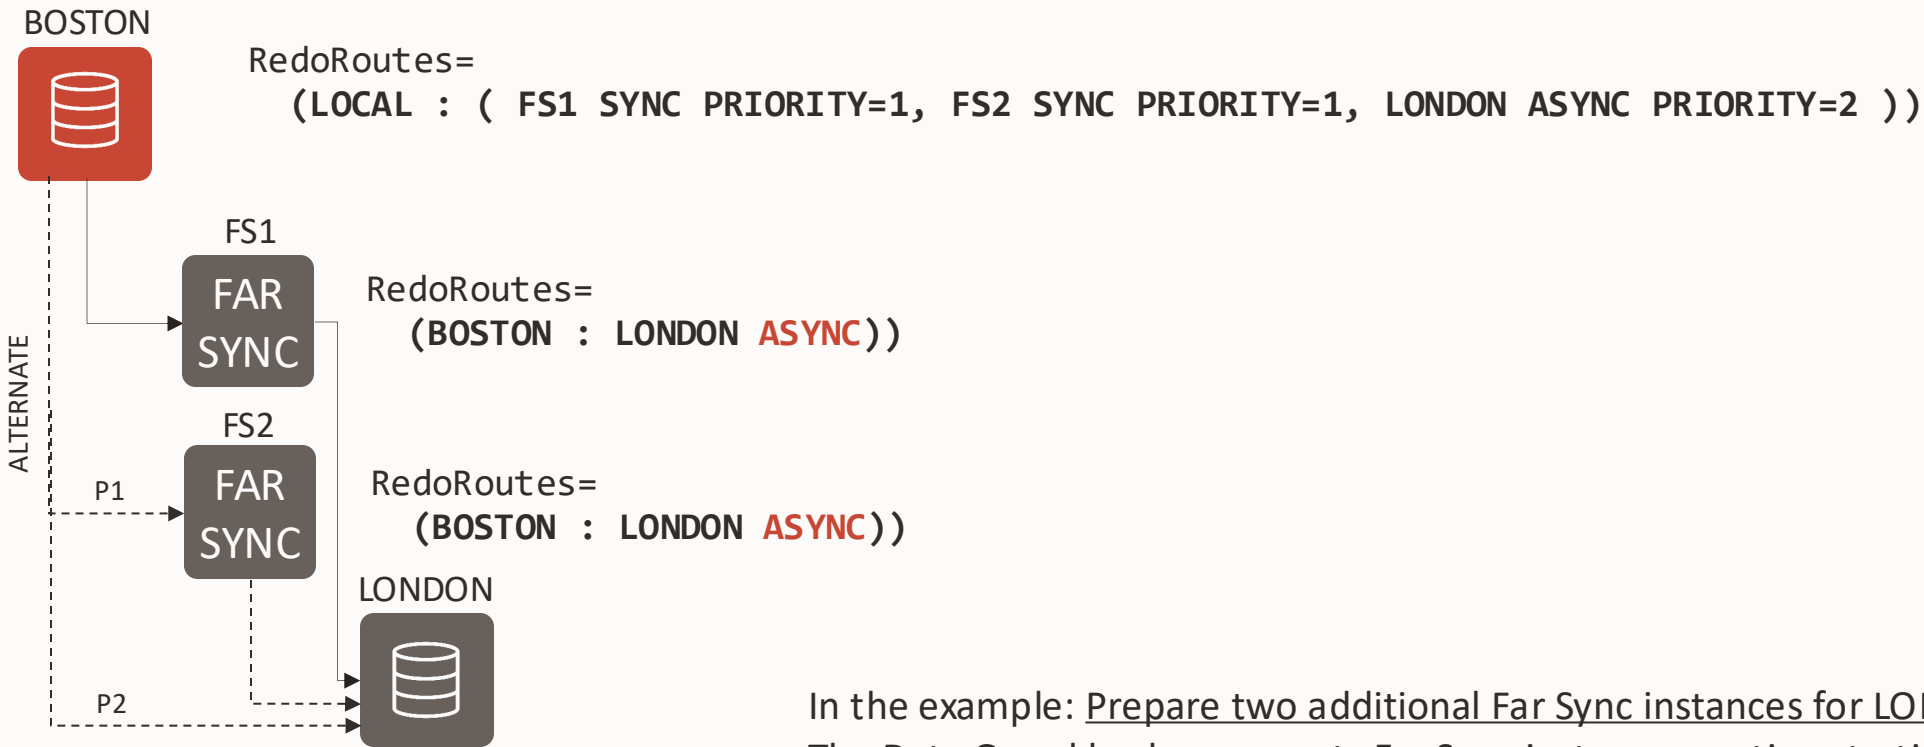
- Special instance:
 - No datafiles
 - No Media Recovery
 - Only control files, archives and standby logs
- Up to 30 direct destinations
- Offload transport compression (Advanced Compression)
- Supports FSFO in MaxAvailability
- Supports FSFO in MaxPerformance (**new in 21c**)

Use different Datacenters or Availability Domains!

- Upon failover, the standby will fetch the very last redo from the Far Sync

Active Data Guard Far Sync

Use RedoRoutes for Far Sync High Availability



In the example: Prepare two additional Far Sync instances for LONDON!
The Data Guard broker supports Far Sync instance creation starting with 21c.



Benefits and Downsides of Far Sync

When to consider Far Sync?

Benefits

- Increased **performance** for existing Sync configurations
- Increased **protection** for existing Async configurations
- Zero Data Loss (Max Availability) across distant regions
- Fully integrated with the broker
- Automatic gap resolution through the Far Sync

Downsides

- Additional server(s) or VM(s) and components
- A Far Sync co-located with the primary might not prevent data loss in case of full site failure



Far Sync and Fast Start Failover

Which Fast Start Failover protection modes are compatible with Far Sync?

FSFO and FAR SYNC	Maximum Performance	Maximum Availability	Maximum Protection
ASync	✓ (21c+)	✗	✗
FAST SYNC	✗	✓	✗
SYNC	✗	✓	✗

FSFO without FAR SYNC	Maximum Performance	Maximum Availability	Maximum Protection
ASync	✓	✗	✗
FAST SYNC	✗	✓	✗
SYNC	✗	✓	✓


FAR SYNC without FSFO	Maximum Performance	Maximum Availability	Maximum Protection
ASync	✓	✗	✗
FAST SYNC	✗	✓	✗
SYNC	✗	✓	✗

<https://docs.oracle.com/en/database/oracle/oracle-database/21/dgbrk/using-data-guard-broker-to-manage-switchovers-failovers.html#GUID-7423C774-27DF-49F9-BB43-7D547BCE7762>



Data Guard Broker Far Sync Instance Creation


One step further automated by the broker



```
DGMGRL> CREATE FAR_SYNC bostonfs
```

```
AS CONNECT IDENTIFIER IS "bostonfs_conn_str"  
PARAMETER_VALUE_CONVERT "boston","bostonfs"  
SET LOG_FILE_NAME_CONVERT "boston","bostonfs"  
SET DB_RECOVERY_FILE_DEST "$ORACLE_HOME/dbs/"  
SET DB_RECOVERY_FILE_DEST_SIZE "100G"  
RESET UNDO_TABLESPACE;
```

- Automated SPFILE and controlfile creation
- The Far Sync is created, started and added to the configuration



Oracle Active Data Guard Rolling Maintenance and Upgrades

Solutions for Database Rolling Maintenance and Upgrades

Manual

Part of Enterprise Edition

Source >= 11.1.0.7 && <= 12.1.0.2

Manual approach

Limited feature support

DBMS_ROLLING

Requires Active Data Guard

Source >= 12.1.0.2

Automated

Comprehensive feature support

GoldenGate

Requires GoldenGate

Source >= 11.2.0.4 (for OCI GG)

Manual approach

Best feature support

Fallback mechanism

Using SQL Apply to Upgrade the Oracle Database

<https://docs.oracle.com/en/database/oracle/oracle-database/19/sbydb/using-sql-apply-to-perform-rolling-upgrade.html>

Using DBMS_ROLLING to Perform a Rolling Upgrade

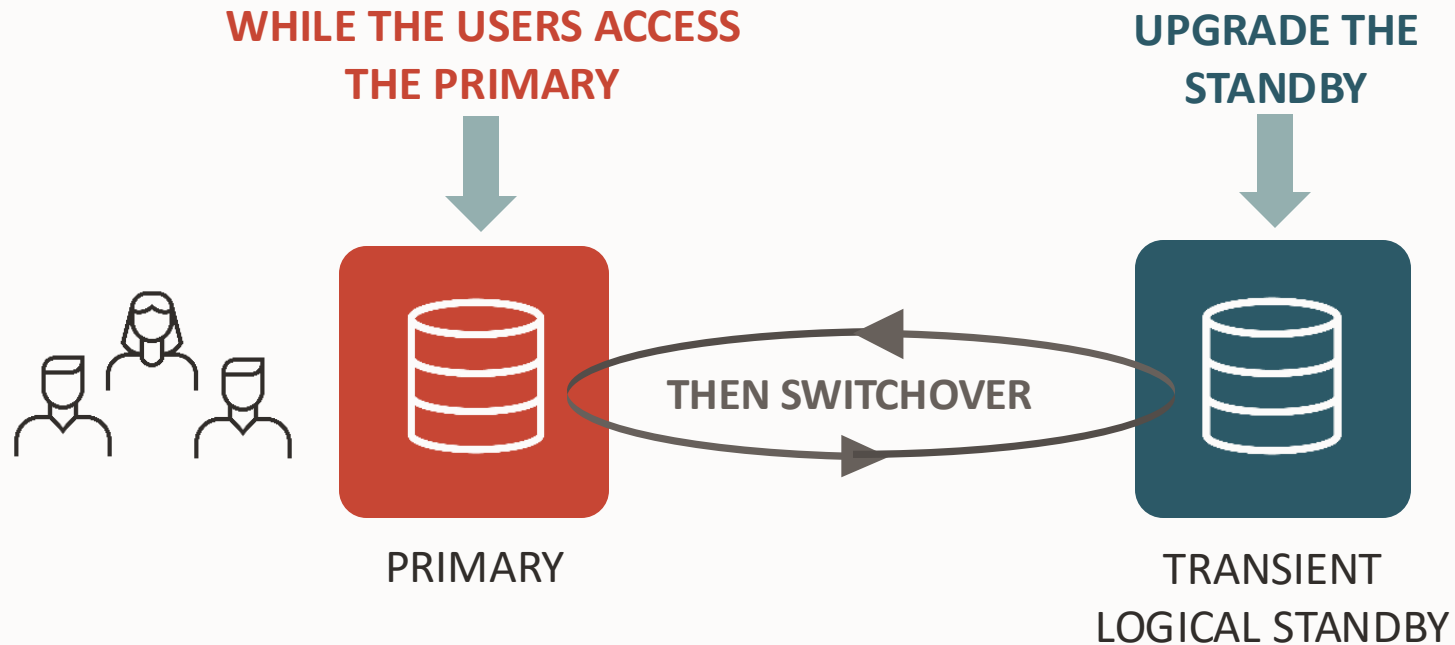
https://docs.oracle.com/en/database/oracle/oracle-database/19/sbydb/using-DBMS_ROLLING-to-perform-rolling-upgrade.html

Overview of Steps for Upgrading Oracle Database Using Oracle GoldenGate

<https://docs.oracle.com/en/database/oracle/oracle-database/19/upgrd/converting-databases-upgrades.html#GUID-8E029631-8265-497C-983B-B8A4ACD47B98>

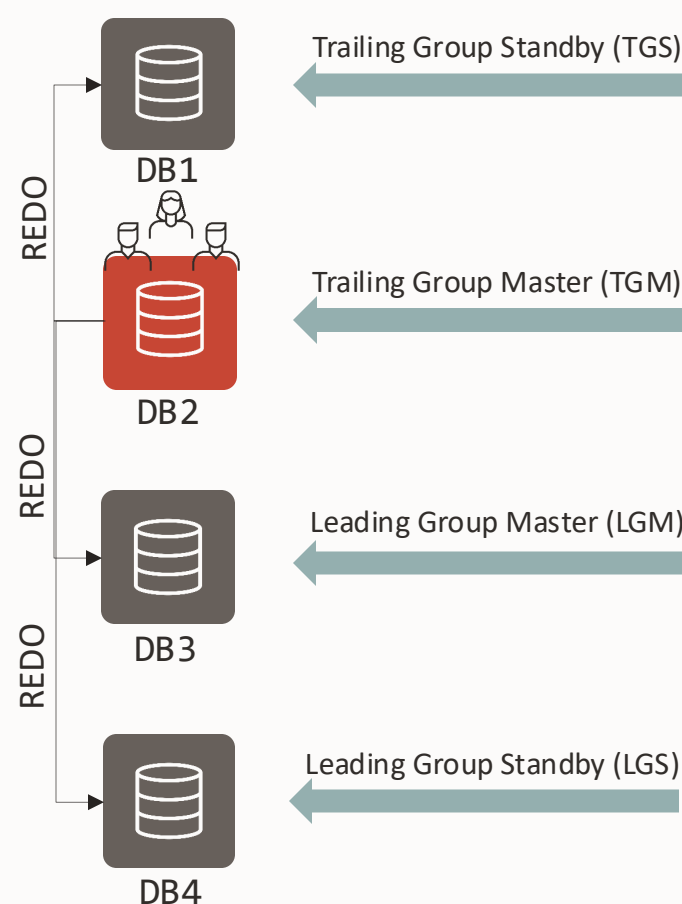
Active Data Guard Rolling Maintenance and Upgrades

Using `DBMS_ROLLING` package



- Use a transient logical standby database to upgrade with very little downtime.
- The only downtime is as little as it takes to perform a switchover.

The DBMS_ROLLING.INIT_PLAN phase



Trailing Group

```
-- check DBA_ROLLING_UNSUPPORTED for incompatible data types
-- initialize the plan and set the future primary
DBMS_ROLLING.INIT_PLAN(future_primary=>'DB3');

-- add the required standbys to the TRAILING GROUP
DBMS_ROLLING.SET_PARAMETER('DB1','MEMBER','TRAILING');

-- add the required standbys to the LEADING GROUP
DBMS_ROLLING.SET_PARAMETER('DB4','MEMBER','LEADING');
```

Leading Group



The DBMS_ROLLING parameters



ACTIVE_SESSIONS_TIMEOUT
ACTIVE_SESSIONS_WAIT
BACKUP_CONTROLFILE
DGBROKER
DICTIONARY_LOAD_TIMEOUT
DICTIONARY_LOAD_WAIT
DICTIONARY_PLS_WAIT_INIT
DICTIONARY_PLS_WAIT_TIMEOUT
EVENT_RECORDS
FAILOVER
GRP_PREFIX
IGNORE_BUILD_WARNINGS
IGNORE_LAST_ERROR
LAD_ENABLED_TIMEOUT
LOG_LEVEL

MEMBER
READY_LGM_LAG_TIME
READY_LGM_LAG_TIMEOUT
READY_LGM_LAG_WAIT
SWITCH_LGM_LAG_TIME
SWITCH_LGM_LAG_TIMEOUT
SWITCH_LGM_LAG_WAIT
SWITCH_LGS_LAG_TIME
SWITCH_LGS_LAG_TIMEOUT
SWITCH_LGS_LAG_WAIT
UPDATED_LGS_TIMEOUT
UPDATED_LGS_WAIT
UPDATED_TGS_TIMEOUT
UPDATED_TGS_WAIT

The DBMS_ROLLING parameters



Example:

```
-- Activate full logging
exec DBMS_ROLLING.SET_PARAMETER (scope=>null, name=>'LOG_LEVEL', value=>'FULL');

-- Wait for the SQL Apply Lag to go below 1 minute before initiating the switchover
exec DBMS_ROLLING.SET_PARAMETER('SWITCH_LGM_LAG_WAIT', '1');
exec DBMS_ROLLING.SET_PARAMETER('SWITCH_LGM_LAG_TIME', '60');
```


Final touches before starting




```
$ # The standby must be mounted  
$ srvctl stop database -d DB3  
$ srvctl start database -d DB3 -o mount
```

```
SQL> -- The PDBs must be open  
SQL> alter pluggable database all open;
```

```
DGMGRL> # no FSFO or MaxProtection  
DGMGRL> disable fast_start failover  
DGMGRL> edit configuration set protection mode as MaxAvailability;
```

The DBMS_ROLLING.BUILD_PLAN phase



```

-- build the plan
DBMS_ROLLING.BUILD_PLAN();

-- check for any errors or warnings
SELECT * FROM DBA_ROLLING_EVENTS;

-- review the plan
SELECT * FROM DBA_ROLLING_PLAN ORDER BY INSTID;
  
```




The DBMS_ROLLING.BUILD_PLAN phase

```
1 START    Notify Data Guard broker that DBMS_ROLLING has started
2 START    Notify Data Guard broker that DBMS_ROLLING has started
3 START    Verify database is a primary
4 START    Verify MAXIMUM PROTECTION is disabled
5 START    Verify database is a physical standby
6 START    Verify physical standby is mounted
7 START    Verify future primary is configured with standby redo logs
8 START    Verify server parameter file exists and is modifiable
9 START    Verify server parameter file exists and is modifiable
10 START   Verify Data Guard broker configuration is enabled
11 START   Verify Data Guard broker configuration is enabled
12 START   Verify Fast-Start Failover is disabled
13 START   Verify Fast-Start Failover is disabled
14 START   Verify fast recovery area is configured
15 START   Verify available flashback restore points
16 START   Verify fast recovery area is configured
17 START   Verify available flashback restore points
18 START   Stop media recovery
19 START   Drop guaranteed restore point DBMSRU_INITIAL
20 START   Create guaranteed restore point DBMSRU_INITIAL
21 START   Drop guaranteed restore point DBMSRU_INITIAL
22 START   Create guaranteed restore point DBMSRU_INITIAL
23 START   Start media recovery
24 START   Verify media recovery is running
25 START   Verify user_dump_dest has been specified
26 START   Backup control file to rolling_change_backup.f
27 START   Verify user_dump_dest has been specified
28 START   Backup control file to rolling_change_backup.f
29 START   Get current supplemental logging on the primary database
30 START   Get current redo branch of the primary database
31 START   Wait until recovery is active on the primary's redo branch
32 START   Reduce to a single instance if database is a RAC
33 START   Verify only a single instance is active if future primary is RAC
34 START   Stop media recovery
35 START   Execute dbms_logstdby.build
36 START   Convert into a transient logical standby
37 START   Open database including instance-peers if RAC
38 START   Verify logical standby is open read/write
39 START   Get redo branch of transient logical standby
40 START   Get reset scn of transient logical redo branch
41 START   Configure logical standby parameters
42 START   Start logical standby apply
43 START   Enable compatibility advance despite presence of GRPs

44 START   Log pre-switchover instructions to events table
45 START   Record start of user upgrade of DB3
46 SWITCH  Verify database is in OPENRW mode
47 SWITCH  Record completion of user upgrade of DB3
48 SWITCH  Scan LADs for presence of DB2 destination
49 SWITCH  Test if DB2 is reachable using configured TNS service
50 SWITCH  Call Data Guard broker to enable redo transport to DB3
51 SWITCH  Archive all current online redo logs
52 SWITCH  Archive all current online redo logs
53 SWITCH  Stop logical standby apply
54 SWITCH  Start logical standby apply
55 SWITCH  Wait until apply lag has fallen below 600 seconds
56 SWITCH  Notify Data Guard broker that switchover to logical standby database is starting
57 SWITCH  Log post-switchover instructions to events table
58 SWITCH  Switch database to a logical standby
59 SWITCH  Notify Data Guard broker that switchover to logical standby database has completed
60 SWITCH  Wait until end-of-redo has been applied
61 SWITCH  Archive all current online redo logs
62 SWITCH  Notify Data Guard broker that switchover to primary is starting
63 SWITCH  Switch database to a primary
64 SWITCH  Notify Data Guard broker that switchover to primary has completed
65 SWITCH  Enable compatibility advance despite presence of GRPs
66 SWITCH  Synchronize plan with new primary
67 FINISH  Reduce to a single instance for FINISH
68 FINISH  Verify only a single instance is active
69 FINISH  Verify database is mounted
70 FINISH  Flashback database
71 FINISH  Convert into a physical standby
72 FINISH  Verify database is open
73 FINISH  Save the DBID of the new primary
74 FINISH  Save the logminer session start scn
75 FINISH  Wait until transient logical redo branch has been registered
76 FINISH  Start media recovery
77 FINISH  Wait until apply/recovery has started on the transient branch
78 FINISH  Wait until upgrade redo has been fully recovered
79 FINISH  Prevent compatibility advance if GRPs are present
80 FINISH  Prevent compatibility advance if GRPs are present
81 FINISH  Drop guaranteed restore point DBMSRU_INITIAL
82 FINISH  Drop guaranteed restore point DBMSRU_INITIAL
83 FINISH  Purge logical standby metadata from database if necessary
84 FINISH  Notify Data Guard broker that DBMS_ROLLING has finished
85 FINISH  Notify Data Guard broker that DBMS_ROLLING has finished
86 FINISH  Restore Supplemental Logging
```





The DBMS_ROLLING.START phase

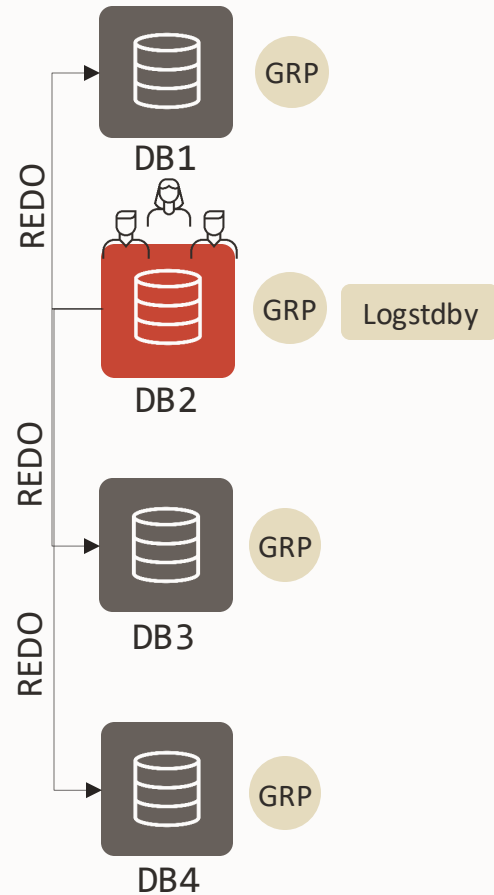
 User sessions

 +1 Upgraded

 Primary

 Physical Standby


 Logical Standby



```
-- start the plan  
DBMS_ROLLING.START_PLAN();
```

- Creates the Guaranteed Restore Point (GRP)
- Builds the logical standby metadata (`dbms_logstdby.build`)

The DBMS_ROLLING.START phase




```
-- start the plan
DBMS_ROLLING.START_PLAN();
```

- Creates the Guaranteed Restore Point (GRP)
- Builds the LogMiner directory (dbms_logstdby.build)
- Converts the LGM to Logical Standby
- Starts SQL Apply
- With a configuration composed of 4 databases, the LGM and TGM are still protected by a physical standby



The DBMS_ROLLING.START phase



```
DGMGRL> show configuration;
```

```
Configuration - geneva
```

```
Protection Mode: MaxAvailability
```

```
Members:
```

```
DB1 - Primary database
```

```
DB3 - Physical standby database
```

```
Warning: ORA-16854: apply lag could not be
determined
```


```
Fast-Start Failover: DISABLED
```

```
Configuration Status:
```

```
ROLLING DATABASE MAINTENANCE IN PROGRESS
```



The DBMS_ROLLING.START phase



```
DGMGRL> show database DB3
```

```
...
```

```
Role: PHYSICAL STANDBY
```

```
Intended State: APPLY-ON
```

```
Transport Lag: 0 seconds (computed 0 seconds ago)
```

```
Apply Lag: 3 minutes 18 seconds (computed 0 seconds ago)
```

```
...
```

```
Database Warning(s):
```


```
ORA-16866: database converted to transient logical standby database for rolling database maintenance
```

```
Database Status:
```

```
WARNING
```



The DBMS_ROLLING.START phase



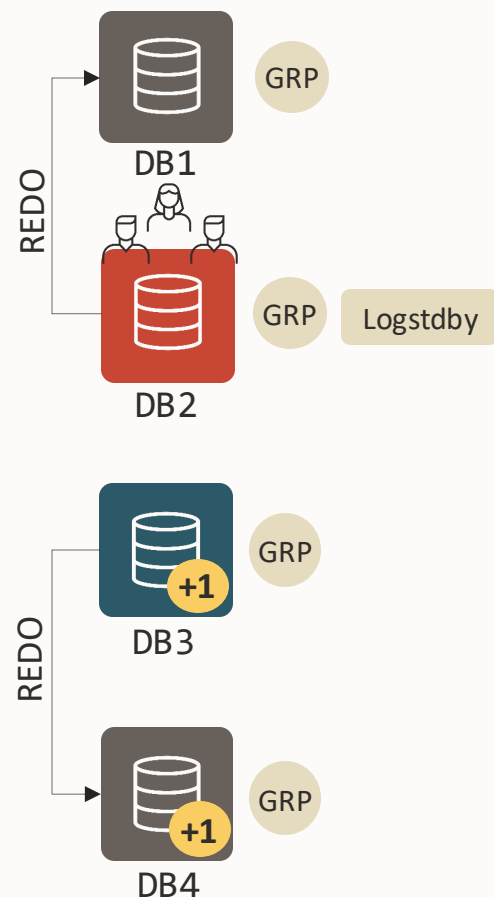
```
-- check the status of the SQL apply:
SQL> select * from V$LOGSTDBY_PROGRESS;
```

```
-- use SQL apply commands if you need
SQL> alter database start logical standby apply immediate;
```

```
-- check for logical standby error messages
SQL> select * from DBA_LOGSTDBY_EVENTS
2>      order by event_timestamp;
```

```
22-NOV-21 06.41.12  DML on "AUDSYS"."AUD$UNIFIED"
ORA-16129: unsupported DML encountered
22-NOV-21 06.41.13  truncate table wri$_adv_addm_pdb$
ORA-16247: DDL skipped on internal schema
```

The Upgrade/Maintenance phase




- Do the maintenance on the Leading Group Master

```
-- e.g. upgrade to a major version with AutoUpgrade
$ java -jar autoupgrade.jar -config CDB1.cfg -mode deploy
```

- This is out of DBMS_ROLLING scope (it is a manual step)
- Don't forget to align the Leading Group Standbys if necessary
- Use it for any major maintenance that requires longer downtimes (change of physical layout, structure changes, offline operations)

The DBMS_ROLLING.SWITCHOVER phase





```
-- switchover to the upgraded database
DBMS_ROLLING.SWITCHOVER()
```

- Depending on the source version and HA configuration, the old connections get FAN notifications and drain automatically
- New connections go to the new primary. Application downtime is minimal.



The DBMS_ROLLING.SWITCHOVER phase

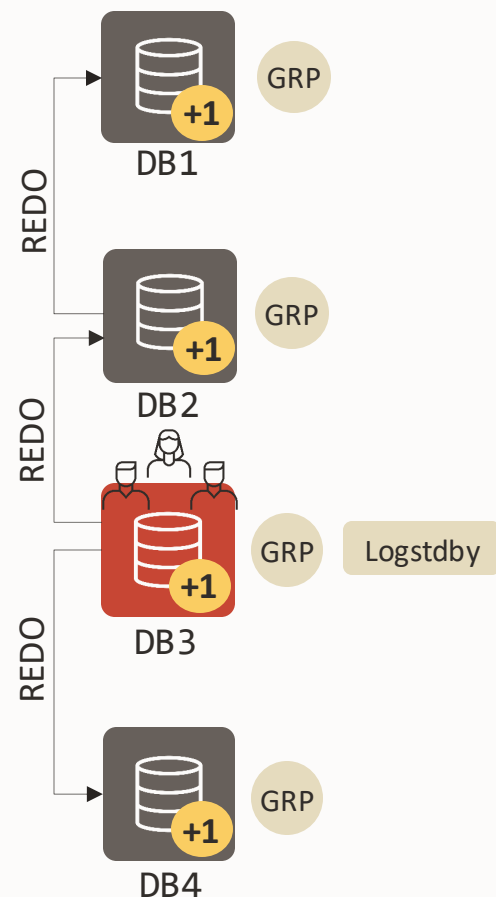


- Start the Trailing Group members with the new binaries (manual)

```
-- run the final part of the plan  
DBMS_ROLLING.FINISH_PLAN()
```

- Flashes back the Trailing Group Master and Standby to the GRP

The DBMS_ROLLING.SWITCHOVER phase




- Start the Trailing Group members with the new binaries (manual)

```
-- run the final part of the plan
DBMS_ROLLING.FINISH_PLAN()
```


- Flashes back the Trailing Group Master and Standby to the GRP
- Converts the Trailing Group Master to a physical standby
- Starts redo apply and catches up with the primary
- Drops the guaranteed restore points and logical standby metadata


The DBMS_ROLLING.SWITCHOVER phase


 User sessions

 +1 Upgraded

 Primary

 Physical Standby

 Logical Standby



```
-- destroy the plan to clean up everything  
DBMS_ROLLING.DESTROY_PLAN()
```

DBMS_ROLLING catalog views



Evaluate	DBA_ROLLING_UNSUPPORTED
Initialize	DBA_ROLLING_PARAMETERS
Build	DBA_ROLLING_DATABASES
	DBA_ROLLING_PLAN
Monitor	DBA_ROLLING_EVENTS
	DBA_ROLLING_STATISTICS
	DBA_ROLLING_STATUS

Check here for unsupported data types!


Get the current parameters before building

Verify the plan before and during the execution

Warning and errors are visible here



DBMS_ROLLING points of attention



Do not create the logical standby on the **same** server as the primary database



Supplemental logging is enabled automatically which introduces an overhead and increases the amount of redo generated



When supplemental logging is enabled all DML cursors are invalidated



Not all data types and partitioning types are supported



For optimal performance all tables should have primary keys or unique keys

Important DBMS_ROLLING milestones

The driver is the SOURCE database!



SOURCE VERSION

12.1

- First version of DBMS_ROLLING for upgrades from 12.1 to higher versions

12.2

- Integration with the **Data Guard broker**
- **FAN events** for Clusterware-backed databases
- Support for **Identity columns**

19c

- Planned in future RU: Support for **Application Continuity** and **Transparent Application Continuity** (backport from 23ai)

21c

- **FAN events without Clusterware**
- Support for **JSON datatype**

23ai

- Support for **Application Continuity** and **Transparent Application Continuity**
- Support for Blockchain tables
- Support for new Boolean data type
- Support for SQL domains

DBMS_ROLLING and client failover

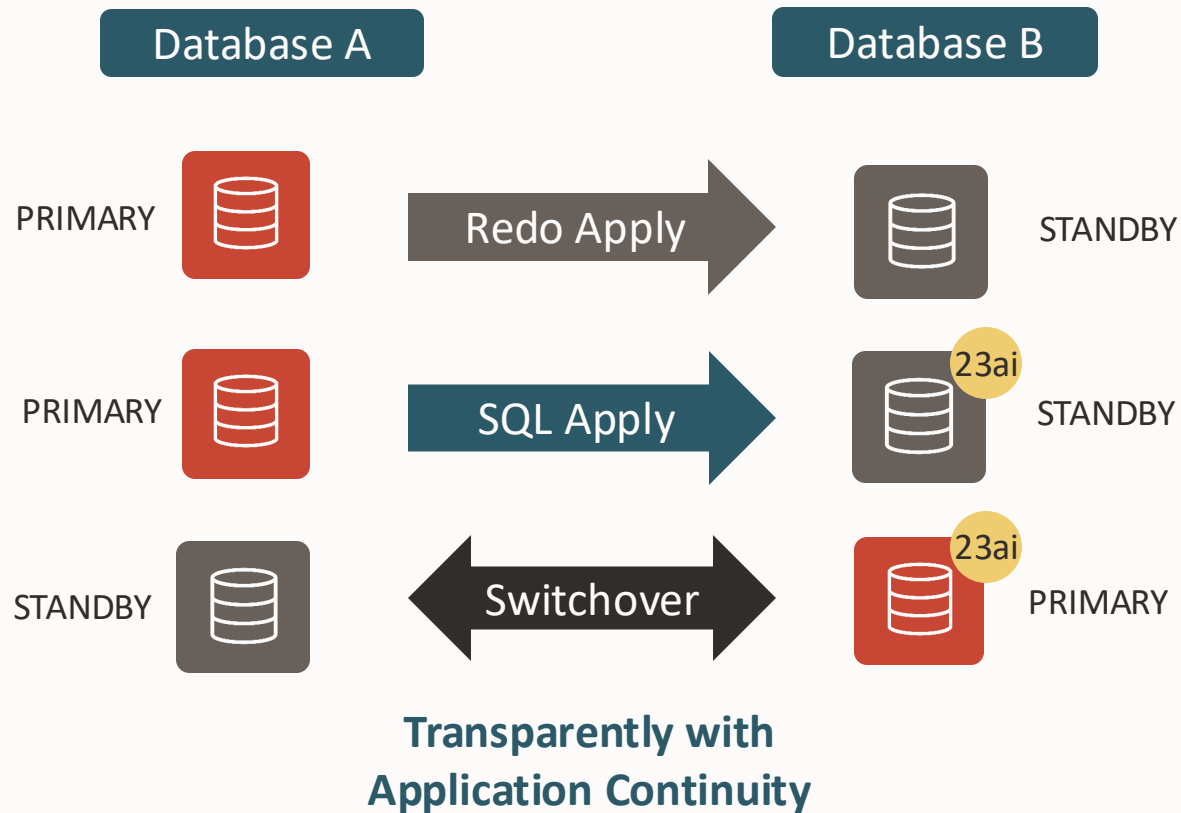


DBMS_ROLLING.SWITCHOVER	Broker + OCW	Broker Only
12.1	Broker Not supported	Broker Not supported
12.2	FAN events	No FAN events
19c	FAN events (AC/TAC backport planned)	No FAN events
21c	FAN events	FAN events
23ai	FAN events + AC/TAC	FAN events + AC/TAC



Zero Application Downtime for Database Release Upgrades

Minimizes application impact throughout the entire database upgrade process



(Transparent) Application Continuity

- Hides database downtime from your users
 - It rebuilds the session state
 - It replays in-flight transactions

DBMS_ROLLING

- Enables the automated rolling application of version-changing upgrades and patch sets.

Together they hide the final switchover needed at the end of the automated process.

DBMS_ROLLING – Read More



Using DBMS_ROLLING to Perform a Rolling Upgrade

https://docs.oracle.com/en/database/oracle/oracle-database/19/sbydb/using-DBMS_ROLLING-to-perform-rolling-upgrade.html

DBMS_ROLLING - PL/SQL Packages and Types Reference

https://docs.oracle.com/en/database/oracle/oracle-database/19/arpls/DBMS_ROLLING.html#GUID-097F1B39-E623-43B5-BA30-DF377BFE05CF

Automated Database Upgrades using Oracle Active Data Guard and DBMS_ROLLING

<https://www.oracle.com/technetwork/database/availability/database-upgrade-dbms-rolling-4126957.pdf>

Oracle Database Rolling Upgrades (without DBMS_ROLLING)

<https://www.oracle.com/technetwork/database/availability/database-rolling-upgrade-3206539.pdf>

DBMS_ROLLING – Read More



MOS Notes:

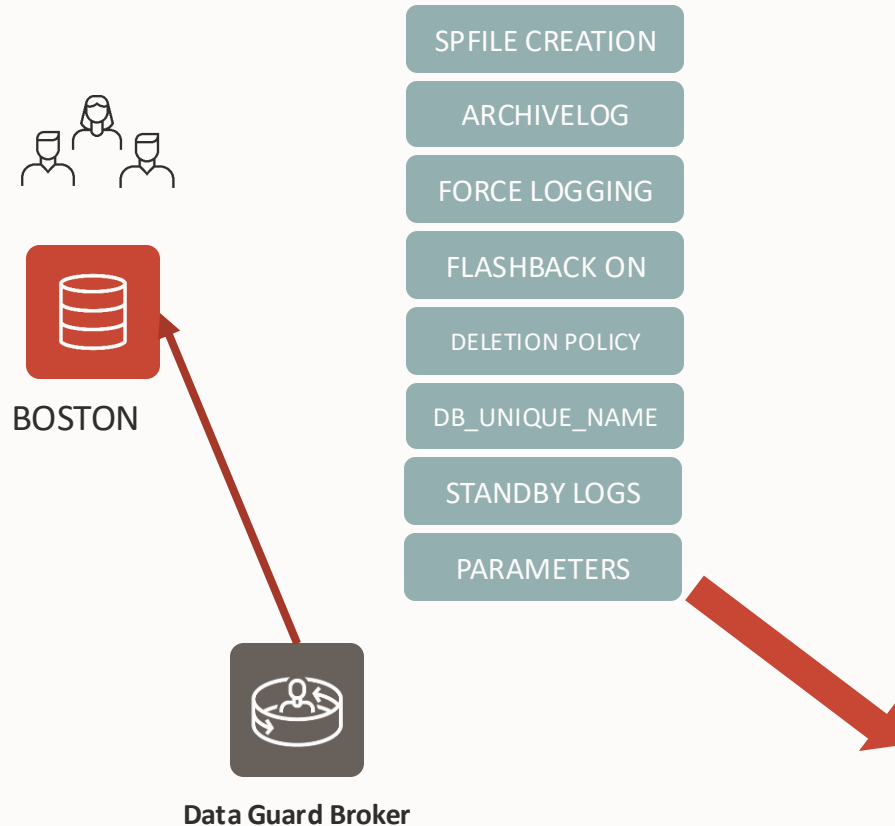
- [Transient Rolling Upgrade Using DBMS_ROLLING - Beginners Guide](#)
- [Rolling upgrade using DBMS_ROLLING - Complete Reference \(Doc ID 2086512.1\)](#)
- [MAA Whitepaper: SQL Apply Best Practices \(Doc ID 1672310.1\)](#)
- [Step by Step How to Do Switchover/Failover on Logical Standby Environment \(Doc ID 2535950.1\)](#)
- [How To Skip A Complete Schema From Application on Logical Standby Database \(Doc ID 741325.1\)](#)
- [How to monitor the progress of the logical standby \(Doc ID 1296954.1\)](#)
- [How To Reduce The Performance Impact Of LogMiner Usage On A Production Database \(Doc ID 1629300.1\)](#)
- [Handling ORA-1403 ora-12801 on logical standby apply \(Doc ID 1178284.1\)](#)
- [Troubleshooting Example - Rolling Upgrade using DBMS_ROLLING \(Doc ID 2535940.1\)](#)
- [DBMS Rolling Upgrade Switchover Fails with ORA-45427: Logical Standby Redo Apply Process Was Not Running \(Doc ID 2696017.1\)](#)
- [SRDC - Collect Logical Standby Database Information \(Doc ID 1910065.1\)](#)
- [MRP fails with ORA-19906 after Flashback of Transient Logical Standby used for Rolling Upgrade \(Doc ID 2069325.1\)](#)
- [What Causes High Redo When Supplemental Logging is Enabled \(Doc ID 1349037.1\)](#)



Other **21c** features for Data Guard and Broker

Automatic Primary Database Preparation

Faster and easier creation of Data Guard environments



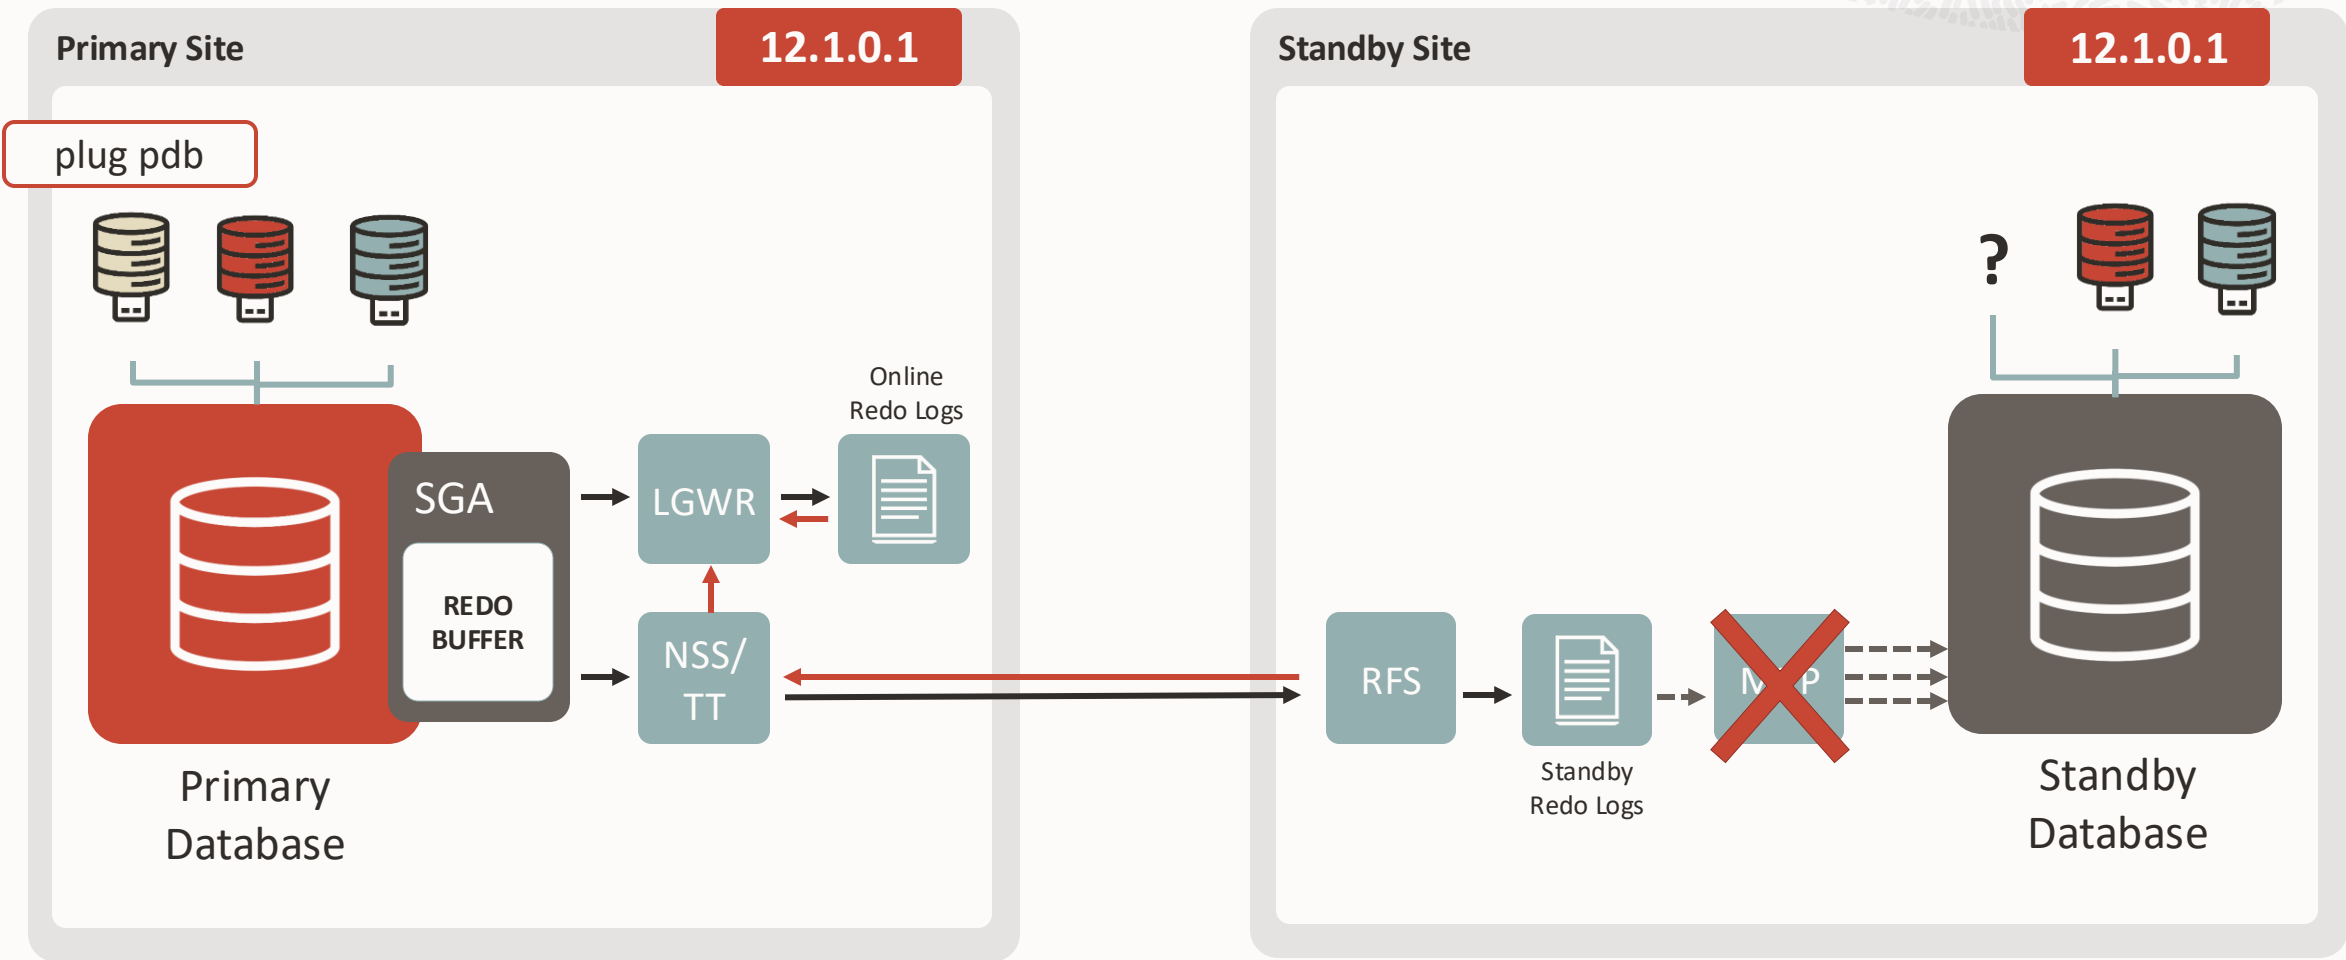
```
DGMGRL> PREPARE DATABASE FOR DATA GUARD
WITH DB_UNIQUE_NAME IS boston
DB_RECOVERY_FILE_DEST IS "+FRA"
DB_RECOVERY_FILE_DEST_SIZE IS "400G"
BROKER_CONFIG_FILE1 IS "+DATA/BOSTON/dg1.dat"
BROKER_CONFIG_FILE2 IS "+FRA/BOSTON/dg2.dat";
```

- If the parameters are good enough, they are not modified
- It restarts the database for:
 - Changes to static parameters
 - Enabling the Archivelog mode

DB_FILES	= 1024
LOG_BUFFER	= 256M
DB_BLOCK_CHECKSUM	= TYPICAL
DB_LOST_WRITE_PROTECT	= TYPICAL
DB_FLASHBACK_RETENTION_TARGET	= 120
PARALLEL_THREADS_PER_CPU	= 1
STANDBY_FILE_MANAGEMENT	= AUTO
DG_BROKER_START	= TRUE

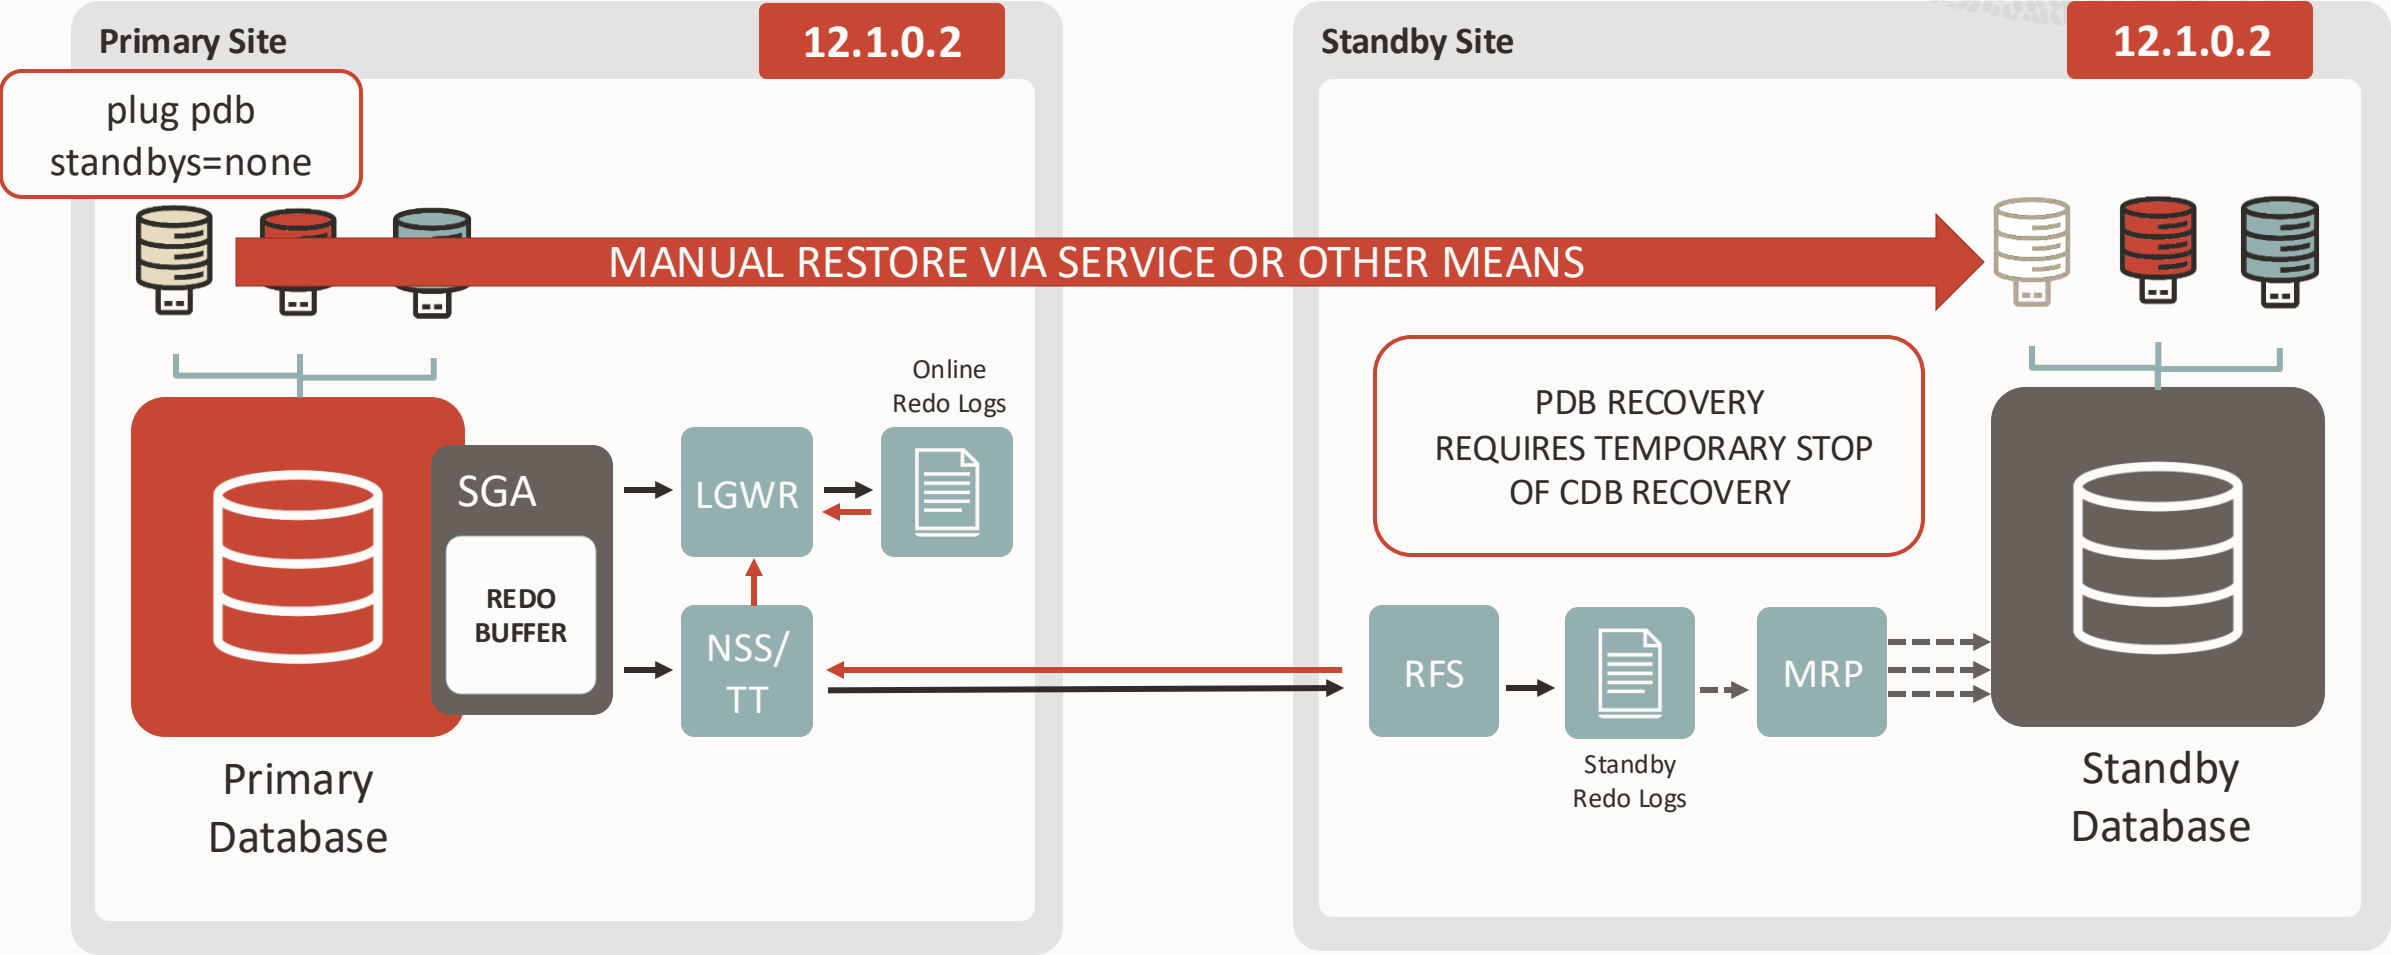
Pluggable Database Recovery Isolation

Simplified PDB cloning in Data Guard configurations



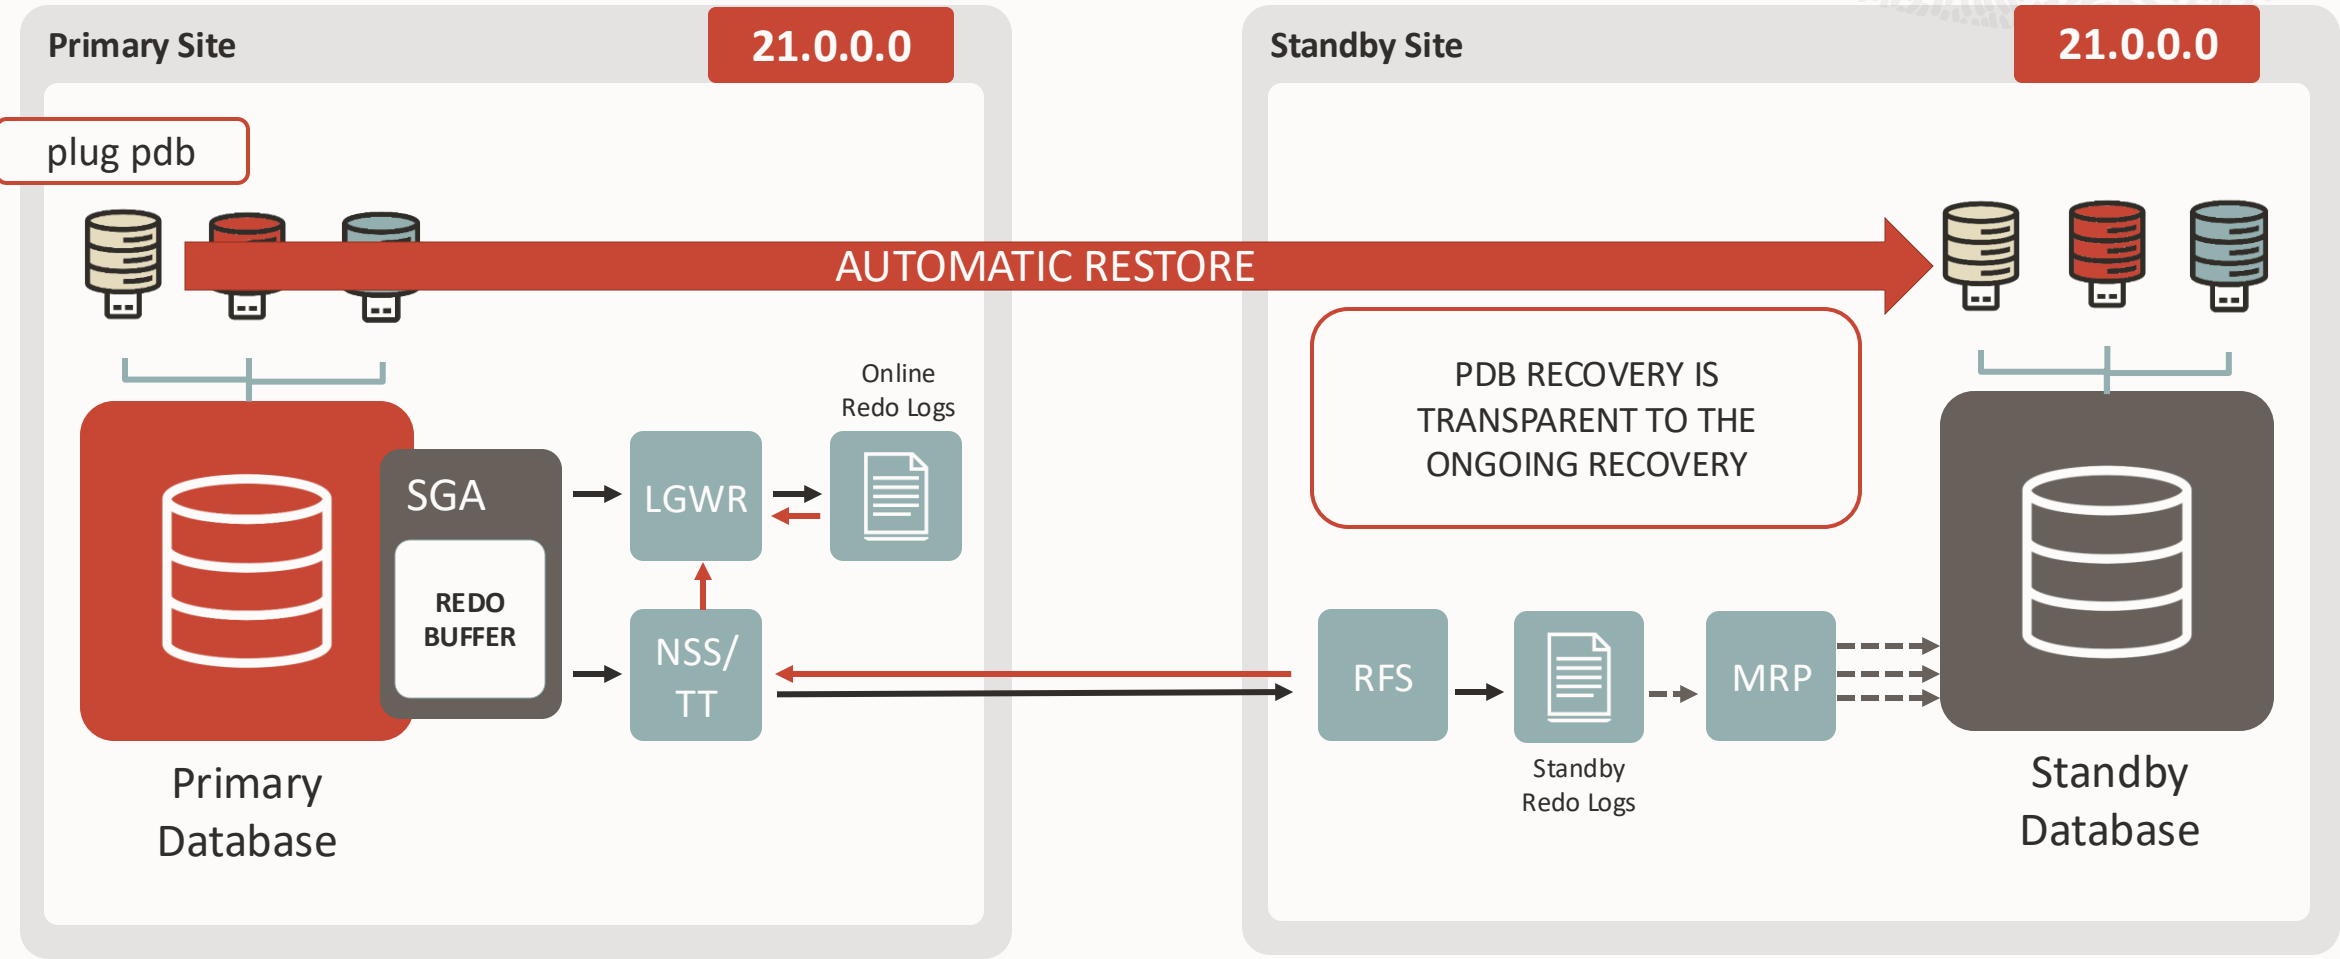
Pluggable Database Recovery Isolation

Simplified PDB cloning in Data Guard configurations



Pluggable Database Recovery Isolation

Simplified PDB cloning in Data Guard configurations



ORDS REST API for Data Guard management

Ready for modern DevOps deployment

New in ORDS 21.4 for 21c databases

```
POST /database/dataguard/configuration/
{
  "primary_connection_identifier": "site1-scan:1521/mydb",
  "primary_database": "mydb_site1"
}
```

Create the configuration

```
POST /database/dataguard/databases/
{
  "connection_identifier": "site2-scan:1521/mydb",
  "database_name": "mydb_site2"
}
```

Add the standby databases

```
PUT /database/dataguard/configuration/
{
  "operation": "ENABLE"
}
```

Enable the configuration

Oracle REST Data Services API - Data Guard REST Endpoints

<https://docs.oracle.com/en/database/oracle/oracle-rest-data-services/21.4/orrst/api-data-guard.html>

Data Guard management from SQLcl

Everything under control with a single command-line tool

New in SQLcl 22.1 for 21c databases

```
SQL> help dg
```

```
DG
```

```
-----
```


```
Run DG commands
```

```
DG ADD DATABASE "<database name>" AS CONNECT IDENTIFIER IS <connect identifier> [ INCLUDE CURRENT DESTINATIONS ];
DG CREATE CONFIGURATION "<config_name>" AS PRIMARY DATABASE IS <database name> CONNECT IDENTIFIER IS <connect_identifier>
    [ INCLUDE CURRENT DESTINATIONS ];
DG DISABLE CONFIGURATION;
DG DISABLE { DATABASE | RECOVERY_APPLIANCE | FAR_SYNC | MEMBER } <member name>;
DG EDIT CONFIGURATION SET PROPERTY <property name> = '<property value>';
DG EDIT { DATABASE | RECOVERY_APPLIANCE | FAR_SYNC | MEMBER } <member name> SET PROPERTY <property name> = '<property value>';
DG ENABLE CONFIGURATION;
DG ENABLE { DATABASE | RECOVERY_APPLIANCE | FAR_SYNC | MEMBER } <member name>;
DG FAILOVER TO <database name> [IMMEDIATE];
DG REINSTATE DATABASE <database name>;
DG REMOVE CONFIGURATION [PRESERVE DESTINATIONS];
DG REMOVE { DATABASE | RECOVERY_APPLIANCE | FAR_SYNC | MEMBER } <name> [PRESERVE DESTINATIONS];
DG SHOW CONFIGURATION [<property name>];
DG SHOW DATABASE <database name> [<property name>];
DG SWITCHOVER TO <database name> [WAIT [<timeout in seconds>]];
```

Other changes in Oracle Data Guard 21c

- **Far Sync can now be used with Fast-Start Failover in Max Performance mode (Active Data Guard)**
Primary can send redo asynchronously to Far Sync.
- **The broker configuration now supports up to four observers**
Before 21c, the limit was three observers.
- **The PreferredObserverHosts property now supports priorities**
Example: PreferredObserverHosts='host-a:1, host-b:2'
- **Properties deprecated in 19c are now **desupported****

ArchiveLagTarget	DbFileNameConvert	LsbyPreserveCommitOrder
DataGuardSyncLatency	LogArchiveFormat	LsbyRecordAppliedDdl
LogArchiveMaxProcesses	LogFileNameConvert	LsbyRecordSkipDdl
LogArchiveMinSucceedDest	LsbyMaxEventsRecorded	LsbyRecordSkipErrors
LogArchiveTrace	LsbyMaxServers	LsbyParameters
StandbyFileManagement	LsbyMaxSga	



Other **23ai** features for Data Guard and Broker

Different Ways to Configure Oracle Data Guard

dgmgrl

```
DGMGRL> create configuration mydb
> as primary database is mydb
> connect identifier is 'clu-scan:1521/mydb'
```



```
SQL> DG create configuration mydb as primary database
is mydb connect identifier is 'clu-scan:1521/mydb'
```

PL/
SQL

```
DECLARE
  severity BINARY_INTEGER;
  retcode  BINARY_INTEGER;
BEGIN
  retcode := DBMS_DG.CREATE_CONFIGURATION (
    config_name      => 'mydb'
    primary_ci       => 'clu-scan:1521/mydb'
    severity         => severity
  );
END;
```



```
POST /database/dataguard/configuration/
{
  "primary_connection_identifier": "clu-
scan:1521/mydb",
  "primary_database": "mydb_site1"
}
```

Easier Integration Thanks to Many SQL Additions

Some useful new views

```
SQL> select member, property, value from V$DG_BROKER_PROPERTY where value is not null;
```

MEMBER	PROPERTY	VALUE
mydb	FastStartFailoverThreshold	180
mydb	OperationTimeout	30
...		
mydb_site1	DGConnectIdentifier	mydb_site1
mydb_site1	FastStartFailoverTarget	mydb_site2
mydb_site1	LogShipping	ON
...		
mydb_site1	StaticConnectIdentifier	(DESCRIPTION=<...>))
mydb_site2	DGConnectIdentifier	mydb_site2
mydb_site2	FastStartFailoverTarget	mydb_site1
...		

66 rows selected.

```
SQL> desc V$FAST_START_FAILOVER_CONFIG;
```

Name	Null?	Type
FSFO_MODE		VARCHAR2(19)
STATUS		VARCHAR2(22)
CURRENT_TARGET		VARCHAR2(30)
THRESHOLD		NUMBER
OBSERVER_PRESENT		VARCHAR2(7)
OBSERVER_HOST		VARCHAR2(512)
PING_INTERVAL		NUMBER
PING_RETRY		NUMBER
PROTECTION_MODE		VARCHAR2(30)
LAG_LIMIT		NUMBER
AUTO_REINSTATE		VARCHAR2(5)
OBSERVER_RECONNECT		NUMBER
OBSERVER_OVERRIDE		VARCHAR2(5)
SHUTDOWN_PRIMARY		VARCHAR2(5)

```
SQL> select * from V$DG_BROKER_ROLE_CHANGE;
```

EVENT	STANDBY_TYPE	OLD_PRIMARY	NEW_PRIMARY	FS_FAILOVER_REASON	BEGIN_TIME	END_TIME
Failover	Physical	mydb1	mydb1b	Manual Failover	30-AUG-2024 19:01:14	30-AUG-2024 19:01:35
Switchover	Physical	mydb1b	mydb1		30-AUG-2024 19:04:53	30-AUG-2024 19:05:15
Switchover	Physical	mydb1	mydb1b		30-AUG-2024 20:51:38	30-AUG-2024 20:52:03
Failover	Physical	mydb1b	mydb1	Manual Failover	30-AUG-2024 20:52:46	30-AUG-2024 20:53:04
Switchover	Logical	mydb1d	mydb1		30-AUG-2024 20:35:27	30-AUG-2024 20:35:48
Fast-Start Failover	Physical	mydb1	mydb1b	Primary Disconnected	30-AUG-2024 20:13:51	30-AUG-2024 20:14:53



Strict Database Validation

More checks, better explanations, increased operational security

(*) available in 23.6

```
VALIDATE DATABASE [VERBOSE] <database>  
[ STRICT { ALL | APPLY_PROPERTY | DATAFILES_OFFLINE (*) | FLASHBACK | FORCE_LOGGING | LOG_FILES_CLEARED  
| LOG_FILE_CONFIGURATION | PDBS_OFFLINE (*) | PDB_SAVE_STATE (*) | TRANSPORT_PROPERTY } ];
```

```
DGMGRL> validate database chicago strict all;  
DGM-17567: Current database session was authenticated using operating system credentials.
```

```
Database Role:      Physical standby database  
Primary Database:   boston
```

```
Ready for Switchover: No  
The primary or standby database does not have flashback database enabled. (*)
```

```
Ready for Failover:   Yes (Primary Running)
```

```
Flashback Database Status:
```

Database	Status	Retention Target
boston	Off	1440
chicago	Off	1440

```
...
```

Switchover and Failover Readiness in 23.5

Checking if the database is ready for a role transition is as easy as selecting a column

Two new columns, SWITCHOVER_READY and FAILOVER_READY, computed every minute by the broker.

```
SQL> select database, dataguard_role, status, severity, switchover_ready, failover_ready, transport_mode
2> from v$dg_broker_config;
```

DATABASE	DATAGUARD_ROLE	STATUS	SEVERITY	SWITCHOVER_READY	FAILOVER_READY	TRANSPORT_MODE
boston	PRIMARY	0	SUCCESS	YES	UNKNOWN	-N/A-
chicago	PHYSICAL STANDBY	0	SUCCESS	YES	YES	ASync

The checks done by the broker are a superset of the ALTER DATABASE SWITCHOVER VERIFY command:

```
SQL> alter database switchover to chicago verify;
alter database switchover to chicago verify
*
ERROR at line 1:
ORA-16470: Redo Apply is not running on switchover target
```



PL/SQL API for Data Guard broker management

Manage Data Guard configurations from any SQL*Net connection

30+ new functions in DBMS_DG PL/SQL package

```
DECLARE
    severity BINARY_INTEGER;
    retcode BINARY_INTEGER;
BEGIN

    retcode := DBMS_DG.CREATE_CONFIGURATION (
        config_name      => 'mydb'
        primary_ci       => 'site1-scan:1521/mydb'
        severity         => severity
    );

    IF retcode != 0 THEN
        /* handle error code */
    END IF;

    retcode := DBMS_DG.ADD_DATABASE (
        database_name    => 'mydb_site2'
        database_ci      => 'site2-scan:1521/mydb'
        severity         => severity
    );

END;
```

Create the configuration

Add the standby databases

PL/SQL Packages and Types Reference - DBMS_DG

https://docs.oracle.com/en/database/oracle/oracle-database/23/arpls/DBMS_DG.html


Easier checking of Data Guard configurations

The new fixed view **V\$DG_BROKER_PROPERTY** contains the properties of the configuration and all the members

```
SQL> select member, property, value from V$DG_BROKER_PROPERTY where value is not null;
```

MEMBER	PROPERTY	VALUE
-----	-----	-----
mydb	FastStartFailoverThreshold	180
mydb	OperationTimeout	30
mydb	TraceLevel	USER
mydb	FastStartFailoverLagLimit	300
mydb	CommunicationTimeout	180
mydb	ObserverReconnect	0
mydb	ObserverPingInterval	0
mydb	ObserverPingRetry	0
mydb	FastStartFailoverAutoReinstate	TRUE
mydb	FastStartFailoverPmyShutdown	TRUE
...		
mydb_site1	DGConnectIdentifier	mydb_site1
mydb_site1	FastStartFailoverTarget	mydb_site2
mydb_site1	LogShipping	ON
mydb_site1	LogXptMode	ASYNC
mydb_site1	DelayMins	0
...		
mydb_site1	StaticConnectIdentifier	(DESCRIPTION=<...>)))
mydb_site1	TopWaitEvents	(monitor)
mydb_site1	SidName	(monitor)
mydb_site2	DGConnectIdentifier	mydb_site2
mydb_site2	FastStartFailoverTarget	mydb_site1
...		

66 rows selected.



New command: **VALIDATE DGConnectIdentifier**

Check network resolution, connectivity, password, service name from the database

```
DGMGRL> validate dgconnectidentifier mydb_site2;
```

```
At instance 'mydb' of member 'mydb_site1'
```

```
Connect Descriptor:
```

```
(DESCRIPTION=(ADDRESS_LIST=(ADDRESS=(PROTOCOL=tcp)(HOST=host2)(PORT=1521)))(CONNECT_DATA=(SERVICE_NAME=mydb_site2.mydomain)(SERVER=DEDICATED)))
```

```
Environment Variables:
```

```
TNS_ADMIN: /u01/app/oracle/product/23.1.0.0/network/admin
```

```
ORACLE_HOME: /u01/app/oracle/product/23.1.0.0
```

```
ORACLE_BASE: /u01/app/oracle
```

```
Initialization Parameters:
```

```
LOCAL_LISTENER: host1:1521
```

```
Connected to instance 'mydb' at member 'mydb_site2'
```

```
At instance 'mydb' of member 'mydb_site2'
```

```
Connect Descriptor:
```

```
(DESCRIPTION=(ADDRESS_LIST=(ADDRESS=(PROTOCOL=tcp)(HOST=host2)(PORT=1521)))(CONNECT_DATA=(SERVICE_NAME=mydb_site2.mydomain)(SERVER=DEDICATED)))
```

```
Environment Variables:
```

```
TNS_ADMIN: /u01/app/oracle/product/23.1.0.0/network/admin
```

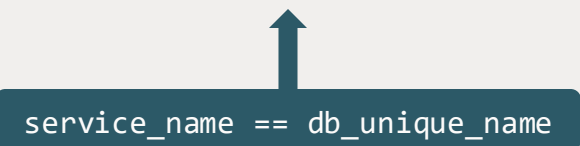
```
ORACLE_HOME: /u01/app/oracle/product/23.1.0.0
```

```
ORACLE_BASE: /u01/app/oracle
```

```
Initialization Parameters:
```

```
LOCAL_LISTENER: host2:1521
```

```
Connected to instance 'mydb' at member 'mydb_site2'
```



service_name == db_unique_name

New commands: SHOW|EDIT ALL MEMBERS

Easier management of member's properties and parameters

-- DATA GUARD PROPERTIES

```
DGMGRL> SHOW ALL MEMBERS logxptmode
```

```
mydb_site1: logxptmode = 'ASYNC'
```

```
mydb_site2: logxptmode = 'ASYNC'
```

```
DGMGRL> EDIT ALL MEMBERS SET PROPERTY logxptmode = 'SYNC';
```

```
Property "logxptmode" updated for member "mydb_site1".
```

```
Property "logxptmode" updated for member "mydb_site2".
```

```
DGMGRL> SHOW ALL MEMBERS logxptmode
```

```
mydb_site1: logxptmode = 'SYNC'
```

```
mydb_site2: logxptmode = 'SYNC'
```

-- DB PARAMETERS

```
DGMGRL> SHOW ALL MEMBERS PARAMETER fast_start_mttr_target
```

```
mydb_site1: fast_start_mttr_target = '0'
```

```
mydb_site2: fast_start_mttr_target = '0'
```

```
DGMGRL> EDIT ALL MEMBERS SET PARAMETER fast_start_mttr_target=15;
```

```
Parameter "fast_start_mttr_target" updated for member "mydb_site1".
```

```
Parameter "fast_start_mttr_target" updated for member "mydb_site2".
```

```
DGMGRL> SHOW ALL MEMBERS PARAMETER fast_start_mttr_target
```

```
mydb_site1: fast_start_mttr_target = '15'
```

```
mydb_site2: fast_start_mttr_target = '15'
```

Get and set Data Guard broker properties

Get and set Database parameters

Automatic tempfile creation on the standby database

Temporary tablespace creation during recovery:



PRIMARY	STANDBY	Non-OMF	OMF
Non-OMF		Standby_file_management AUTO: ✓ Creates one tempfile with the default size using db_file_name_convert. Standby_file_management MANUAL: ✗ Does not create a tempfile.	Standby_file_management AUTO: ✓ Creates one tempfile with the default size with OMF naming. Standby_file_management MANUAL: ✗ Does not create a tempfile.(?)
OMF		✗ Does not create a tempfile.	✓ Creates one tempfile with the default size with OMF naming.



When the standby opens and a temporary tablespace has no tempfiles:

Non-OMF	✗ Does not create a tempfile.
OMF	✓ Creates one tempfile with the default size with OMF naming.





Questions & Answers

Thank you



ORACLE